

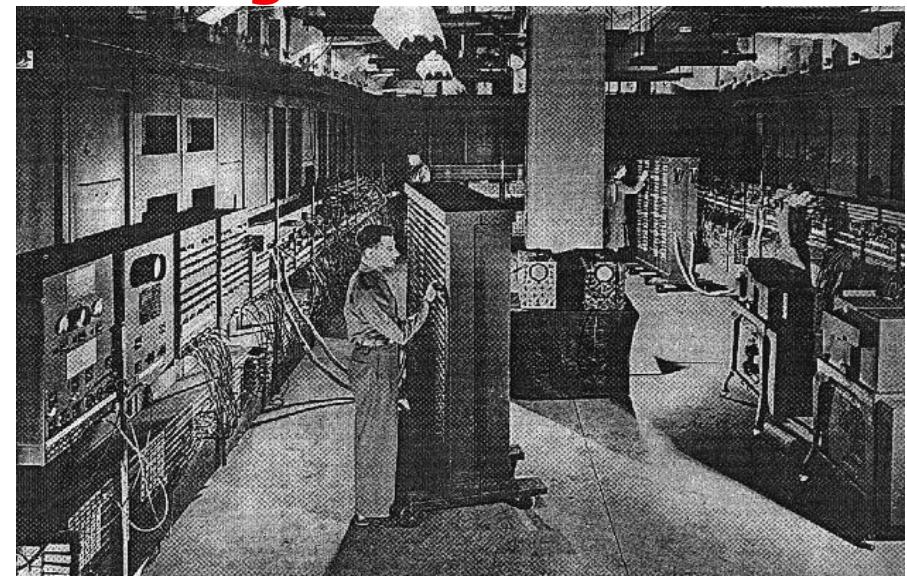
Course Overview

Why Study Computer Architecture?

- **Understand where computers are going**
 - Future capabilities drive the (computing) world
 - Real world-impact: no computer architecture → no computers!
- **Understand high-level design concepts**
 - The best architects understand all the levels
 - Devices, circuits, architecture, compiler, applications
- **Understand computer performance**
 - learn valid experimental methodologies
- **Write better software**
 - The best software designers also understand hardware
 - Need to understand hardware to write fast software
- **Design hardware**
 - At Intel, AMD, IBM, ARM, Qualcomm, Oracle, NVIDIA, Samsung

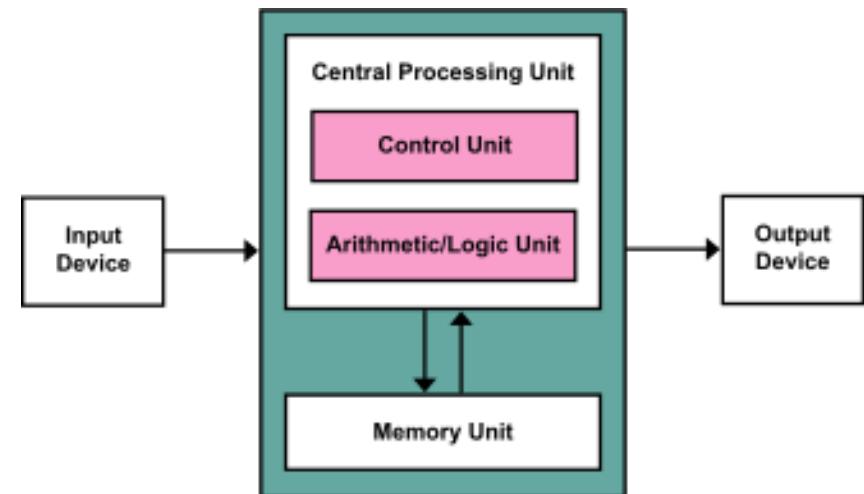
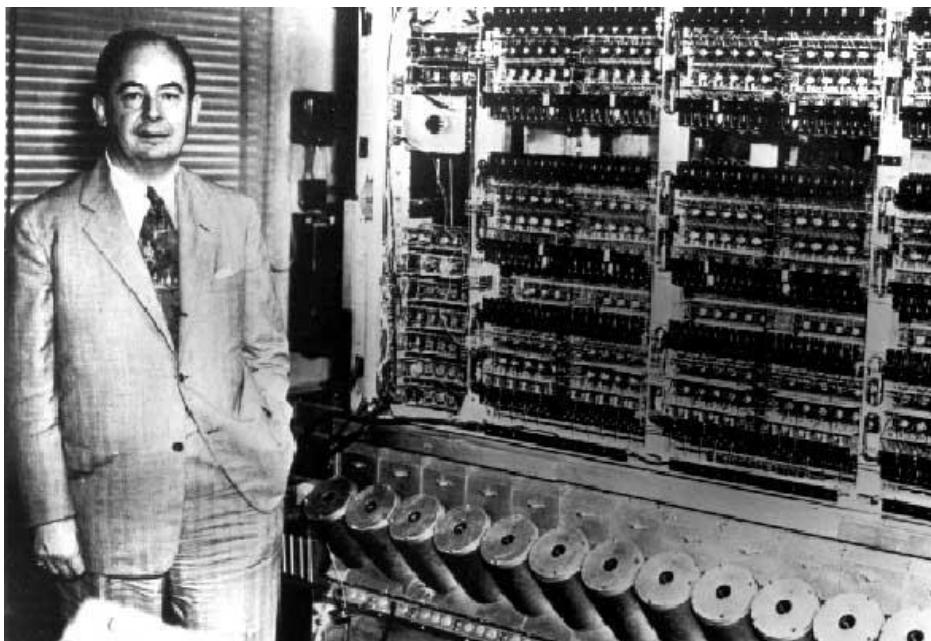
ENIAC

- **ENIAC:** electronic numerical integrator and calculator
 - First operational general-purpose stored-program computer
 - Designed and built here by Eckert and Mauchly from 1943-6
 - ENIAC was designed to calculate artillery firing tables for the US Army's Ballistic Research Laboratory (BRL).
 - Von Neumann architecture? (No)
- **First seminars on computer design**
 - Moore School Lectures, 1946
 - "*Theory and Techniques for Design of Electronic Digital Computers*"

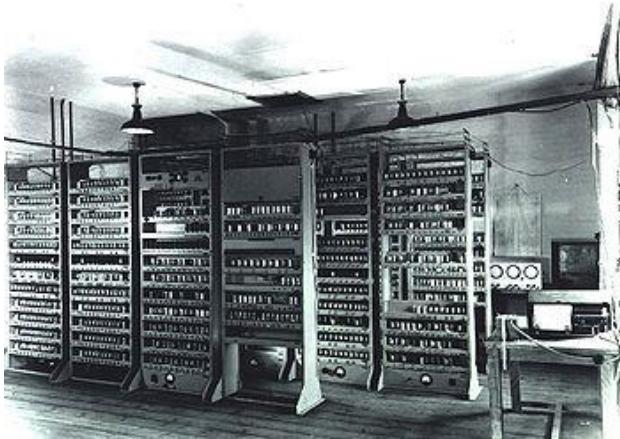


Von Neumann

- The Von Neumann architecture, also known as the Von Neumann model and Princeton architecture.



Von Neumann architecture computer



- **EDSAC (Electronic Delay Storage Automatic Calculator)** was the first practical stored-program electronic computer (May 1949)
- **EDVAC (Electronic Discrete Variable Automatic Computer)**
 - EDVAC's construction in August 1944. EDVAC was delivered to the Ballistics Research Laboratory in August 1949

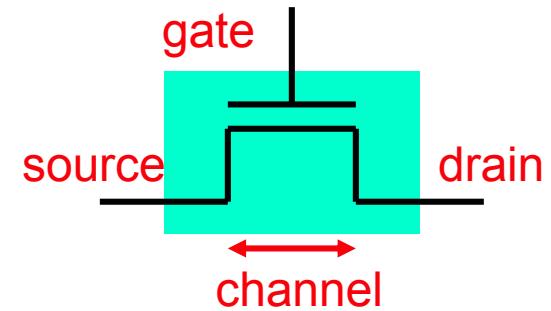
Application Specific Designs

- This class is about **general-purpose CPUs**
 - Processor that can do anything, run a full OS, etc.
 - E.g., Intel Core i7, AMD Athlon, IBM Power, ARM A-series
- In contrast to **application-specific chips**
 - Or **ASICs** (Application specific integrated circuits)
 - Also application-domain specific processors
 - Implement critical domain-specific functionality in hardware
 - Examples: video encoding, 3D graphics
 - General rules
 - Hardware is less flexible than software
 - + Hardware more effective (speed, power, cost) than software
 - + Domain specific more “parallel” than general purpose
 - But general mainstream processors becoming more parallel
- Trend: from **specific** to **general** (for a specific domain)

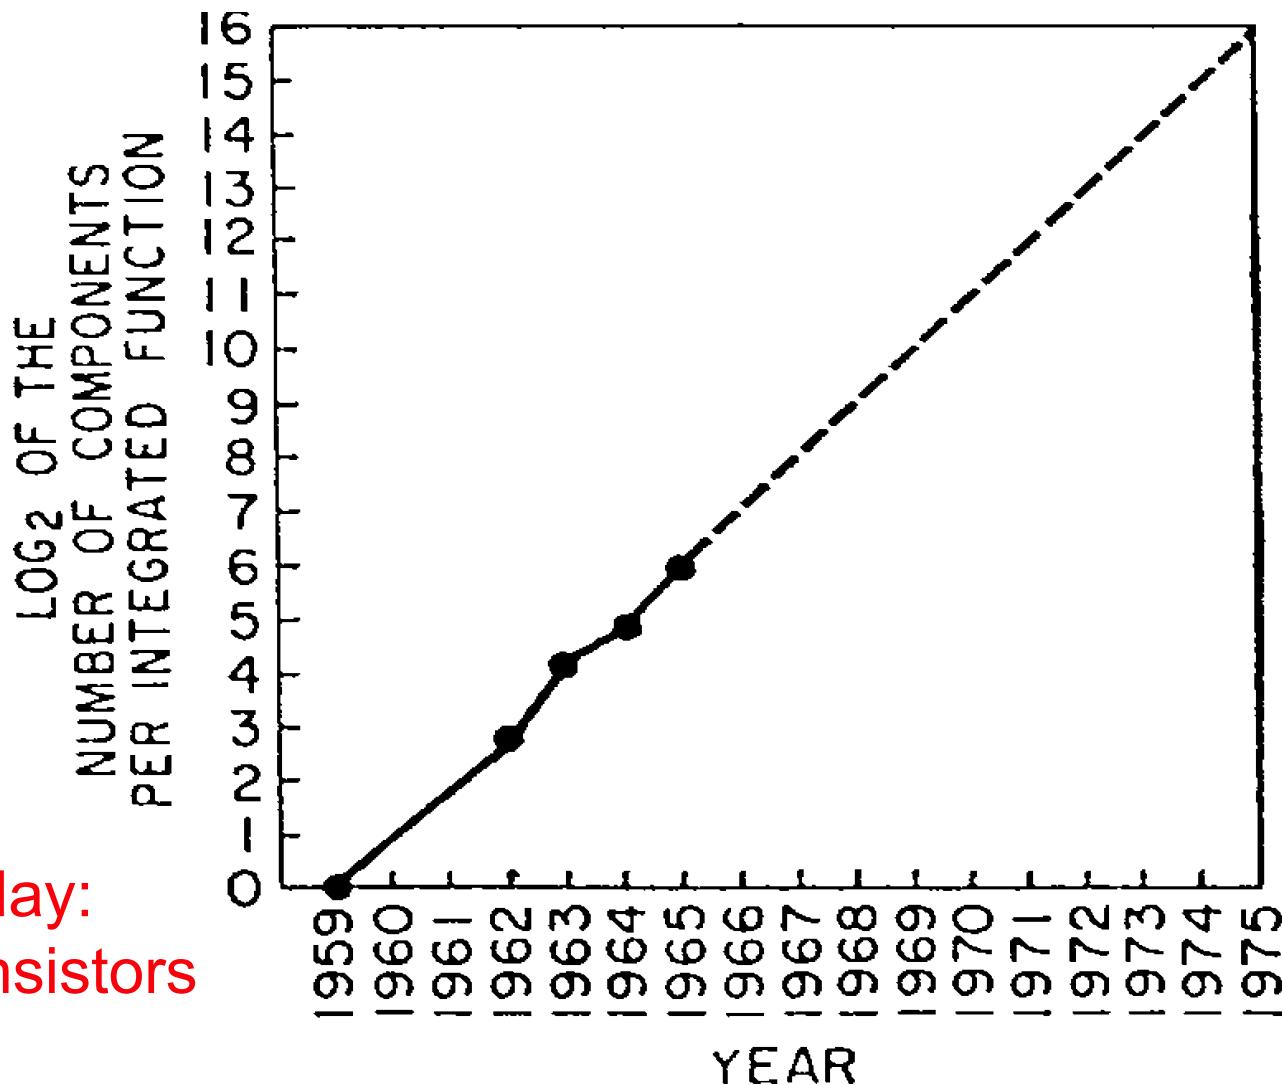
Technology Trends

“Technology”

- Basic element
 - Solid-state **transistor**
 - Building block of **integrated circuits (ICs)**
- What's so great about ICs? Everything
 - + High performance, high reliability, low cost, low power
 - + Lever of mass production
- Several kinds of integrated circuit families
 - **SRAM/logic**: optimized for speed (used for processors)
 - **DRAM**: optimized for density, cost, power (used for memory)
 - **Flash**: optimized for density, cost (used for disks)
 - Increasing opportunities for integrating multiple technologies
- Non-transistor storage and inter-connection technologies
 - Disk, optical storage, ethernet, fiber optics, wireless



Moore's Law - 1965



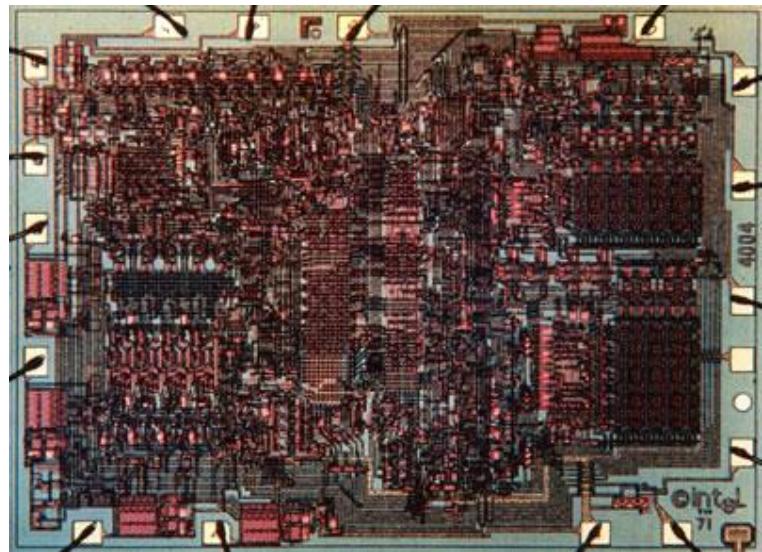
Today:
 2^{32} transistors

Revolution I: The Microprocessor

- **Microprocessor revolution**
 - One significant technology threshold was crossed in 1970s
 - Enough transistors ($\sim 25K$) to put a 16-bit processor on one chip
 - Huge performance advantages: fewer slow chip-crossings
 - Even bigger cost advantages: one “stamped-out” component
- Microprocessors have allowed new market segments
 - Desktops, CD/DVD players, laptops, game consoles, set-top boxes, mobile phones, digital camera, mp3 players, GPS, automotive
- And replaced incumbents in existing segments
 - Microprocessor-based system replaced supercomputers, “mainframes”, “minicomputers”, etc.

First Microprocessor

- Intel 4004 (1971)
 - Application: calculators
 - Technology: 10000 nm
 - 2300 transistors
 - 13 mm²
 - 108 KHz
 - 12 Volts
 - 4-bit data
 - Single-cycle datapath



Tracing the Microprocessor Revolution

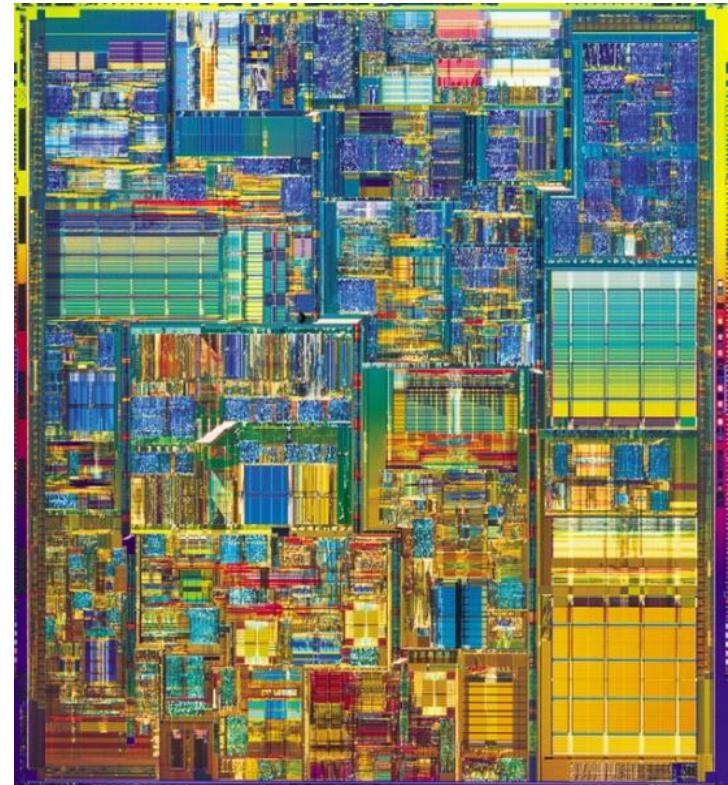
- How were growing transistor counts used?
- Initially to widen the datapath
 - 4004: 4 bits → AMD Opteron (2003): 64 bits
- ... and also to add more powerful instructions
 - To amortize overhead of fetch and decode
 - To simplify assembly programming (which was done by hand then)

Revolution II: Implicit Parallelism

- Then to **extract implicit instruction-level parallelism**
 - Hardware provides parallel resources, figures out how to use them
 - Software is oblivious
- Initially using pipelining ...
 - Which also enabled increased clock frequency
- ... caches ...
 - Which became necessary as processor clock frequency increased
- ... and integrated floating-point
- Then deeper pipelines and branch speculation
- Then multiple instructions per cycle (superscalar)
- Then dynamic scheduling (out-of-order execution)
- We will discuss this in detail

Pinnacle of Single-Core Microprocessors

- Intel Pentium4 (2003)
 - Application: desktop/server
 - Technology: 90nm (1/100th of 4004)
 - 55M transistors (20,000x)
 - 101 mm² (10x)
 - 3.4 GHz (10,000x)
 - 1.2 Volts (1/10th)
 - 32/64-bit data (16x)
 - 22-stage pipelined datapath
 - 3 instructions per cycle (superscalar)
 - Two levels of on-chip cache
 - data-parallel vector (SIMD) instructions, hyperthreading

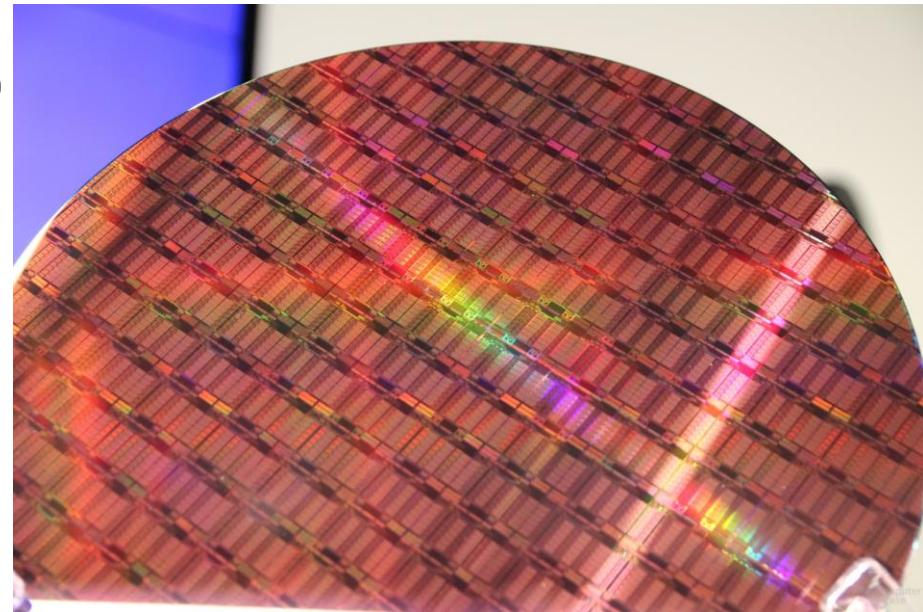


Revolution III: Explicit Parallelism

- Then to support **explicit data & thread level parallelism**
 - Hardware provides parallel resources, software specifies usage
 - Why? diminishing returns on instruction-level-parallelism
- First using (subword) vector instructions..., Intel's SSE
 - One instruction does four parallel multiplies
- ... and general support for multi-threaded programs
 - Coherent caches, hardware synchronization primitives
- Then using support for multiple concurrent threads on chip
 - First with single-core multi-threading, now with multi-core

Modern Multicore Processor

- Intel Xeon E5-2699 V4 (2016)
 - Application: server
 - Technology: 14nm (16% of P4)
 - 7.2B transistors (130x)
 - 456 mm² (4.5x)
 - 2.4 to 3.6 Ghz (~1x)
 - 1.8 Volts (~1x)
 - 256-bit data (2x)
 - 14-stage pipelined datapath (0.5x)
 - 4 instructions per cycle (1x)
 - Three levels of on-chip cache
 - data-parallel vector (SIMD) instructions, hyperthreading
 - **22-core multicore** (22x)



Revolution IV: Heterogeneous Processing

- Combining **multiple** kinds of compute engines in one die
 - not just homogenous collection of cores
 - System-on-Chip (SoC) is one common example in mobile space
- Lots of stuff on the chip beyond just CPUs
 - Graphics Processing Units (GPUs)
 - throughput-oriented specialized multicore processors
 - good for gaming, machine learning, computer vision, ...
 - Special-purpose logic
 - media codecs, cell phone radio, encryption, compression
- Excellent energy efficiency and performance
 - extremely complicated to program!

Technology Disruptions

- Classic examples:
 - The transistor
 - Microprocessor
- More recent examples:
 - Multicore processors
 - Flash-based solid-state storage
- Near-term potentially disruptive technologies:
 - Non-volatile memory (e.g., phase-change memory)
 - Chip stacking (aka “3D die stacking”)
- Disruptive “end-of-scaling”
 - “If something can’t go on forever, it must stop eventually”
 - Can we continue to shrink transistors for ever?
 - Transistor energy efficiency is not improving as rapidly

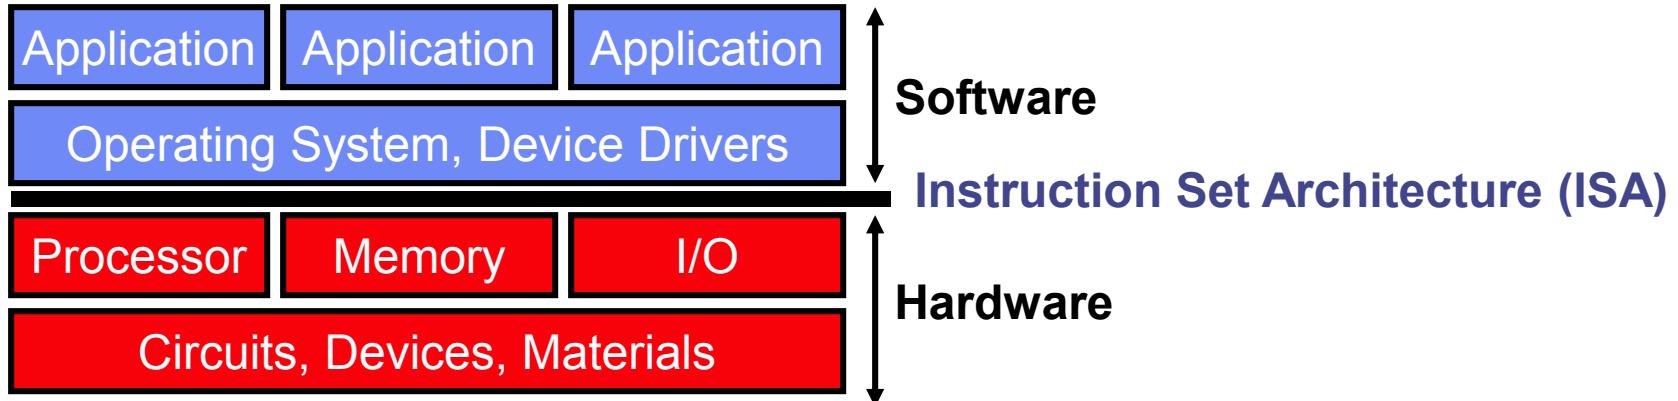
Managing This Mess

- Architect must consider all factors
 - Goals/constraints, applications, implementation technology
- Questions
 - How to deal with all of these inputs?
 - How to manage changes?
- Answers
 - Accrued institutional knowledge (stand on each other's shoulders)
 - Experience, rules of thumb
 - Discipline: clearly defined end state, keep your eyes on the ball
 - **Abstraction and layering**

Pervasive Ideas: Abstraction and Layering

- **Abstraction**: only way of dealing with complex systems
 - Divide world into objects, each with an...
 - **Interface**: knob(s)
 - **Implementation**: “black box”
 - Only specialists deal with implementation, rest of us with interface
- **Layering**: repeated abstraction makes life even simpler
 - Divide objects in system into layers, objects at layer n :
 - are implemented using interfaces of layer $n - 1$
 - don't need to know interfaces of layer $n - 2$ (sometimes helps)
- **Inertia**: the dark side of layering
 - Layer interfaces become entrenched over time (“standards”)
 - Very difficult to change even if benefit is clear (example: Digital TV)
- **Opacity**: hard to reason about performance across layers

Abstraction, Layering, and Computers



- **Computer architecture**
 - Definition of **ISA** to facilitate implementation of software layers
- This course mostly on **computer micro-architecture**
 - Design of chip to implement **ISA**
- Touch on compilers & OS ($n + 1$), circuits ($n - 1$) as well

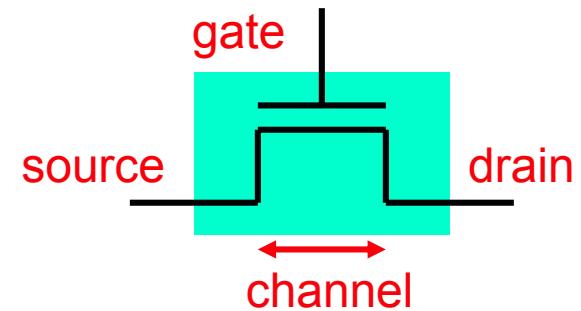
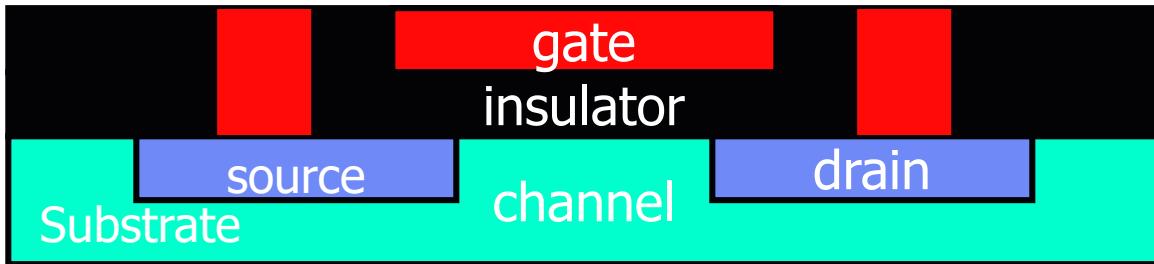
Technology & Energy

Technology & Energy

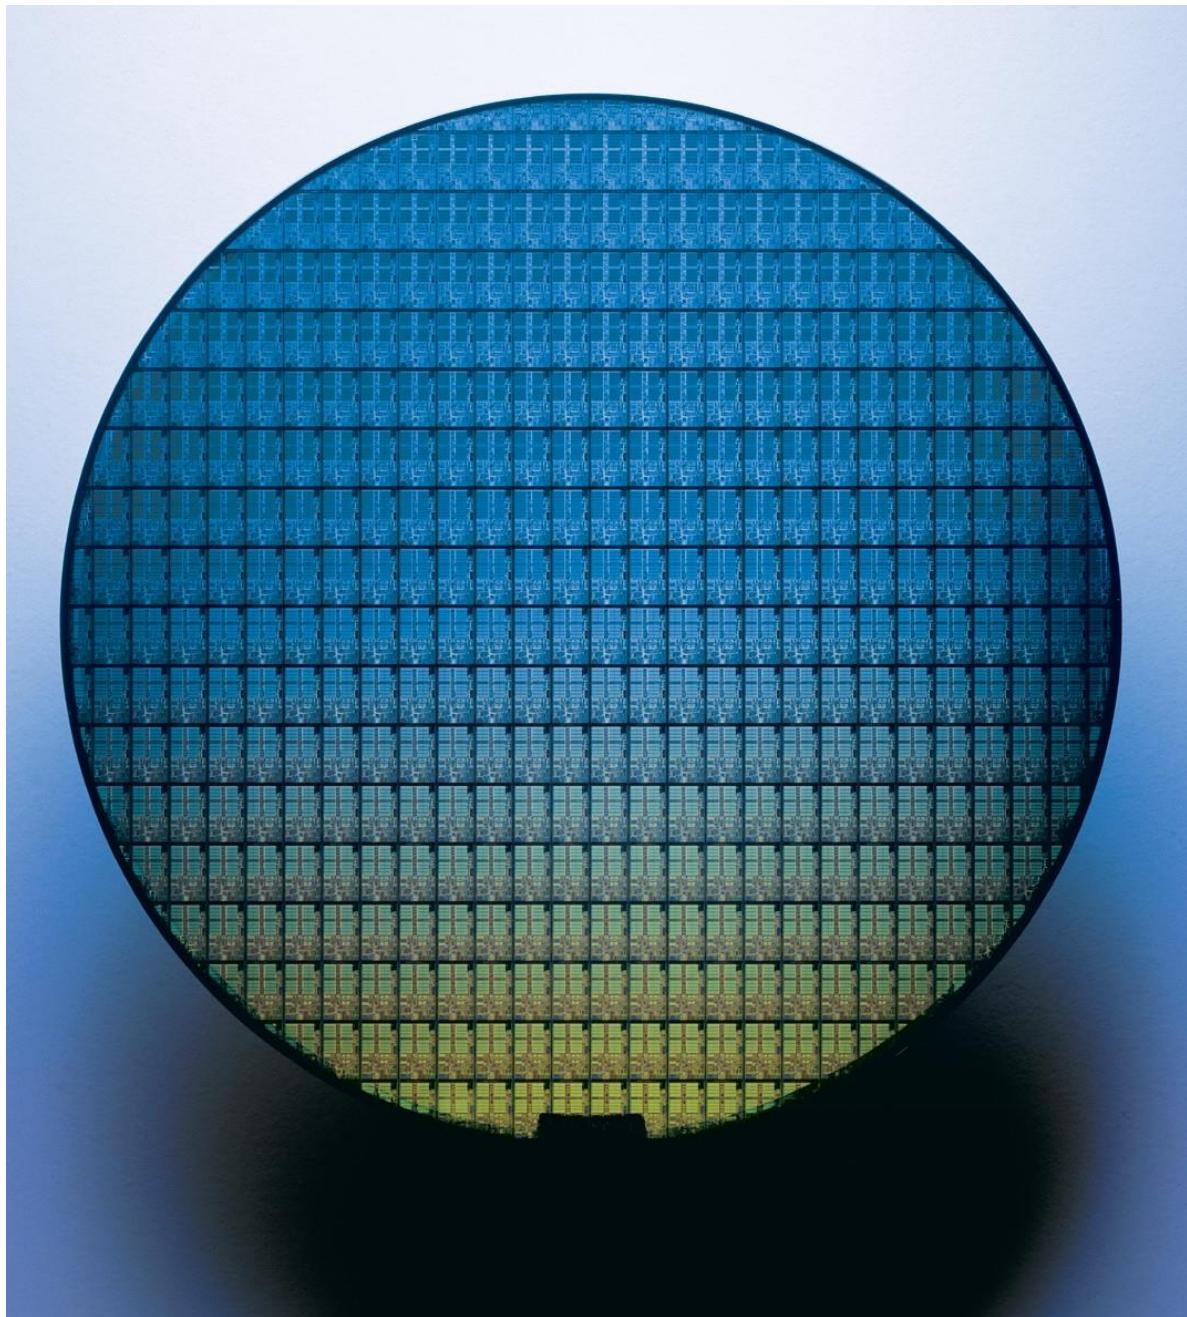
- Technology basis
 - Fabrication (manufacturing) & cost
 - Transistors & wires
 - Implications of transistor scaling (Moore's Law)
- Energy & power

Technology & Fabrication

Semiconductor Technology



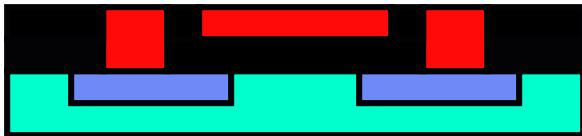
- Basic technology element: **MOSFET**
 - Solid-state component acts like electrical switch
 - **MOS**: metal-oxide-semiconductor
 - Conductor, insulator, semi-conductor
- **FET**: field-effect transistor
 - Channel conducts source→drain only when voltage applied to gate
- **Channel length**: characteristic parameter (short → fast)
 - Aka “feature size” or “technology node”
 - Currently: 14 nanometers (nm)
 - Continued miniaturization (scaling) known as “**Moore’s Law**”
 - Won’t last forever, physical limits approaching (or are they?)



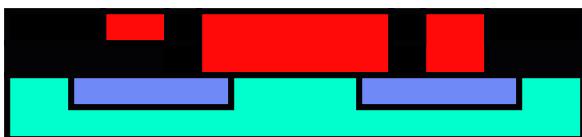
Intel
Pentium M
Wafer

Manufacturing Defects

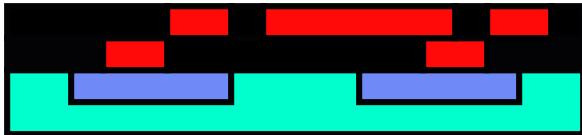
Correct:



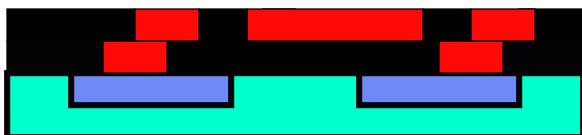
Defective:



Defective:



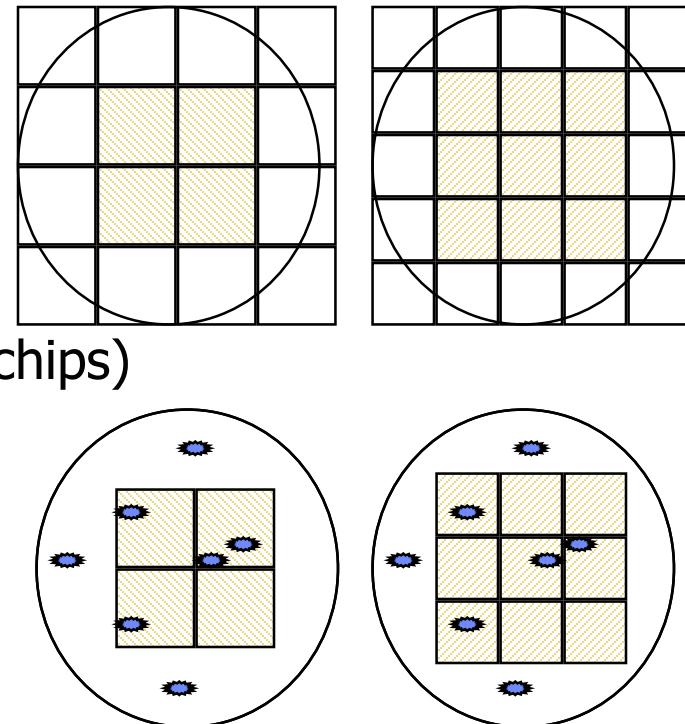
Slow:



- Defects can arise
 - Under-/over-doping
 - Over-/under-dissolved insulator
 - Mask mis-alignment
 - Particle contaminants
- Try to minimize defects
 - Process margins
 - Design rules
 - Minimal transistor size, separation
- Or, tolerate defects
 - Redundant or “spare” memory cells
 - Can substantially improve yield

Cost Implications of Defects

- Chips built in multi-step chemical processes on **wafers**
 - Cost / wafer is constant, $f(\text{wafer size, number of steps})$
- Chip (die) cost is related to **area**
 - Larger chips means fewer of them
- Cost is **superlinear** in area
 - Why? random defects
 - **Larger chip, more chance of defect**
 - Result: lower “yield” (fewer working chips)
- **Wafer yield**: % wafer that is chips
- **Die yield**: % chips that work
- Yield is increasingly non-binary - fast vs slow chips



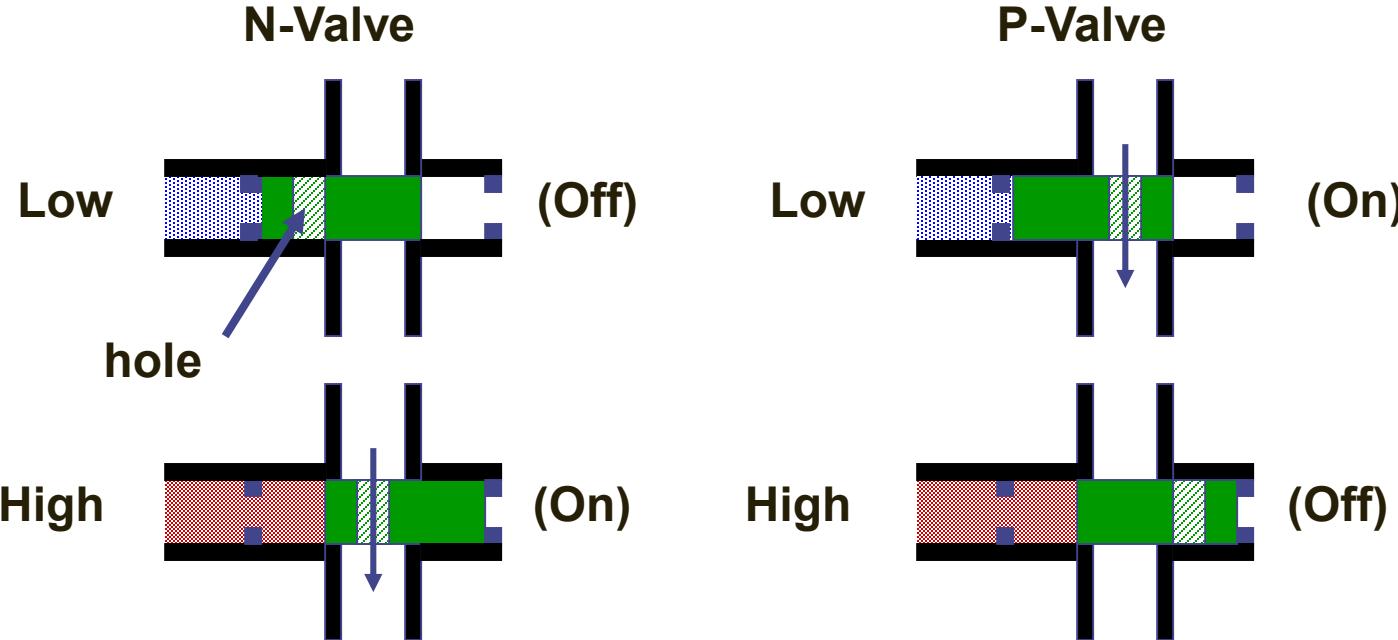
Manufacturing Cost

- **Chip cost vs system cost**
 - Cost of memory, storage, display, battery, etc.
- **Cost vs price**
 - Relationship complicated; **microprocessors not commodities**
 - Specialization, compatibility, different cost/performance/power
 - Economies of scale
- **Unit costs:** die manufacturing, testing, packaging, burn-in
 - Die cost based on area & defect rate (yield)
 - Package cost related to heat dissipation & number of pins
- **Fixed costs:** design & verification, fab cost
 - Amortized over “proliferations”, e.g., Core i3, i5, i7 variants
 - Building new “fab” costs billions of dollars today
 - Both getting worse; trend toward “foundry” & “fabless” models

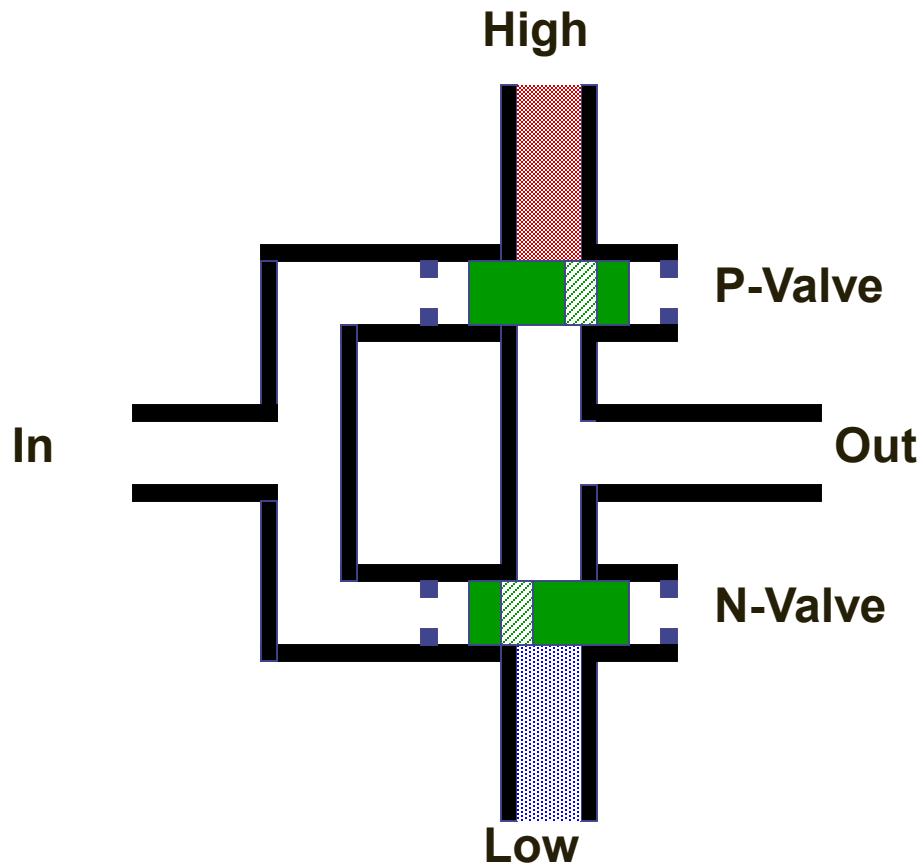
Transistor Switching Speed

A Transistor Analogy: Computing with Air

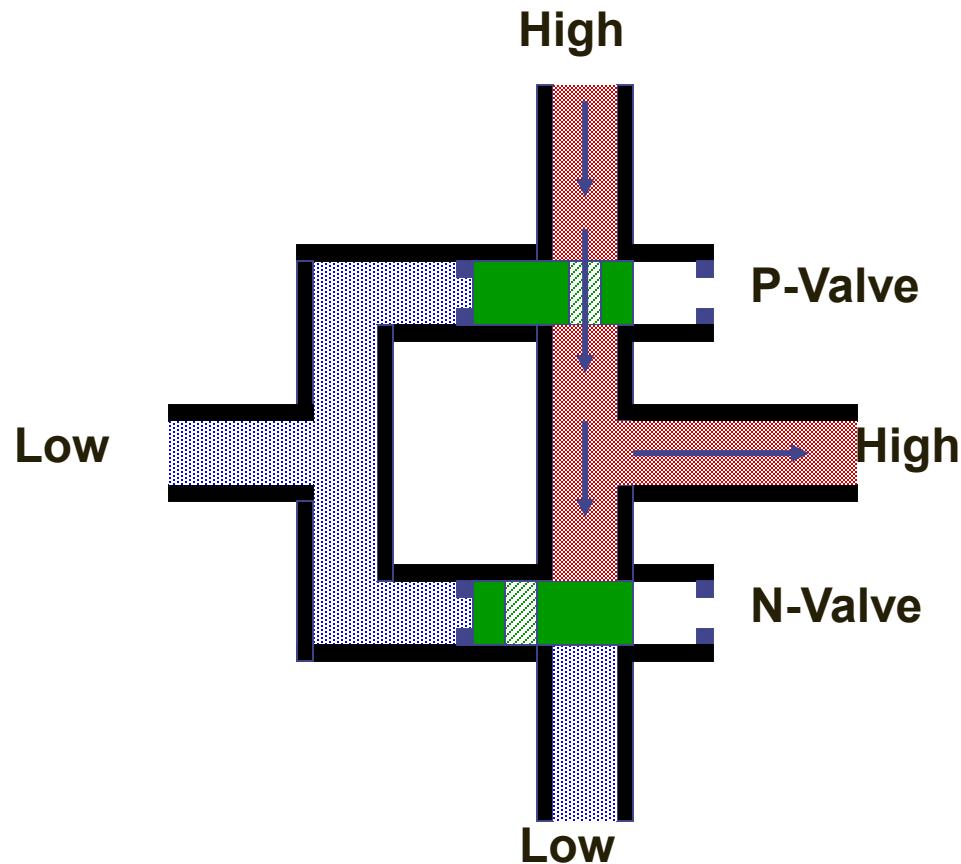
- Use air pressure to encode values
 - High pressure represents a “1” (blow)
 - Low pressure represents a “0” (suck)
- Valve can allow or disallow the flow of air
 - Two types of valves



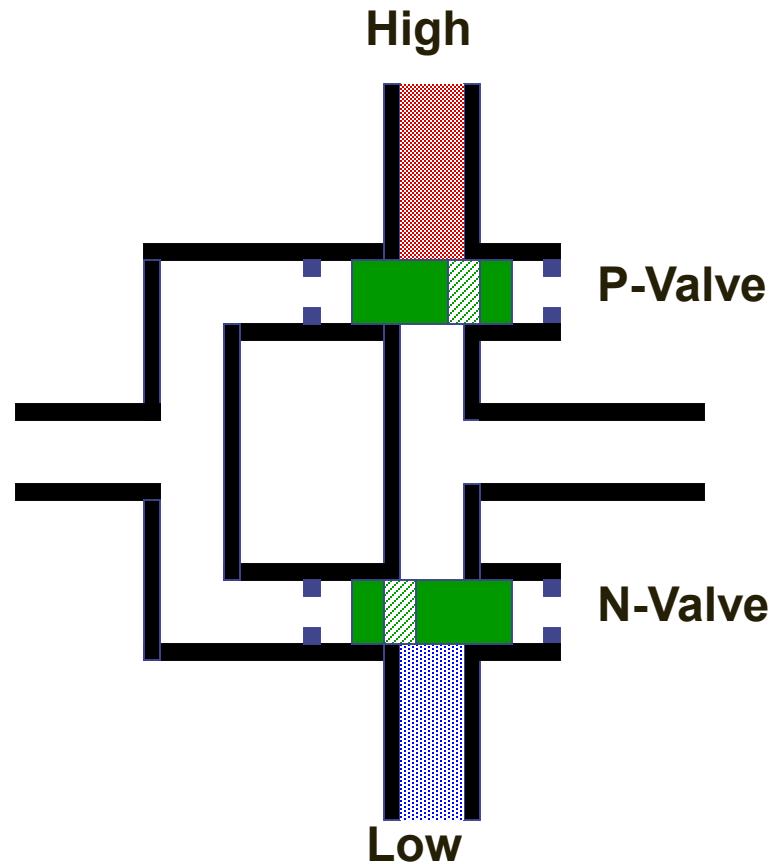
Pressure Inverter



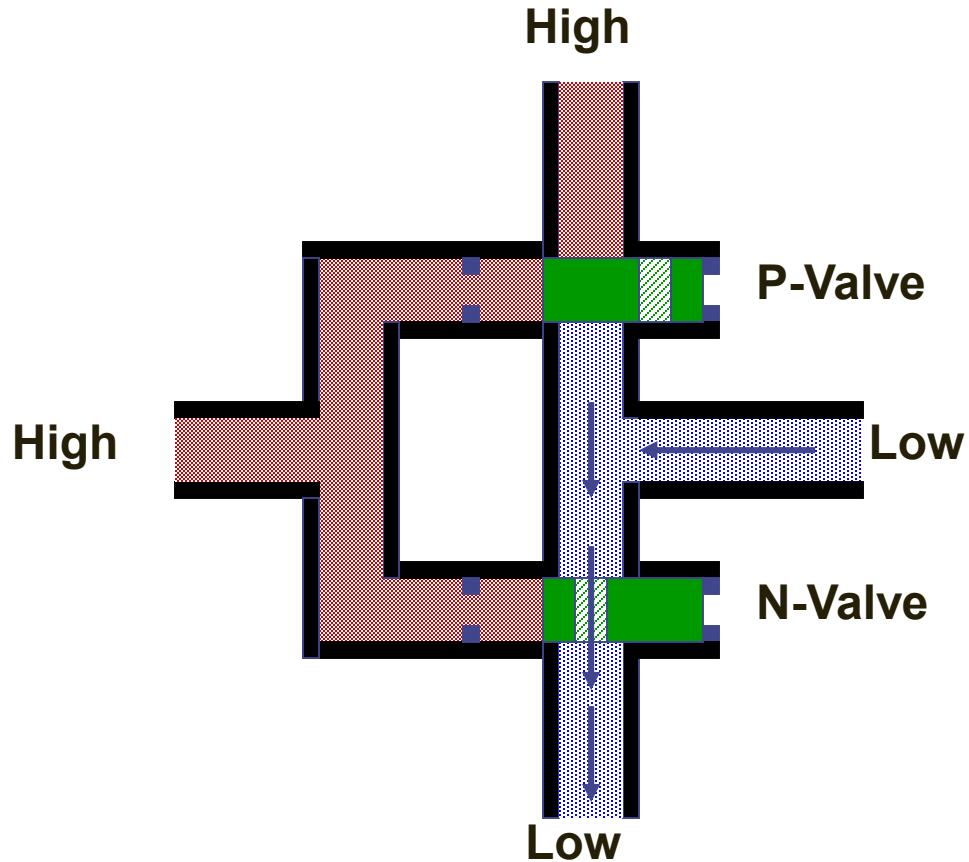
Pressure Inverter (Low to High)



Pressure Inverter



Pressure Inverter (High to Low)

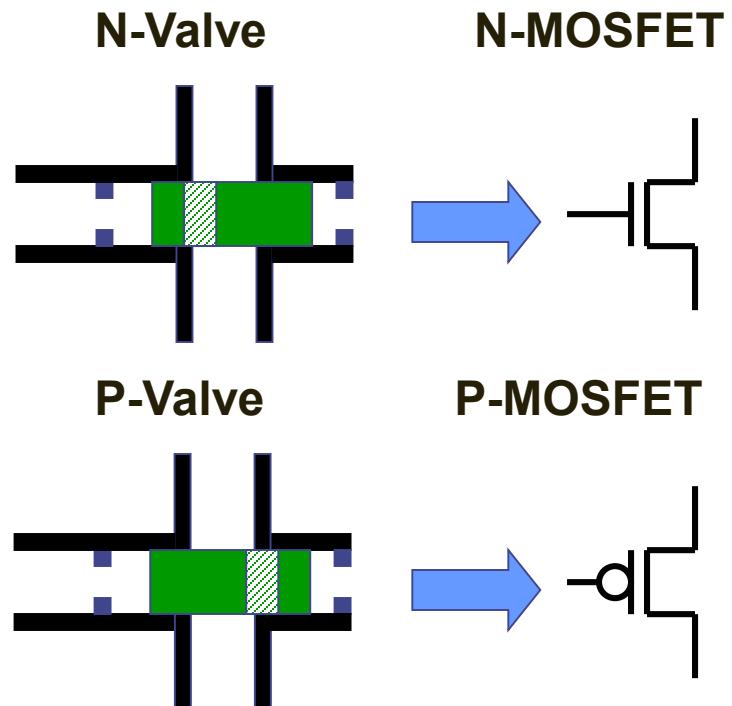


Analogy Explained

- Pressure differential → electrical potential (voltage)
 - Air molecules → electrons
 - Pressure (molecules per volume) → voltage
 - High pressure → high voltage
 - Low pressure → low voltage
- Air flow → electrical current
 - Pipes → wires
 - Air only flows from high to low pressure
 - Electrons only flow from high to low voltage
 - Flow only occurs when changing from 1 to 0 or 0 to 1
- Valve → transistor
 - The transistor: one of the century's most important inventions

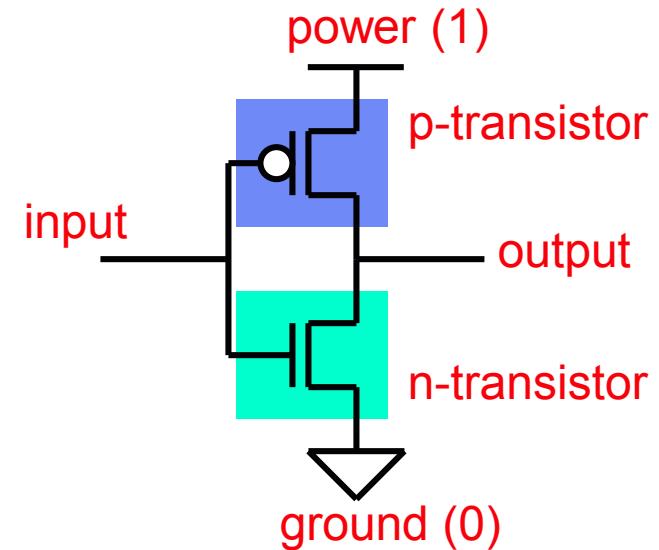
Transistors as Switches

- Two types
 - N-type
 - P-type
- Properties
 - Solid state (no moving parts)
 - Reliable (low failure rate)
 - Small (14nm channel length)
 - Fast (<0.1ns switch latency)



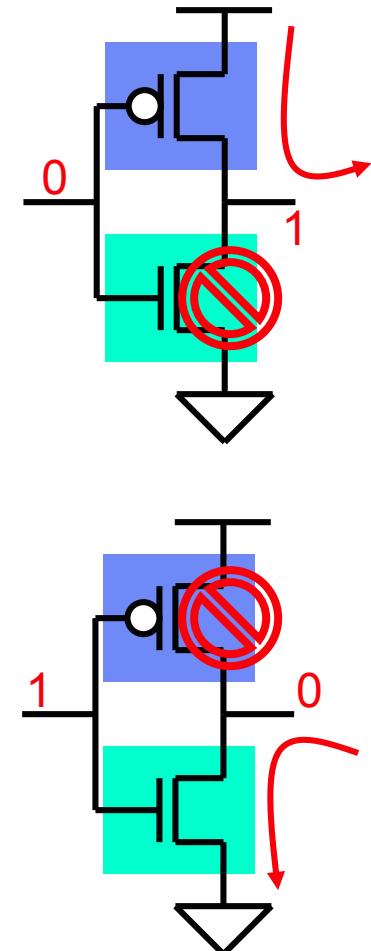
Complementary MOS (CMOS)

- Voltages as values
 - Power (V_{DD}) = "1", Ground = "0"
- Two kinds of MOSFETs
 - **N-transistors**
 - Conduct when gate voltage is 1
 - Good at passing 0s
 - **P-transistors**
 - Conduct when gate voltage is 0
 - Good at passing 1s
- **CMOS**
 - Complementary n/p- networks form boolean logic (i.e., gates)
 - And some non-gate elements too (important example: RAMs)



Basic CMOS Logic Gate

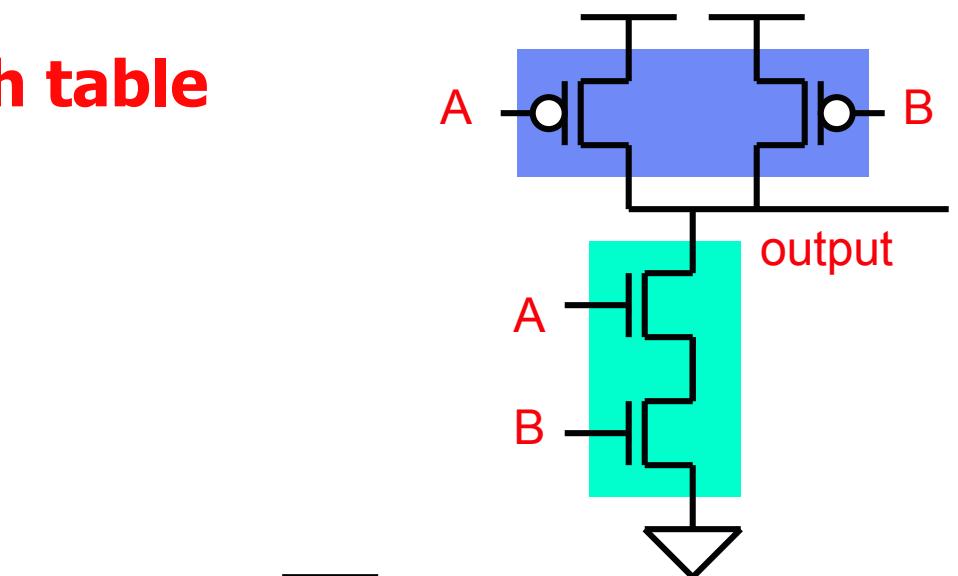
- **Inverter:** NOT gate
 - One p-transistor, one n-transistor
 - Basic operation
 - Input = 0
 - P-transistor closed, n-transistor open
 - Power charges output (1)
 - Input = 1
 - P-transistor open, n-transistor closed
 - Output discharges to ground (0)



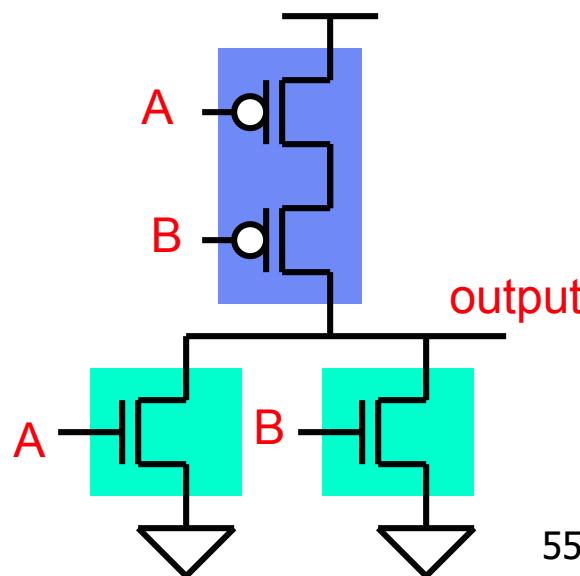
Another CMOS Gate Example

- What is this? Look at **truth table**

- 0, 0 → 1
- 0, 1 → 1
- 1, 0 → 1
- 1, 1 → 0
- Result: **NAND** (NOT AND)
- NAND is “universal”



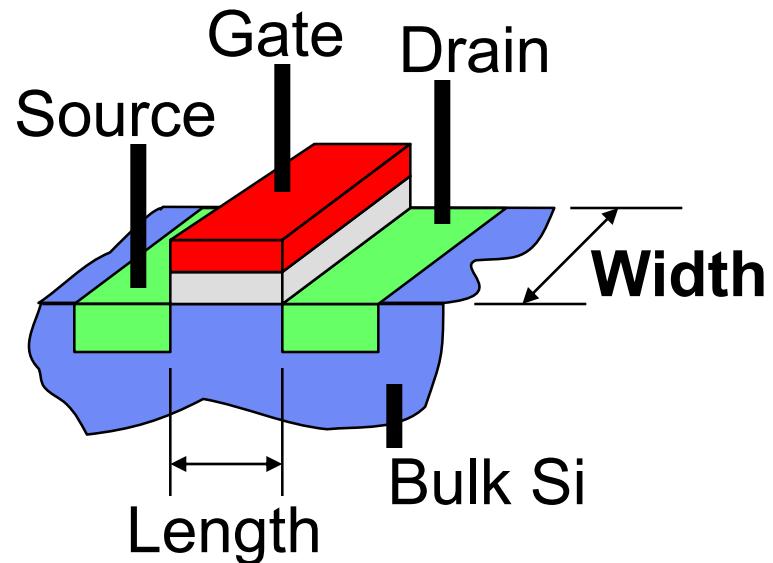
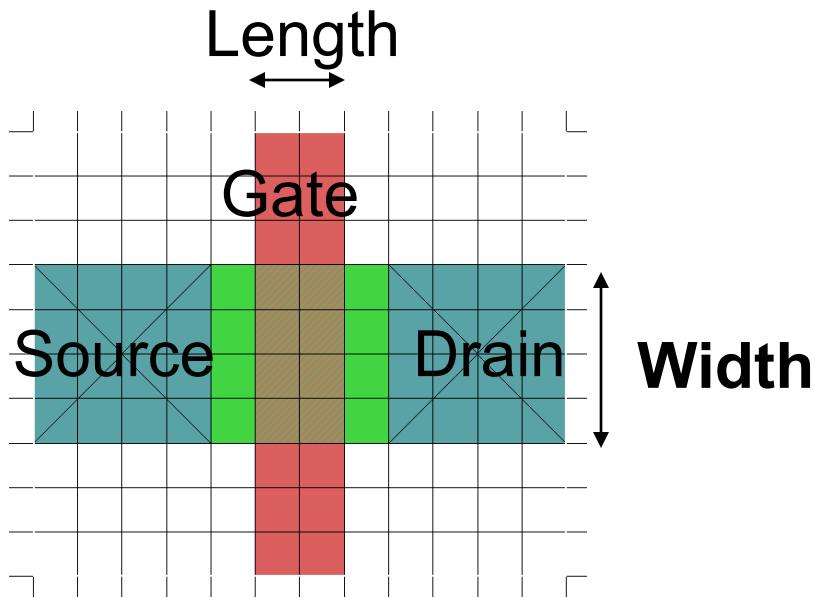
- What function is this?



Technology Basis of Transistor Speed

- Physics 101: delay through an electrical component $\propto \text{RC}$
 - **Resistance (R)**  \propto length / cross-section area
 - Slows rate of charge flow
 - **Capacitance (C)**  \propto length * area / distance-to-other-plate
 - Stores charge
 - **Voltage (V)**
 - Electrical pressure
 - **Threshold Voltage (V_t)**
 - Voltage at which a transistor turns “on”
 - Property of transistor based on fabrication technology
 - **Switching time** $\propto (\mathbf{R} * \mathbf{C}) / (\mathbf{V} - \mathbf{V}_t)$
- Two kinds of electrical components
 - CMOS transistors (gates, sources, drains)
 - Wires

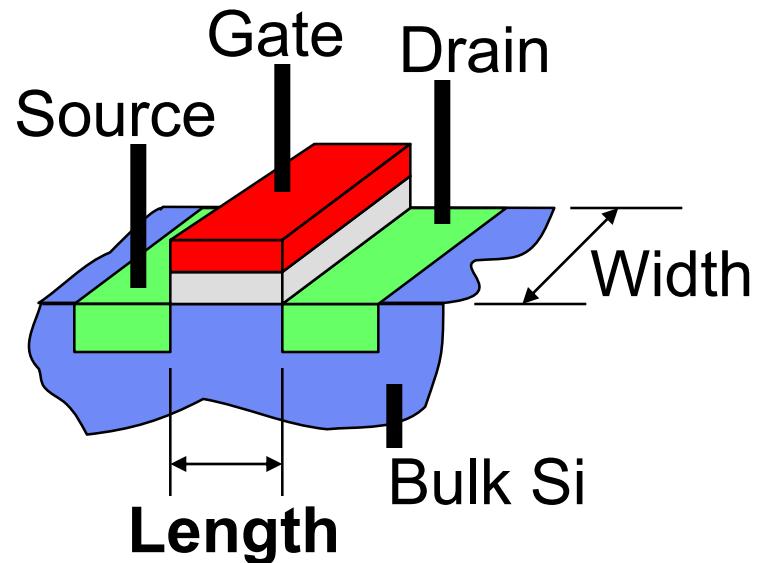
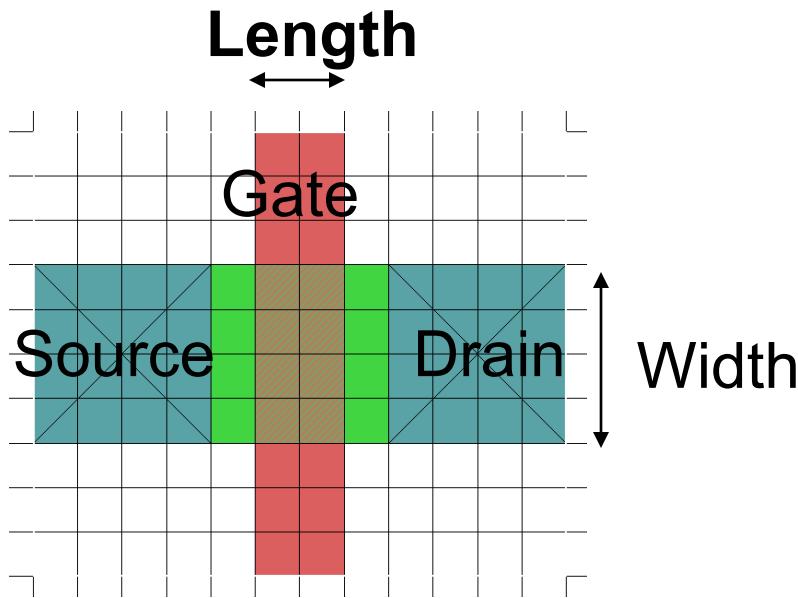
Transistor Geometry: Width



Diagrams © Krste Asanovic, MIT

- **Transistor width**, set by designer for each transistor
- Wider transistors:
 - **Lower resistance** of channel (increases drive strength) – good!
 - But, **increases capacitance** of gate/source/drain – bad!
- Result: set width to balance these conflicting effects

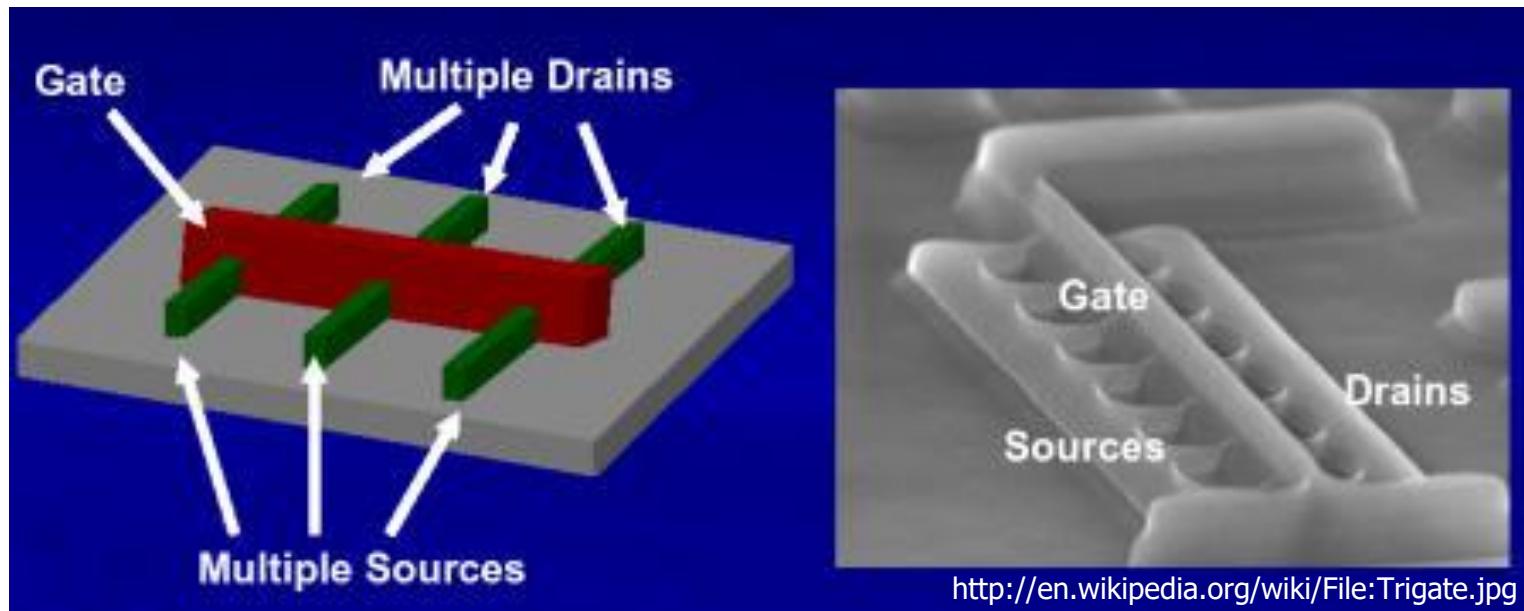
Transistor Geometry: Length & Scaling



Diagrams © Krste Asanovic, MIT

- **Transistor length:** characteristic of “process generation”
 - “22nm” refers to the transistor gate length
- Each process generation shrinks transistor length by 1.4x
 - “Moore’s law” -> roughly 2x improvement in transistor density
 - Roughly linear improvement in switching speeds (lower resistance)

Trigate FinFET Transistors



- nonplanar (or “3D”) transistors
 - trigate: multiple sources/drains/gates
 - FinFET: gate is wrapped around the channel
- lower leakage, faster switching times
- Intel’s trigate design released in mid-2012 (Ivy Bridge)
 - TSMC FinFET in mid-2015, Samsung FinFET in early 2016

Technology Scaling Trends

Dennard Scaling

“Design of ion-implanted MOSFET’s with very small physical dimensions”
Robert H. Dennard, Fritz H. Gaensslen, Hwa-Nien Yu, V. Leo Rideout, Ernest
Bassous, and Andre R. LeBlanc
IEEE Journal of Solid-State Circuits, October 1974

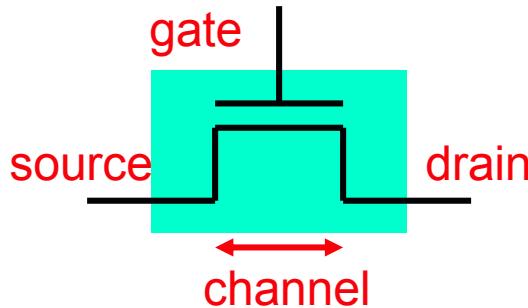
Table 1
Scaling Results for Circuit Performance

Device or Circuit Parameter	Scaling Factor
Device dimension t_{ox} , L , W	$1/\kappa$
Doping concentration N_a	κ
Voltage V	$1/\kappa$
Current I	$1/\kappa$
Capacitance $\epsilon A/t$	$1/\kappa$
Delay time/circuit VC/I	$1/\kappa$
Power dissipation/circuit VI	$1/\kappa^2$
Power density VI/A	1

Dennard Scaling

- stopped in ~2005 due to leakage
 - V close to V_t , transistors never really “on” or “off”
 - gate-oxide leakage due to very small oxide thickness
 - quantum-mechanical electron tunneling
- Moore’s Law still in effect!
 - but power usage goes up with increasing transistor counts

Moore's Law: Technology Scaling



- **Moore's Law:** aka “technology scaling”
 - Continued miniaturization (esp. reduction in channel length)
 - + Improves switching speed, power/transistor, area(cost)/transistor
 - Reduces transistor reliability
 - Literally: DRAM density (transistors/area) doubles every 18 months
 - Public interpretation: performance doubles every 18 months
 - Not quite right, but helps performance in several ways...

Moore's Law in the Future

- Won't last forever, approaching physical limits
 - "If something must eventually stop, it can't go on forever"
 - But betting against it has proved foolish in the past
 - Perhaps will "slow" rather than stop abruptly
- Transistor count will likely continue to scale
 - 3D "die stacking" is arriving
- But transistor performance scaling?
 - Running into physical limits
 - Example: gate oxide < 10 silicon atoms!
 - Can't decrease it much further
 - Power is becoming the limiting factor

The image shows a screenshot of a news article from AnandTech. The header features the AnandTech logo and navigation icons (menu, home, search). Below the header, the breadcrumb navigation shows "Home > CPUs". The main title of the article is "Tick Tock On The Rocks: Intel Delays 10nm, Adds 3rd Gen 14nm Core Product 'Kaby Lake'". A byline indicates the article was written by Brett Howse & Ryan Smith on July 16, 2015 at 10:15 AM EST. There are 136 comments available. The bottom right corner of the image displays the large blue Intel logo.

Power & Energy

Power/Energy Are Increasingly Important

- **Battery life** for mobile devices
 - Laptops, phones, cameras
- **Tolerable temperature** for devices without active cooling
 - Power means temperature, active cooling means **cost**
 - No room for a fan in a cell phone, no market for a hot cell phone
- **Electric bill** for compute/data centers
 - Pay for power twice: once in, once out (to cool)
- **Environmental concerns**
 - IT accounts for growing fraction of electricity consumption

Energy & Power

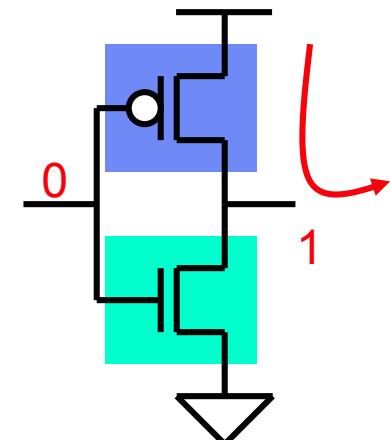
- **Energy:** measured in Joules or Watt-seconds
 - Total amount of energy stored/used
 - Battery life, electric bill, environmental impact
 - Instructions per Joule (car analogy: miles per gallon)
- **Power:** energy per unit time (measured in Watts)
 - Related to “performance” (which is also a “per unit time” metric)
 - Power impacts power supply and cooling requirements (cost)
 - Power-density (Watt/mm²): important related metric
 - Peak power vs average power
 - E.g., camera: power “spikes” when you actually take a picture
 - Joules per second (car analogy: gallons per hour)
- Two sources:
 - **Dynamic power:** active switching of transistors
 - **Static power:** leakage of transistors even while inactive

Recall: Tech. Basis of Transistor Speed

- Physics 101: delay through an electrical component $\propto \text{RC}$
 - **Resistance (R)**  \propto length / cross-section area
 - Slows rate of charge flow
 - **Capacitance (C)**  \propto length * area / distance-to-other-plate
 - Stores charge
 - **Voltage (V)**
 - Electrical pressure
 - **Threshold Voltage (V_t)**
 - Voltage at which a transistor turns “on”
 - Property of transistor based on fabrication technology
 - **Switching time** $\propto (\mathbf{R} * \mathbf{C}) / (\mathbf{V} - \mathbf{V}_t)$

Dynamic Power

- **Dynamic power ($P_{dynamic}$)**: aka switching or active power
 - Energy to switch a gate (0 to 1, 1 to 0)
 - Each gate has capacitance (C)
 - Charge stored $\propto C * V$
 - Energy to charge/discharge a capacitor $\propto C * V^2$
 - Time to charge/discharge a capacitor $\propto V$
 - Result: frequency $\propto V$
 - $P_{dynamic} \approx N * C * V^2 * f * A$
 - N: number of transistors
 - C: capacitance per transistor (size of transistors)
 - V: voltage (supply voltage for gate)
 - f: frequency (transistor switching freq. \propto clock freq.)
 - A: activity factor (not all transistors may switch this cycle)

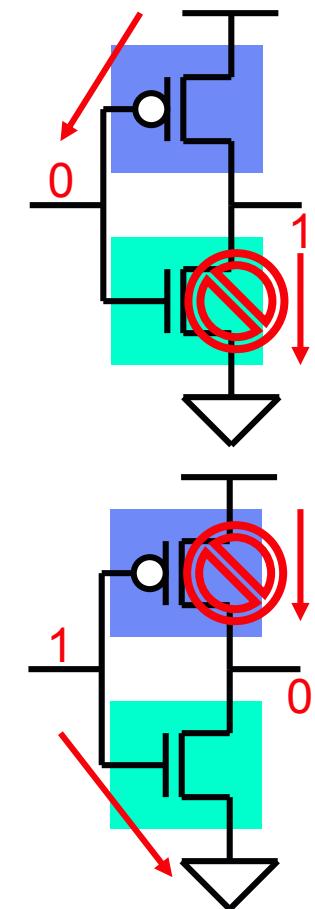


Reducing Dynamic Power

- Target each component: $P_{\text{dynamic}} \approx N * C * V^2 * f * A$
- **Reduce number of transistors (N)**
 - Use fewer transistors and gates
- **Reduce capacitance (C)**
 - Smaller transistors (Moore's law)
- **Reduce voltage (V)**
 - Quadratic reduction in energy consumption!
 - But also slows transistors (transistor speed $\propto V$)
- **Reduce frequency (f)**
 - Slower clock frequency (reduces power but not energy) Why?
- **Reduce activity (A)**
 - “Clock gating” disable clocks to unused parts of chip
 - Don’t switch gates unnecessarily

Static Power

- **Static power (P_{static})**: aka idle or leakage power
 - Transistors don't turn off all the way
 - Transistors are "leaky valves"
 - $P_{\text{static}} \approx N * V * e^{-V_t}$
 - N: number of transistors
 - V: voltage
 - **V_t (threshold voltage)**: voltage at which transistor conducts (begins to switch)
- Switching speed vs leakage trade-off
- The lower the V_t :
 - Faster transistors (linear)
 - Transistor speed $\propto V - V_t$
 - Leakier transistors (exponential)



Reducing Static Power

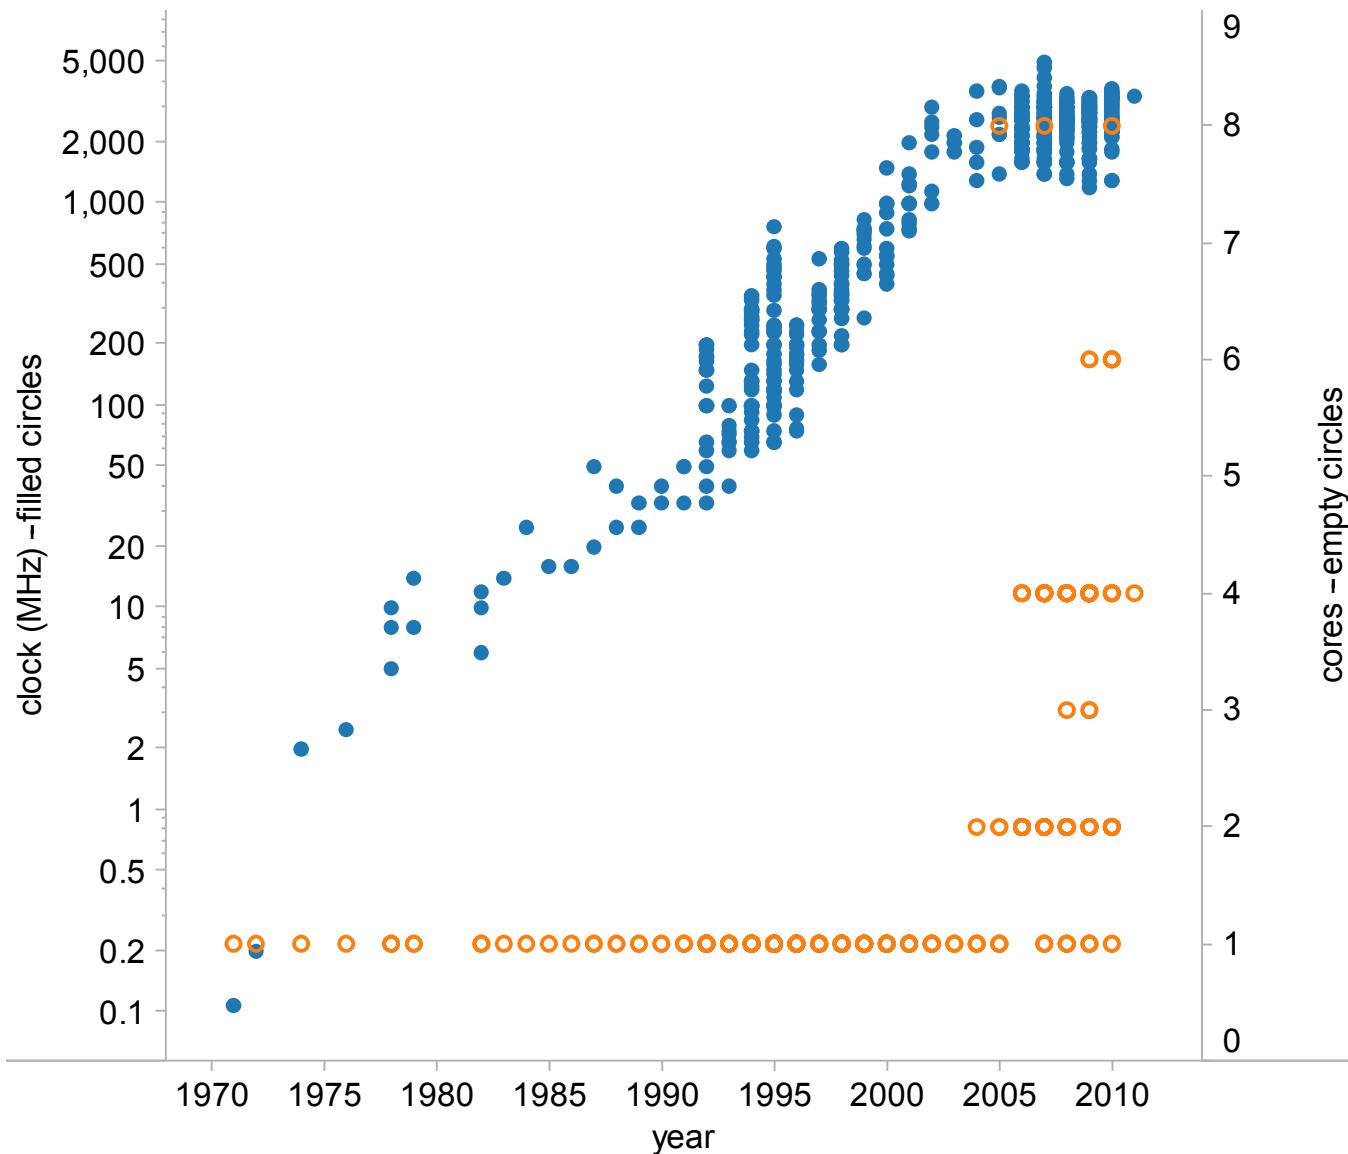
- Target each component: $P_{\text{static}} \approx N * V * e^{-Vt}$
- **Reduce number of transistors (N)**
 - Use fewer transistors/gates
- **Disable transistors** (also targets N)
 - “Power gating” disable power to unused parts (long latency to power up)
 - Power down units (or entire cores) not being used
- **Reduce voltage (V)**
 - Linear reduction in static energy consumption
 - But also slows transistors ($\text{transistor speed} \propto V$)
- **Dual V_t** – use a mixture of high and low V_t transistors
 - Use slow, low-leak transistors in SRAM arrays
 - Requires extra fabrication steps (cost)
- **Low-leakage transistors**
 - High-K/Metal-Gates in Intel’s 45nm process, “tri-gate” in Intel’s 22nm
 - Reducing frequency can hurt energy efficiency due to leakage power

Dynamic Voltage/Frequency Scaling

- **Dynamically trade-off power for performance**
 - Change the voltage and frequency at runtime
 - Under control of operating system and/or hardware
- Recall: $P_{\text{dynamic}} \approx N * C * V^2 * f * A$
 - Because frequency \propto to $V - V_t$...
 - $P_{\text{dynamic}} \propto$ to $V^2(V - V_t) \approx V^3$
- Reduce both voltage and frequency linearly
 - **Cubic decrease in dynamic power**
 - Linear decrease in performance (actually sub-linear)
 - Thus, only about quadratic decrease in energy
 - Linear decrease in static power
 - Thus, static energy can become dominant
- Newer chips can adjust frequency on a per-core basis

Frequency and Core Count

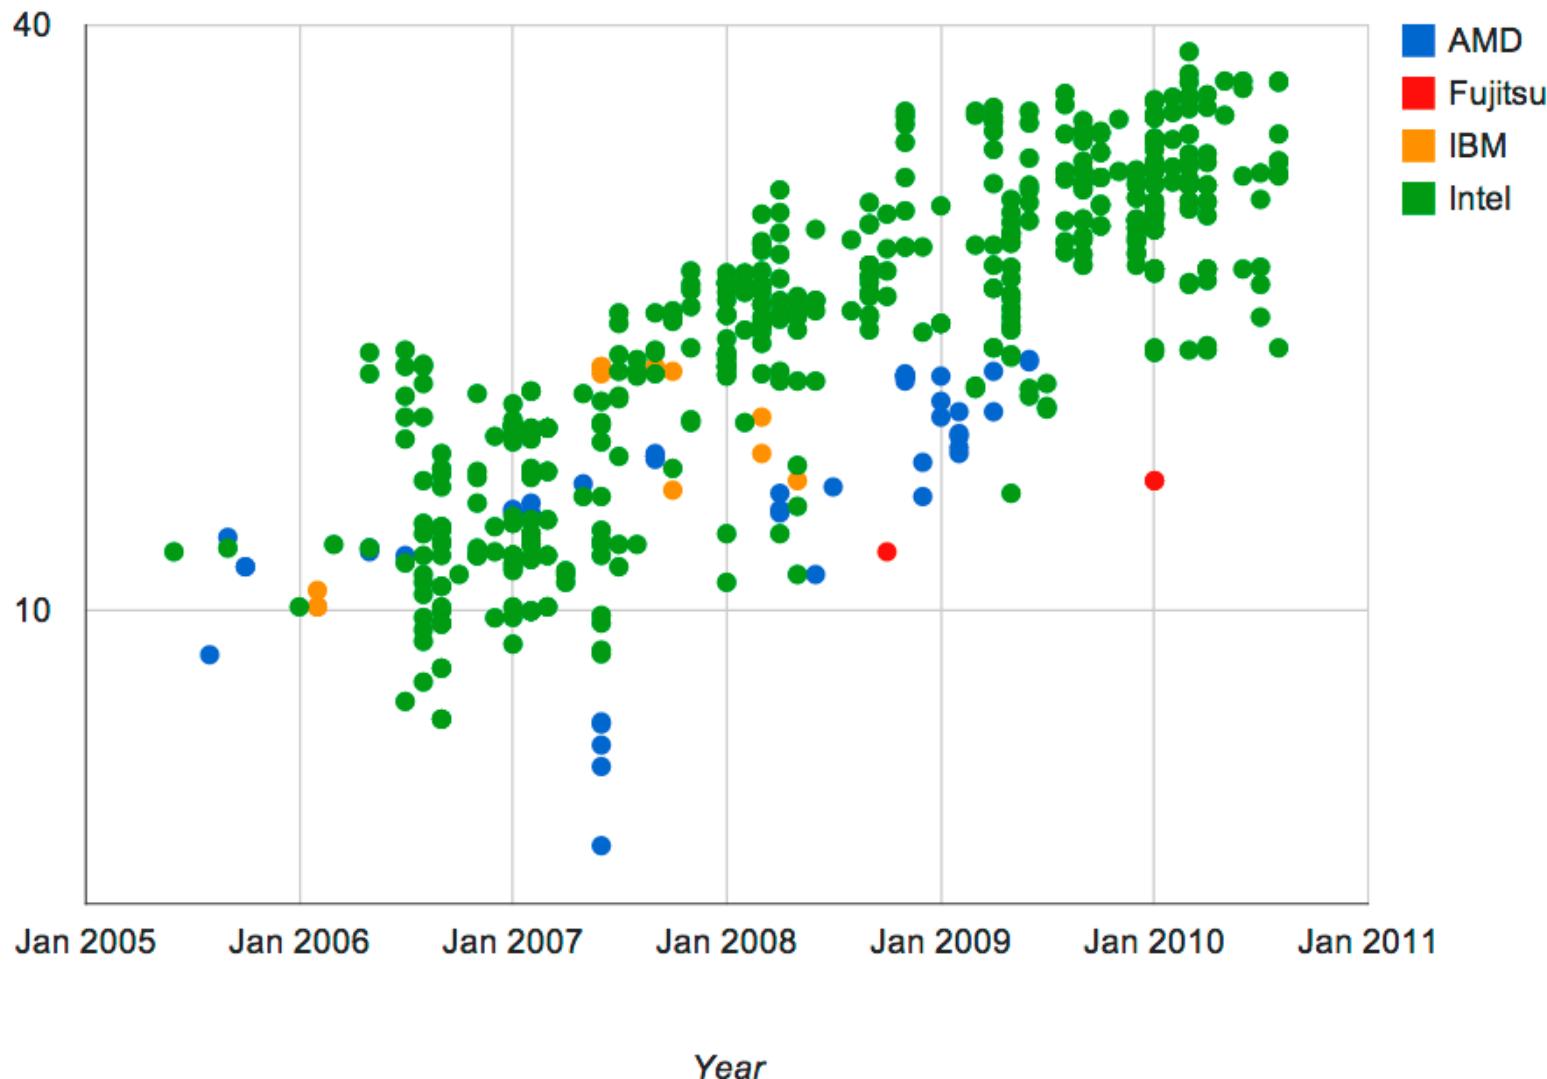
data from
<http://cpudb.stanford.edu>



SpecINT 2006 performance

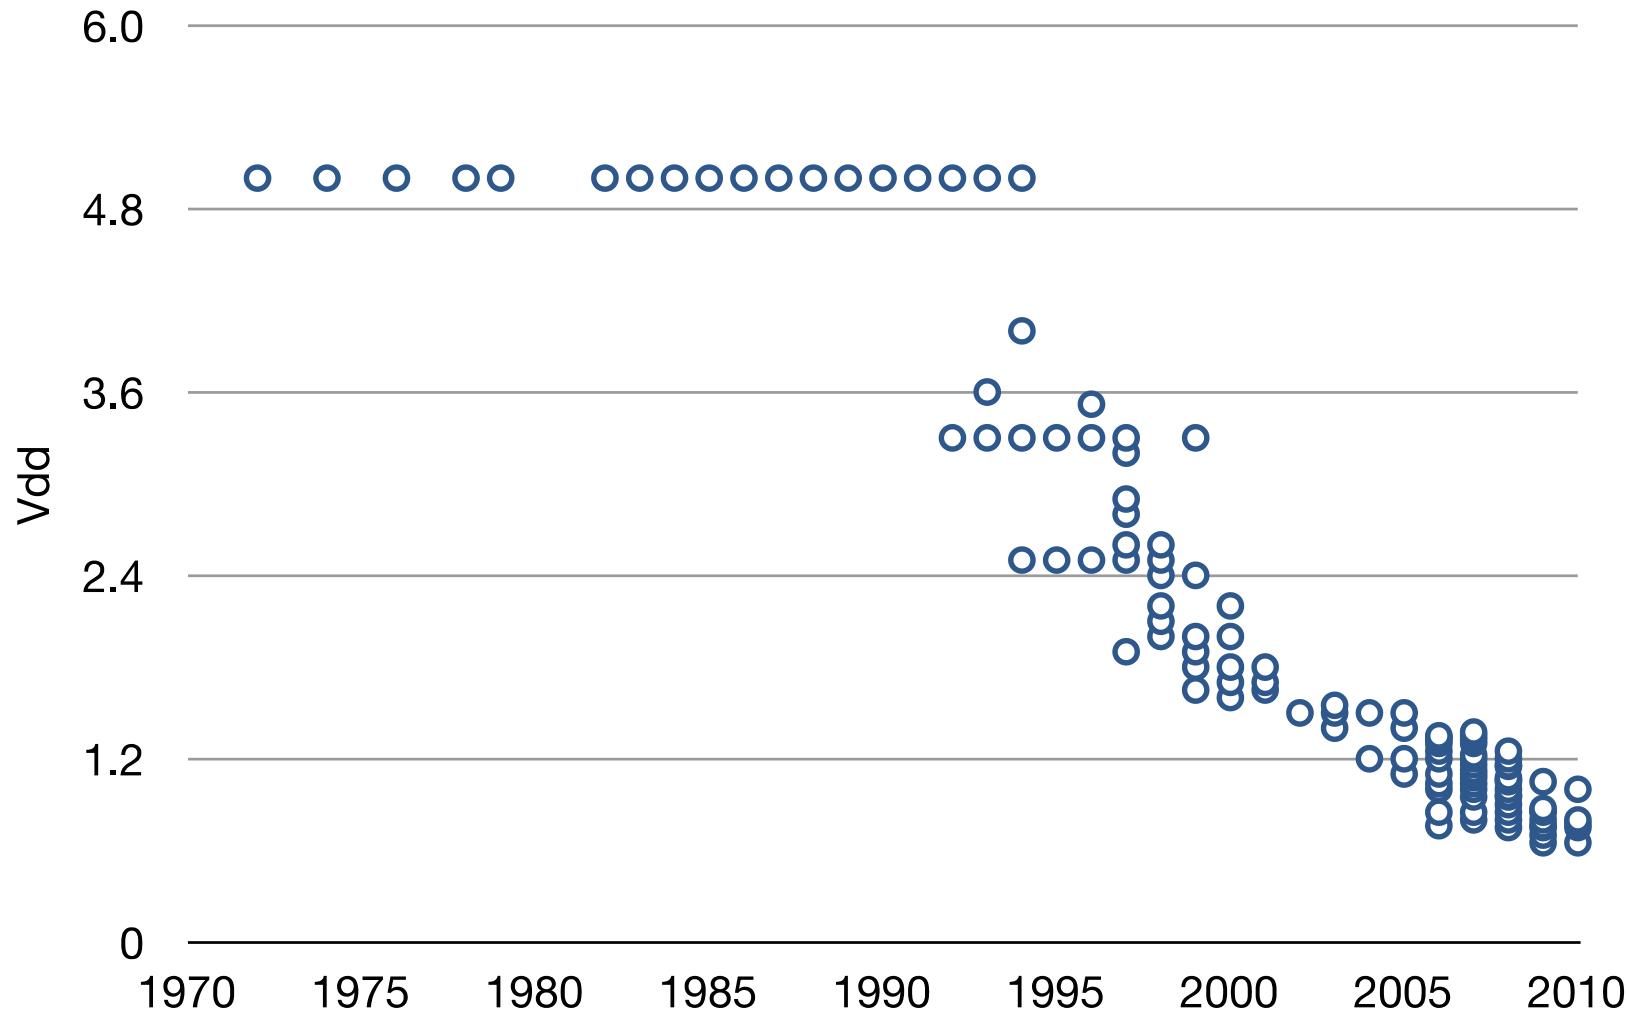
graph from
<http://cpudb.stanford.edu>

Spec Int 2006



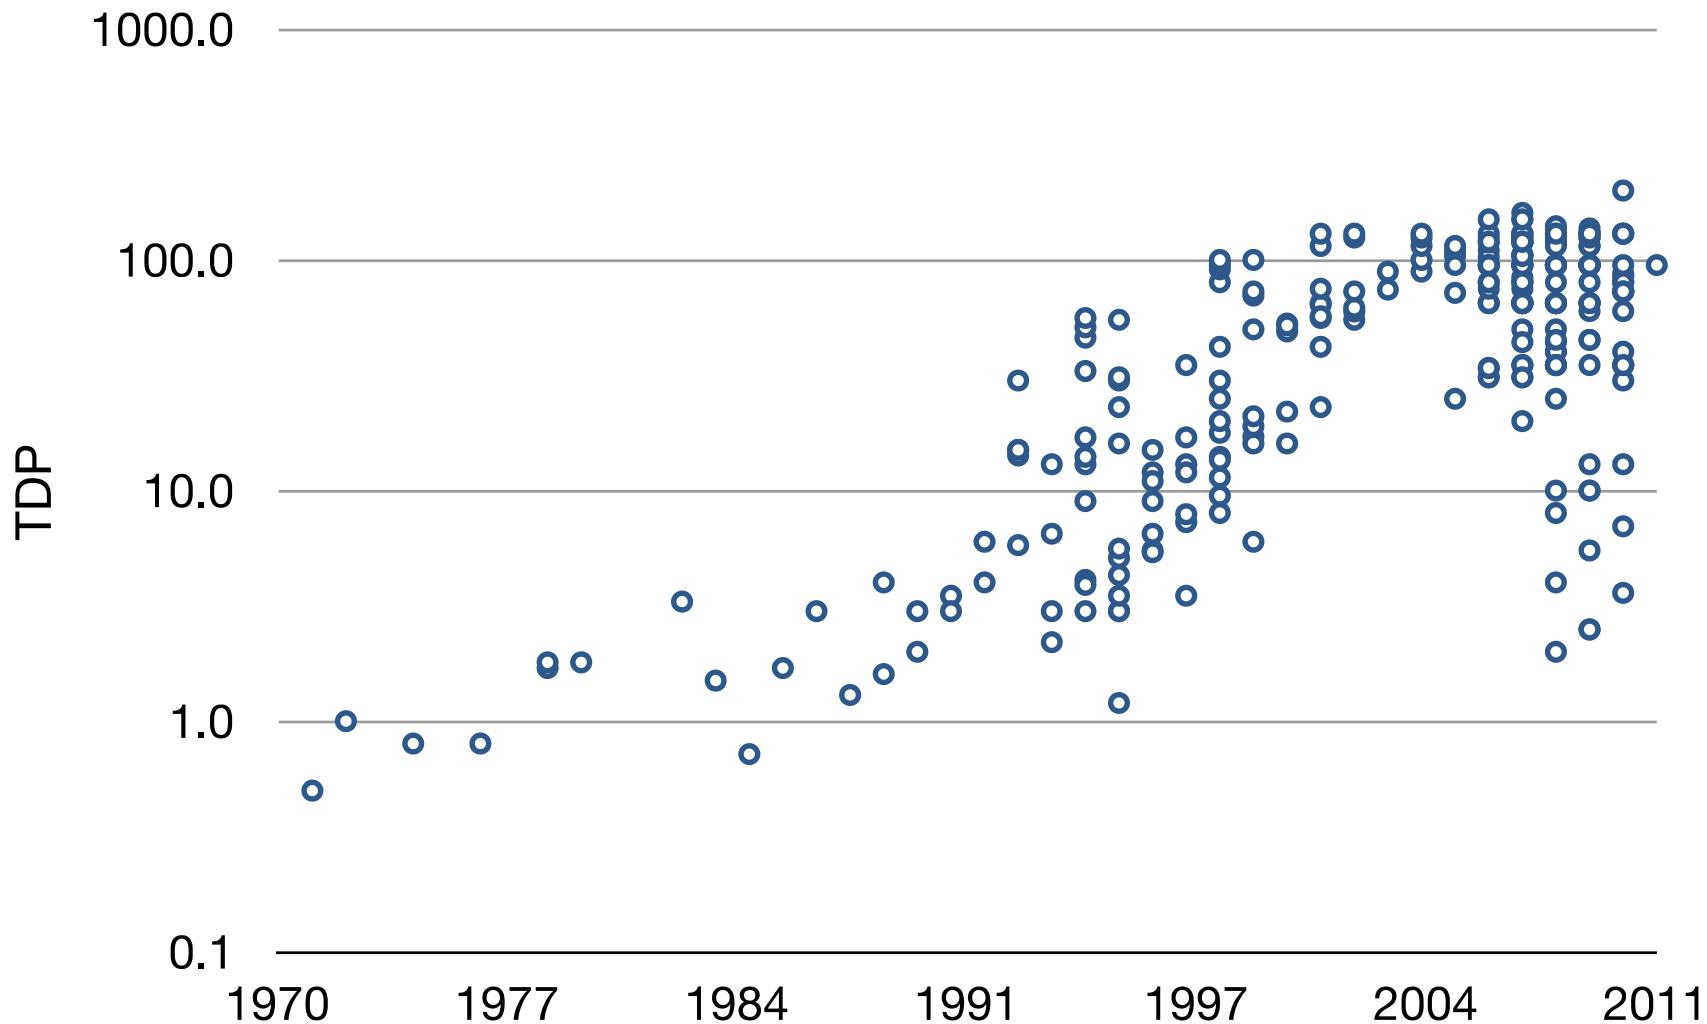
Supply Voltage

data from
<http://cpudb.stanford.edu>



Thermal Design Power

data from
<http://cpudb.stanford.edu>



Moore without Dennard

- + Dennard scaling reduced power/transistor...
 - Required reducing V , which requires a trade-off:
 - Keeping V_t the same and reducing frequency (f)
 - Lowering V_t and increasing leakage exponentially
- + Moore's Law still gives more transistors
 - + Use techniques like high-K/metal gate, dual- V_T , tri-gate
- The end of voltage scaling & "dark silicon"
 - Current projections: power per transistor reduced by 25-35% per technology node
 - What are the implications?

Implications on Software

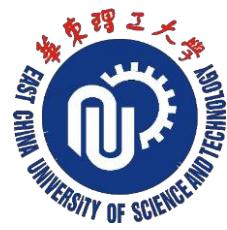
- Software-controlled dynamic voltage/frequency scaling
 - Example: video decoding
 - Too high a clock frequency – wasted energy (battery life)
 - Too low a clock frequency – quality of video suffers
 - “Race to sleep” versus “slow and steady” approaches
- Managing low-power modes
 - Don’t want to “wake up” the processor every millisecond
- Tuning software
 - Faster algorithms can be converted to lower-power algorithms
 - Via dynamic voltage/frequency scaling
- Exploiting parallelism & heterogeneous cores
 - ARM BIG.little design (a few “big” cores & 1 “low power” core)
- Specialized hardware accelerators

Summary

Technology Summary

- Has a first-order impact on computer architecture
 - Performance (transistor delay, wire delay)
 - Cost (die area & defects)
 - **Changing rapidly**
- Most significant trends for architects
 - More and more transistors
 - What to do with them? → integration → **parallelism**
 - Logic is improving faster than memory & cross-chip wires
 - “Memory wall” → caches, more integration
- Power and energy
 - Voltage vs frequency, parallelism, special-purpose hardware
- This unit: a quick overview, just scratching the surface

Rest of course



Thanks!
Q&A

