

# Course Overview

# Why Study Computer Architecture?

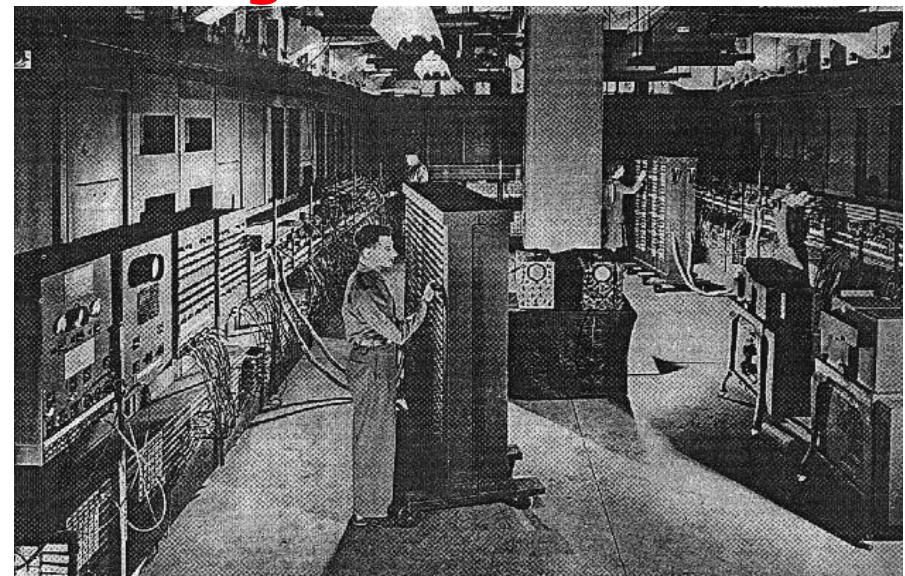
---

- **Understand where computers are going**
  - Future capabilities drive the (computing) world
  - Real world-impact: no computer architecture → no computers!
- **Understand high-level design concepts**
  - The best architects understand all the levels
    - Devices, circuits, architecture, compiler, applications
- **Understand computer performance**
  - learn valid experimental methodologies
- **Write better software**
  - The best software designers also understand hardware
  - Need to understand hardware to write fast software
- **Design hardware**
  - At Intel, AMD, IBM, ARM, Qualcomm, Oracle, NVIDIA, Samsung

# ENIAC

---

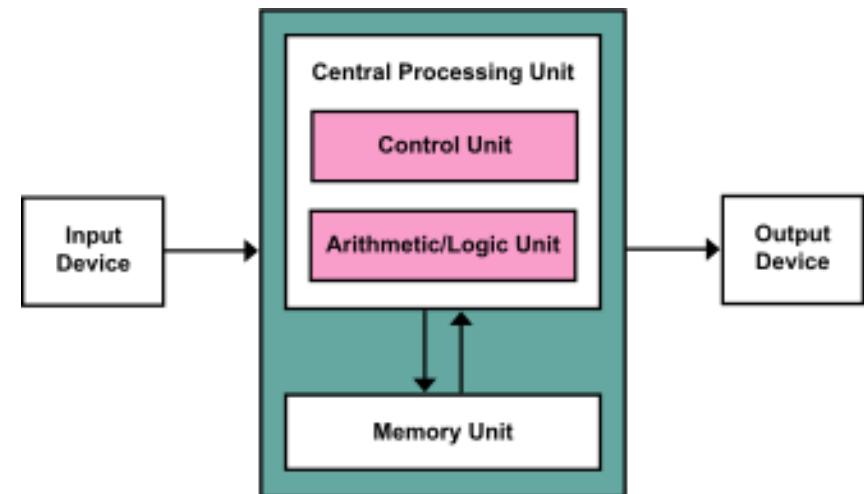
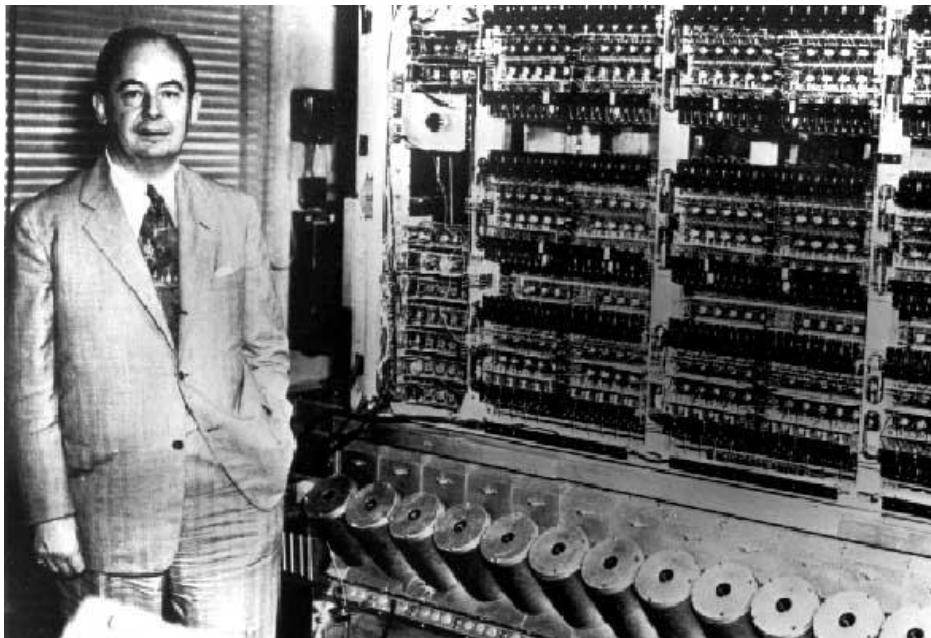
- **ENIAC**: electronic numerical integrator and calculator
  - First operational general-purpose stored-program computer
  - Designed and built here by Eckert and Mauchly from 1943-6
  - ENIAC was designed to calculate artillery firing tables for the US Army's Ballistic Research Laboratory (BRL).
  - Von Neumann architecture? (No)
- **First seminars on computer design**
  - Moore School Lectures, 1946
  - *"Theory and Techniques for Design of Electronic Digital Computers"*



# Von Neumann

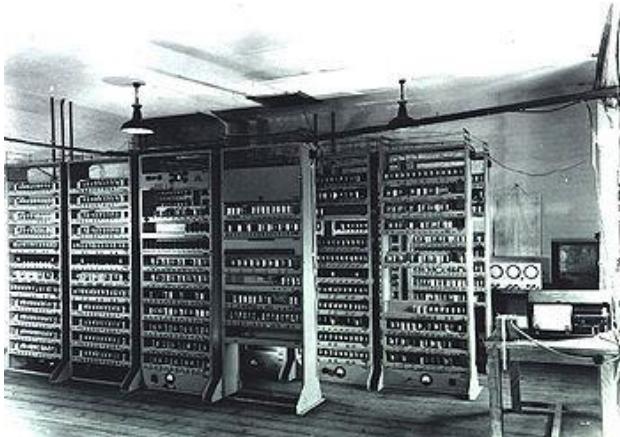
---

- The Von Neumann architecture, also known as the Von Neumann model and Princeton architecture.



# Von Neumann architecture computer

---



- **EDSAC (Electronic Delay Storage Automatic Calculator)** was the first practical stored-program electronic computer (May 1949)
- **EDVAC (Electronic Discrete Variable Automatic Computer)**
  - EDVAC's construction in August 1944. EDVAC was delivered to the Ballistics Research Laboratory in August 1949

# Application Specific Designs

---

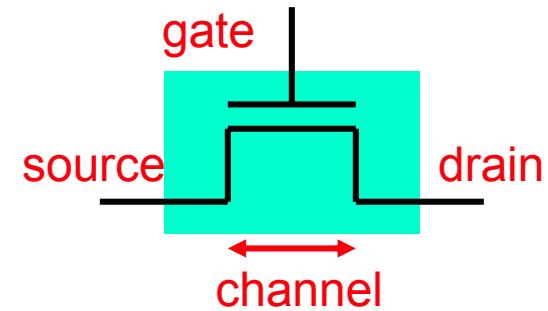
- This class is about **general-purpose CPUs**
  - Processor that can do anything, run a full OS, etc.
  - E.g., Intel Core i7, AMD Athlon, IBM Power, ARM A-series
- In contrast to **application-specific chips**
  - Or **ASICs** (Application specific integrated circuits)
    - Also application-domain specific processors
  - Implement critical domain-specific functionality in hardware
    - Examples: video encoding, 3D graphics
  - General rules
    - Hardware is less flexible than software
    - + Hardware more effective (speed, power, cost) than software
    - + Domain specific more “parallel” than general purpose
      - But general mainstream processors becoming more parallel
- Trend: from **specific** to **general** (for a specific domain)

# Technology Trends

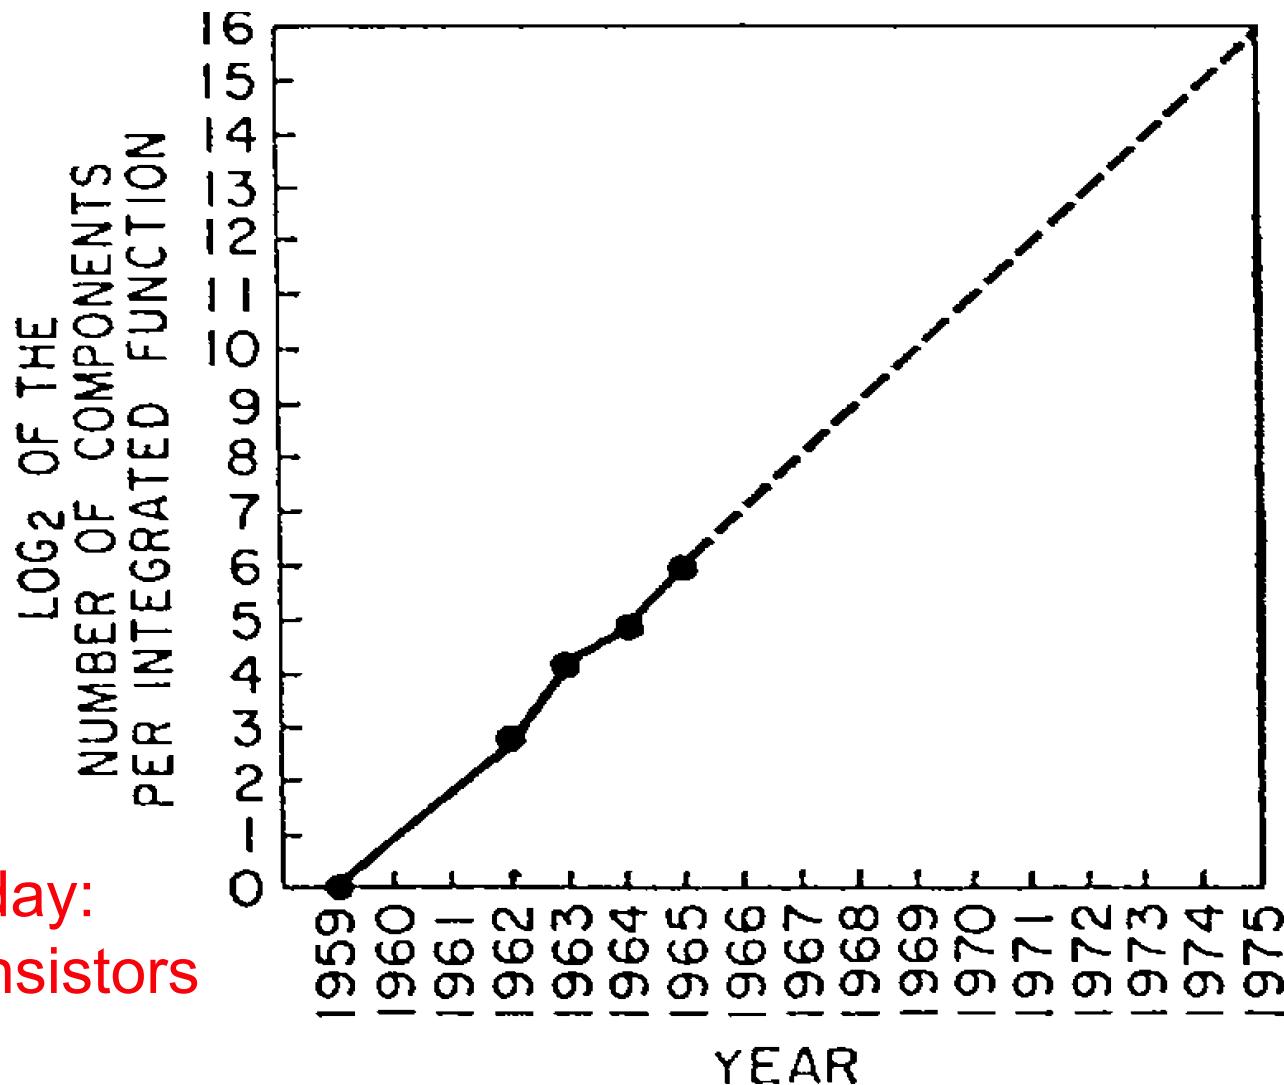
# “Technology”

---

- Basic element
  - Solid-state **transistor**
  - Building block of **integrated circuits (ICs)**
- What's so great about ICs? Everything
  - + High performance, high reliability, low cost, low power
  - + Lever of mass production
- Several kinds of integrated circuit families
  - **SRAM/logic**: optimized for speed (used for processors)
  - **DRAM**: optimized for density, cost, power (used for memory)
  - **Flash**: optimized for density, cost (used for disks)
  - Increasing opportunities for integrating multiple technologies
- Non-transistor storage and inter-connection technologies
  - Disk, optical storage, ethernet, fiber optics, wireless



# Moore's Law - 1965



# Revolution I: The Microprocessor

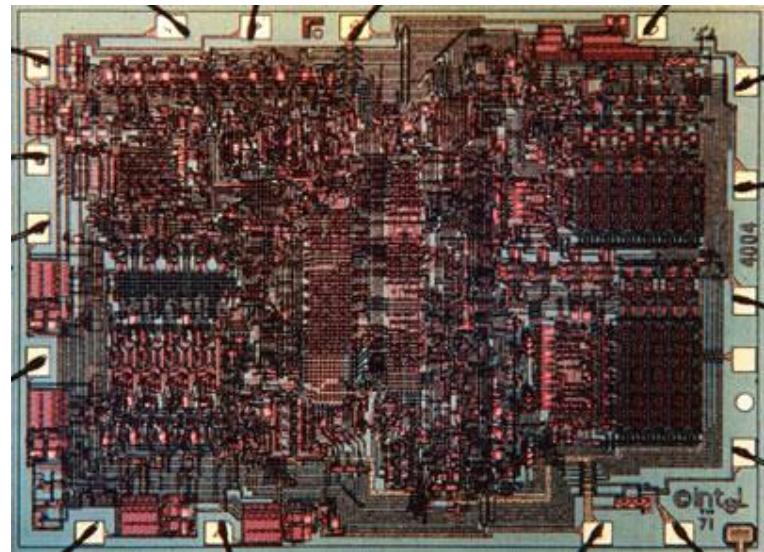
---

- **Microprocessor revolution**
  - One significant technology threshold was crossed in 1970s
  - Enough transistors ( $\sim 25K$ ) to put a 16-bit processor on one chip
  - Huge performance advantages: fewer slow chip-crossings
  - Even bigger cost advantages: one “stamped-out” component
- Microprocessors have allowed new market segments
  - Desktops, CD/DVD players, laptops, game consoles, set-top boxes, mobile phones, digital camera, mp3 players, GPS, automotive
- And replaced incumbents in existing segments
  - Microprocessor-based system replaced supercomputers, “mainframes”, “minicomputers”, etc.

# First Microprocessor

---

- Intel 4004 (1971)
  - Application: calculators
  - Technology: 10000 nm
  - 2300 transistors
  - 13 mm<sup>2</sup>
  - 108 KHz
  - 12 Volts
  - 4-bit data
  - Single-cycle datapath



# Tracing the Microprocessor Revolution

---

- How were growing transistor counts used?
- Initially to widen the datapath
  - 4004: 4 bits → AMD Opteron (2003): 64 bits
- ... and also to add more powerful instructions
  - To amortize overhead of fetch and decode
  - To simplify assembly programming (which was done by hand then)

# Revolution II: Implicit Parallelism

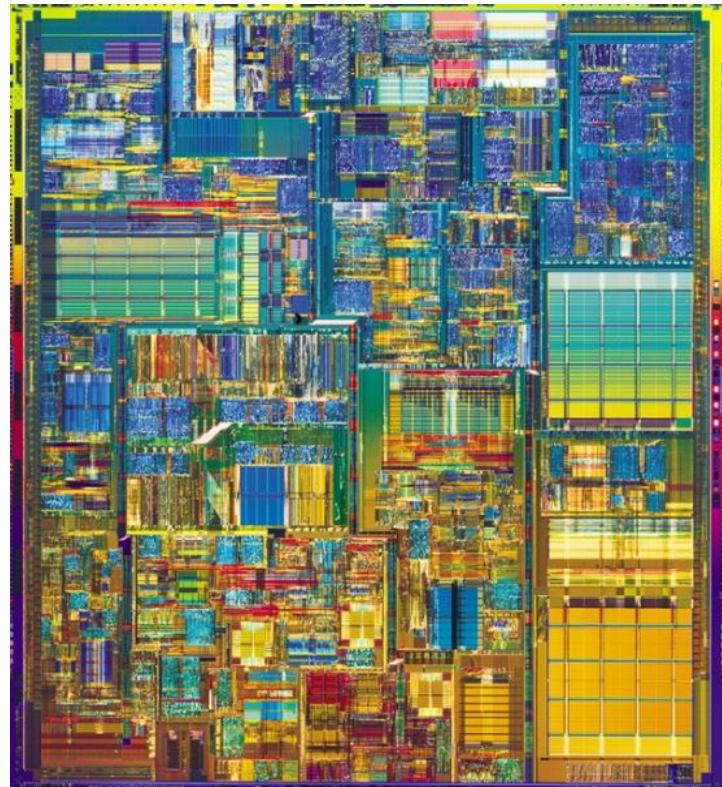
---

- Then to **extract implicit instruction-level parallelism**
  - Hardware provides parallel resources, figures out how to use them
  - Software is oblivious
- Initially using pipelining ...
  - Which also enabled increased clock frequency
- ... caches ...
  - Which became necessary as processor clock frequency increased
- ... and integrated floating-point
- Then deeper pipelines and branch speculation
- Then multiple instructions per cycle (superscalar)
- Then dynamic scheduling (out-of-order execution)
- We will discuss this in detail

# Pinnacle of Single-Core Microprocessors

---

- Intel Pentium4 (2003)
  - Application: desktop/server
  - Technology: 90nm (1/100th of 4004)
  - 55M transistors (20,000x)
  - 101 mm<sup>2</sup> (10x)
  - 3.4 GHz (10,000x)
  - 1.2 Volts (1/10th)
  - 32/64-bit data (16x)
  - 22-stage pipelined datapath
  - 3 instructions per cycle (superscalar)
  - Two levels of on-chip cache
  - data-parallel vector (SIMD) instructions, hyperthreading



# Revolution III: Explicit Parallelism

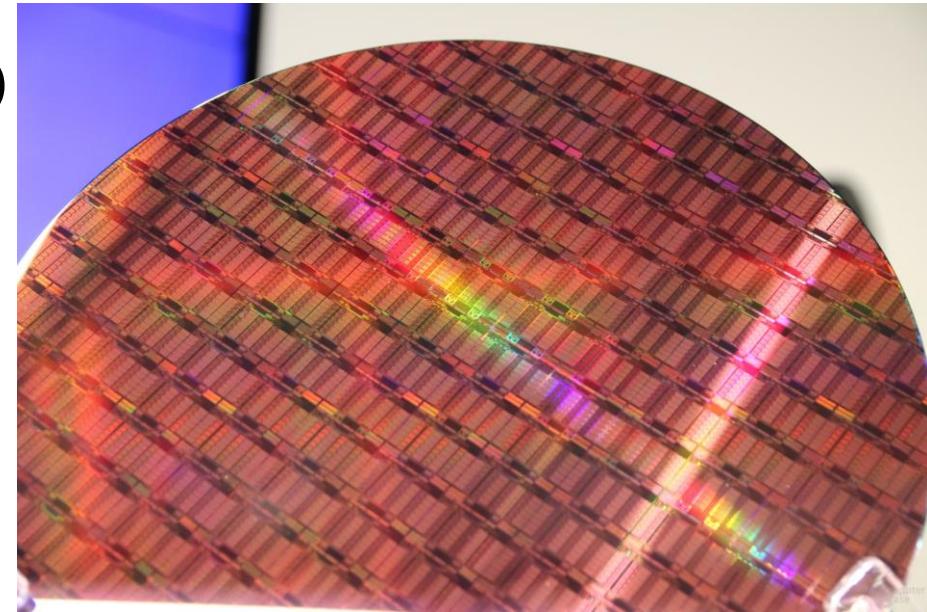
---

- Then to support **explicit data & thread level parallelism**
  - Hardware provides parallel resources, software specifies usage
  - Why? diminishing returns on instruction-level-parallelism
- First using (subword) vector instructions..., Intel's SSE
  - One instruction does four parallel multiplies
- ... and general support for multi-threaded programs
  - Coherent caches, hardware synchronization primitives
- Then using support for multiple concurrent threads on chip
  - First with single-core multi-threading, now with multi-core

# Modern Multicore Processor

---

- Intel Xeon E5-2699 V4 (2016)
  - Application: server
  - Technology: 14nm (16% of P4)
  - 7.2B transistors (130x)
  - 456 mm<sup>2</sup> (4.5x)
  - 2.4 to 3.6 Ghz (~1x)
  - 1.8 Volts (~1x)
  - 256-bit data (2x)
  - 14-stage pipelined datapath (0.5x)
  - 4 instructions per cycle (1x)
  - Three levels of on-chip cache
  - data-parallel vector (SIMD) instructions, hyperthreading
  - **22-core multicore** (22x)



# Revolution IV: Heterogeneous Processing

---

- Combining **multiple** kinds of compute engines in one die
  - not just homogenous collection of cores
  - System-on-Chip (SoC) is one common example in mobile space
- Lots of stuff on the chip beyond just CPUs
  - Graphics Processing Units (GPUs)
    - throughput-oriented specialized multicore processors
    - good for gaming, machine learning, computer vision, ...
  - Special-purpose logic
    - media codecs, cell phone radio, encryption, compression
- Excellent energy efficiency and performance
  - extremely complicated to program!

# Technology Disruptions

---

- Classic examples:
  - The transistor
  - Microprocessor
- More recent examples:
  - Multicore processors
  - Flash-based solid-state storage
- Near-term potentially disruptive technologies:
  - Non-volatile memory (e.g., phase-change memory)
  - Chip stacking (aka “3D die stacking”)
- Disruptive “end-of-scaling”
  - “If something can’t go on forever, it must stop eventually”
  - Can we continue to shrink transistors for ever?
  - Transistor energy efficiency is not improving as rapidly

# Managing This Mess

---

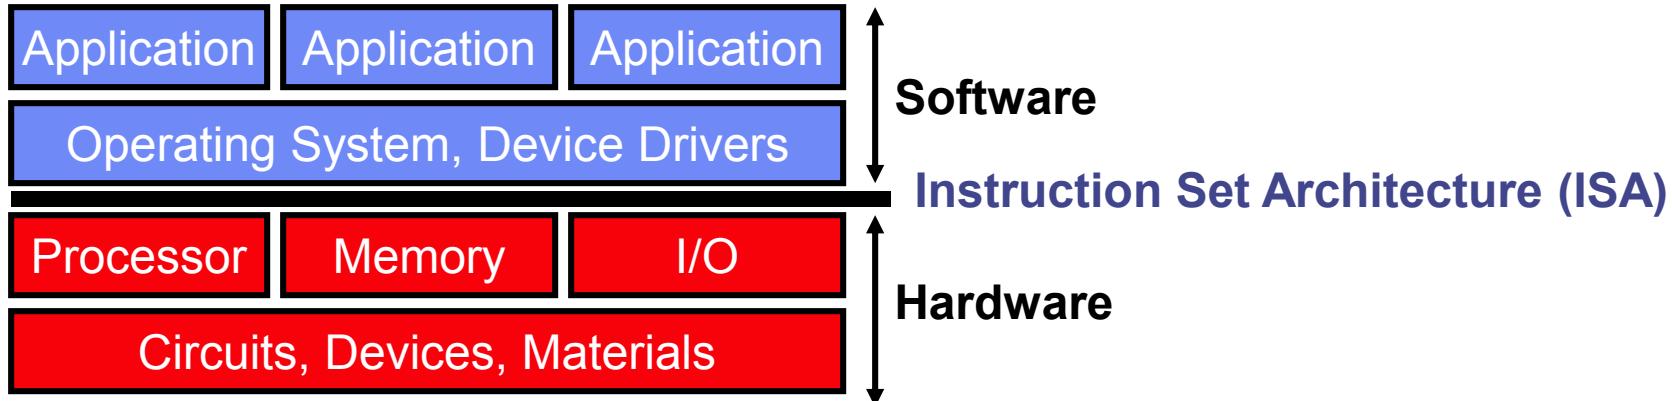
- Architect must consider all factors
  - Goals/constraints, applications, implementation technology
- Questions
  - How to deal with all of these inputs?
  - How to manage changes?
- Answers
  - Accrued institutional knowledge (stand on each other's shoulders)
  - Experience, rules of thumb
  - Discipline: clearly defined end state, keep your eyes on the ball
  - **Abstraction and layering**

# Pervasive Ideas: Abstraction and Layering

---

- **Abstraction**: only way of dealing with complex systems
  - Divide world into objects, each with an...
    - **Interface**: knob(s)
    - **Implementation**: “black box”
  - Only specialists deal with implementation, rest of us with interface
- **Layering**: repeated abstraction makes life even simpler
  - Divide objects in system into layers, objects at layer  $n$ :
    - are implemented using interfaces of layer  $n - 1$
    - don't need to know interfaces of layer  $n - 2$  (sometimes helps)
- **Inertia**: the dark side of layering
  - Layer interfaces become entrenched over time (“standards”)
  - Very difficult to change even if benefit is clear (example: Digital TV)
- **Opacity**: hard to reason about performance across layers

# Abstraction, Layering, and Computers



- **Computer architecture**
  - Definition of **ISA** to facilitate implementation of software layers
- This course mostly on **computer micro-architecture**
  - Design of chip to implement **ISA**
- Touch on compilers & OS ( $n + 1$ ), circuits ( $n - 1$ ) as well

# Technology & Energy

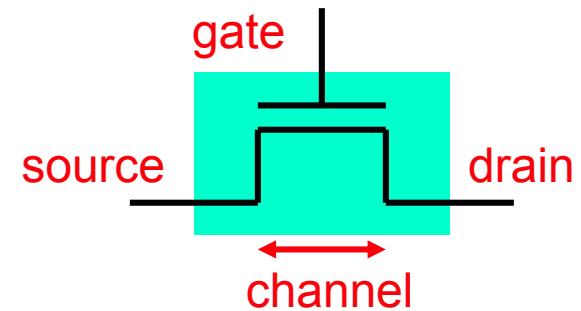
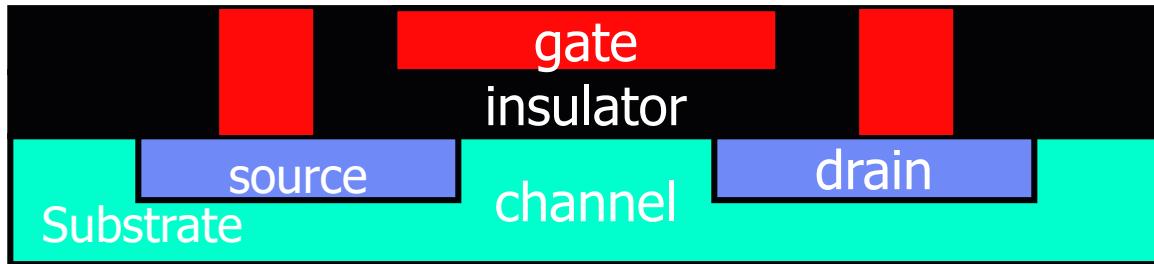
# Technology & Energy

---

- Technology basis
  - Fabrication (manufacturing) & cost
  - Transistors & wires
  - Implications of transistor scaling (Moore's Law)
- Energy & power

# Technology & Fabrication

# Semiconductor Technology

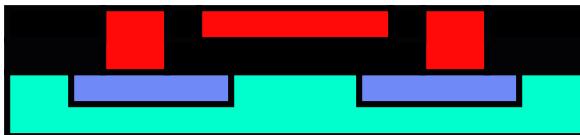


- Basic technology element: **MOSFET**
  - Solid-state component acts like electrical switch
  - **MOS**: metal-oxide-semiconductor
    - Conductor, insulator, semi-conductor
- **FET**: field-effect transistor
  - Channel conducts source→drain only when voltage applied to gate
- **Channel length**: characteristic parameter (short → fast)
  - Aka “feature size” or “technology node”
  - Currently: 14 nanometers (nm)
  - Continued miniaturization (scaling) known as “**Moore’s Law**”
    - Won’t last forever, physical limits approaching (or are they?)

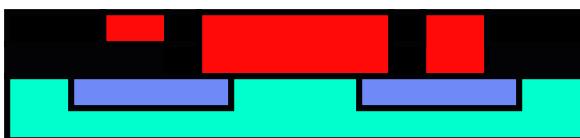
# Manufacturing Defects

---

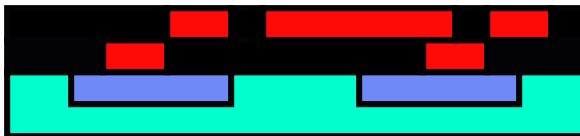
Correct:



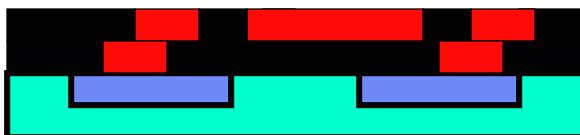
Defective:



Defective:



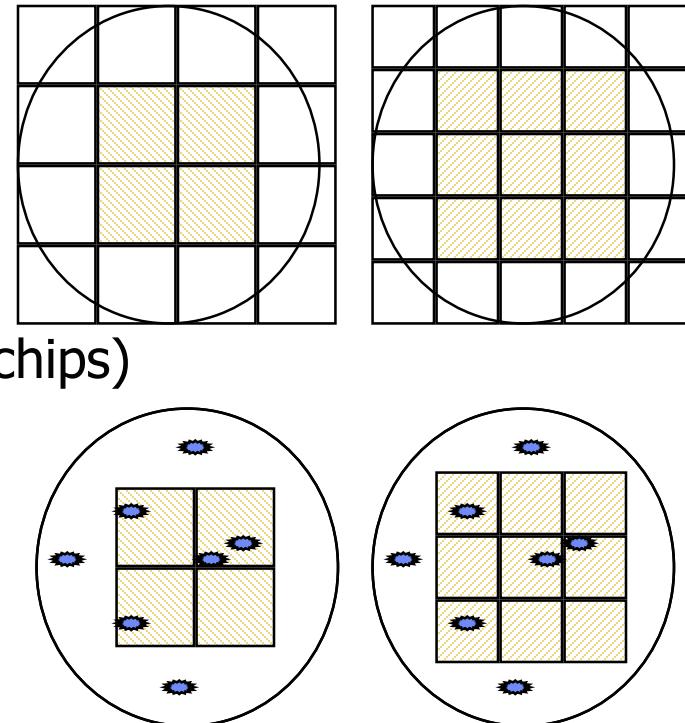
Slow:



- Defects can arise
  - Under-/over-doping
  - Over-/under-dissolved insulator
  - Mask mis-alignment
  - Particle contaminants
- Try to minimize defects
  - Process margins
  - Design rules
    - Minimal transistor size, separation
- Or, tolerate defects
  - Redundant or “spare” memory cells
  - Can substantially improve yield

# Cost Implications of Defects

- Chips built in multi-step chemical processes on **wafers**
  - Cost / wafer is constant,  $f(\text{wafer size, number of steps})$
- Chip (die) cost is related to **area**
  - Larger chips means fewer of them
- Cost is **superlinear** in area
  - Why? random defects
  - **Larger chip, more chance of defect**
    - Result: lower “yield” (fewer working chips)
- **Wafer yield**: % wafer that is chips
- **Die yield**: % chips that work
- Yield is increasingly non-binary - fast vs slow chips



# Manufacturing Cost

---

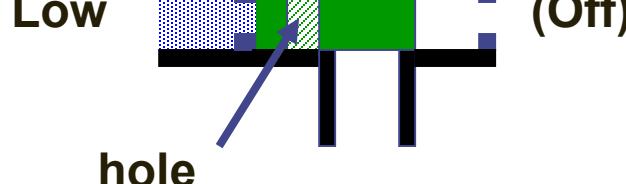
- **Chip cost vs system cost**
  - Cost of memory, storage, display, battery, etc.
- **Cost vs price**
  - Relationship complicated; **microprocessors not commodities**
  - Specialization, compatibility, different cost/performance/power
  - Economies of scale
- **Unit costs: die manufacturing, testing, packaging, burn-in**
  - Die cost based on area & defect rate (yield)
  - Package cost related to heat dissipation & number of pins
- **Fixed costs: design & verification, fab cost**
  - Amortized over “proliferations”, e.g., Core i3, i5, i7 variants
  - Building new “fab” costs billions of dollars today
  - Both getting worse; trend toward “foundry” & “fabless” models

# **Transistor Switching Speed**

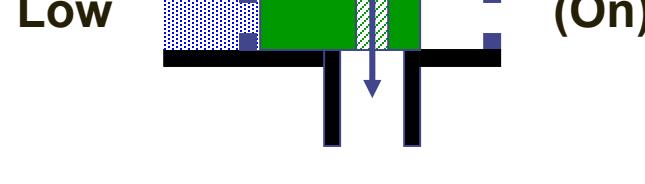
# A Transistor Analogy: Computing with Air

- Use air pressure to encode values
  - High pressure represents a “1” (blow)
  - Low pressure represents a “0” (suck)
- Valve can allow or disallow the flow of air
  - Two types of valves

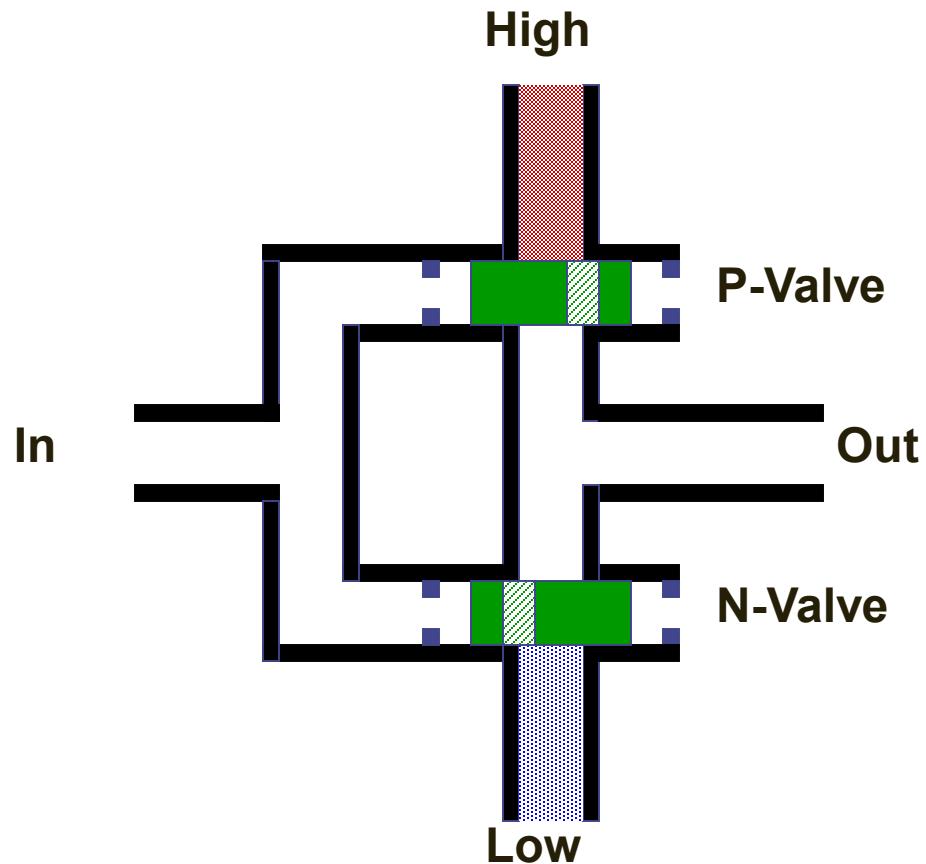
**N-Valve**



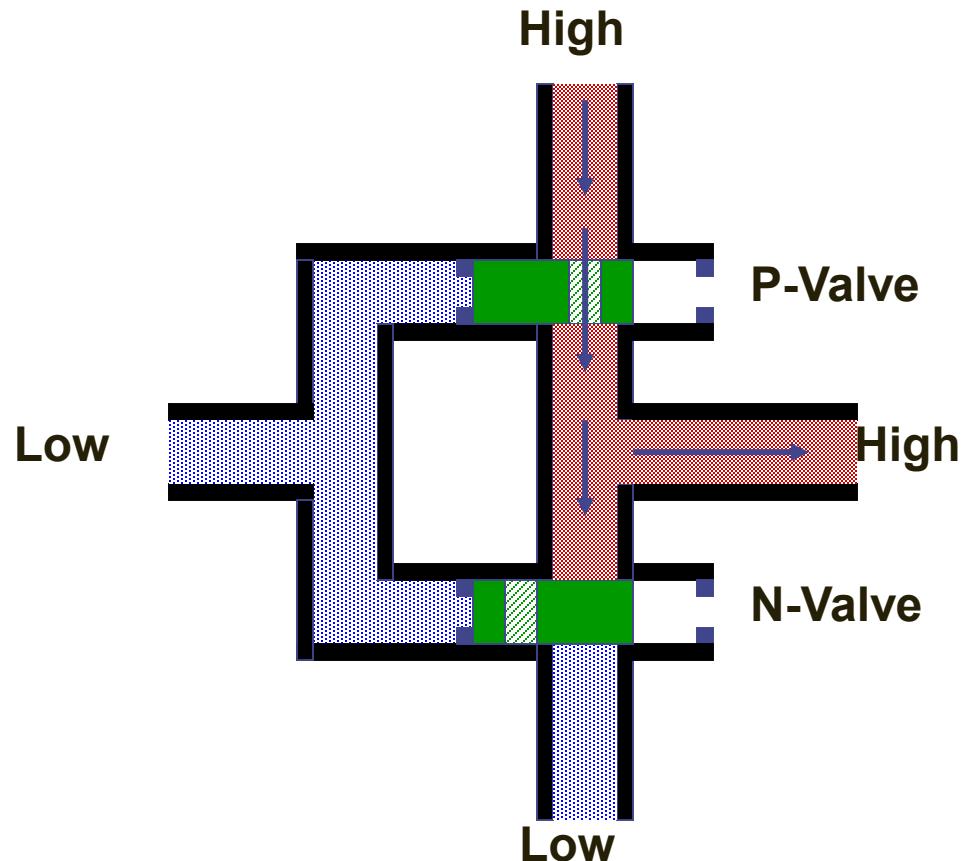
**P-Valve**



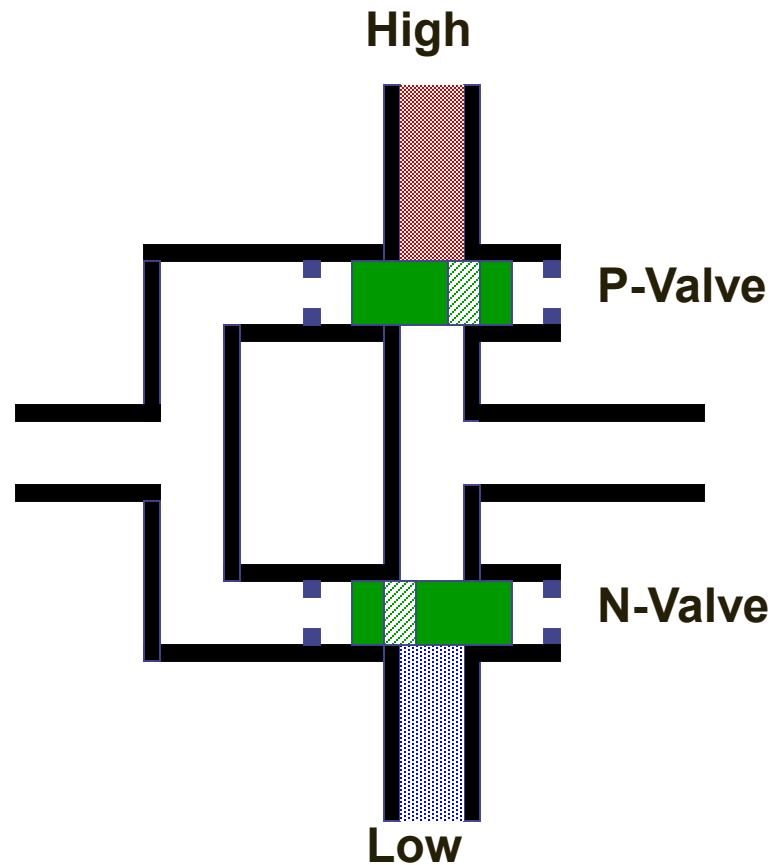
# Pressure Inverter



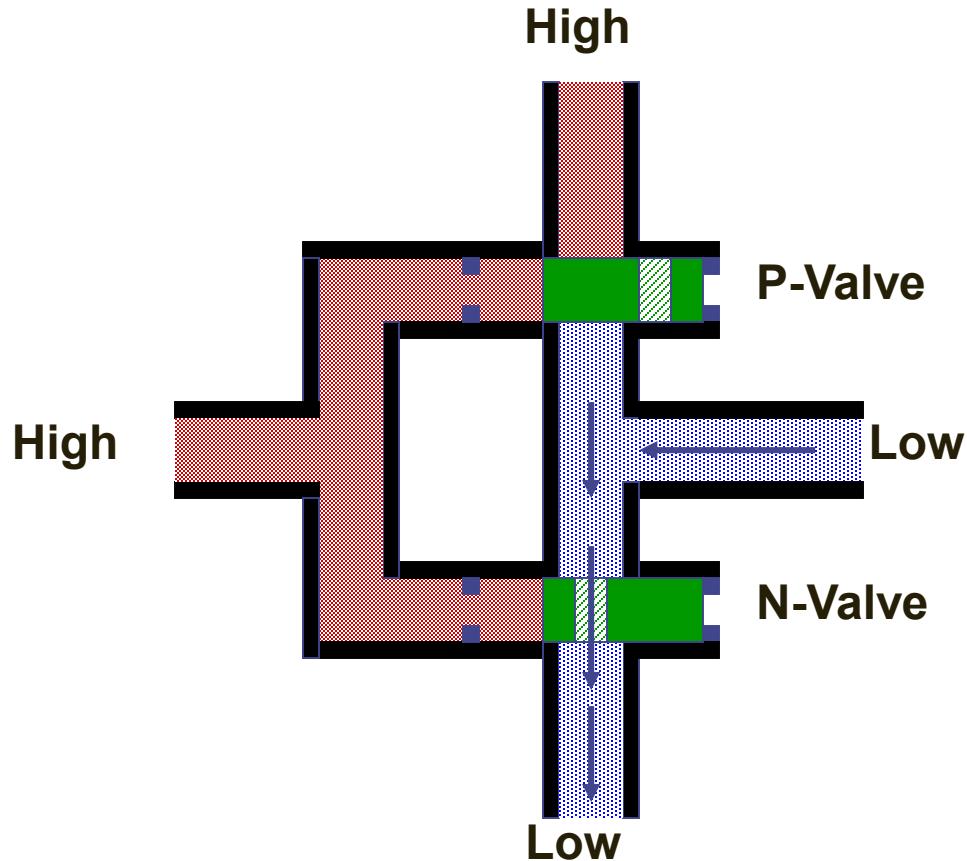
# Pressure Inverter (Low to High)



# Pressure Inverter



# Pressure Inverter (High to Low)



# Analogy Explained

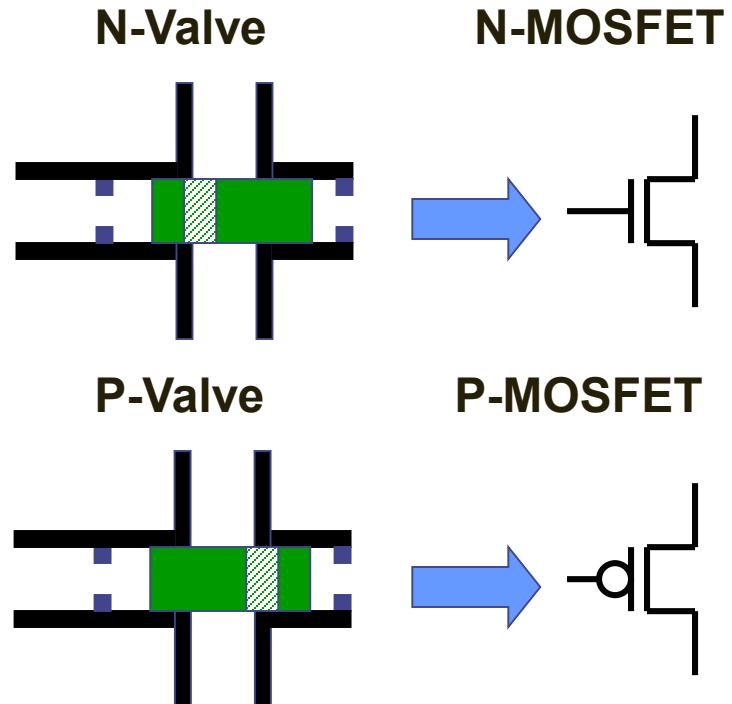
---

- Pressure differential → electrical potential (voltage)
  - Air molecules → electrons
  - Pressure (molecules per volume) → voltage
  - High pressure → high voltage
  - Low pressure → low voltage
- Air flow → electrical current
  - Pipes → wires
  - Air only flows from high to low pressure
  - Electrons only flow from high to low voltage
  - Flow only occurs when changing from 1 to 0 or 0 to 1
- Valve → transistor
  - The transistor: one of the century's most important inventions

# Transistors as Switches

---

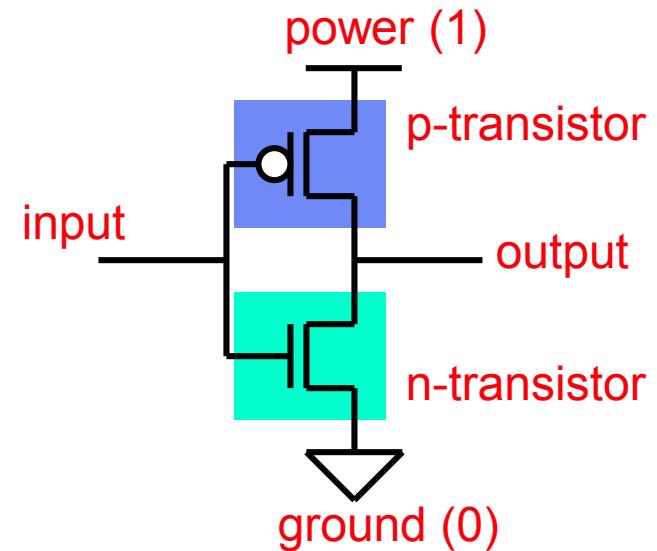
- Two types
  - N-type
  - P-type
- Properties
  - Solid state (no moving parts)
  - Reliable (low failure rate)
  - Small (14nm channel length)
  - Fast (<0.1ns switch latency)



# Complementary MOS (CMOS)

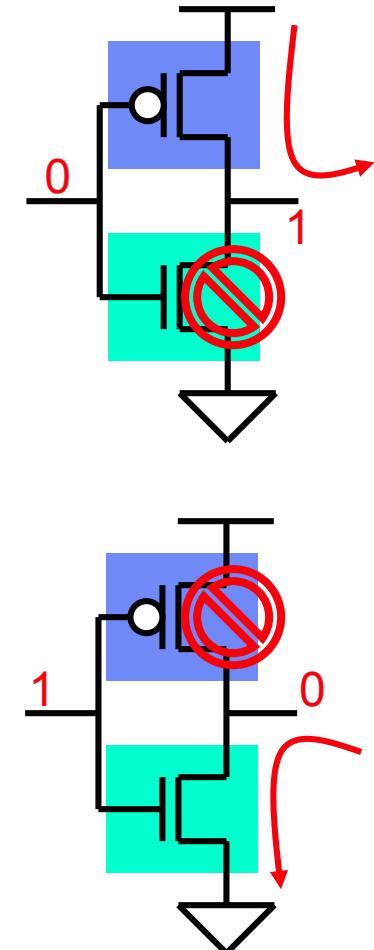
---

- Voltages as values
  - Power ( $V_{DD}$ ) = "1", Ground = "0"
- Two kinds of MOSFETs
  - **N-transistors**
    - Conduct when gate voltage is 1
    - Good at passing 0s
  - **P-transistors**
    - Conduct when gate voltage is 0
    - Good at passing 1s
- **CMOS**
  - Complementary n/p- networks form boolean logic (i.e., gates)
  - And some non-gate elements too (important example: RAMs)



# Basic CMOS Logic Gate

- **Inverter**: NOT gate
  - One p-transistor, one n-transistor
  - Basic operation
  - Input = 0
    - P-transistor closed, n-transistor open
    - Power charges output (1)
  - Input = 1
    - P-transistor open, n-transistor closed
    - Output discharges to ground (0)

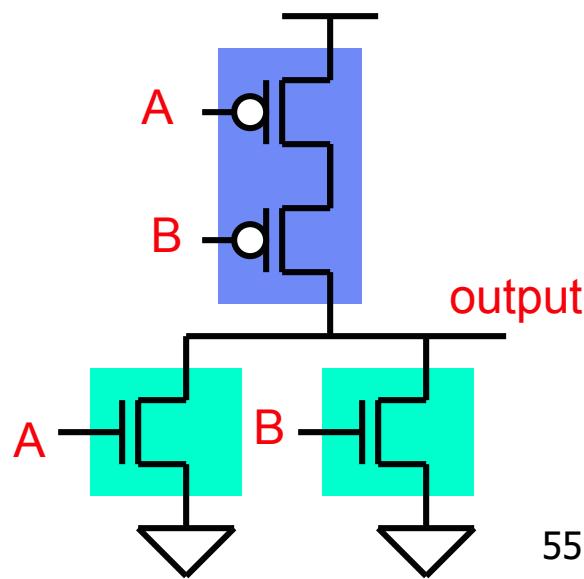
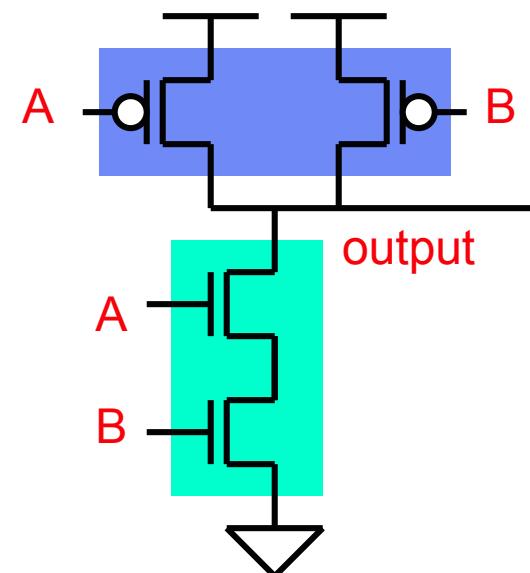


# Another CMOS Gate Example

- What is this? Look at **truth table**

- $0, 0 \rightarrow 1$
- $0, 1 \rightarrow 1$
- $1, 0 \rightarrow 1$
- $1, 1 \rightarrow 0$
- Result: **NAND** (NOT AND)
- NAND is “universal”

- What function is this?

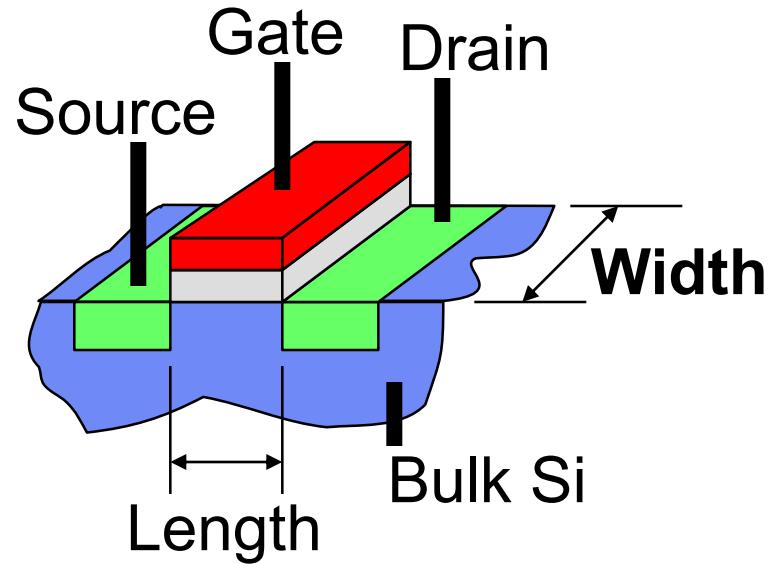
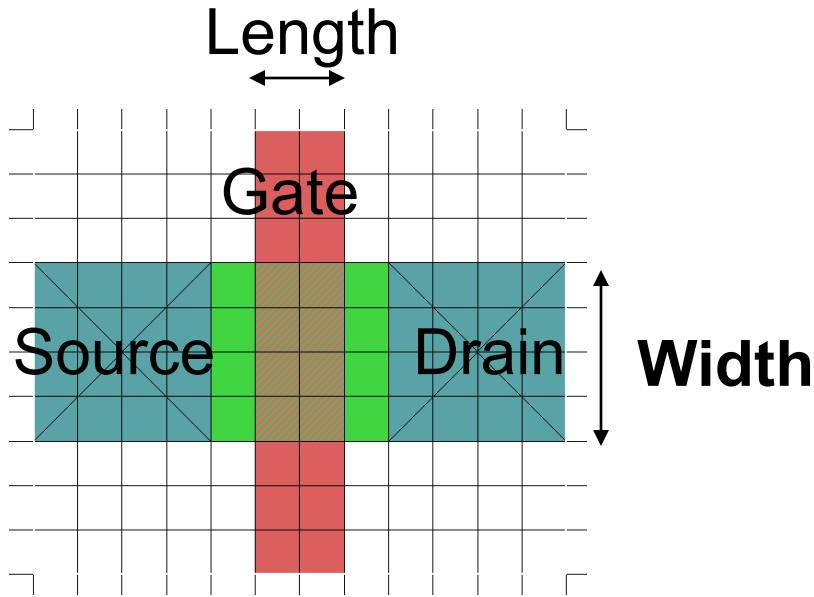


# Technology Basis of Transistor Speed

---

- Physics 101: delay through an electrical component  $\propto \text{RC}$ 
  - **Resistance (R)**   $\propto$  length / cross-section area
    - Slows rate of charge flow
  - **Capacitance (C)**   $\propto$  length \* area / distance-to-other-plate
    - Stores charge
  - **Voltage (V)**
    - Electrical pressure
  - **Threshold Voltage ( $V_t$ )**
    - Voltage at which a transistor turns “on”
    - Property of transistor based on fabrication technology
  - **Switching time**  $\propto (R * C) / (V - V_t)$
- Two kinds of electrical components
  - CMOS transistors (gates, sources, drains)
  - Wires

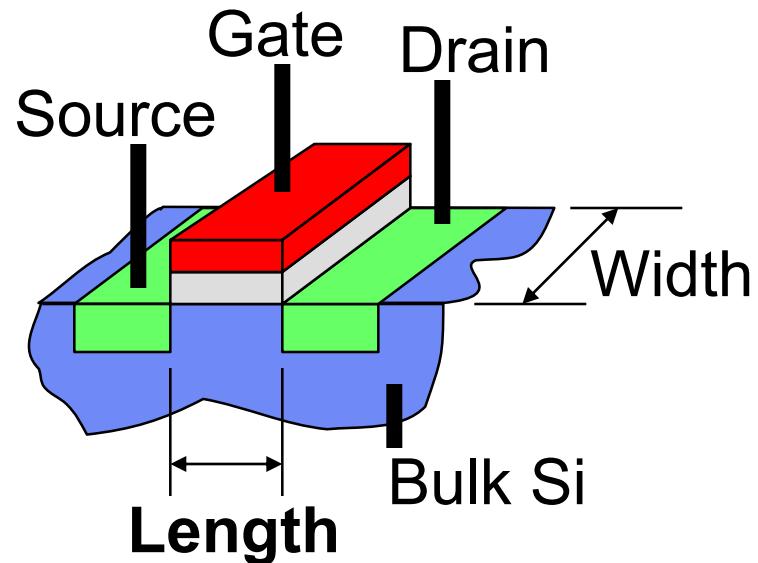
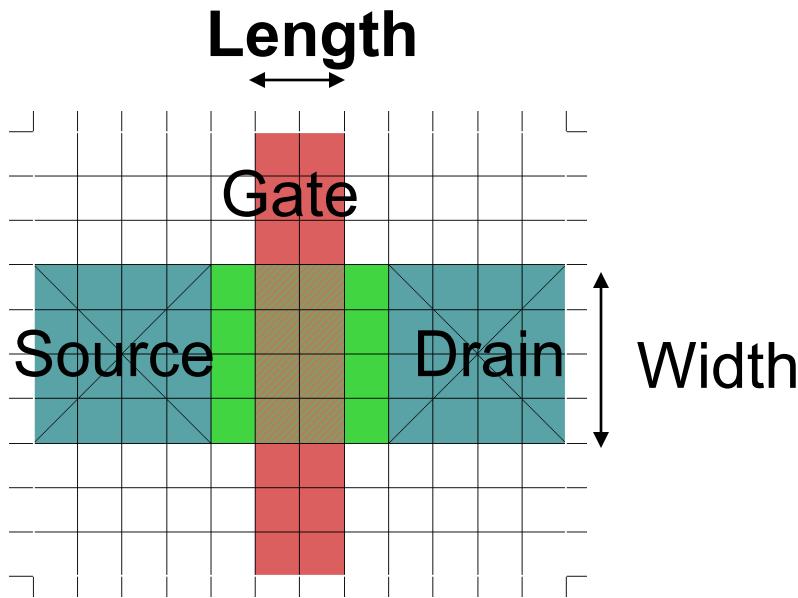
# Transistor Geometry: Width



Diagrams © Krste Asanovic, MIT

- **Transistor width**, set by designer for each transistor
- Wider transistors:
  - **Lower resistance** of channel (increases drive strength) – good!
  - But, **increases capacitance** of gate/source/drain – bad!
- Result: set width to balance these conflicting effects

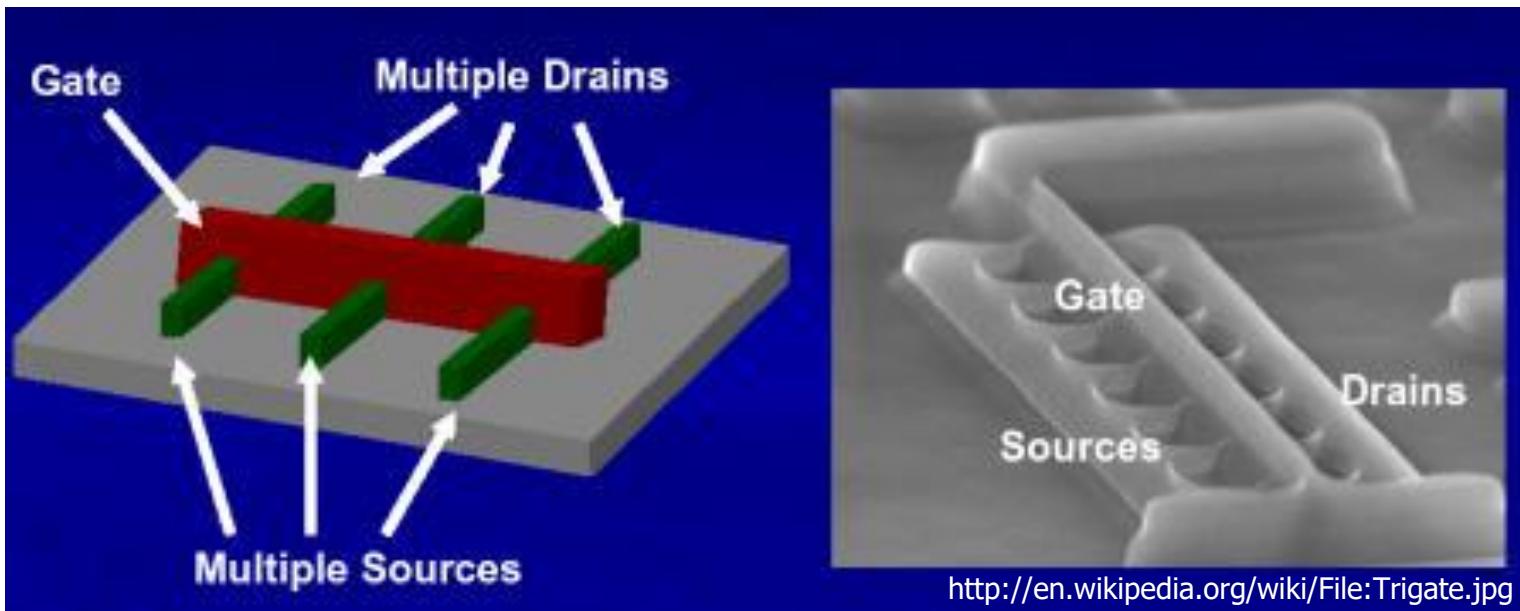
# Transistor Geometry: Length & Scaling



Diagrams © Krste Asanovic, MIT

- **Transistor length:** characteristic of “process generation”
  - “22nm” refers to the transistor gate length
- Each process generation shrinks transistor length by 1.4x
  - “Moore’s law” -> roughly 2x improvement in transistor density
  - Roughly linear improvement in switching speeds (lower resistance)

# Trigate FinFET Transistors



- nonplanar (or “3D”) transistors
  - trigate: multiple sources/drains/gates
  - FinFET: gate is wrapped around the channel
- lower leakage, faster switching times
- Intel’s trigate design released in mid-2012 (Ivy Bridge)
  - TSMC FinFET in mid-2015, Samsung FinFET in early 2016

# Technology Scaling Trends

# Dennard Scaling

“Design of ion-implanted MOSFET’s with very small physical dimensions”  
Robert H. Dennard, Fritz H. Gaensslen, Hwa-Nien Yu, V. Leo Rideout, Ernest  
Bassous, and Andre R. LeBlanc  
IEEE Journal of Solid-State Circuits, October 1974

**Table 1**  
Scaling Results for Circuit Performance

Device or Circuit Parameter	Scaling Factor
Device dimension $t_{ox}$ , $L$ , $W$	$1/\kappa$
Doping concentration $N_a$	$\kappa$
Voltage $V$	$1/\kappa$
Current $I$	$1/\kappa$
Capacitance $\epsilon A/t$	$1/\kappa$
Delay time/circuit $VC/I$	$1/\kappa$
Power dissipation/circuit $VI$	$1/\kappa^2$
Power density $VI/A$	1

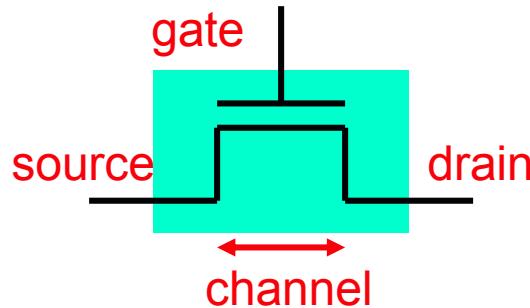
# Dennard Scaling

---

- stopped in ~2005 due to leakage
  - $V$  close to  $V_t$ , transistors never really “on” or “off”
  - gate-oxide leakage due to very small oxide thickness
    - quantum-mechanical electron tunneling
- Moore’s Law still in effect!
  - but power usage goes up with increasing transistor counts

# Moore's Law: Technology Scaling

---



- **Moore's Law:** aka “technology scaling”
  - Continued miniaturization (esp. reduction in channel length)
    - + Improves switching speed, power/transistor, area(cost)/transistor
    - Reduces transistor reliability
  - Literally: DRAM density (transistors/area) doubles every 18 months
  - Public interpretation: performance doubles every 18 months
    - Not quite right, but helps performance in several ways...

# Moore's Law in the Future

---

- Won't last forever, approaching physical limits
  - "If something must eventually stop, it can't go on forever"
  - But betting against it has proved foolish in the past
  - Perhaps will "slow" rather than stop abruptly
- Transistor count will likely continue to scale
  - 3D "die stacking" is arriving
- But transistor performance scaling?
  - Running into physical limits
  - Example: gate oxide < 10 silicon atoms!
    - Can't decrease it much further
  - Power is becoming the limiting factor

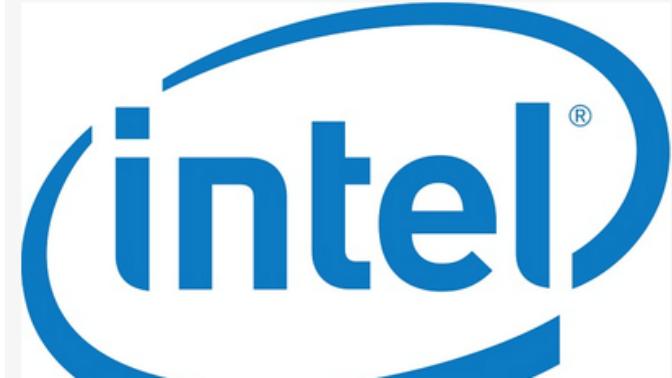


The image shows a screenshot of an AnandTech news article. The header features the AnandTech logo and navigation icons for menu, home, and search. The main title is "Tick Tock On The Rocks: Intel Delays 10nm, Adds 3rd Gen 14nm Core Product 'Kaby Lake'". Below the title is the author information "by Brett Howse & Ryan Smith on July 16, 2015 10:15 AM EST" and a link to "136 Comments | Add A Comment". The bottom of the image features the large, stylized Intel logo.

## Tick Tock On The Rocks: Intel Delays 10nm, Adds 3rd Gen 14nm Core Product "Kaby Lake"

by Brett Howse & Ryan Smith on July 16, 2015 10:15 AM EST

136 Comments | Add A Comment



# Power & Energy

# Power/Energy Are Increasingly Important

---

- **Battery life** for mobile devices
  - Laptops, phones, cameras
- **Tolerable temperature** for devices without active cooling
  - Power means temperature, active cooling means **cost**
  - No room for a fan in a cell phone, no market for a hot cell phone
- **Electric bill** for compute/data centers
  - Pay for power twice: once in, once out (to cool)
- **Environmental concerns**
  - IT accounts for growing fraction of electricity consumption

# Energy & Power

---

- **Energy**: measured in Joules or Watt-seconds
  - Total amount of energy stored/used
  - Battery life, electric bill, environmental impact
  - Instructions per Joule (car analogy: miles per gallon)
- **Power**: energy per unit time (measured in Watts)
  - Related to “performance” (which is also a “per unit time” metric)
  - Power impacts power supply and cooling requirements (cost)
    - Power-density (Watt/mm<sup>2</sup>): important related metric
  - Peak power vs average power
    - E.g., camera: power “spikes” when you actually take a picture
  - Joules per second (car analogy: gallons per hour)
- Two sources:
  - **Dynamic power**: active switching of transistors
  - **Static power**: leakage of transistors even while inactive

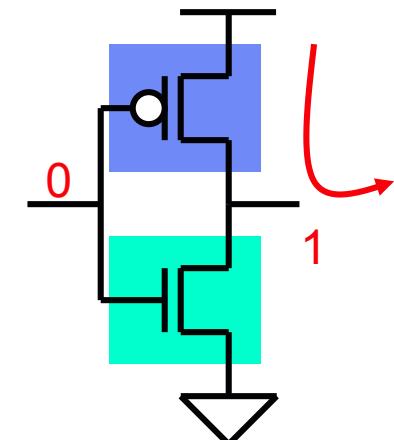
# Recall: Tech. Basis of Transistor Speed

---

- Physics 101: delay through an electrical component  $\propto \text{RC}$ 
  - **Resistance (R)**   $\propto$  length / cross-section area
    - Slows rate of charge flow
  - **Capacitance (C)**   $\propto$  length \* area / distance-to-other-plate
    - Stores charge
  - **Voltage (V)**
    - Electrical pressure
  - **Threshold Voltage ( $V_t$ )**
    - Voltage at which a transistor turns “on”
    - Property of transistor based on fabrication technology
  - **Switching time**  $\propto (R * C) / (V - V_t)$

# Dynamic Power

- **Dynamic power ( $P_{dynamic}$ )**: aka switching or active power
  - Energy to switch a gate (0 to 1, 1 to 0)
  - Each gate has capacitance (C)
    - Charge stored  $\propto C * V$
    - Energy to charge/discharge a capacitor  $\propto C * V^2$
    - Time to charge/discharge a capacitor  $\propto V$ 
      - Result: frequency  $\propto V$
  - $P_{dynamic} \approx N * C * V^2 * f * A$ 
    - N: number of transistors
    - C: capacitance per transistor (size of transistors)
    - V: voltage (supply voltage for gate)
    - f: frequency (transistor switching freq.  $\propto$  clock freq.)
    - A: activity factor (not all transistors may switch this cycle)



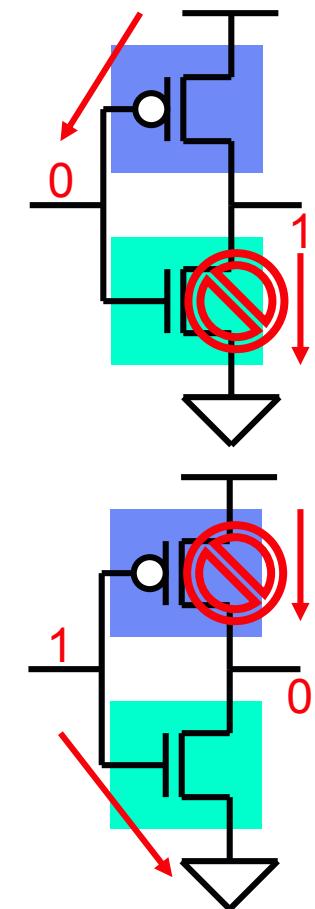
# Reducing Dynamic Power

---

- Target each component:  $P_{\text{dynamic}} \approx N * C * V^2 * f * A$
- **Reduce number of transistors (N)**
  - Use fewer transistors and gates
- **Reduce capacitance (C)**
  - Smaller transistors (Moore's law)
- **Reduce voltage (V)**
  - Quadratic reduction in energy consumption!
  - But also slows transistors (transistor speed  $\propto V$ )
- **Reduce frequency (f)**
  - Slower clock frequency (reduces power but not energy) Why?
- **Reduce activity (A)**
  - “Clock gating” disable clocks to unused parts of chip
  - Don’t switch gates unnecessarily

# Static Power

- **Static power ( $P_{\text{static}}$ )**: aka idle or leakage power
  - Transistors don't turn off all the way
  - Transistors are “leaky valves”
  - $P_{\text{static}} \approx N * V * e^{-V_t}$
  - N: number of transistors
  - V: voltage
  - **$V_t$  (threshold voltage)**: voltage at which transistor conducts (begins to switch)
- Switching speed vs leakage trade-off
- The lower the  $V_t$ :
  - Faster transistors (linear)
    - Transistor speed  $\propto V - V_t$
  - Leakier transistors (exponential)



# Reducing Static Power

---

- Target each component:  $P_{\text{static}} \approx N * V * e^{-Vt}$
- **Reduce number of transistors (N)**
  - Use fewer transistors/gates
- **Disable transistors** (also targets N)
  - “Power gating” disable power to unused parts (long latency to power up)
  - Power down units (or entire cores) not being used
- **Reduce voltage (V)**
  - Linear reduction in static energy consumption
  - But also slows transistors (transistor speed  $\propto V$ )
- **Dual  $V_t$**  – use a mixture of high and low  $V_t$  transistors
  - Use slow, low-leak transistors in SRAM arrays
  - Requires extra fabrication steps (cost)
- **Low-leakage transistors**
  - High-K/Metal-Gates in Intel’s 45nm process, “tri-gate” in Intel’s 22nm
  - Reducing frequency can hurt energy efficiency due to leakage power

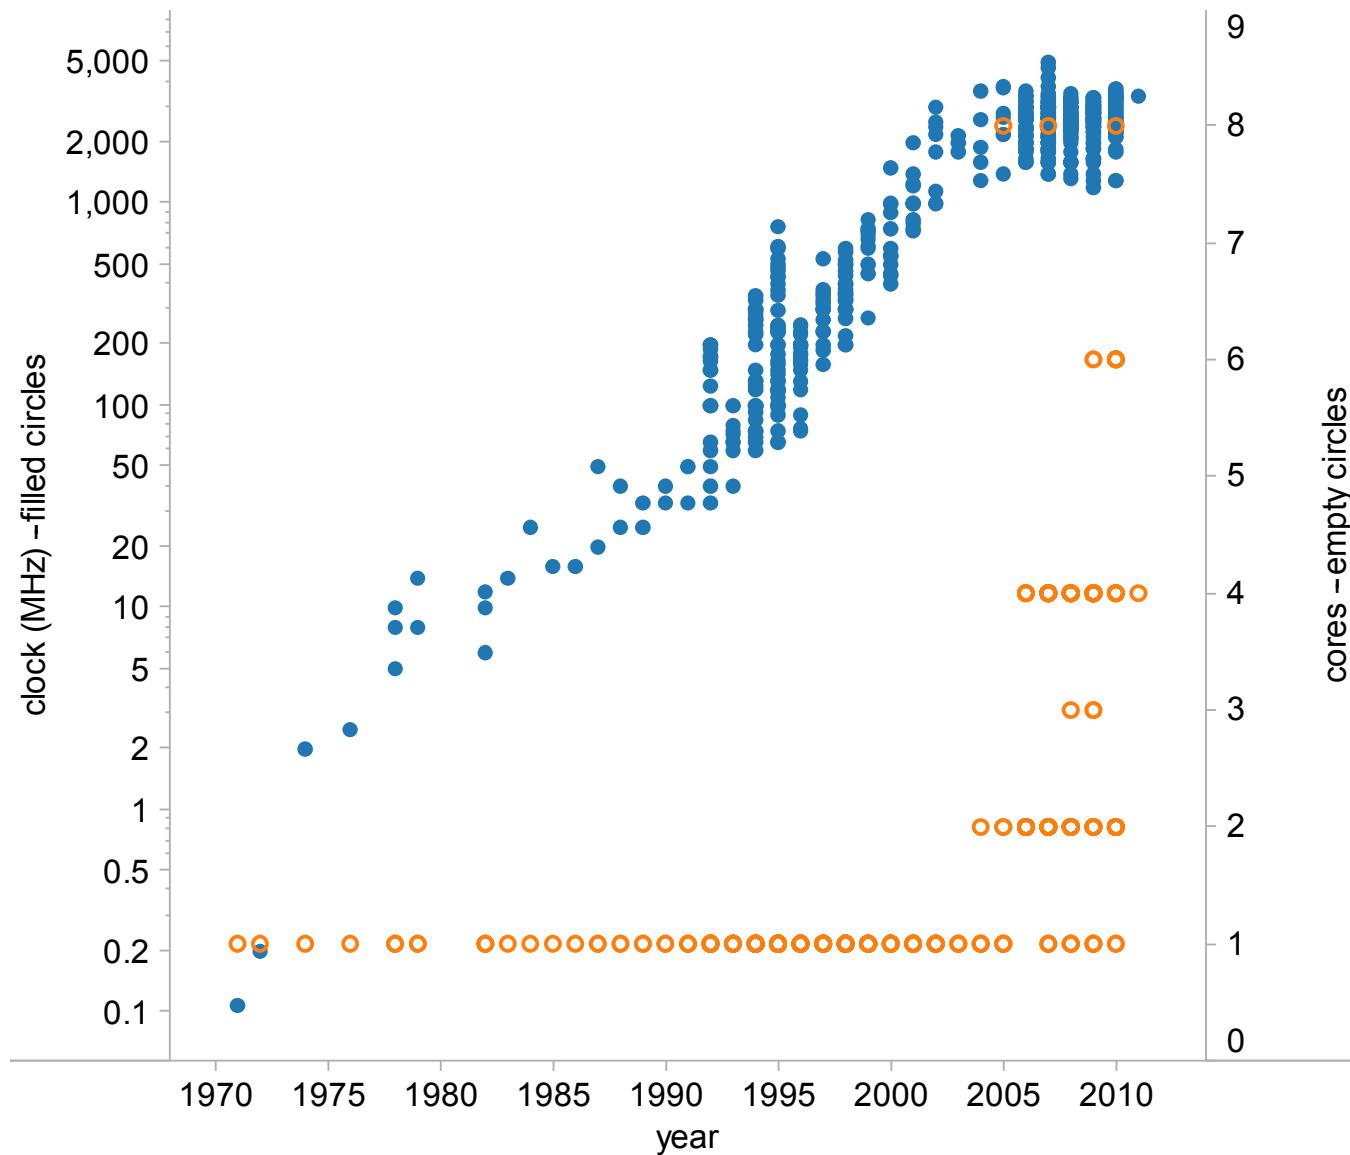
# Dynamic Voltage/Frequency Scaling

---

- **Dynamically trade-off power for performance**
  - Change the voltage and frequency at runtime
  - Under control of operating system and/or hardware
- Recall:  $P_{\text{dynamic}} \approx N * C * V^2 * f * A$ 
  - Because frequency  $\propto$  to  $V - V_t$ ...
  - $P_{\text{dynamic}} \propto$  to  $V^2(V - V_t) \approx V^3$
- Reduce both voltage and frequency linearly
  - **Cubic decrease in dynamic power**
  - Linear decrease in performance (actually sub-linear)
    - Thus, only about quadratic decrease in energy
  - Linear decrease in static power
    - Thus, static energy can become dominant
- Newer chips can adjust frequency on a per-core basis

# Frequency and Core Count

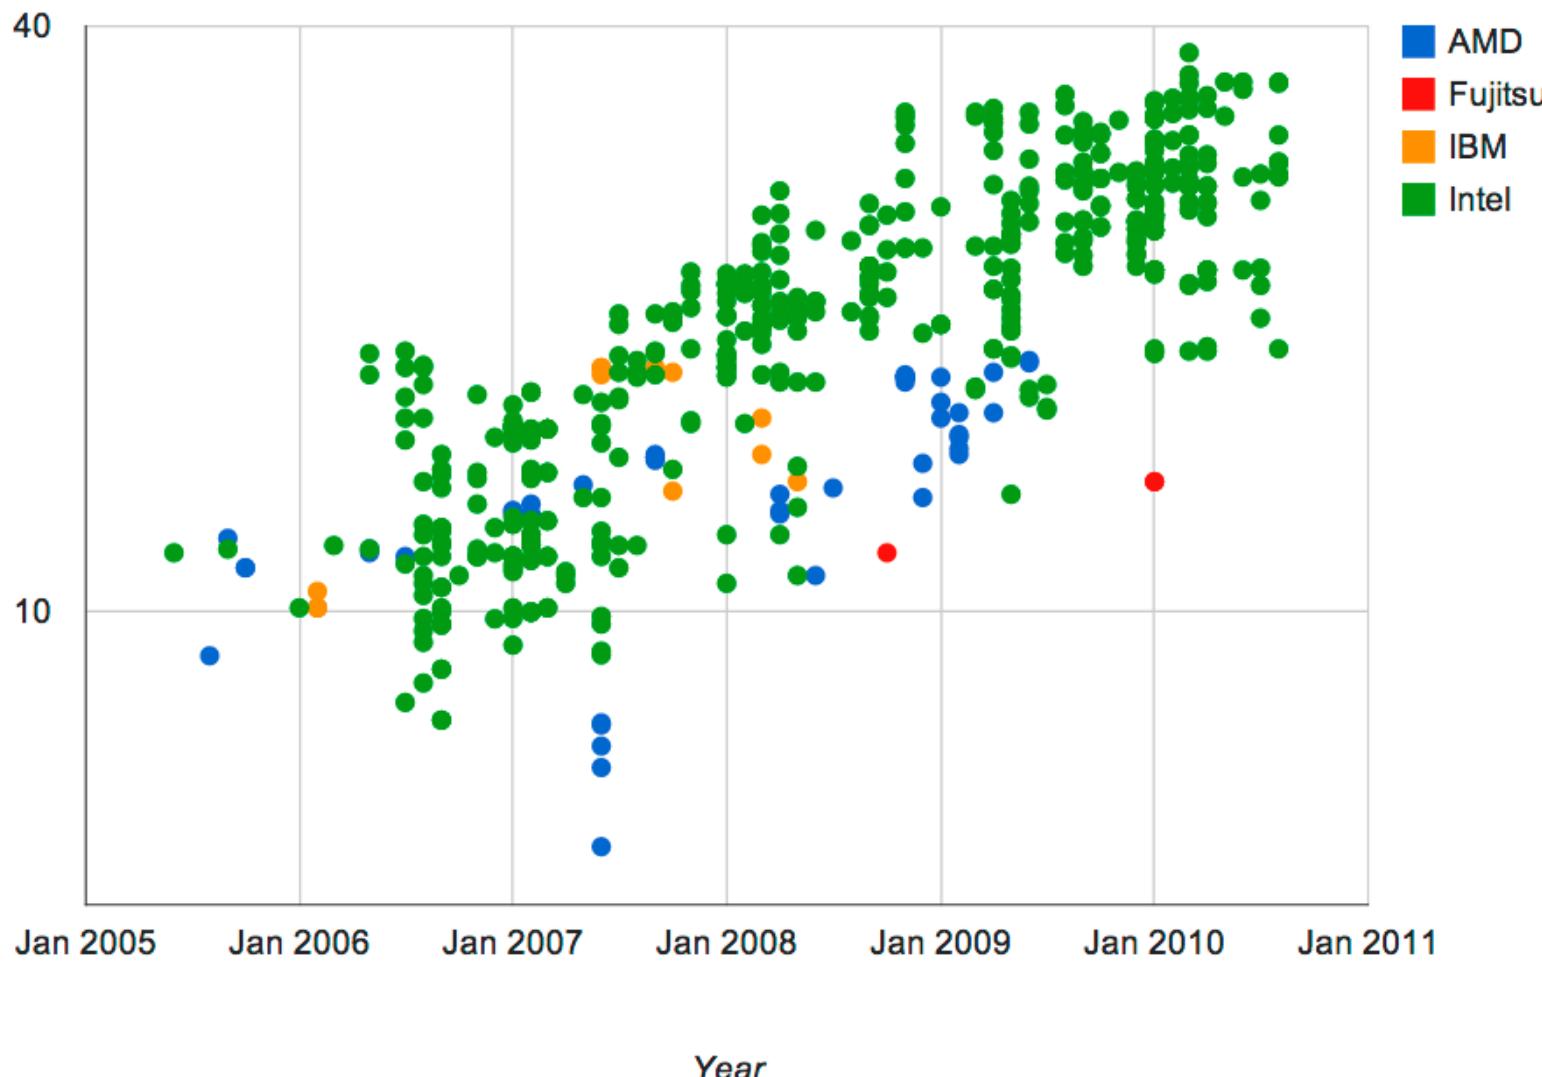
data from  
<http://cpudb.stanford.edu>



# SpecINT 2006 performance

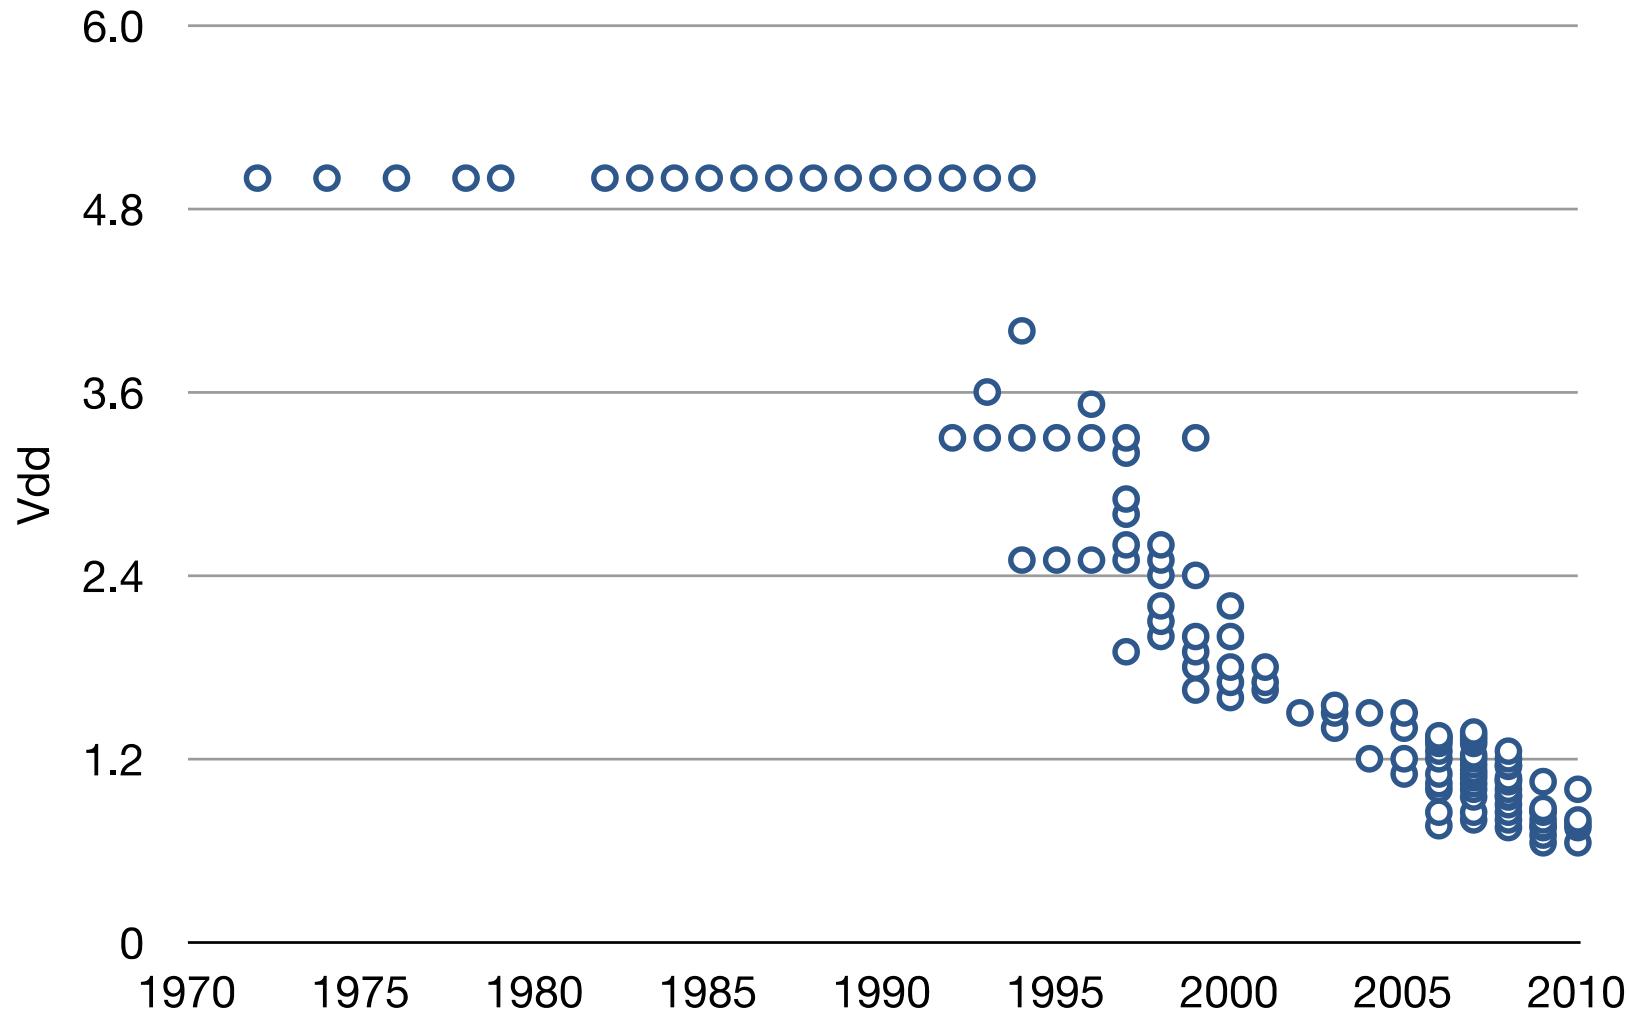
graph from  
<http://cpudb.stanford.edu>

Spec Int 2006



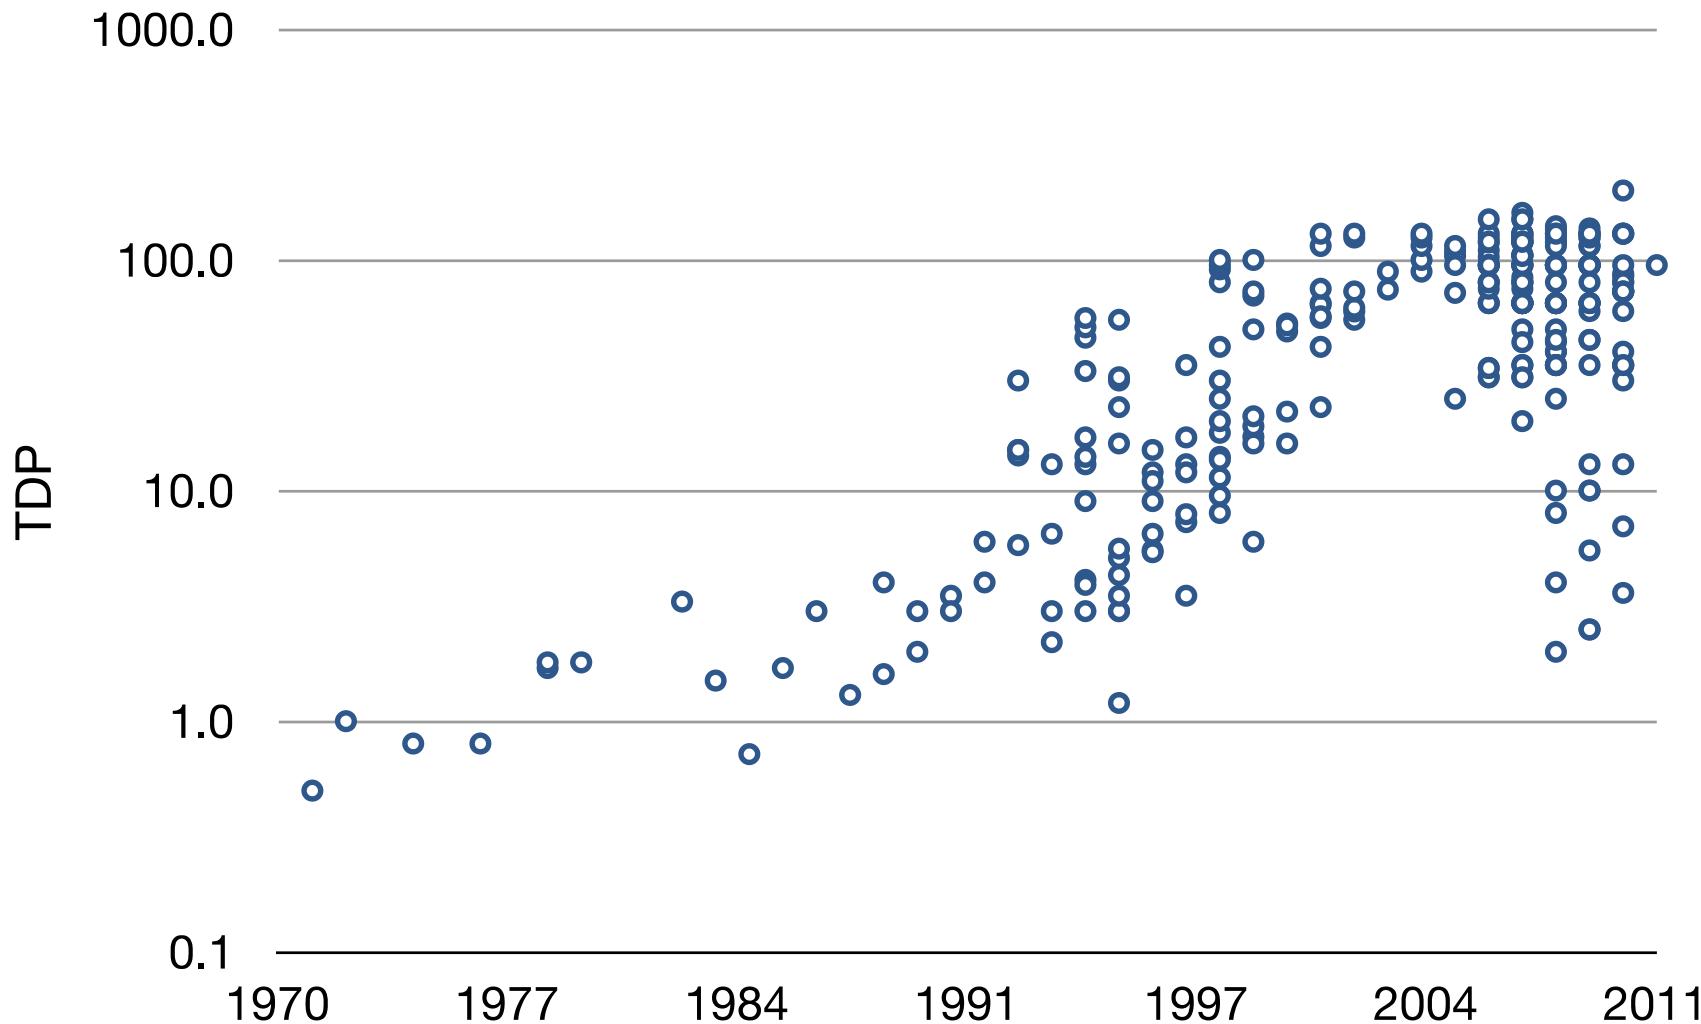
# Supply Voltage

data from  
<http://cpudb.stanford.edu>



# Thermal Design Power

data from  
<http://cpudb.stanford.edu>



# Moore without Dennard

---

- + Dennard scaling reduced power/transistor...
  - Required reducing  $V$ , which requires a trade-off:
    - Keeping  $V_t$  the same and reducing frequency ( $f$ )
    - Lowering  $V_t$  and increasing leakage exponentially
- + Moore's Law still gives more transistors
  - + Use techniques like high-K/metal gate, dual- $V_T$ , tri-gate
- The end of voltage scaling & "dark silicon"
  - Current projections: power per transistor reduced by 25-35% per technology node
  - What are the implications?

# Implications on Software

---

- Software-controlled dynamic voltage/frequency scaling
  - Example: video decoding
    - Too high a clock frequency – wasted energy (battery life)
    - Too low a clock frequency – quality of video suffers
  - “Race to sleep” versus “slow and steady” approaches
- Managing low-power modes
  - Don’t want to “wake up” the processor every millisecond
- Tuning software
  - Faster algorithms can be converted to lower-power algorithms
  - Via dynamic voltage/frequency scaling
- Exploiting parallelism & heterogeneous cores
  - ARM BIG.little design (a few “big” cores & 1 “low power” core)
- Specialized hardware accelerators

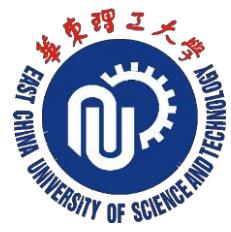
# Summary

# Technology Summary

---

- Has a first-order impact on computer architecture
  - Performance (transistor delay, wire delay)
  - Cost (die area & defects)
  - **Changing rapidly**
- Most significant trends for architects
  - More and more transistors
    - What to do with them? → integration → **parallelism**
  - Logic is improving faster than memory & cross-chip wires
    - “Memory wall” → caches, more integration
- Power and energy
  - Voltage vs frequency, parallelism, special-purpose hardware
- This unit: a quick overview, just scratching the surface

Rest of course



Thanks!  
Q&A

