Modern Computer Architecture

with sources that included UPenn & University of Wisconsin slides
by Joe Devietti, Milo Martin & Amir Roth ,   Mark Hill, Guri Sohi, Jim Smith, and David Wood
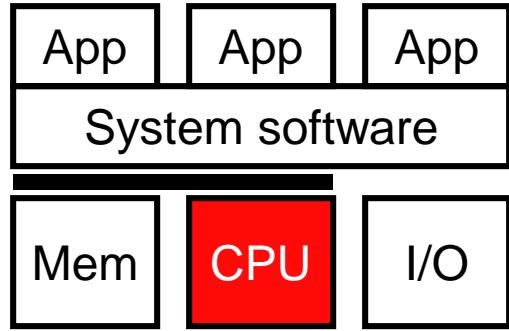
Schedule

- Introduction and Transistors
- Parallel computing (Isaac D. Scherson, University of California, Irvine, in October)
- ISAs
- Performance
- Pipelining Basic
- Branch Prediction
- Caches
- Virtual Memory
- Out-of-Order Execution
- Multicore multi-thread
- Vectors/GPUs for data parallelism.

This Unit: Pipelining
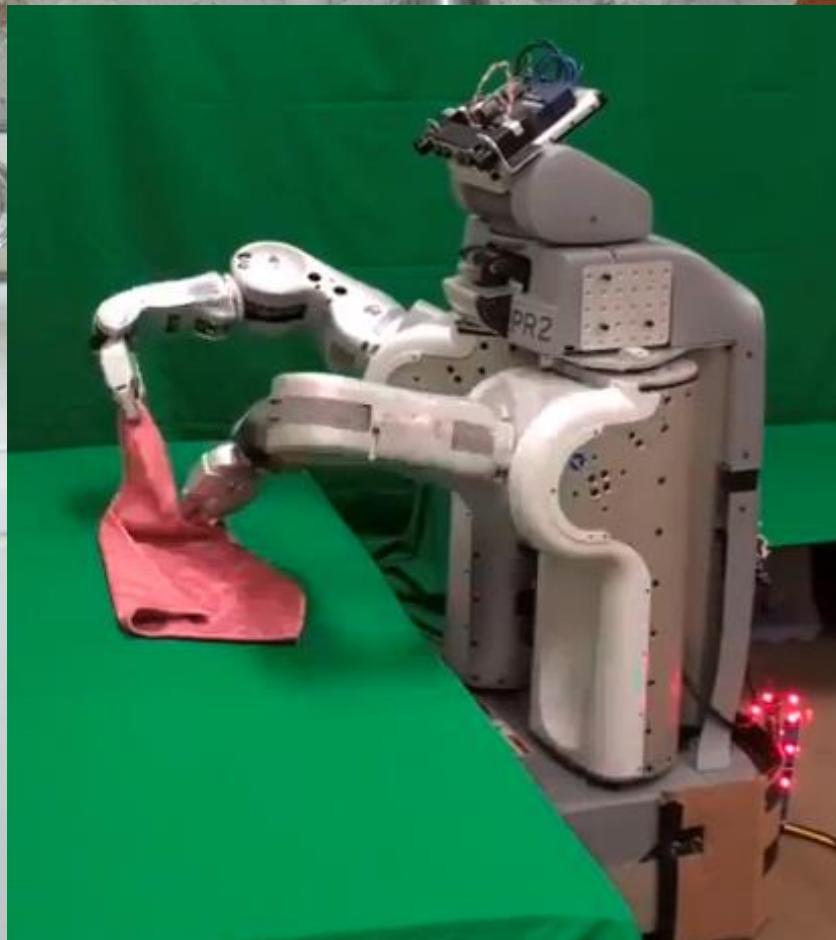
| App	App	App
System software		

Mem	CPU	I/O

- Single-cycle datapaths
- Latency vs. throughput & performance
- Basic pipelining
- Data hazards
    - Bypassing
    - Load-use stalling
- Pipelined multi-cycle operations
- Control hazards
    - Branch prediction

Main Concept

- Instructions broken down into finite set of assembly language instructions

- Instructions executed sequentially

- Pipelining method speeds up sequential execution of these instructions

The eternal pipelining metaphor

In-Class Exercise

- You have a washer, dryer, and "folding robot"
  - Each takes 1 unit of time per load
  - How long for one load in total?
  - How long for two loads of laundry?
  - How long for 100 loads of laundry?

- Now assume:
  - Washing takes 30 minutes, drying 60 minutes, and folding 15 min
  - How long for one load in total?
  - How long for two loads of laundry?
  - How long for 100 loads of laundry?

In-Class Exercise Answers
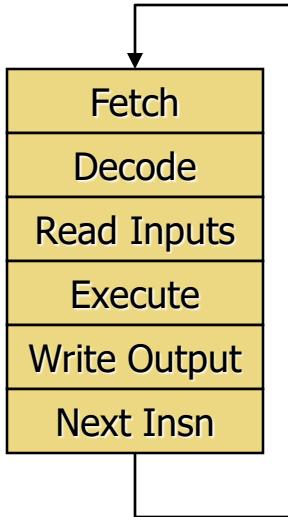
- You have a washer, dryer, and "folding robot"
  - Each takes 1 unit of time per load
  - How long for one load in total?
  - How long for two loads of laundry?
  - How long for 100 loads of laundry?

- Now assume:
  - Washing takes 30 minutes, drying 60 minutes, and folding 15 min
  - How long for one load in total?   30+60+15=**1h45m**
  - How long for two loads of laundry?   30+(60*2)+15 = **2h45m**
  - How long for 100 loads of laundry?   30+(60*100)+15= **100h45m**

Datapath Background

Recall: The Sequential Model

| Fetch |
| Decode |
| Read Inputs |
| Execute |
| Write Output |
| Next Insn |

- **Basic structure of all modern ISAs**
  - Often called Von Neumann, but in ENIAC before

- **Program order**: total order on dynamic insns
  - Order and **named storage** define computation

- Convenient feature: **program counter (PC)**
  - Insn itself stored in memory at location pointed to by PC
  - Next PC is next insn unless insn says otherwise

- Processor logically executes loop at left

- **Atomic**: insn finishes before next insn starts
  - Implementations can break this constraint physically
  - **But must maintain illusion to preserve correctness**

Recall: Maximizing Performance

**Execution time =
(instructions/program) * (seconds/cycle) * (cycles/instruction)**

**(1 billion instructions) * (1ns per cycle) * (1 cycle per insn)
= 1 second**

- Instructions per program:
  - Determined by program, compiler, instruction set architecture (ISA)
- **Cycles per instruction: "CPI"**
  - Typical range today: 2 to 0.5
  - Determined by program, compiler, ISA, micro-architecture
- **Seconds per cycle: "clock period" - same each cycle**
  - Typical range today: 2ns to 0.25ns
  - Reciprocal is frequency: 0.5 Ghz to 4 Ghz (1 Hz = 1 cycle per sec)
  - Determined by micro-architecture, technology parameters
- For minimum execution time, minimize each term
 - Difficult: *often pull against one another*

Modern Computer Architecture

Recall: Latency vs. Throughput

- **Latency (execution time)**: time to finish a fixed task
- **Throughput (bandwidth)**: number of tasks in fixed time
  - Different: exploit parallelism for throughput, not latency (e.g., bread)
  - Often contradictory (latency **vs.** throughput)
    - Will see many examples of this
  - Choose definition of performance that matches your goals
    - Scientific program? Latency, web server: throughput?
- Example: move people 10 miles
  - Car: capacity = 5, speed = 60 miles/hour
  - Bus: capacity = 60, speed = 20 miles/hour
  - Latency: **car = 10 min**, bus = 30 min
  - Throughput: car = 15 PPH (count return trip), **bus = 60 PPH**
- Fastest way to send 1PB of data?

Pipelined Datapath

Latency versus Throughput
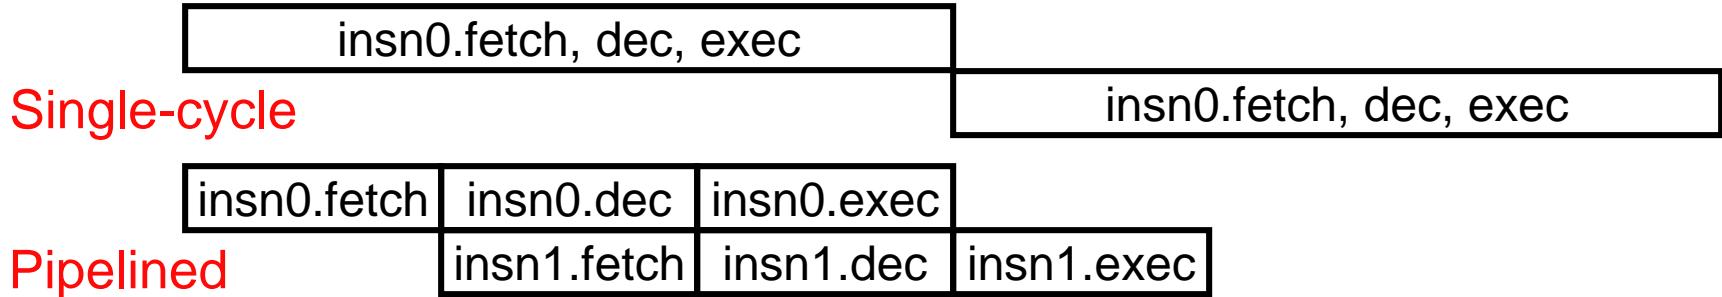
| insn0.fetch, dec, exec	
	insn1.fetch, dec, exec

Single-cycle

- Can we have both low CPI and short clock period?
  - Not if datapath executes only one insn at a time
- Latency and throughput: two views of performance …
  - (1) at the program level and (2) at the instruction level
- Single instruction latency
  - Doesn't matter: programs comprised of billions of instructions
  - Difficult to reduce anyway
- Goal is to make programs, not individual insns, go faster
  - Instruction throughput $\rightarrow$ program latency
  - Key: **exploit inter-insn parallelism**

Pipelining

insn0.fetch, dec, exec	

Single-cycle

	insn0.fetch, dec, exec

insn0.fetch	insn0.dec	insn0.exec
	insn1.fetch	insn1.dec	insn1.exec
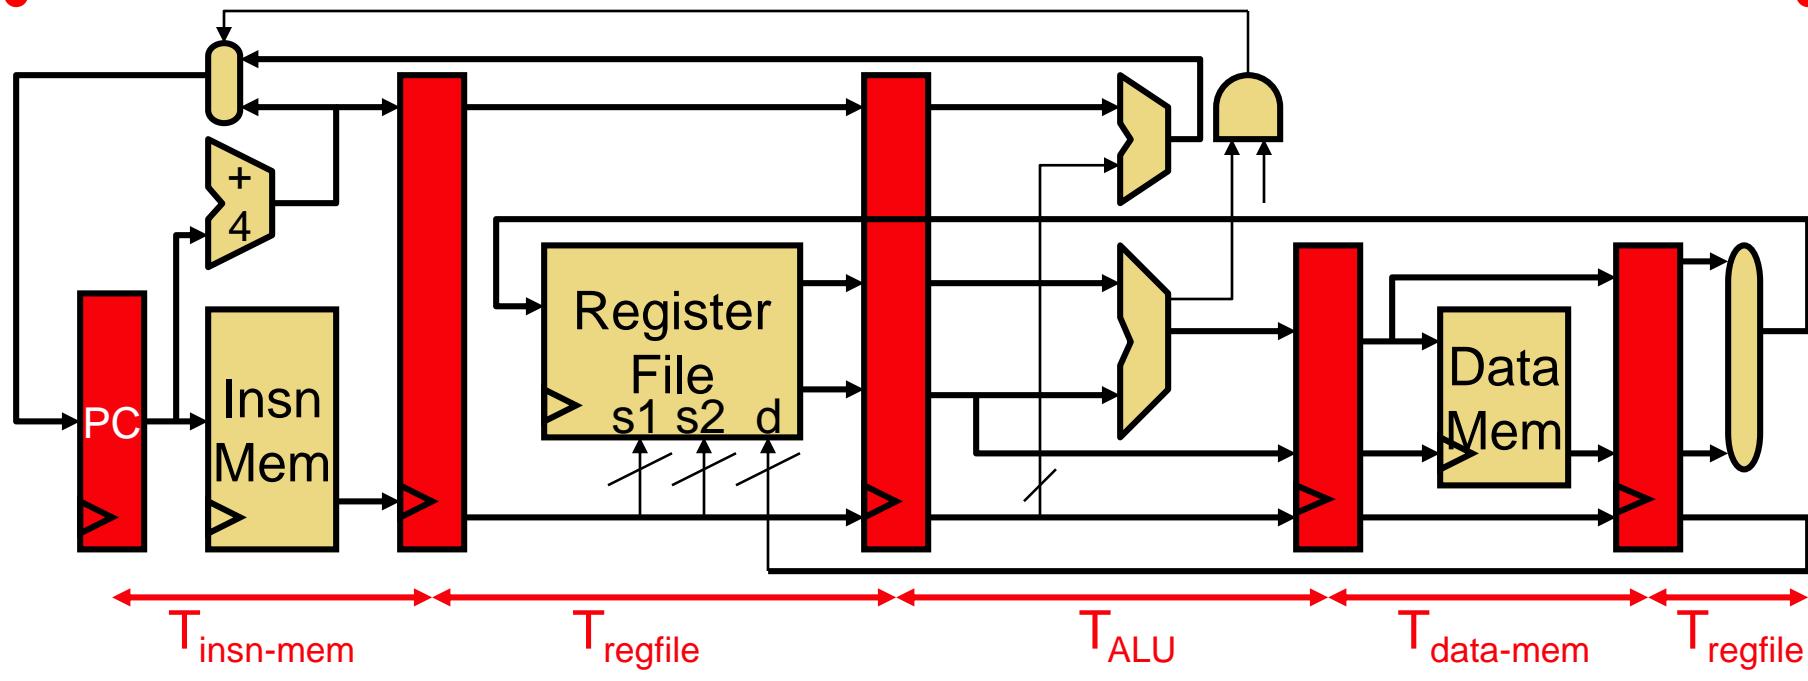
Pipelined

- Important performance technique
  - **Improves instruction throughput, not instruction latency**
- Begin with cycle design
  - When insn advances from stage 1 to 2, next insn enters at stage 1
  - Form of parallelism: "insn-stage parallelism"
  - Maintains illusion of sequential fetch/execute loop
  - Individual instruction takes the same number of stages
  - + **But instructions enter and leave at a much faster rate**
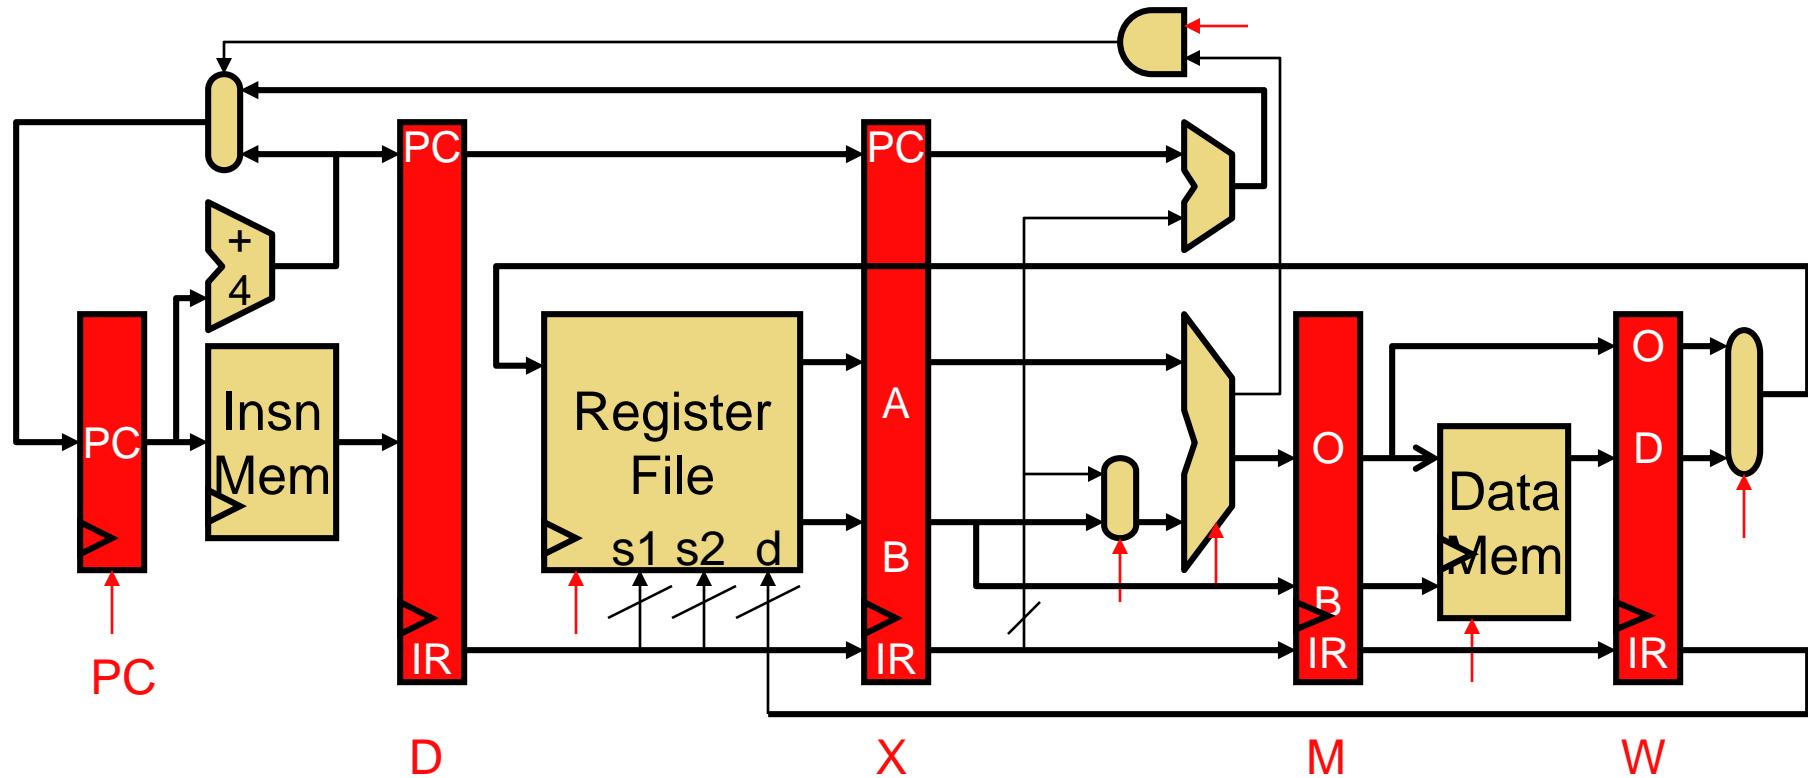
5 Stage Pipeline: Inter-Insn Parallelism

- **Pipelining**: cut datapath into N stages (here 5)
  - One insn in each stage in each cycle
  - + Clock period = MAX($T_{insn-mem}$, $T_{regfile}$, T_{ALU}, $T_{data-mem}$)
  - + Base CPI = 1: insn enters and leaves every cycle
  - – Actual CPI > 1: pipeline must often "stall"
  - Individual insn latency increases (pipeline overhead)

5 Stage Pipelined Datapath

- Five stage: **F**etch, **D**ecode, e**X**ecute, **M**emory, **W**riteback
  - Nothing magical about 5 stages (Pentium 4 had 22 stages!)
- **Latches** (pipeline registers) named by stages they begin(connect)
  - **PC**, **D(**IF/ID**)**, **X(**ID/EX**)**, **M(**EX/MEM**)**, **W(**MEM/WB**)**

What does each stage do?

- **Fetch**
  - read insn bytes from insn memory
- **Decode**
  - determine opcode, register operands, read from register file
- **Execute**
  - perform ALU operation
- **Memory**
  - read/write data memory
- **Writeback**
  - write destination to register file

Works in the MIPS 5 stage pipeline

- IF (Instruction fetch cycle)
    - IR←Mem[PC];
    - NPC ←PC=PC+4;
- ID (Instruction decode/register fetch cycle)
    - A ←Regs[rs];
    - B ←Regs[rt];
    - Imm ←sign-extended immediate field of IR;

- Note: The first two stages of MIPS pipeline do the same functions for all kinds of instructions.

PC, **D(**IF/ID**)**, **X(**ID/EX**)**, **M(**EX/MEM**)**, **W(**MEM/WB**)**

The third stage of MIPS pipeline

- EX (Execution/effective address cycle)
  - Memory reference:
    - ALUoutput ←A+Imm
  - Register-Register ALU instruction:
    - ALUoutput ←A func B;
  - Register-Immediate ALU instruction:
    - ALUoutput ←A op Imm;
  - Branch:
    - ALUoutput ←NPC+(Imm <<2);
    - Cond ←(A==0)

PC, **D(**IF/ID**)**, **X(**ID/EX**)**, **M(**EX/MEM**)**, **W(**MEM/WB**)**

The last two stages of MIPS pipeline

- ## MEM(Memory acces/branch completion cycle)
  - ### Memory reference(for Load/Store):
    - LMD ←Mem[ALUoutput] or   (LMD—load memory data register)
    - Mem[ALUoutput] ←B
  - ### Branch:
    - If (cond) PC ←ALUoutput
- ## WB (Write back cycle)
  - ### Register-Register ALU instruction
    - Regs[rd] ← ALUoutput;
  - ### Register-Immediate ALU instruction
    - Regs[rt] ←ALUoutput;
  - ### Load Instruction:
    - Regs[rt] ←LMD;

PC, **D(**IF/ID**)**, **X(**ID/EX**)**, **M(**EX/MEM**)**, **W(**MEM/WB**)**

Table: Events on every stage

| Stage	Any instruction		
IF	IF/ID.IR←Mem[PC]; IF/ID.NPC, PC ←(if ((EX/MEM.opcode==branch)&EX/MEM.cond) { EX/MEM.ALUoutput} else {PC+4});		
ID	ID/EX.A ←Regs[IF/ID.IR[rs]]; ID/EX.B ←Regs[IF/ID.IR[rt]]; **ID/EX.NPC ←IF/ID.NPC;** ID/EX.IR ←IF/ID.IR; ID/EX.Imm ←sign-extend(IF/ID.IR[immediate field]);		
	ALU instruction	Ld/st instruction	Branch instruction
EX	EX/MEM.IR ←ID/EX.IR; EX/MEM.ALUoutput ←ID/EX.A func ID/EX.B; or EX/MEM.ALUoutput ←ID/EX.A op ID/EX.Imm;	EX/MEM.IR ←ID/EX.IR; EX/MEM.ALUoutput ←ID/EX.A + ID/EX.Imm; EX/MEM.B ←ID/EX.B;	EX/MEM.ALUoutput ←ID/EX.NPC + (ID/EX.Imm<<2); EX/MEM.cond ←(ID/EX.A==0);
MEM	MEM/WB.IR ← EX/MEM.IR; MEM/WB.ALUoutput ←EX/MEM.ALUoutput;	MEM/WB.IR ← EX/MEM.IR; MEM/WB.LMD ←Mem[EX/MEM.ALUoutput]; Or MEM/WB.LMD ←Mem[EX/MEM.ALUoutput];	
WB	Regs[MEM/WB.IR[rd]] ←MEM/WB.ALUoutput; or Regs[MEM/WB.IR[rt]] ←MEM/WB.ALUoutput;	For Load only; Regs[MEM/WB.IR[rt]] ←MEM/WB.LMD	

More Terminology & Foreshadowing
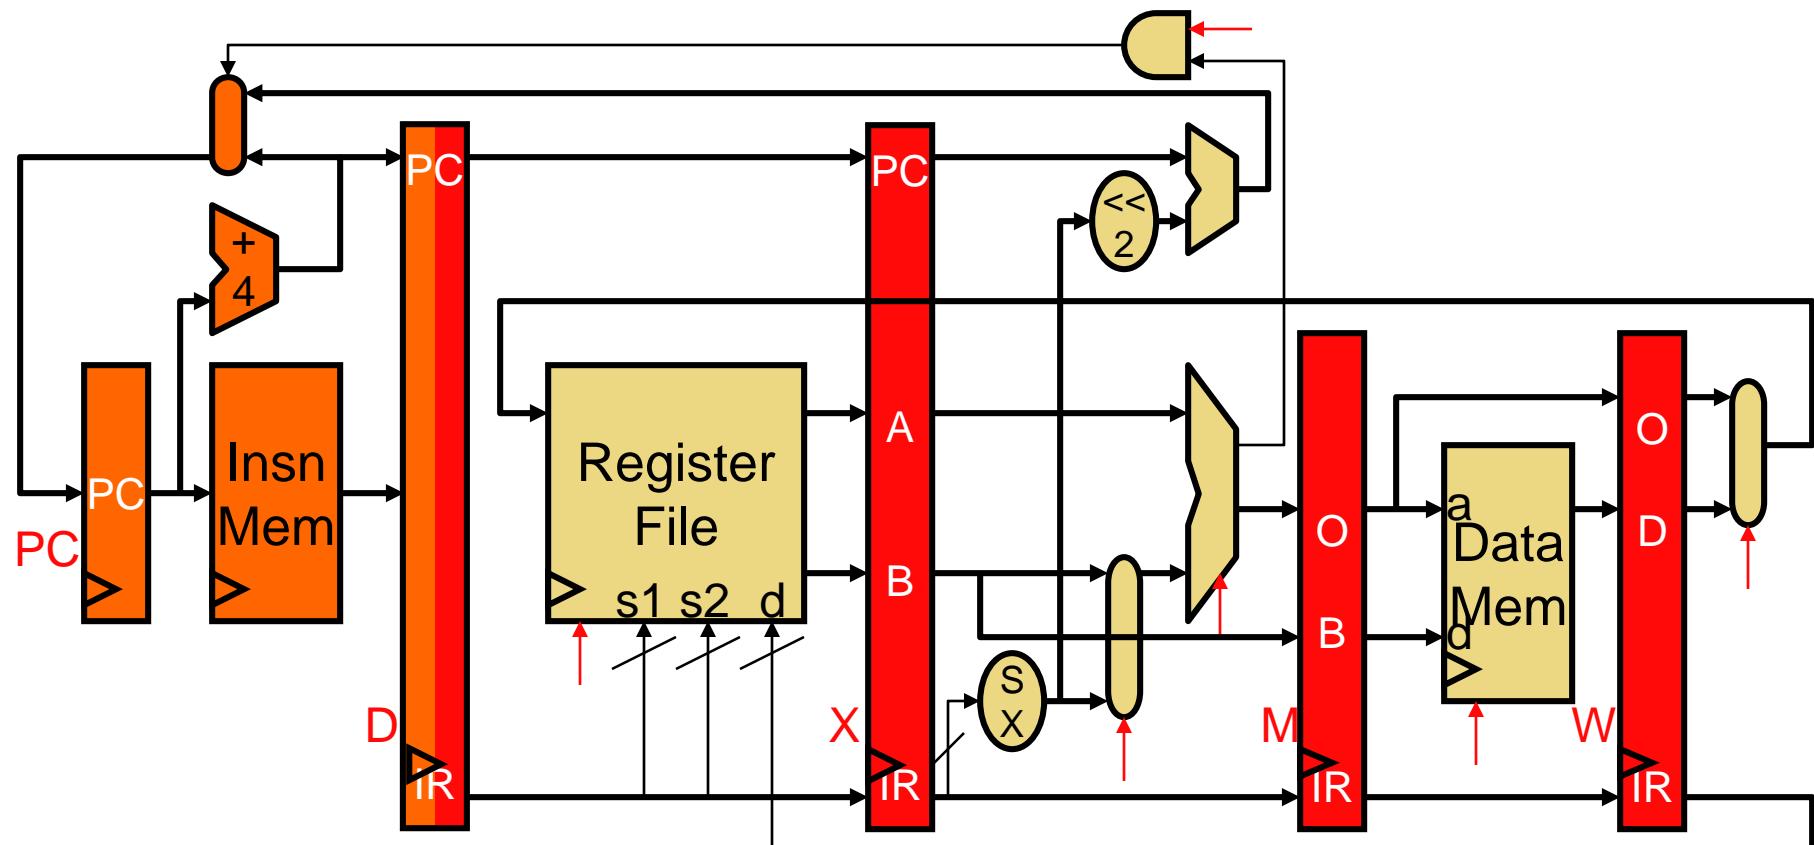
- **Scalar pipeline**: one insn per stage per cycle
  - Alternative: "superscalar" (>1 insns/cycle, later)

- **In-order pipeline**: insns enter execute stage in order
  - Alternative: "out-of-order" (later)

- **Pipeline depth**: number of pipeline stages
  - Nothing magical about five
  - Contemporary high-performance cores have ~15 stage pipelines

Instruction Convention

- Different ISAs use inconsistent register orders

- Some ISAs (for example MIPS)
  - Instruction destination (i.e., output) **on the left**
  - add $1, $2, $3 means $1 ← $2+$3

- Other ISAs
  - Instruction destination (i.e., output) **on the right**
  - `add r1,r2,r3` means `r1+r2➜r3`
  - `ld 8(r5),r4` means `mem[r5+8]➜r4`
  - `st r4,8(r5)` means `r4➜mem[r5+8]`

- Will try to specify to avoid confusion, next slides MIPS style
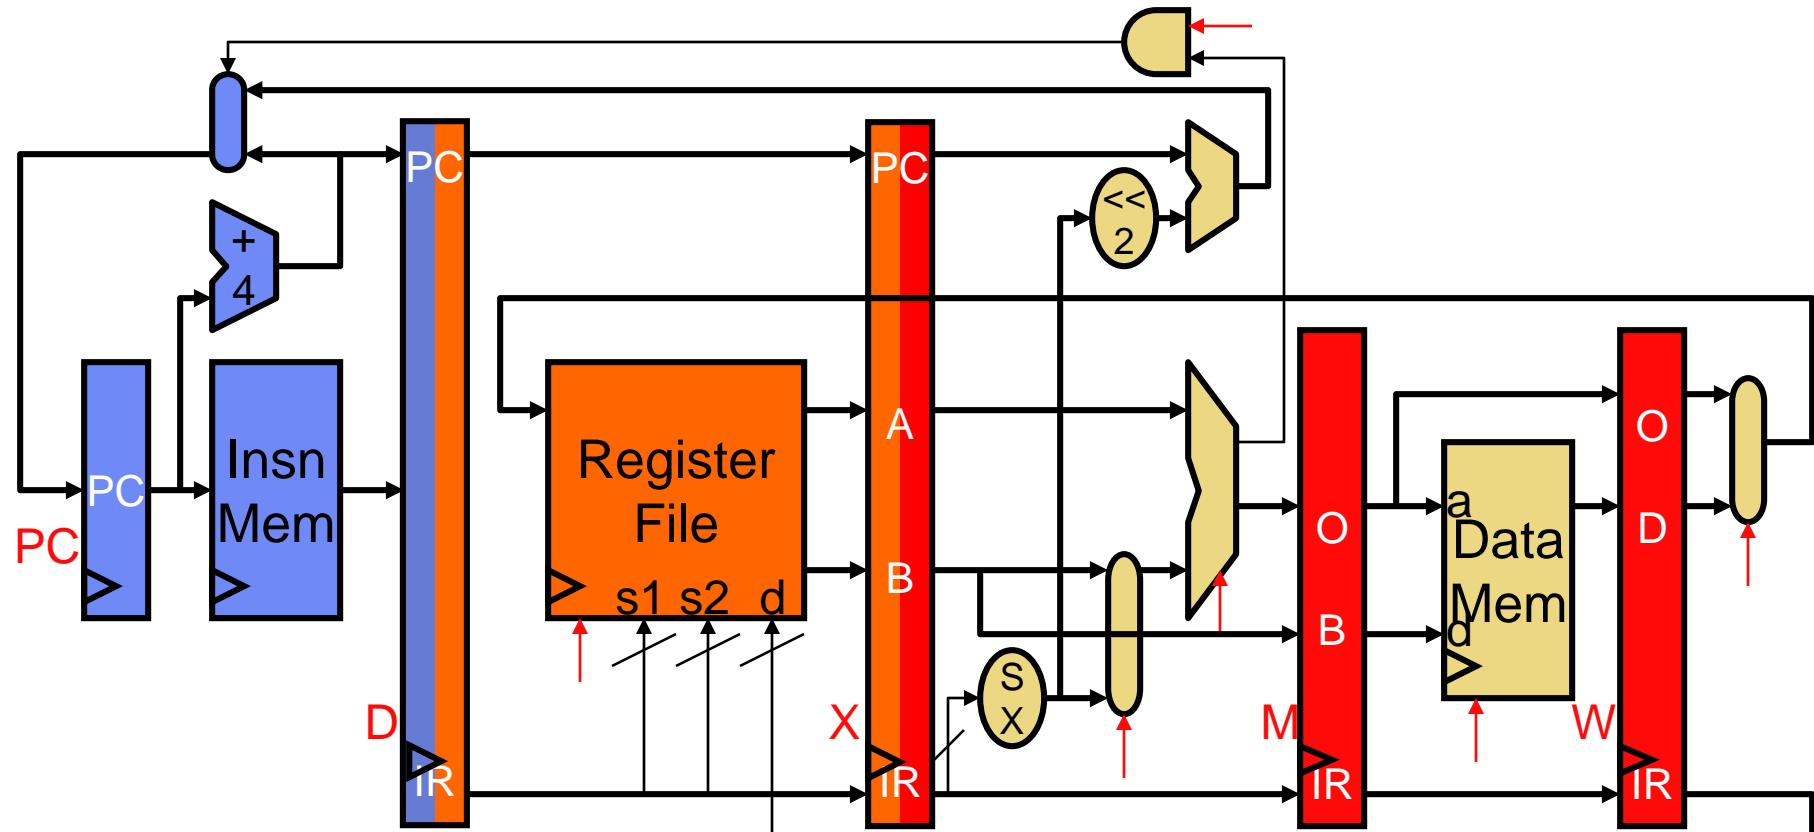
Pipeline Example: Cycle 1

add $3<-$2,$1

- 3 instructions
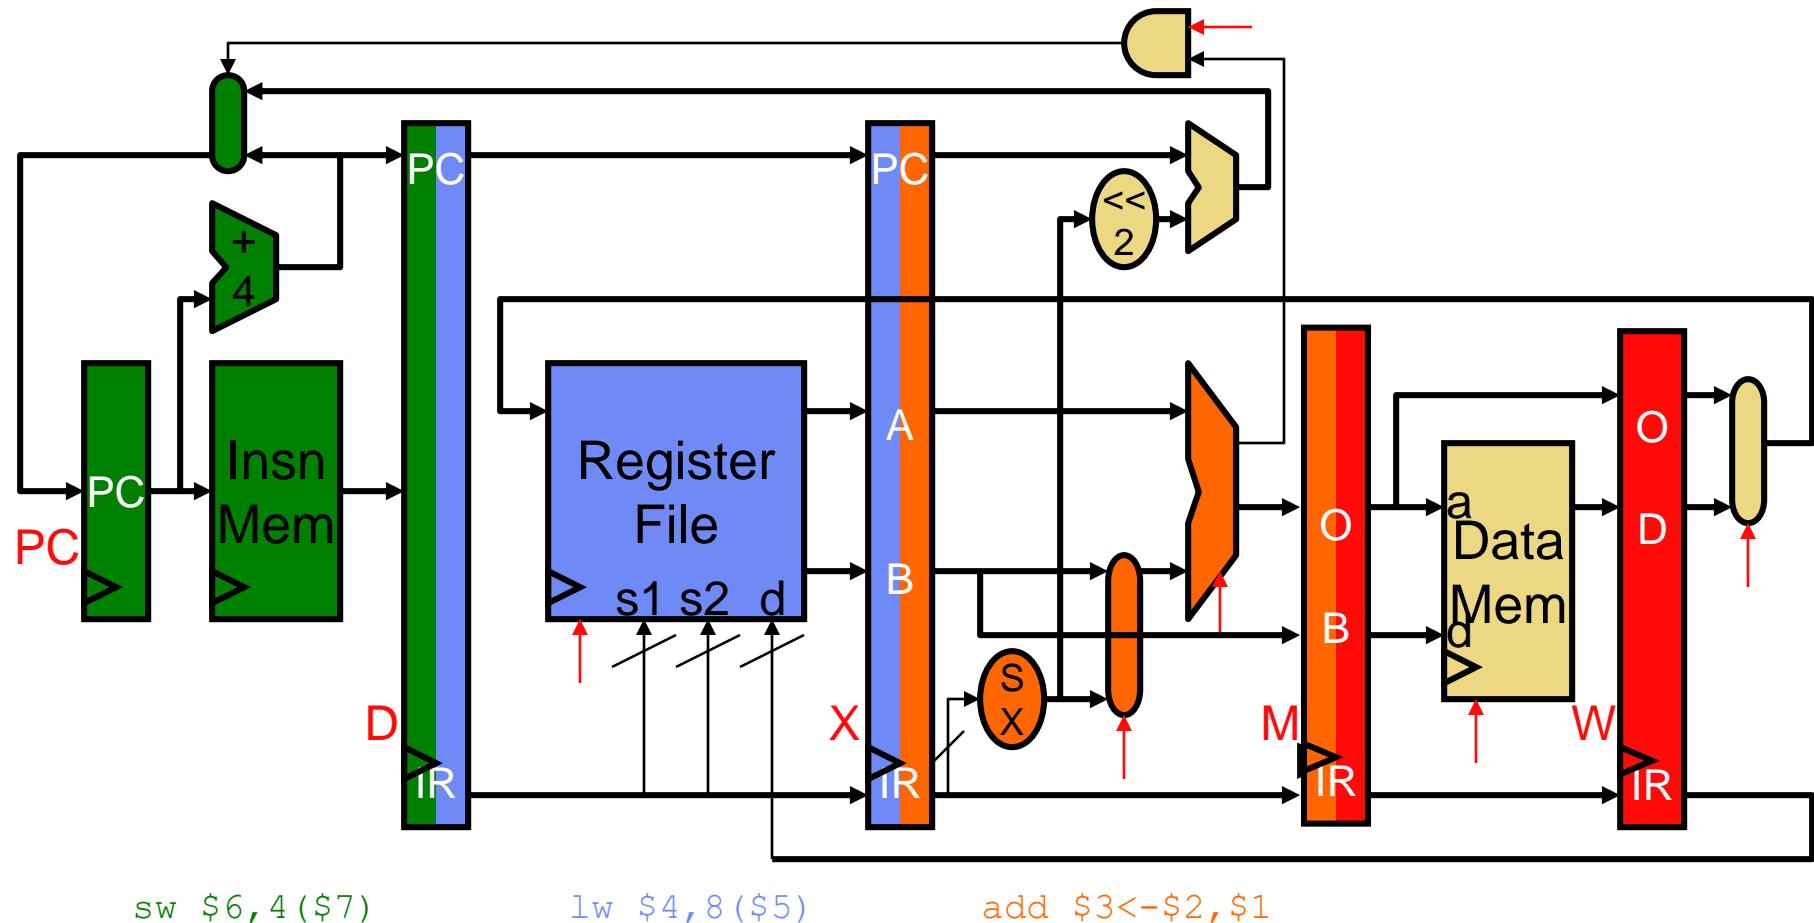
Pipeline Example: Cycle 2

lw $4,8($5)   add $3<-$2,$1

Pipeline Example: Cycle 3

PC

Insn Mem

PC

Register File
s1 s2 d

D

X

M

W

SX

Data Mem

sw $6,4($7)   lw $4,8($5)   add $3<-$2,$1

Pipeline Example: Cycle 4

sw $6,4($7)   lw $4,8($5)   add $3<-$2,$1

- 3 instructions

Modern Computer Architecture
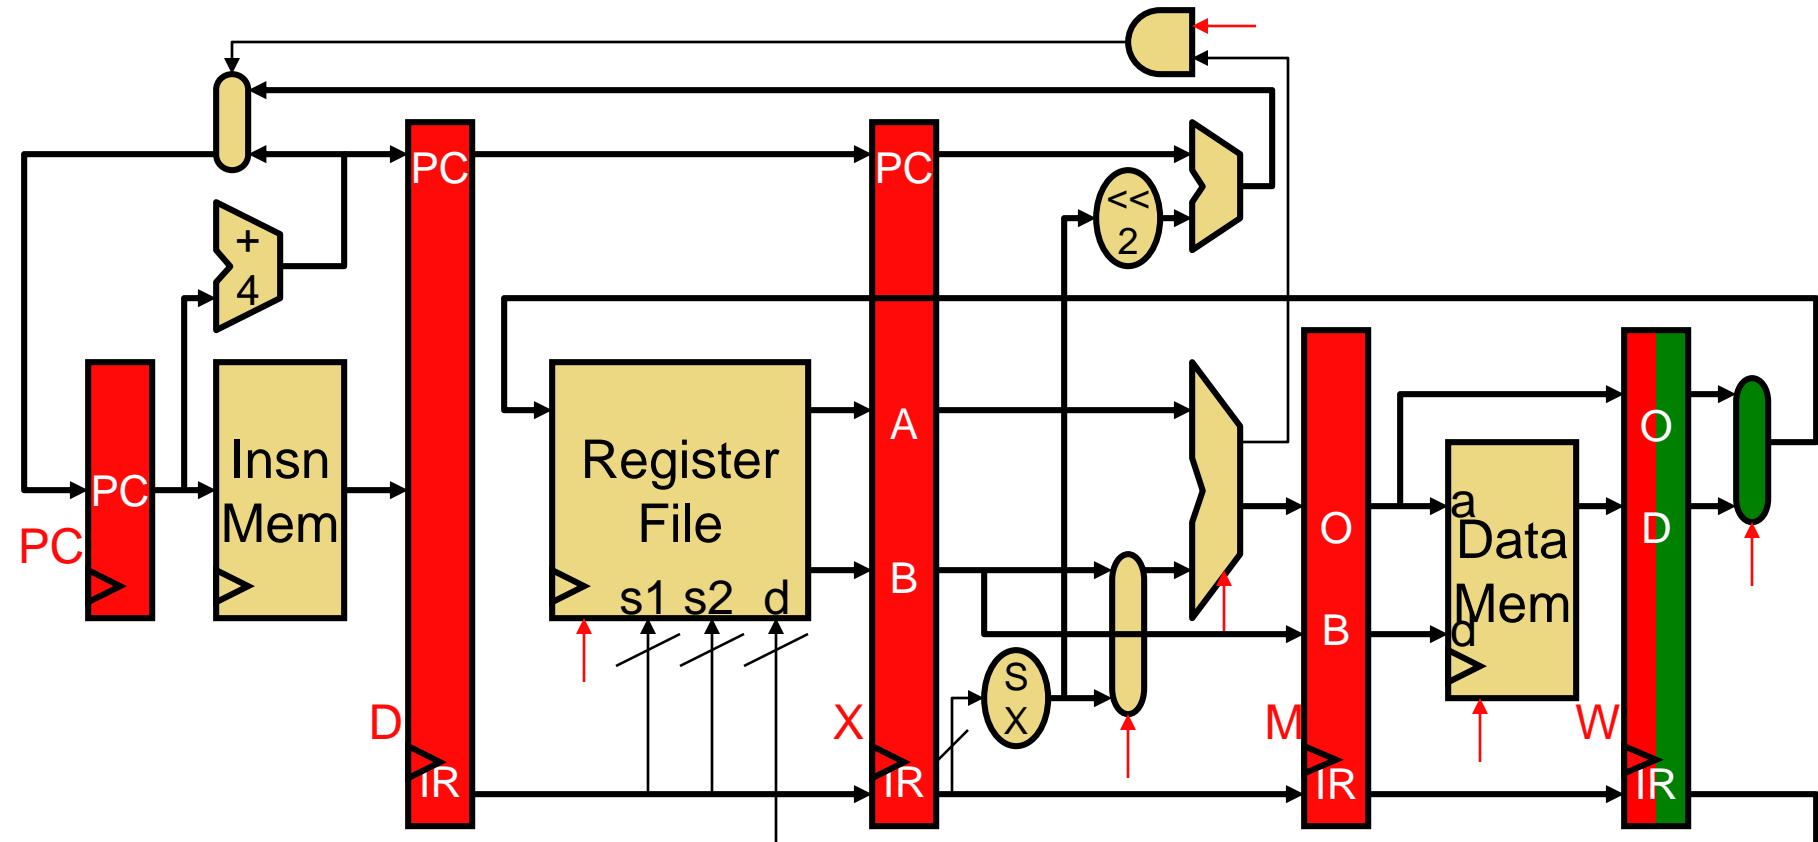
Pipeline Example: Cycle 5

sw $6,4($7)   lw $4,8($5)   add

sw $6,4(7)   lw

Pipeline Example: Cycle 7

sw

Pipeline Diagram

- **Pipeline diagram**: shorthand for what we just saw
  - Across: cycles
  - Down: insns
  - Convention: **X** means `lw $4,8($5)` finishes e**X**ecute stage and writes into M latch at end of cycle 4

	1	2	3	4	5	6	7	8	9
`add $3<-$2,$1`	F	D	X	M	W				
`lw $4,8($5)`		F	D	**X**	M	W			
`sw $6,4($7)`			F	D	X	M	W		

Example Pipeline Perf. Calculation

- Single-cycle
  - Clock period = 50ns, CPI = 1
  - Performance = 50ns/insn

- 5-stage pipeline
  - Clock period = **12ns**   approx. (50ns / 5 stages) + overheads
  + CPI = **1** (each insn takes 5 cycles, but 1 completes each cycle)
    + Performance = **12ns/insn**
  - Well actually … CPI = 1 + some penalty for pipelining (next)
    - CPI = **1.5** (on average insn completes every 1.5 cycles)
    - Performance = **18ns/insn**
    - Much higher performance than single-cycle

Question

Q1: Why Is Pipeline Clock Period > (delay thru datapath) / (number of pipeline stages)?

- Three reasons:
  - Latches add delay
  - Pipeline stages have different delays, clock period is max delay
  - Extra datapaths for pipelining (bypassing paths)

- These factors have implications for ideal number pipeline stages
  - Diminishing clock frequency gains for longer (deeper) pipelines

Question

Q2: Why Is Pipeline CPI > 1?
- CPI for scalar in-order pipeline is 1 **+ stall penalties**
- Stalls used to *resolve* hazards
  - **Hazard**: condition that jeopardizes sequential illusion
  - **Stall**: pipeline delay introduced to restore sequential illusion

- Calculating pipeline CPI
  - **Frequency of stall** * **stall cycles**
  - Penalties add (stalls generally don't overlap in in-order pipelines)
  - $1 + (\text{stall-freq}_1 * \text{stall-cyc}_1) + (\text{stall-freq}_2 * \text{stall-cyc}_2) + \ldots$

- Correctness vs. performance tradeoff
  - make common case fast
  - Long penalties OK if they are rare, e.g., $1 + (0.01 * 10) = 1.1$
  - Stalls also have implications for ideal number of pipeline stages
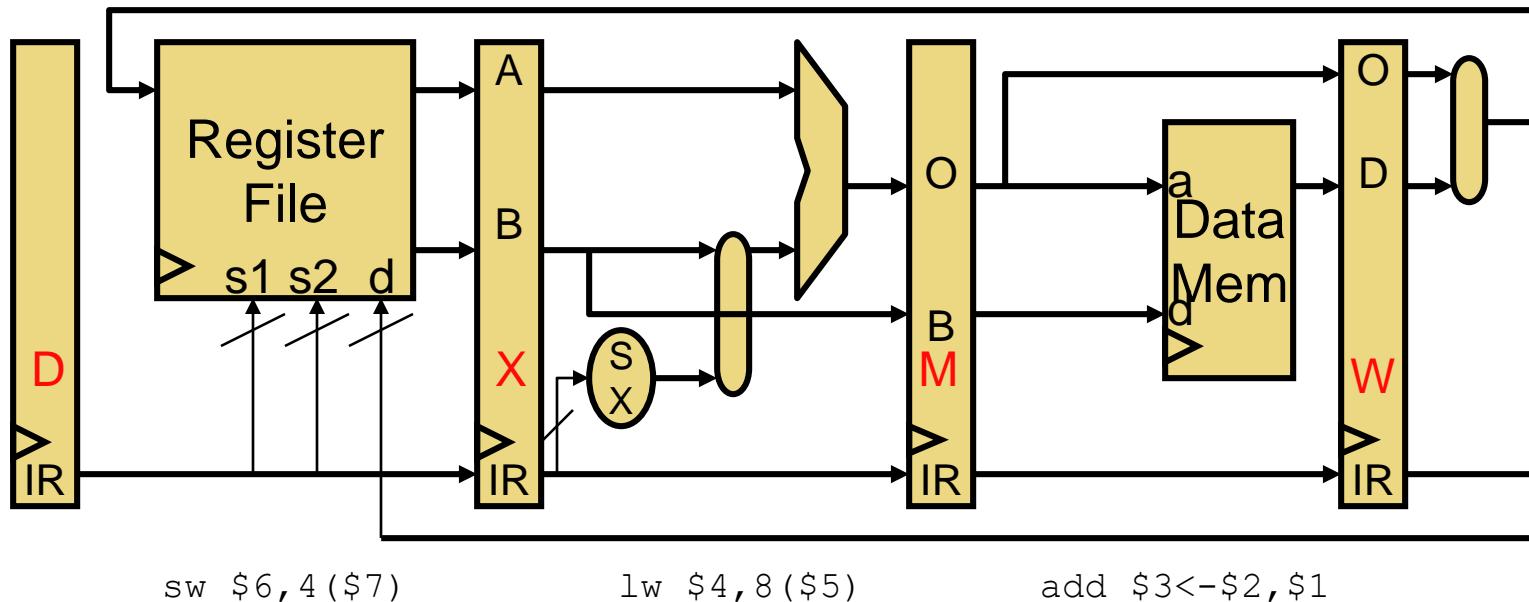
Data Dependences, Pipeline Hazards, and Bypassing

Dependences and Hazards

- **Dependence**: relationship between two insns
  - **Data dep.**: two insns use same storage location
  - **Control dep.**: one insn effects whether another executes at all
  - Not a bad thing, programs would be boring without them
  - **Enforced** by making older insn go before younger one
    - Happens naturally in single-cycle designs
    - But not in a pipeline!

- **Hazard**: dependence & possibility of wrong insn order
  - Effects of wrong insn order cannot be externally visible
    - **Stall**: for order by keeping younger insn in same stage
  - Hazards are a bad thing: stalls reduce performance

Data Hazards

sw $6,4($7)   lw $4,8($5)   add $3<-$2,$1

- Let's forget about branches and control flow for now
- The three insn sequence we saw earlier executed fine…
  - But it wasn't a real program
  - Real programs have **data dependences**
    - They pass values via registers and memory

Dependent Operations

- Independent operations

```
add $3<-$2,$1
add $6<-$5,$4
```
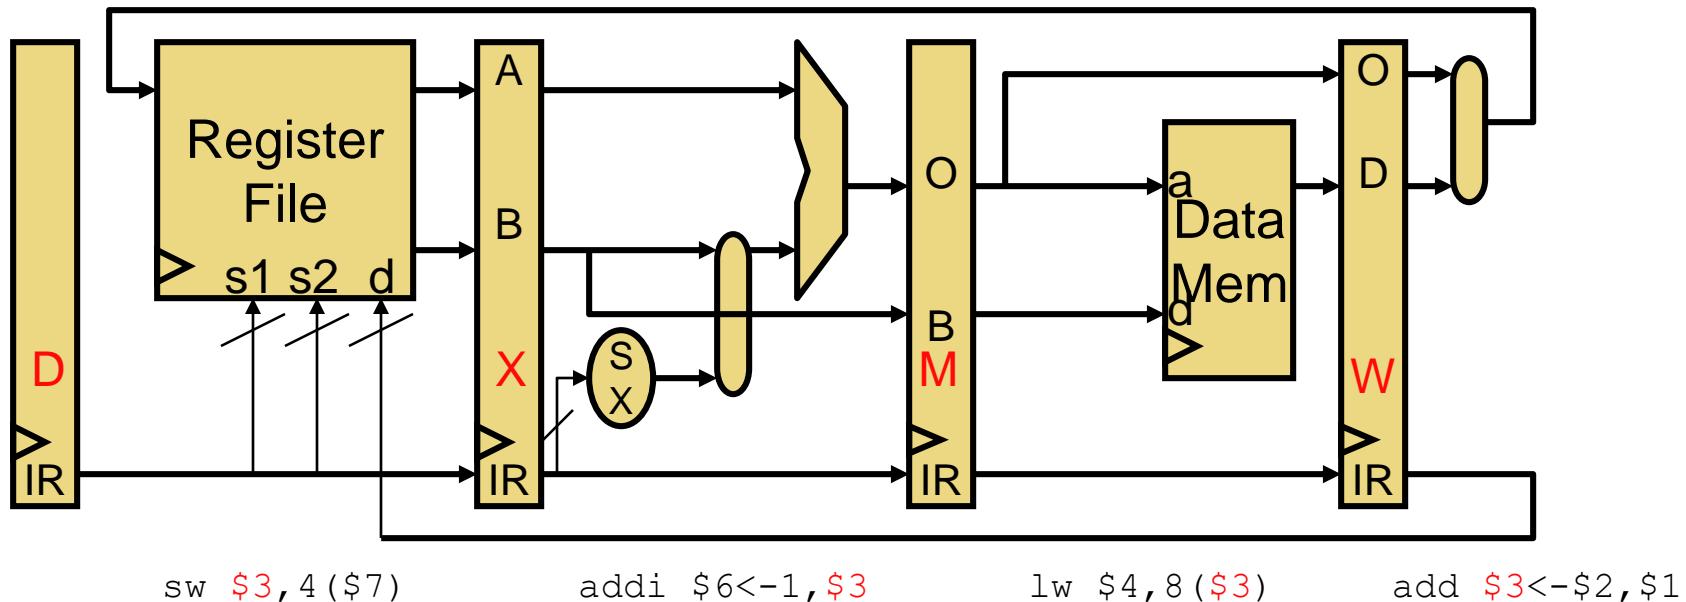
- Would this program execute correctly on a pipeline?

```
add $3<-$2,$1
add $6<-$5,$3
```

- What about this program?
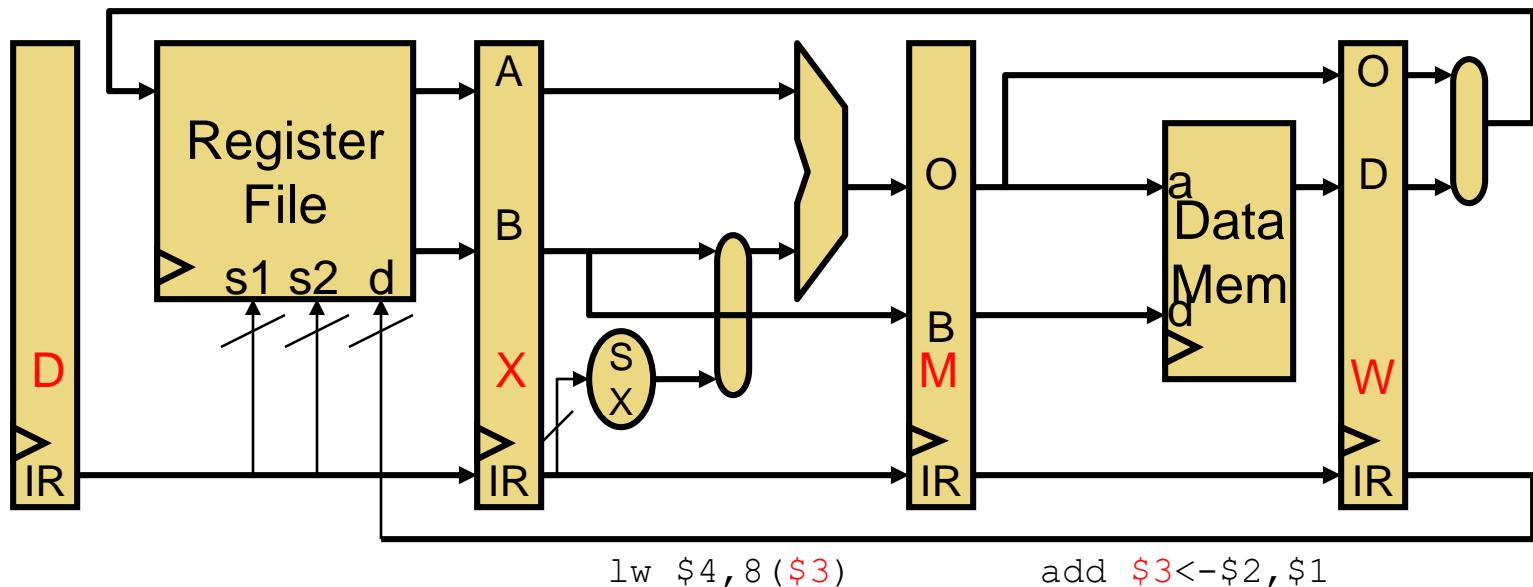
```
add $3<-$2,$1
lw $4,8($3)
addi $6<-1,$3
sw $3,8($7)
```

Data Hazards

sw $3,4($7)   addi $6<-1,$3   lw $4,8($3)   add $3<-$2,$1
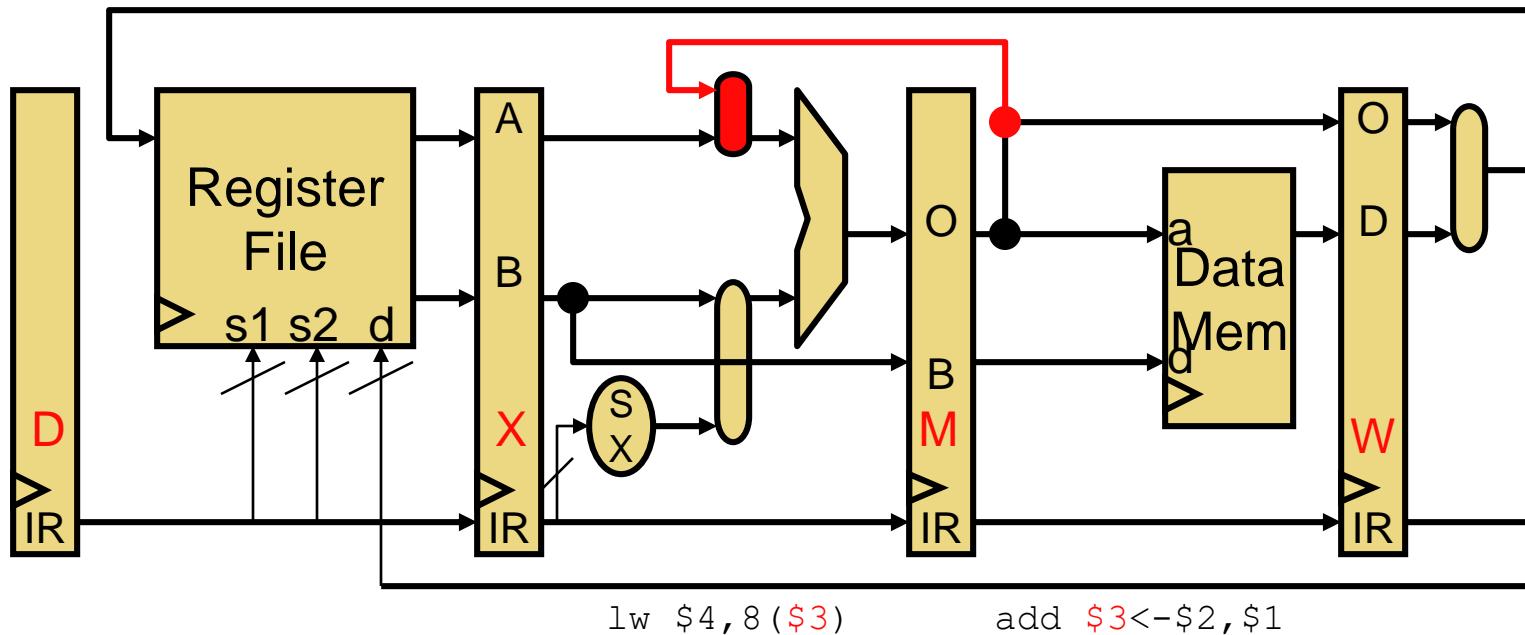
- Would this "program" execute correctly on this pipeline?
  - Which insns would execute with correct inputs?
  - **add** is writing its result into **$3** in current cycle
  - **lw** read **$3** two cycles ago → got wrong value
  - **addi** read **$3** one cycle ago → got wrong value
  - **sw** is reading **$3** this cycle → maybe ok (depends on regfile design)

Observation!


```
lw $4,8($3)          add $3<-$2,$1
```

- Technically, this situation is broken
    - `lw $4,8($3)` has already read $3 from regfile
    - `add $3<-$2,$1` hasn't yet written $3 to regfile
- But fundamentally, everything is OK
    - `lw $4,8($3)` hasn't actually used $3 yet
    - `add $3<-$2,$1` has already computed $3

Bypassing

lw $4,8($3)   add $3<-$2,$1
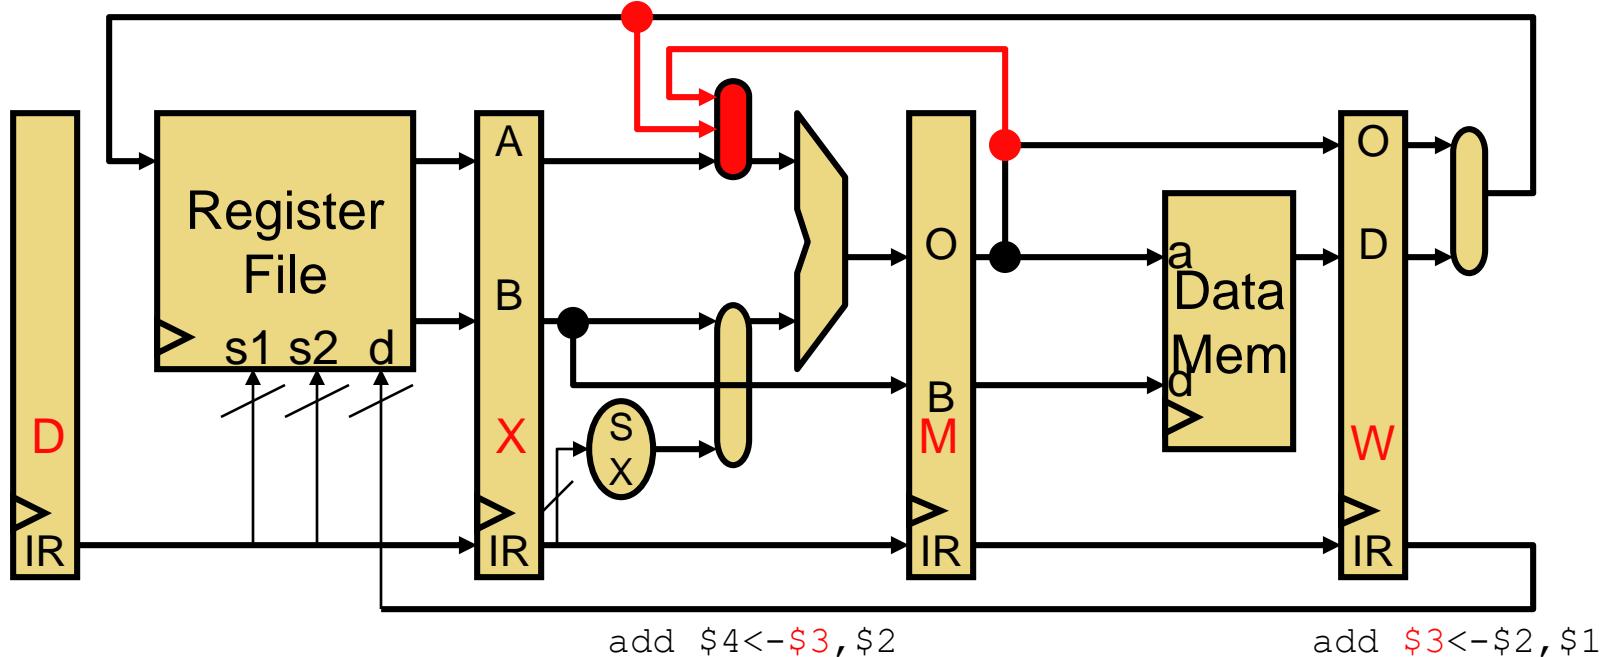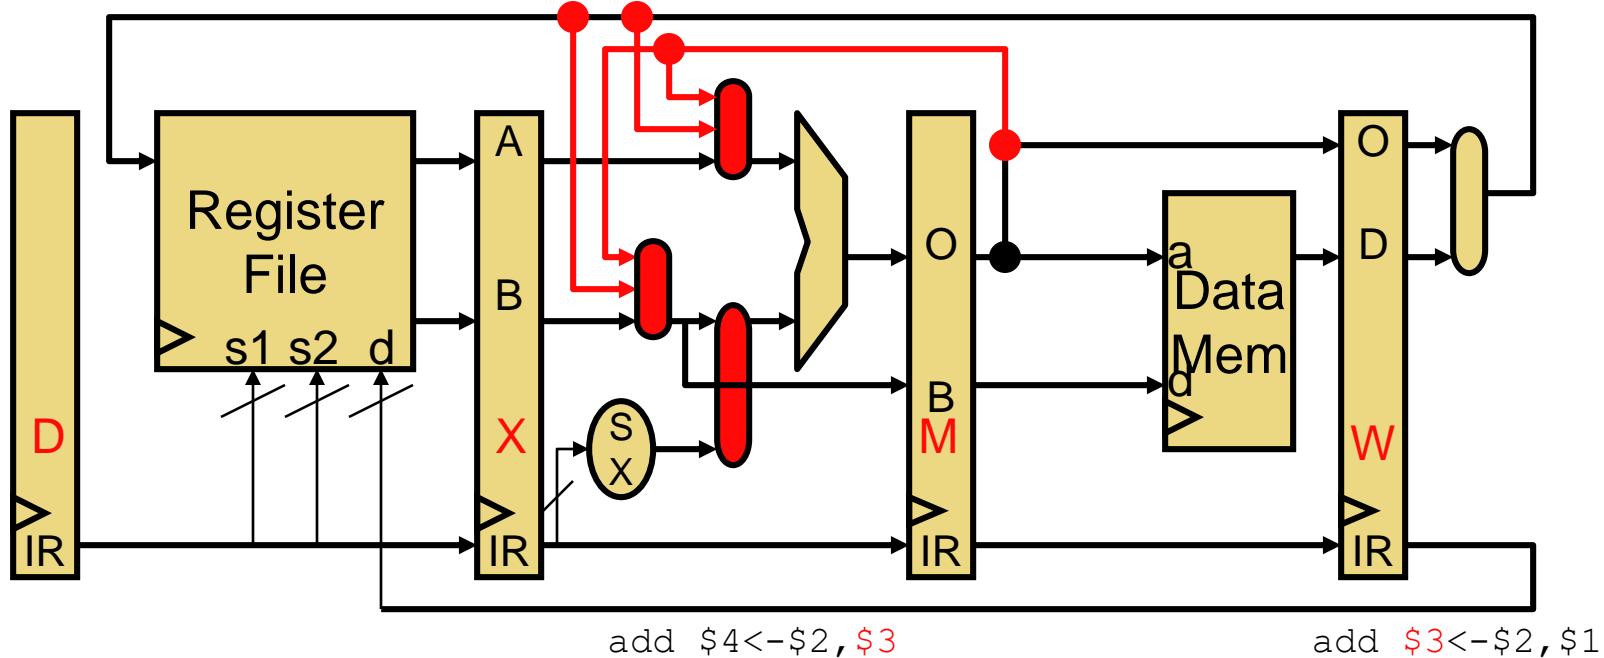
- **Bypassing**
  - Reading a value from an intermediate (μarchitectural) source
  - Not waiting until it is available from primary source
  - Here, we are bypassing the register file
  - Also called **forwarding**

WX Bypassing

add $4<-$3,$2   add $3<-$2,$1
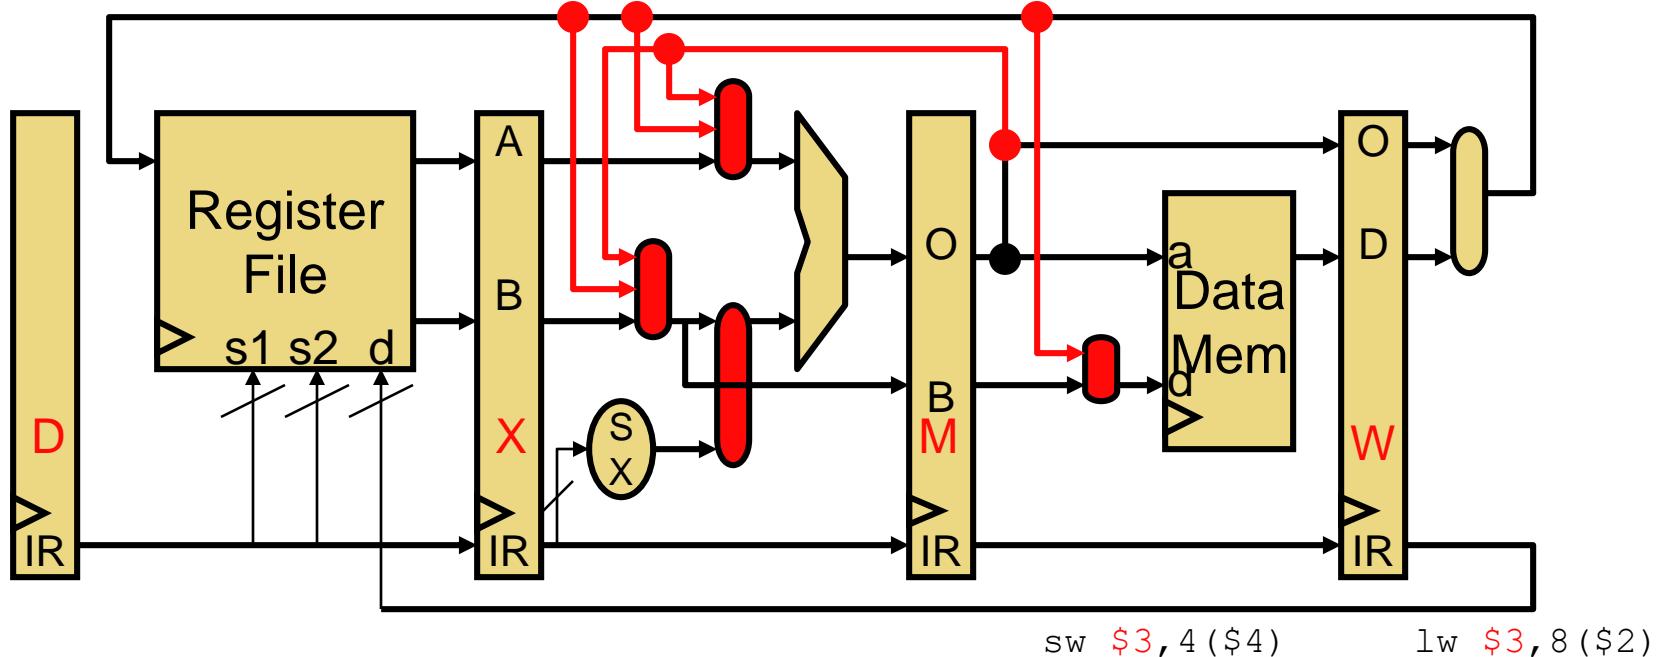
- What about this combination?
  - Add another bypass path and MUX (multiplexor) input
  - First one was an **MX** bypass
  - This one is a **WX** bypass

ALUinB Bypassing

add $4<-$2,$3   add $3<-$2,$1

- Can also bypass to ALU input B

WM Bypassing?


```
sw $3,4($4)          lw $3,8($2)
```

- Does WM bypassing work?
  - Not to the address input (why not?)
    ```
    sw $4,4($3)      lw $3,8($2)
    ```
    X
  - But to the store data input, yes
    ```
    sw $3,4($4)      lw $3,8($2)
    ```
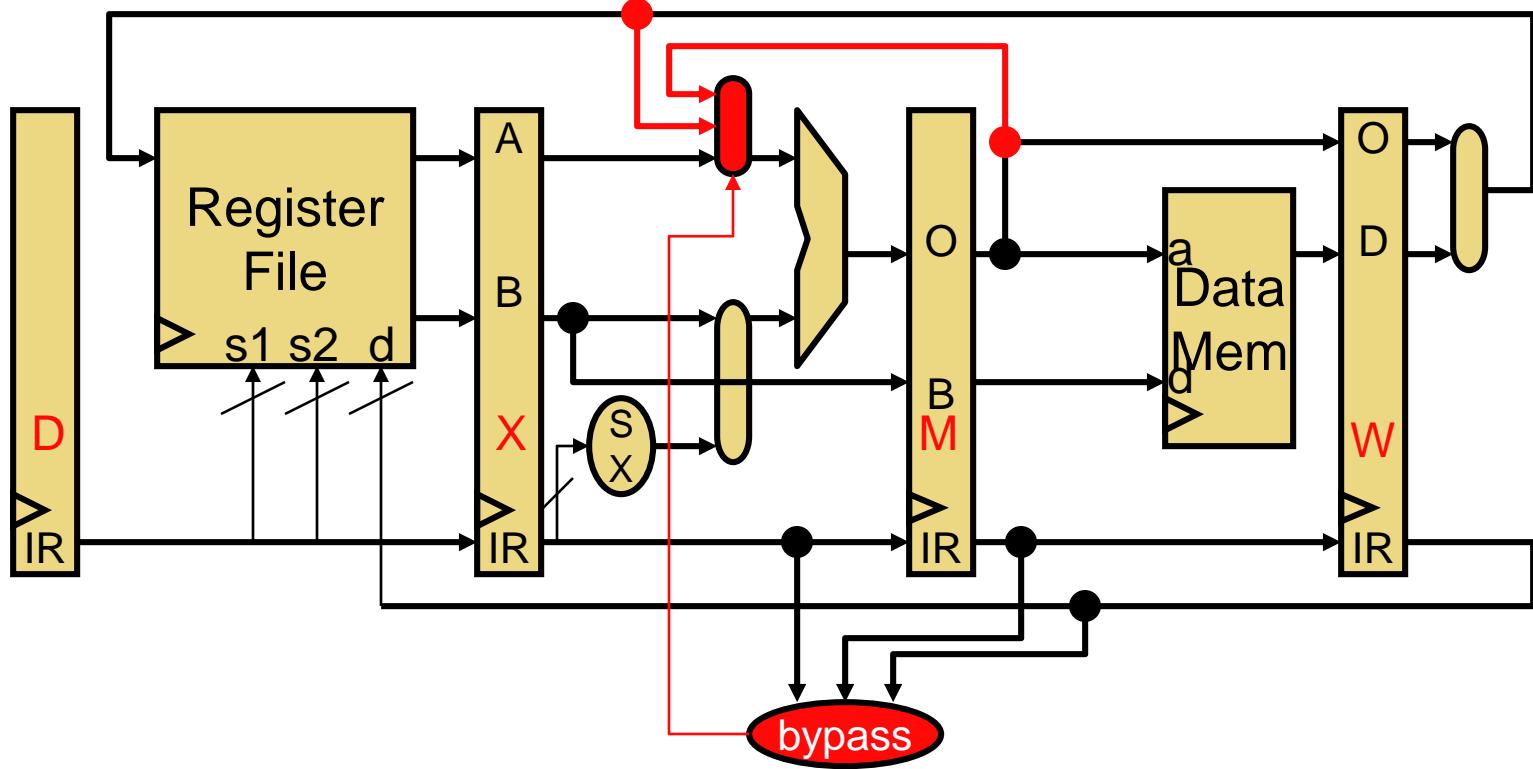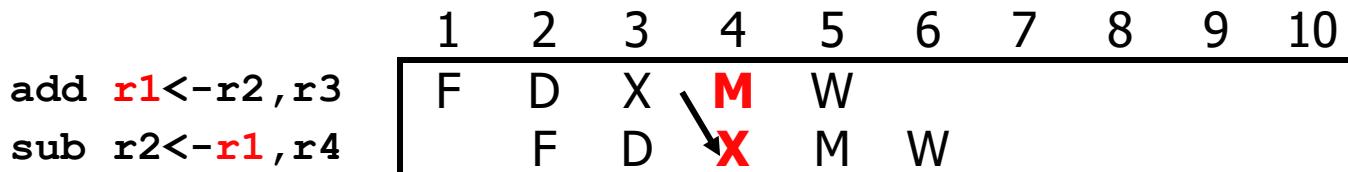
Bypass Logic
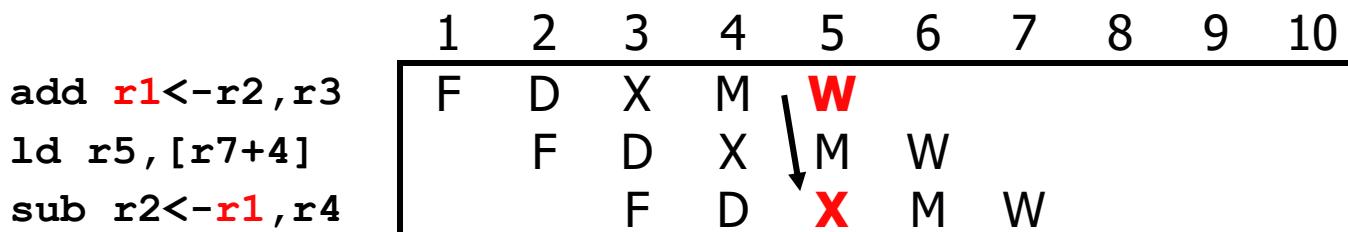
- Each multiplexor has its own logic, here it is for "ALUinA"

  (X.IR.RegSrc1 == M.IR.RegDest) => 0

  (X.IR.RegSrc1 == W.IR.RegDest) => 1

  Else => 2

Pipeline Diagrams with Bypassing
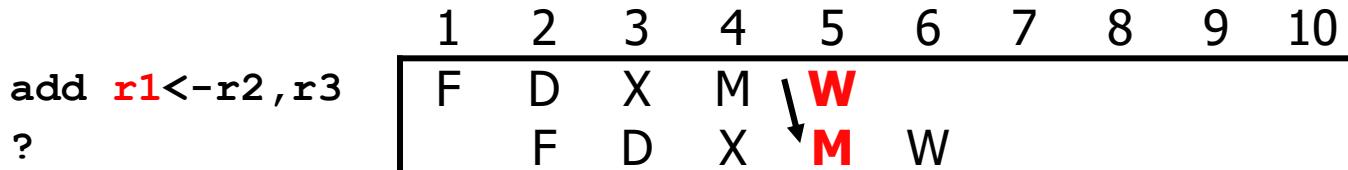
- If bypass exists, "from"/"to" stages execute in same cycle
  - Example: MX bypass

	1	2	3	4	5	6	7	8	9	10
add r1<-r2,r3	F	D	X	**M**	W					
sub r2<-r1,r4		F	D	**X**	M	W				

  - Example: WX bypass

|	1	2	3	4	5	6	7	8	9	10
add r1<-r2,r3	F	D	X	M	**W**					
ld r5,[r7+4]		F	D	X	M	W				
sub r2<-r1,r4			F	D	**X**	M	W			

  - Example: WM bypass

	1	2	3	4	5	6	7	8	9	10
add r1<-r2,r3	F	D	X	M	**W**					
?		F	D	X	**M**	W				
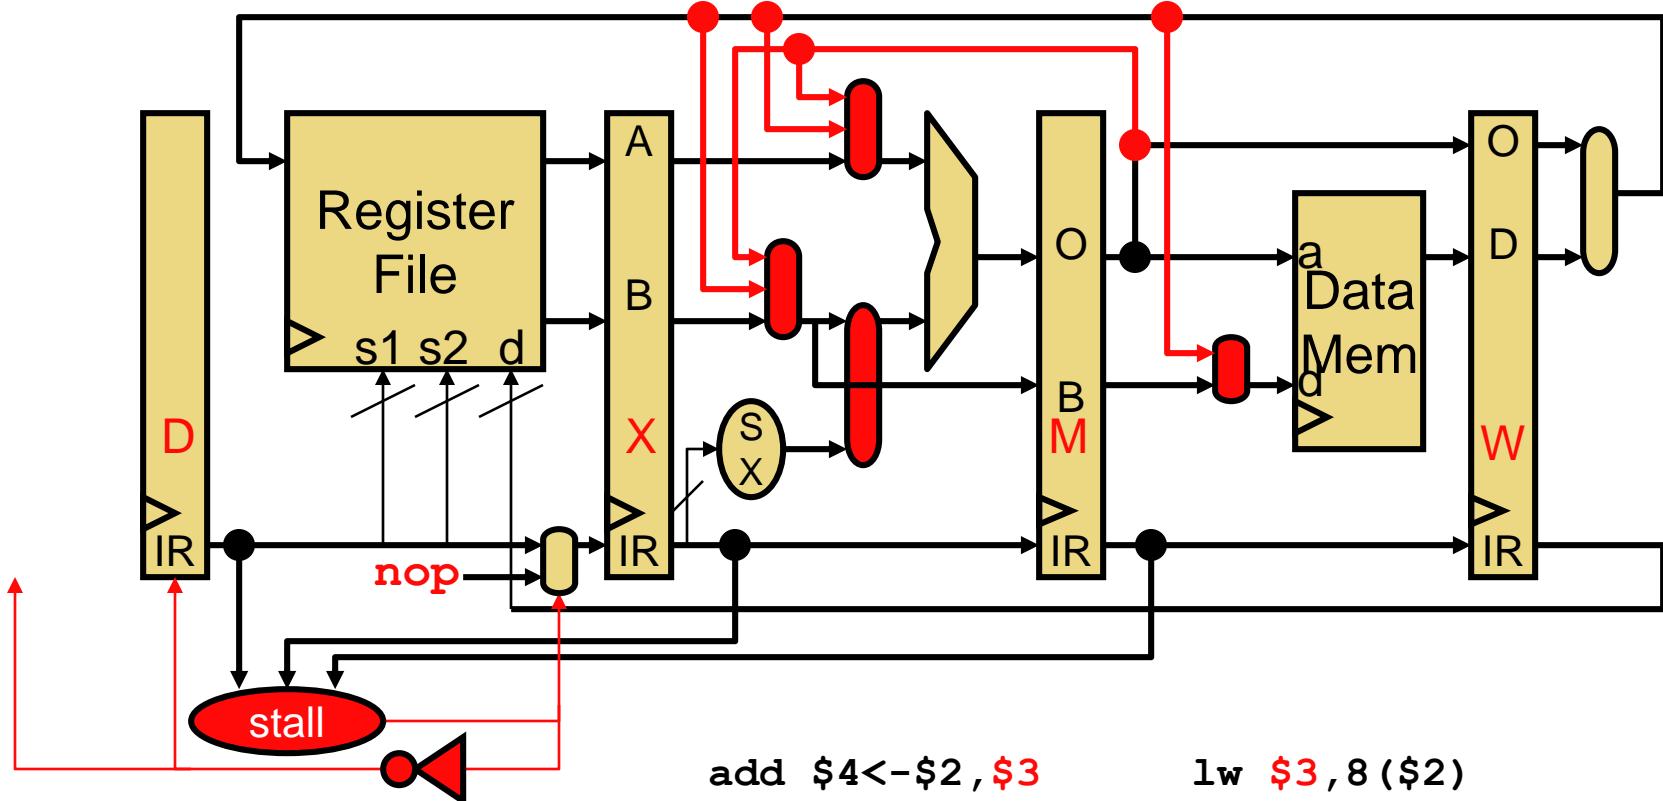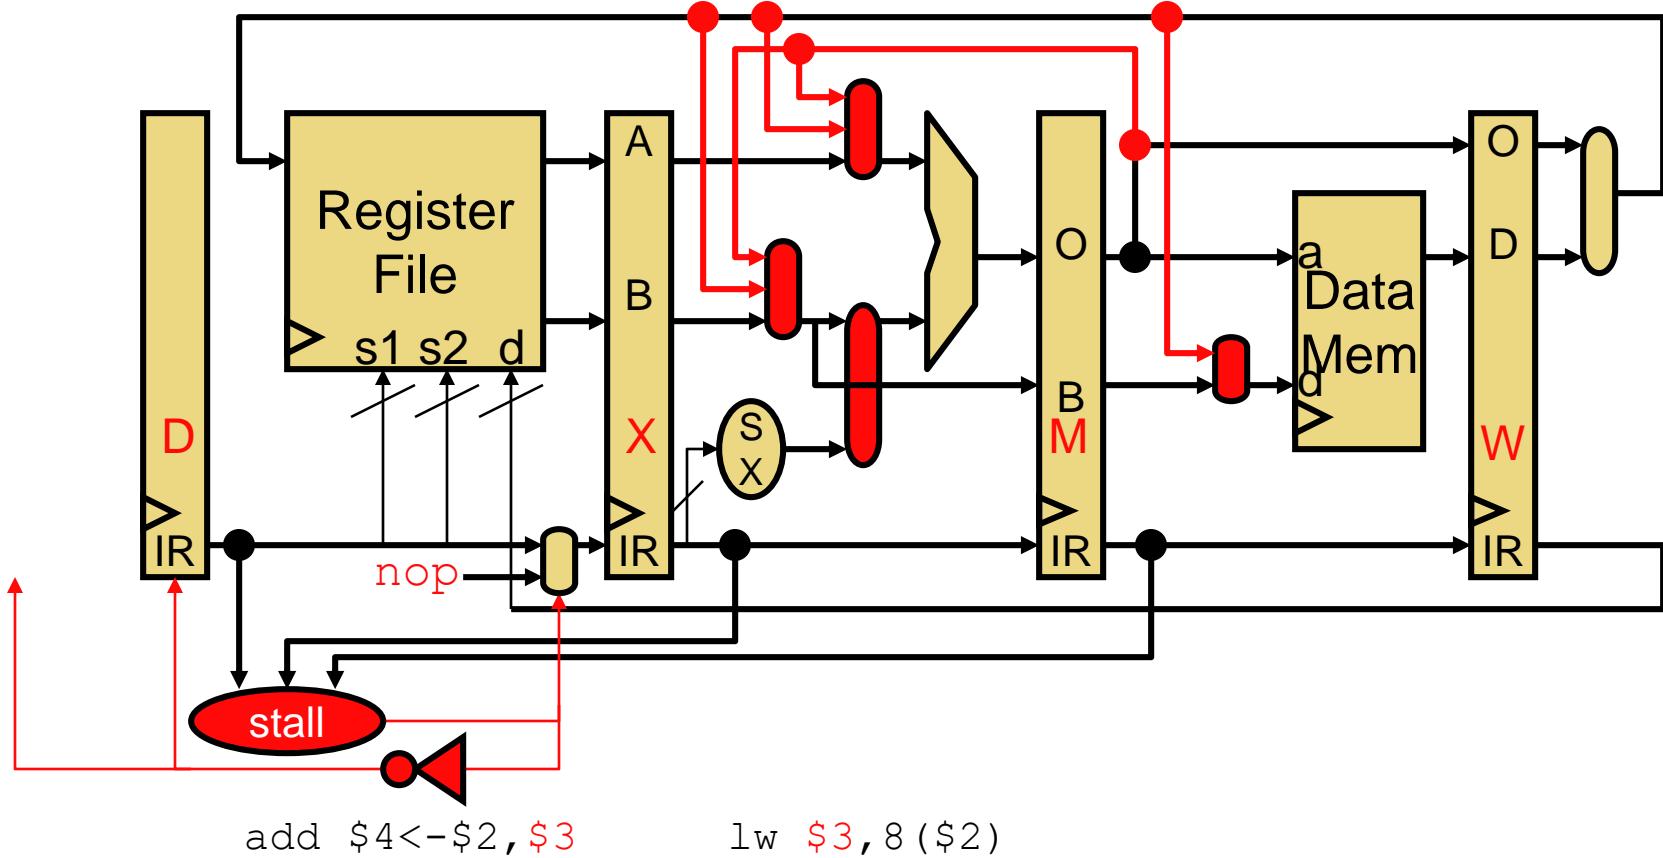
- Can you think of a code example that uses the WM bypass?

Have We Prevented All Data Hazards?

add $4<-$2,$3   lw $3,8($2)
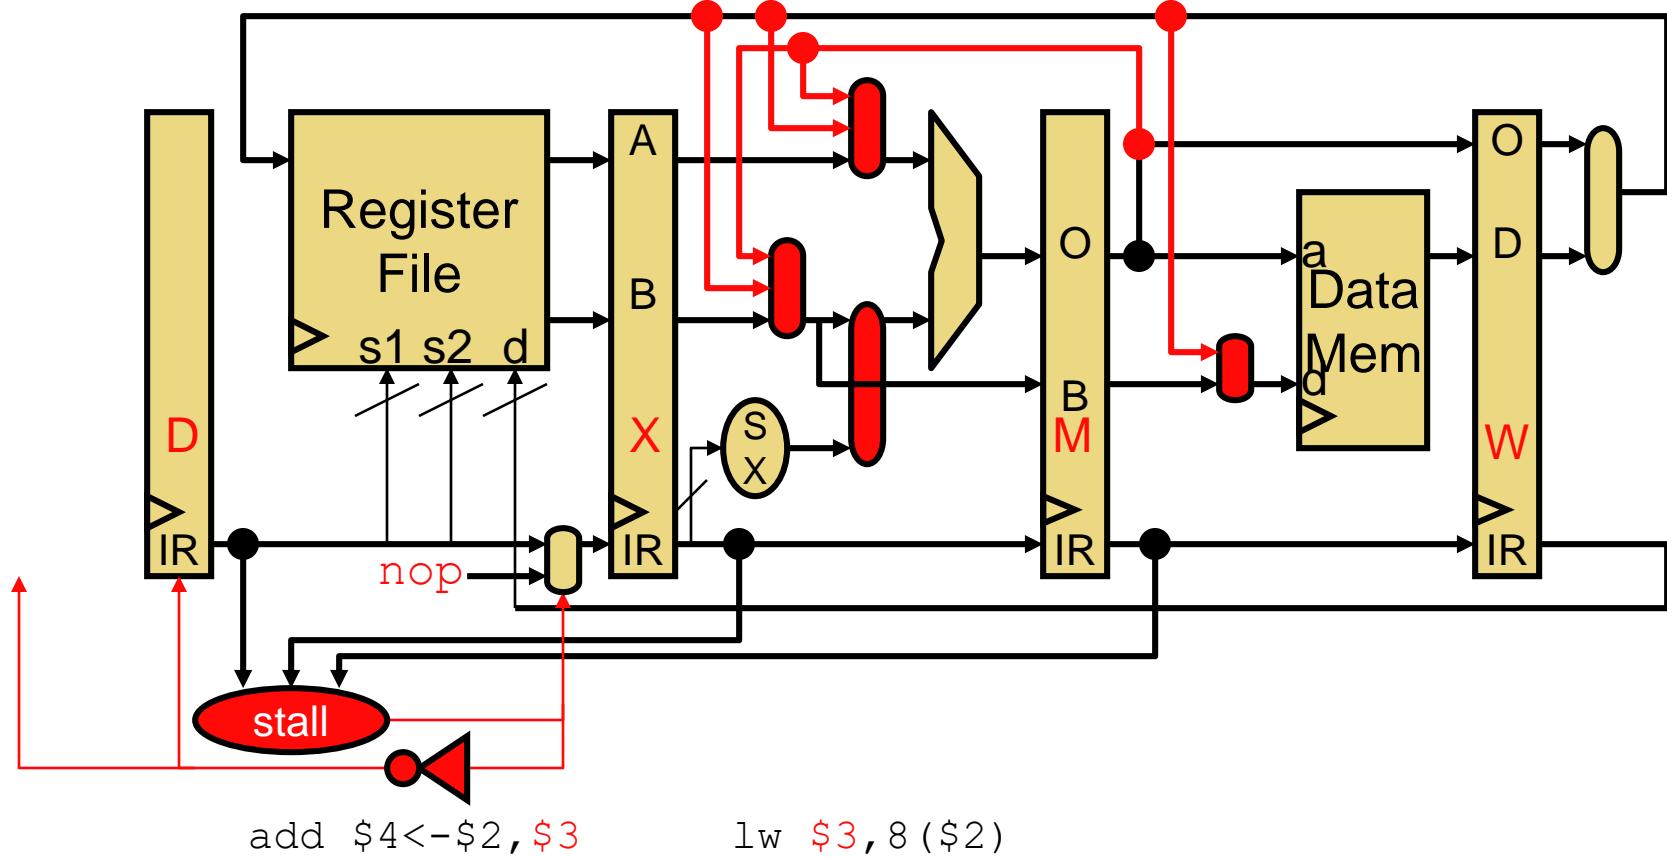
- No. Consider a "load" followed by a dependent "add" insn
- Bypassing alone isn't sufficient!
- Hardware solution: detect this situation and inject a stall cycle
- Software solution: ensure compiler doesn't generate such code

Stalling on Load-To-Use Dependences


```
add $4<-$2,$3        lw $3,8($2)
```
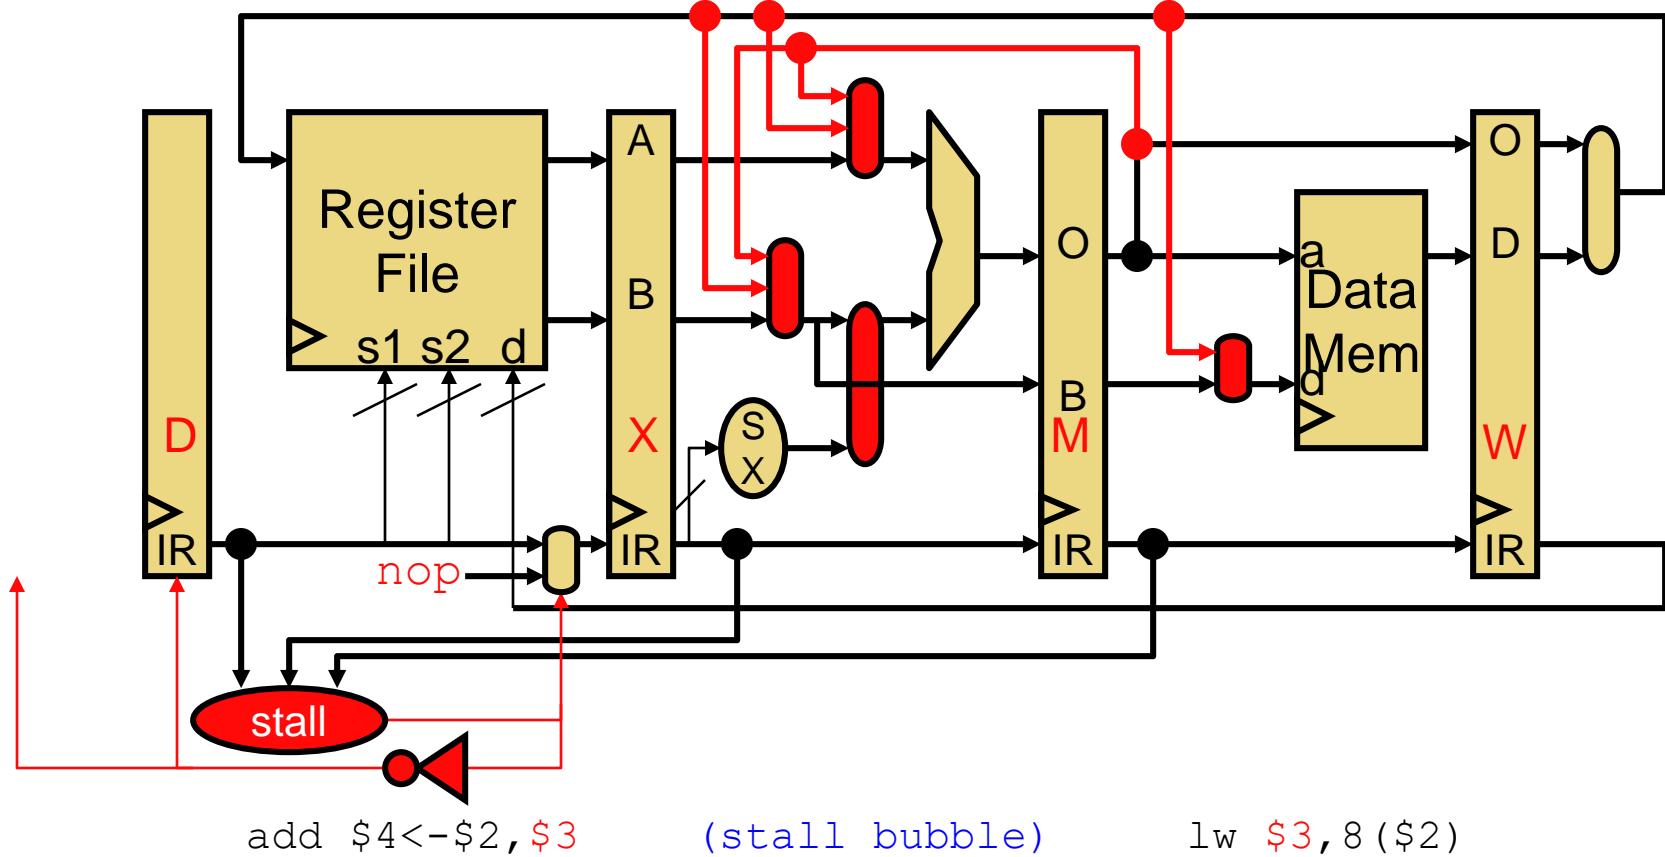
- Prevent "D insn" from advancing this cycle
  - Write **nop** into X.IR (effectively, insert **nop** in hardware)
  - Keep same "D insn", same PC next cycle
- Re-evaluate situation next cycle

Stalling on Load-To-Use Dependences
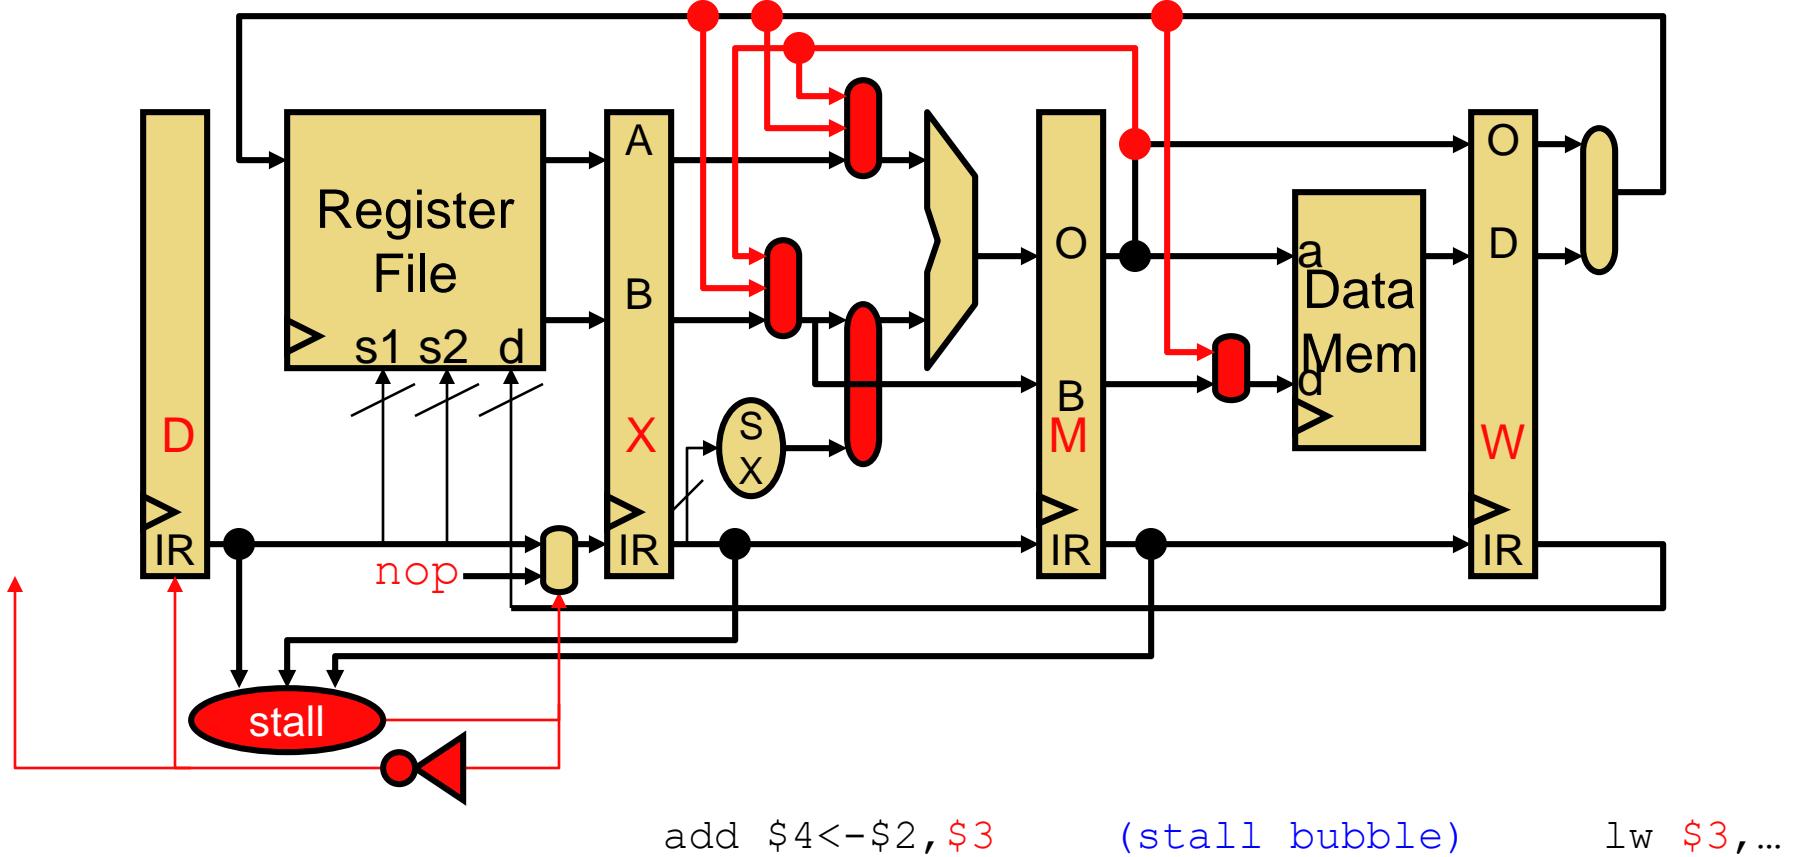

```
add $4<-$2,$3        lw $3,8($2)
```

Stall = (X.IR.Operation == LOAD) &&

    ((D.IR.RegSrc2 == X.IR.RegDest) ||

      ((D.IR.RegSrc1 == X.IR.RegDest) && (D.IR.Op != STORE))

)

Stalling on Load-To-Use Dependences

add $4<-$2,$3   (stall bubble)   lw $3,8($2)

Stall = (X.IR.Operation == LOAD) &&

    ((D.IR.RegSrc2 == X.IR.RegDest) ||

    ((D.IR.RegSrc1 == X.IR.RegDest) && (D.IR.Op != STORE))

)

Stalling on Load-To-Use Dependences

add $4<-$2,$3   (stall bubble)   lw $3,…

Stall = (X.IR.Operation == LOAD) &&

     ((D.IR.RegSrc2 == X.IR.RegDest) ||

      ((D.IR.RegSrc1 == X.IR.RegDest) && (D.IR.Op != STORE))

)

Modern Computer Architecture

Performance Impact of Load/Use Penalty

- Assume
    - Branch: 20%, load: 20%, store: 10%, other: 50%
    - 50% of loads are followed by dependent instruction
        - require 1 cycle stall (I.e., insertion of 1 `nop`)

- Calculate CPI
    - CPI = 1 + (1 * 20% * 50%) = **1.1**

Reducing Load-Use Stall Frequency
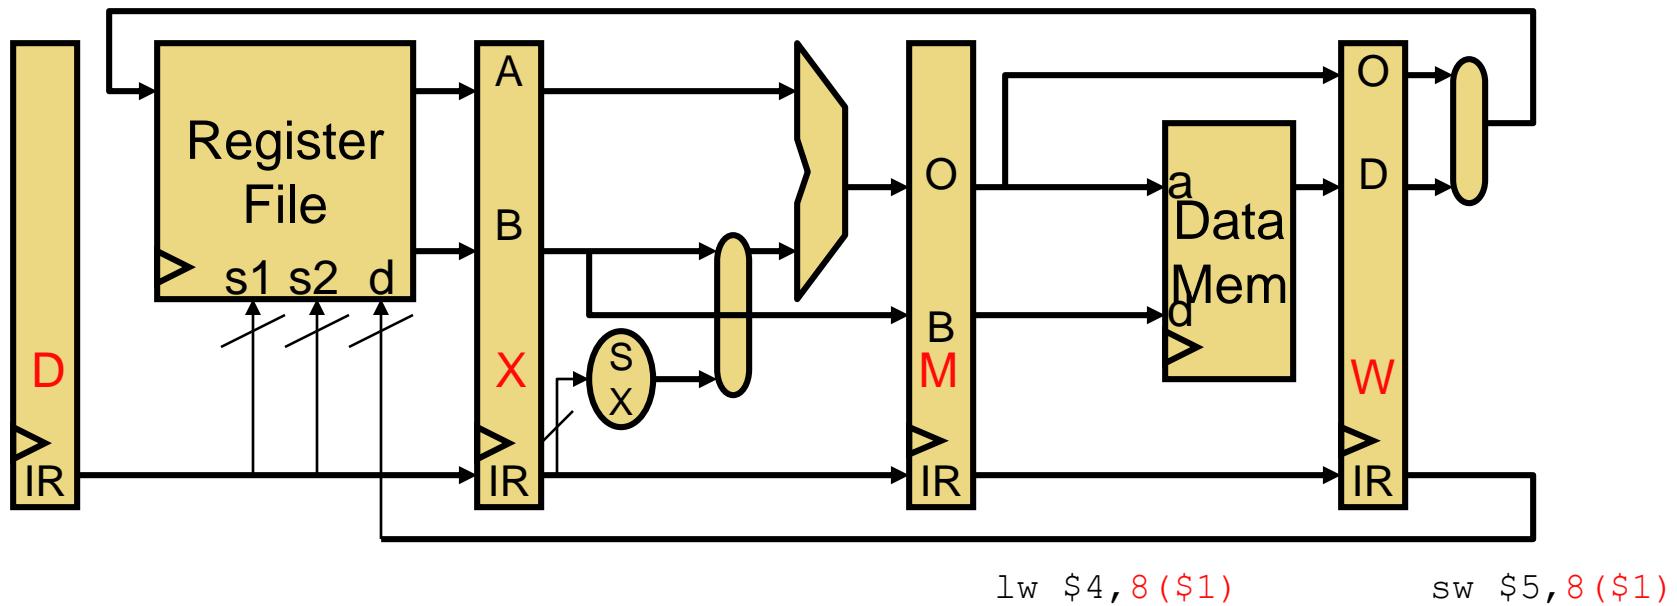
- d* = stall due to data hazard

	1	2	3	4	5	6	7	8	9
add $3<-$2,$1	F	D	X	M	W				
lw $4,4($3)		F	D	X	M	W			
addi $6<-$4,1			F	D	d*	X	M	W	
sub $8<-$3,$1				F	D	X	M	W

- Use compiler scheduling to reduce load-use stall frequency

	1	2	3	4	5	6	7	8	9
add $3<-$2,$1	F	D	X	M	W				
lw $4,4($3)		F	D	X	M	W			
sub $8<-$3,$1			F	D	X	M	W		
addi $6<-$4,1				F	D	X	M	W	

Dependencies Through Memory


```
lw $4,8($1)          sw $5,8($1)
```
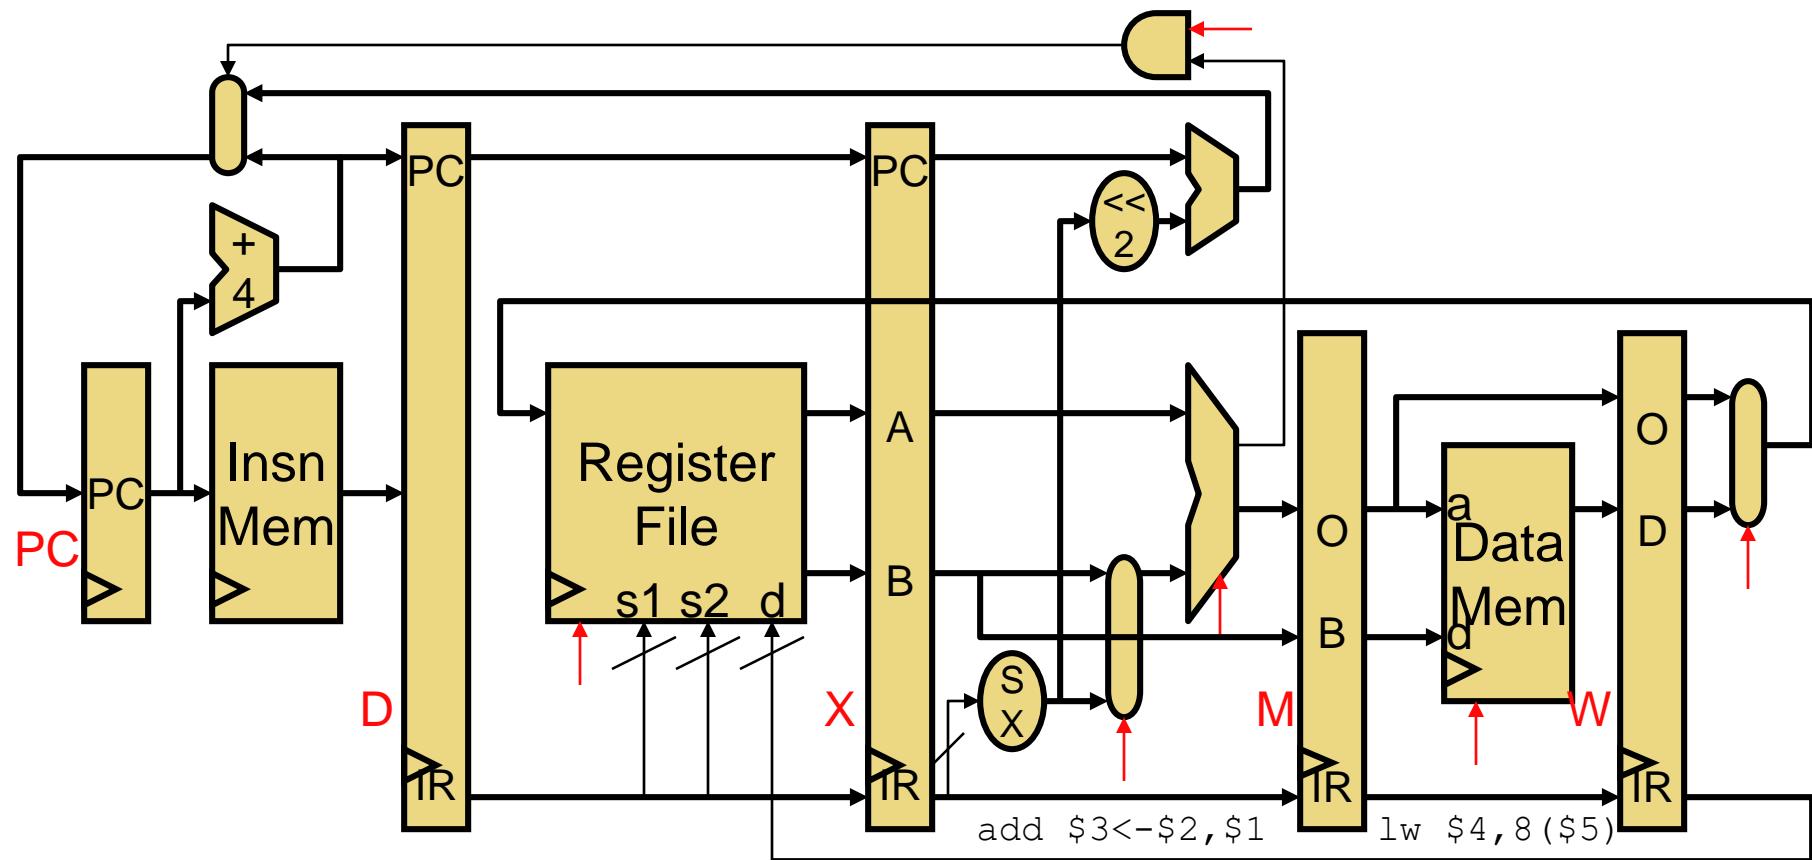
- Are "load to store" memory dependencies a problem?
  - No, `lw` following `sw` to same address in next cycle, gets right value
  - Why? Data mem read/write always take place in same stage

- Are there any other sort of hazards to worry about?

Structural Hazards

- **Structural hazards**
  - Two insns trying to use same circuit at same time
    - E.g., structural hazard on register file write port
- **To avoid structural hazards**
  - Avoided if:
    - Each insn uses every structure exactly once
    - For at most one cycle
    - All instructions travel through all stages
  - Add more resources:
    - Example: two memory accesses per cycle (Fetch & Memory)
    - Split instruction & data memories allows simultaneous access
- **Tolerate structure hazards**
  - Add stall logic to stall pipeline when hazards occur
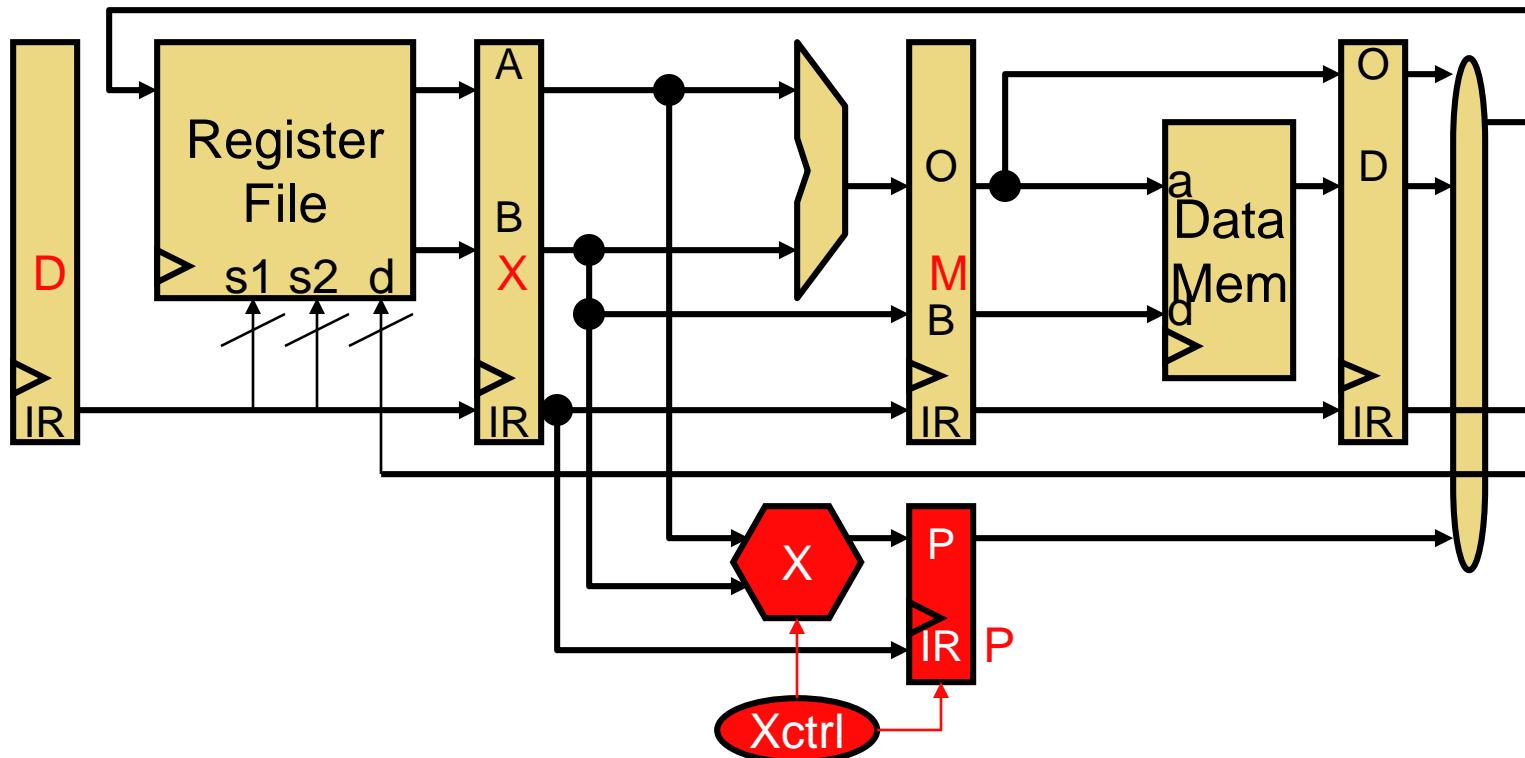
Why Does Every Insn Take 5 Cycles?


```
add $3<-$2,$1        lw $4,8($5)
```

- Could/should we allow **add** to skip M and go to W? No
  - It wouldn't help: peak fetch still only 1 insn per cycle
  - **Structural hazards**: imagine **add** after **lw** (only 1 reg. write port)
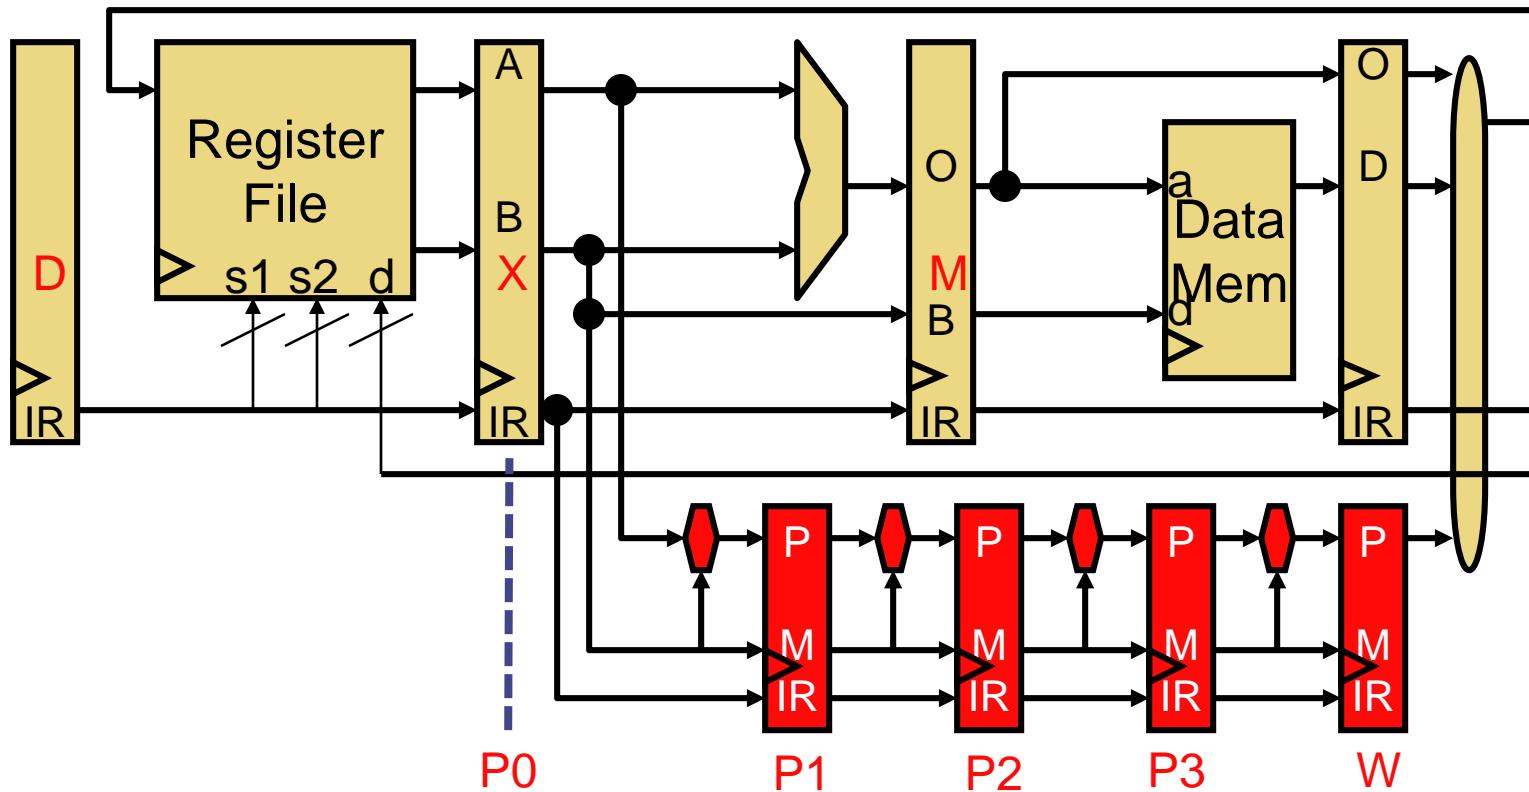
Multi-Cycle Operations

Pipelining and Multi-Cycle Operations

- What if you wanted to add an operation that takes multiple cycles to execute?
  - E.g., 4-cycle multiply
  - **P**: separate output latch connects to W stage
  - Controlled by pipeline control finite state machine (FSM)

A Pipelined Multiplier

- Multiplier itself is often pipelined, what does this mean?
  - Product/multiplicand register/ALUs/latches replicated
  - Can start different multiply operations in consecutive cycles
  - **But still takes 4 cycles to generate output value**

Pipeline Diagram with Multiplier

- Allow **independent** instructions

	1	2	3	4	5	6	7	8	9
`mul $4<-$3,$5`	F	D	P0	P1	P2	P3	W		
`addi $6<-$7,1`		F	D	X	M	W			
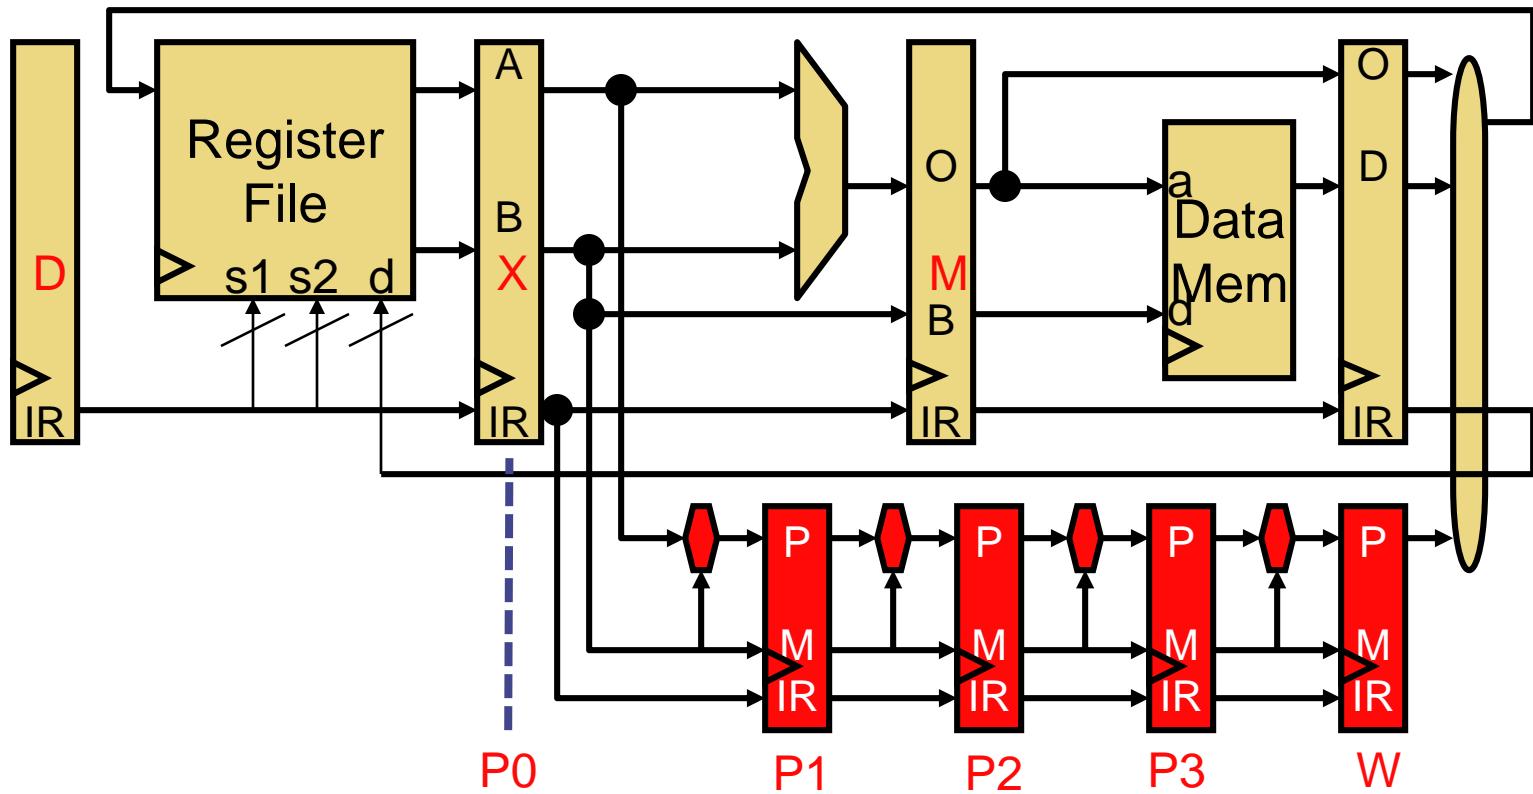
- Even allow **independent multiply** instructions

	1	2	3	4	5	6	7	8	9
`mul $4<-$3,$5`	F	D	P0	P1	P2	P3	W		
`mul $6<-$7,$8`		F	D	P0	P1	P2	P3	W	

- But must stall subsequent **dependent** instructions:
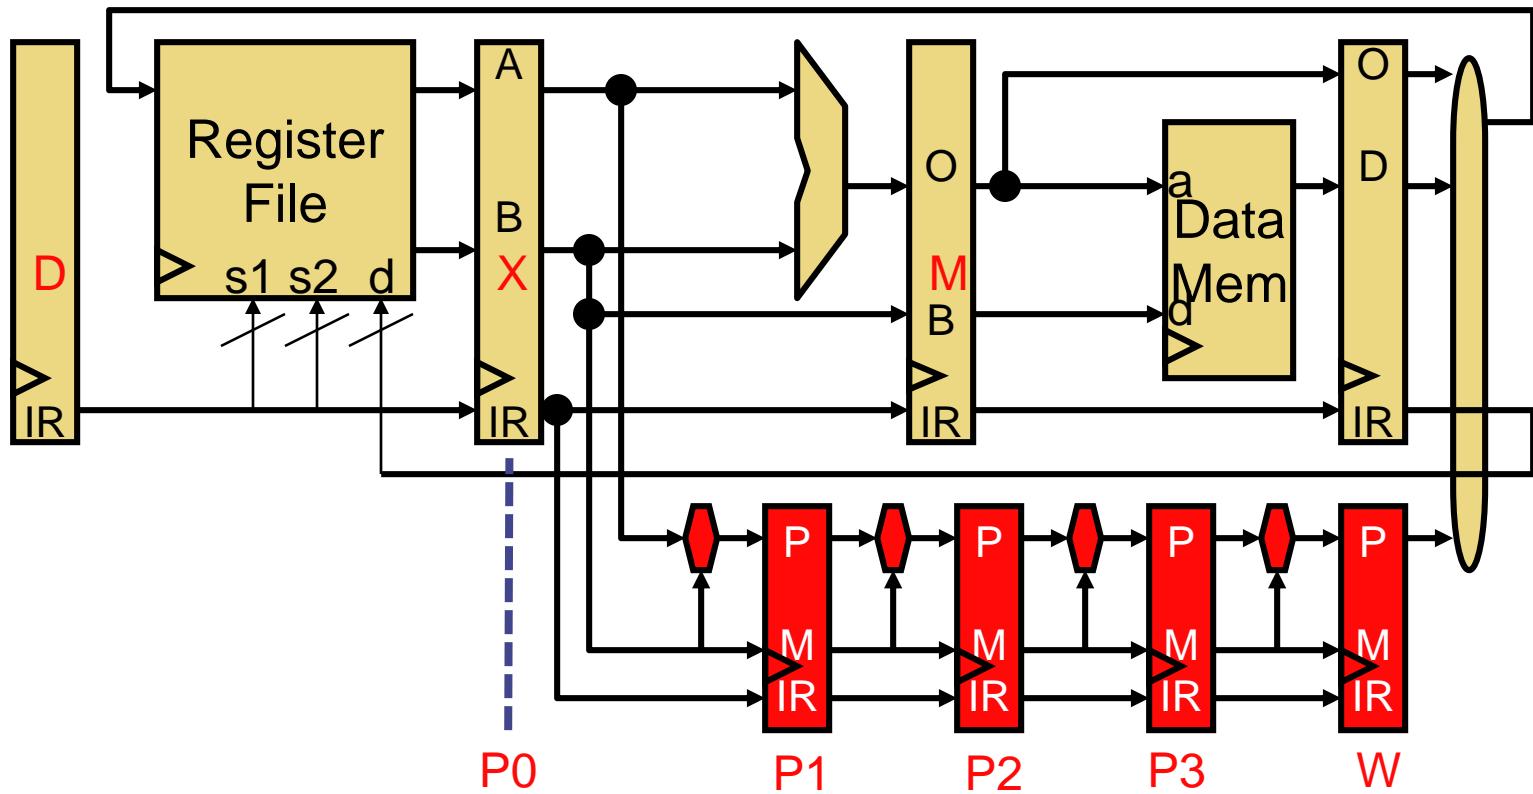
	1	2	3	4	5	6	7	8	9
`mul $4<-$3,$5`	F	D	P0	P1	P2	P3	W		
`addi $6<-$4,1`		F	D	d*	d*	d*	X	M	W

What about Stall Logic?

	1	2	3	4	5	6	7	8	9
`mul $4<-$3,$5`	F	D	P0	P1	P2	P3	W		
`addi $6<-$4,1`		F	D	**d***	**d***	**d***	X	M	W

What about Stall Logic?

Stall = (OldStallLogic) ||
   (D.IR.RegSrc1 == P0.IR.RegDest) || (D.IR.RegSrc2 == P0.IR.RegDest) ||
   (D.IR.RegSrc1 == P1.IR.RegDest) || (D.IR.RegSrc2 == P1.IR.RegDest) ||
   (D.IR.RegSrc1 == P2.IR.RegDest) || (D.IR.RegSrc2 == P2.IR.RegDest)

Multiplier Write Port Structural Hazard
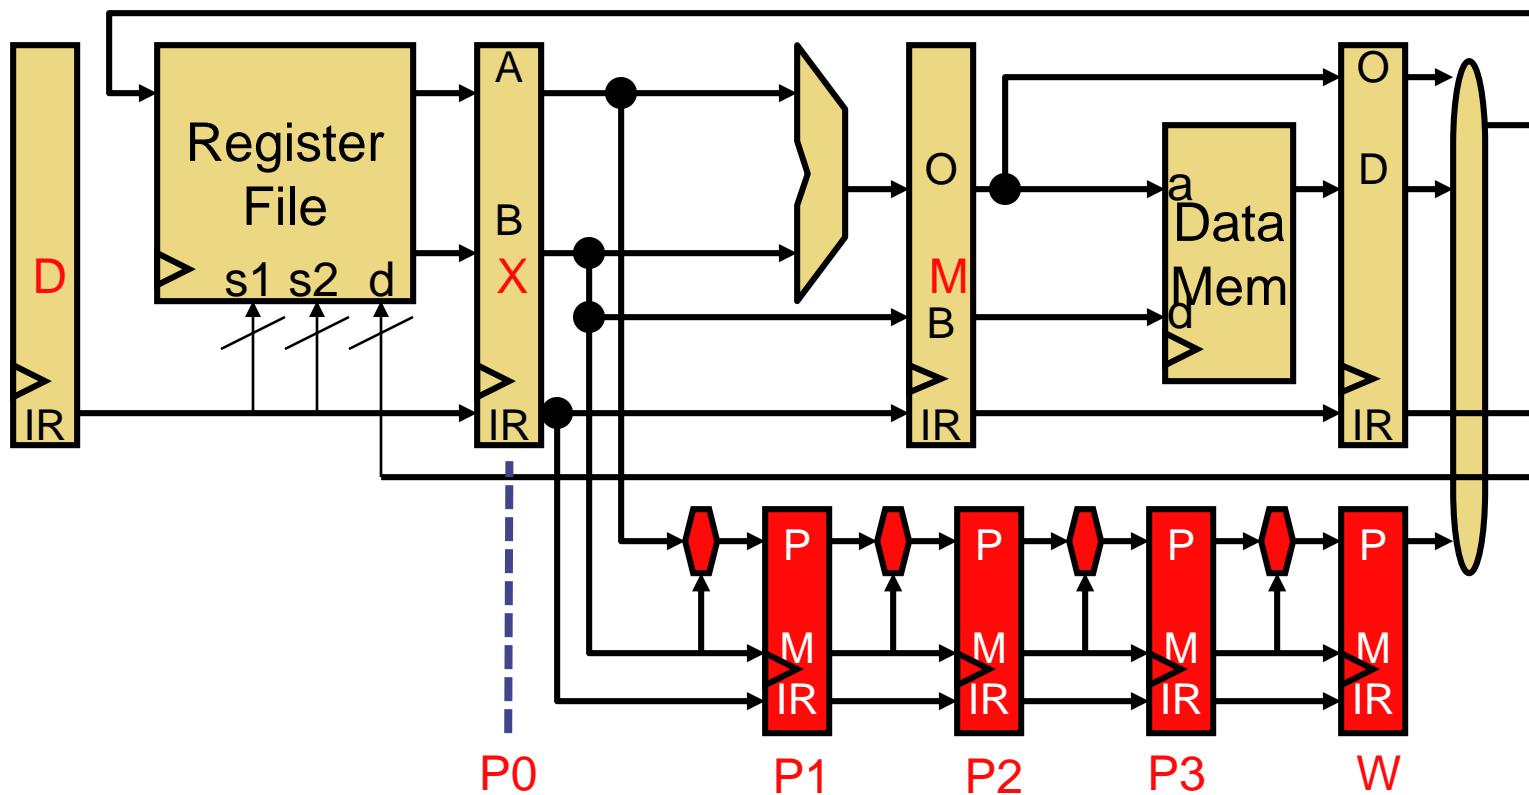
- What about...
  - Two instructions trying to write register file in same cycle?
  - Structural hazard!

- Must prevent:

	1	2	3	4	5	6	7	8	9
`mul $4<-$3,$5`	F	D	P0	P1	P2	P3	W		
`addi $6<-$1,1`		F	D	X	M	W			
`add $5<-$6,$10`			F	D	X	M	**W**		

- Solution? stall the subsequent instruction

	1	2	3	4	5	6	7	8	9
`mul $4<-$3,$5`	F	D	P0	P1	P2	P3	W		
`addi $6<-$1,1`		F	D	X	M	W			
`add $5<-$6,$10`			F	D	**d***	X	M	**W**	

Preventing Structural Hazard

- Fix to problem on previous slide:

  Stall = (OldStallLogic) ||

  (**D.IR.RegDest "is valid" &&**

  D.IR.Operation != MULT && P1.IR.RegDest "is valid")

More Multiplier Nasties
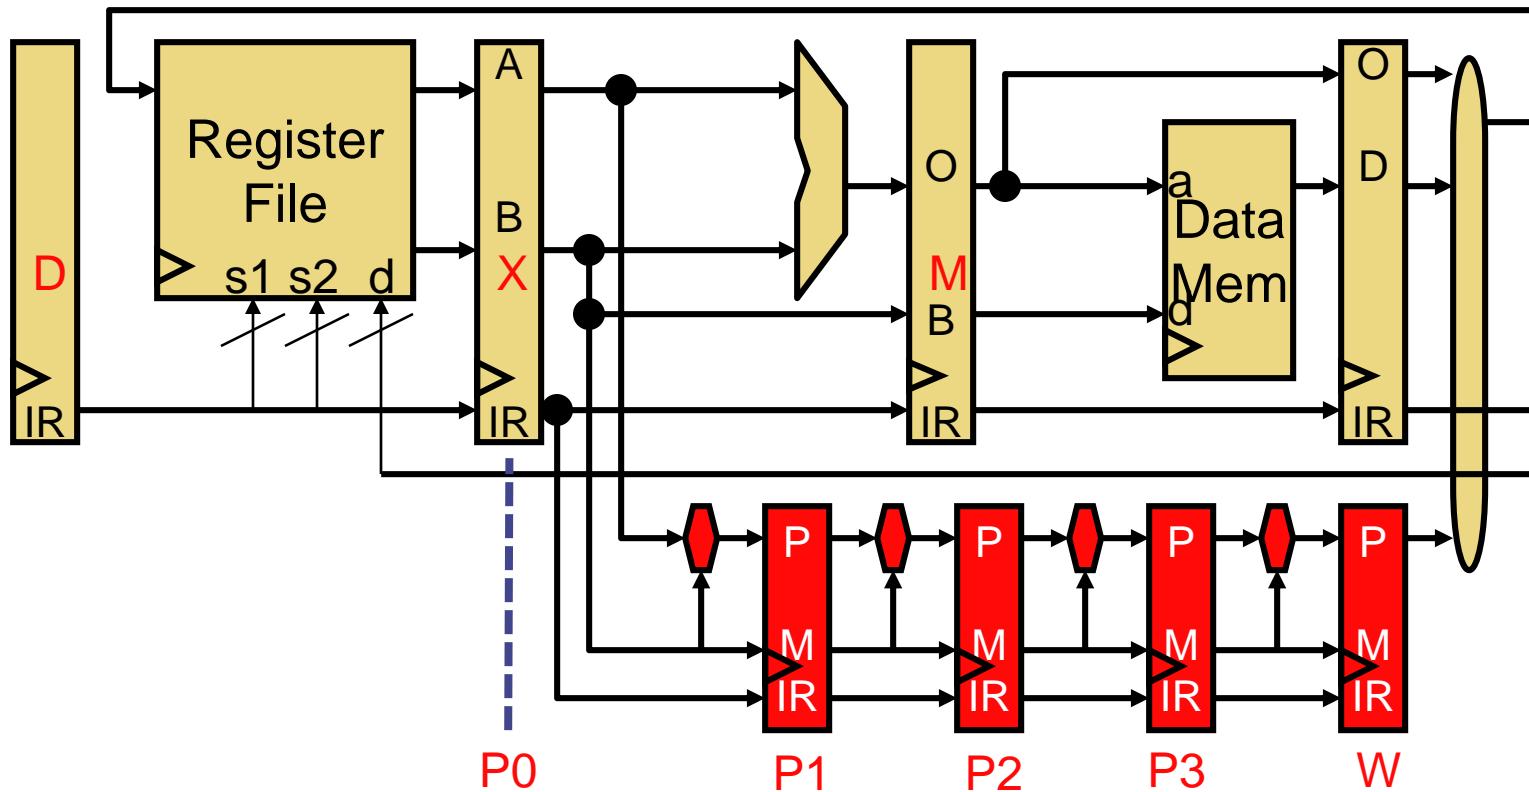
- ## What about…
  - Mis-ordered writes to the same register
  - Software thinks `add` gets `$4` from `addi`, actually gets it from `mul`

	1	2	3	4	5	6	7	8	9
`mul $4,$3,$5`	F	D	P0	P1	P2	P3	W		
`addi $4,$1,1`		F	D	X	M	**W**			
…									
…									
`add $10,$4,$6`									

- ## Common? Not for a 4-cycle multiply with 5-stage pipeline
  - More common with deeper pipelines
  - In any case, must be correct

Preventing Mis-Ordered Reg. Write

- Fix to problem on previous slide:

  Stall = (OldStallLogic) ||

  ((**D.IR.RegDest == P0.IR.RegDest) && (P0.IR.Operation == MULT)
  && (D.IR.Operation != MULT)**)

Corrected Pipeline Diagram

- With the correct stall logic
  - Prevent mis-ordered writes to the same register
  - Why two cycles of delay?

	1	2	3	4	5	6	7	8	9
`mul $4,$3,$5`	F	D	P0	P1	P2	P3	W		
`addi $4,$1,1`		F	D	**d***	**d***	X	M	**W**	
…									
…									
`add $10,$4,$6`									

- **Multi-cycle operations complicate pipeline logic**

Pipelined Functional Units

- Almost all multi-cycle functional units are pipelined
  - Each operation takes *N* cycles
  - But can start initiate a new (independent) operation every cycle
  - Requires internal latching and some hardware replication
  - + A cheaper way to add bandwidth than multiple non-pipelined units

	1	2	3	4	5	6	7	8	9	10	11
`mulf f0,f1,f2`	F	D	E*	E*	E*	E*	W				
`mulf f3,f4,f5`		F	D	E*	E*	E*	E*	W			
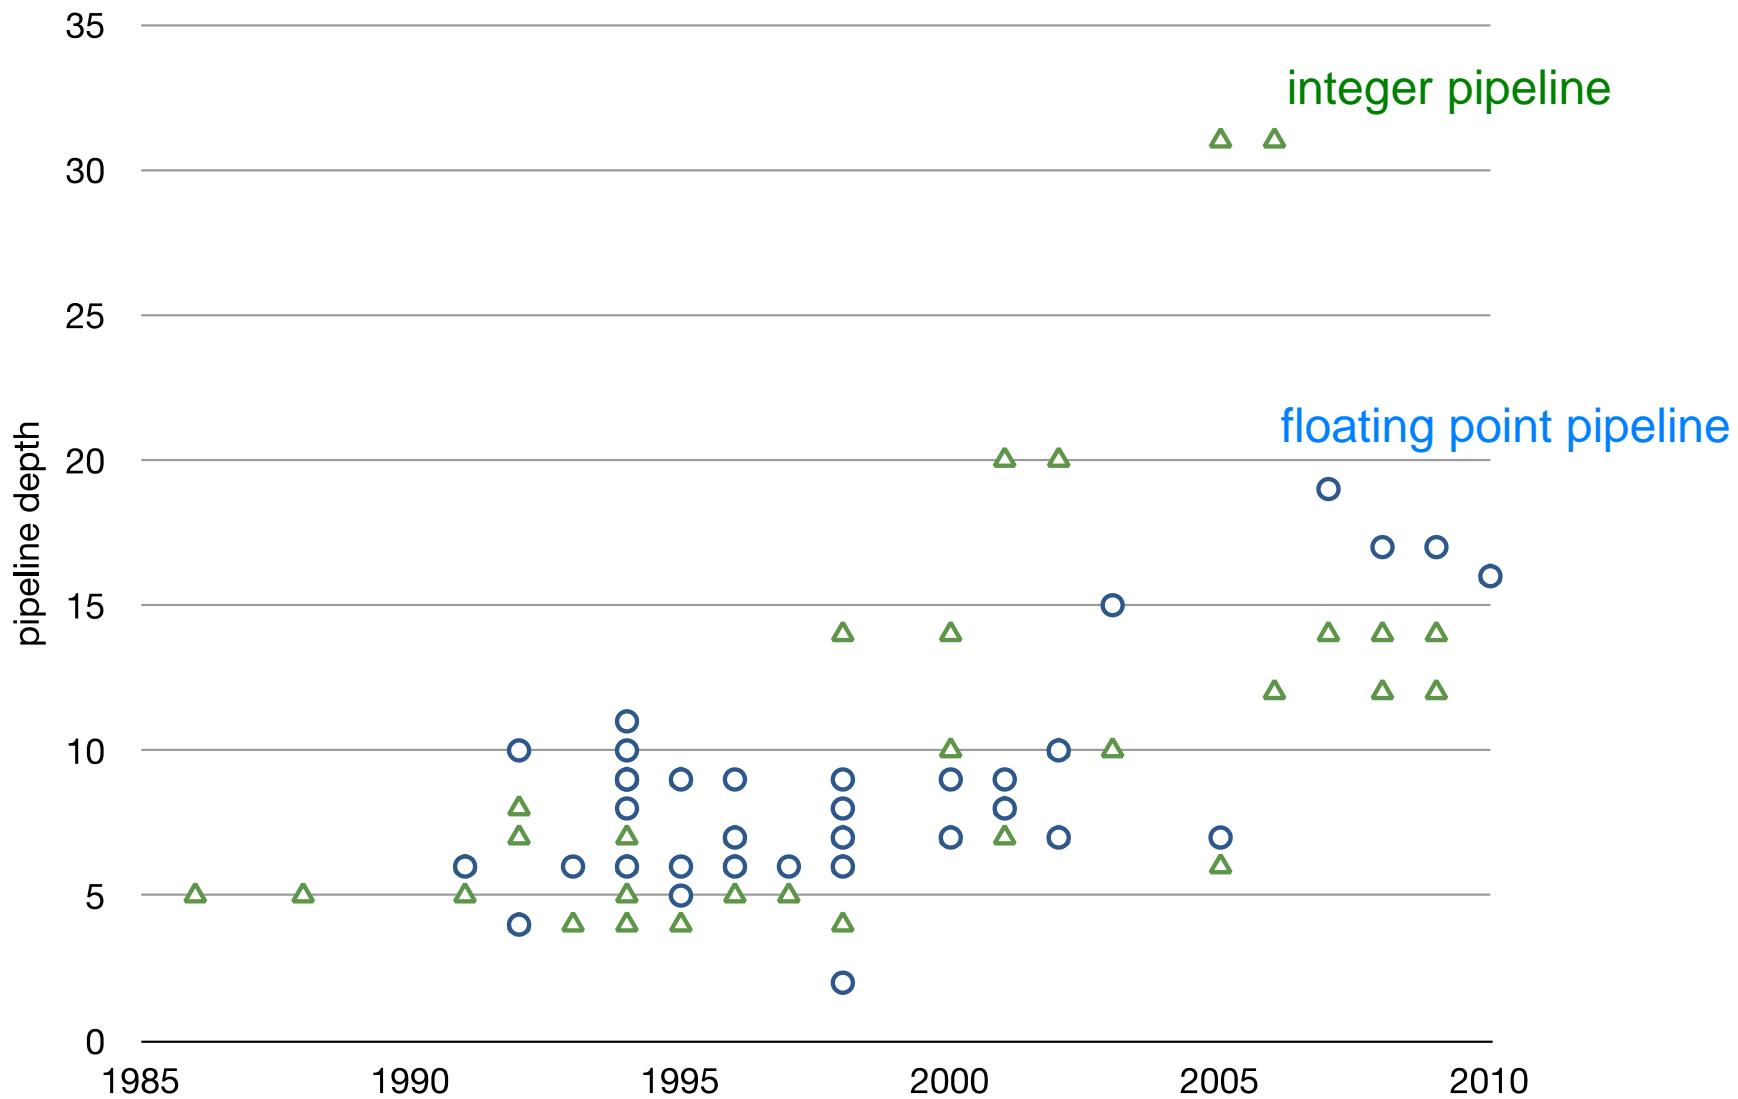
- One exception: int/FP divide: difficult to pipeline and not worth it

	1	2	3	4	5	6	7	8	9	10	11
`divf f0,f1,f2`	F	D	E/	E/	E/	E/	W				
`divf f3,f4,f5`		F	D	**s***	**s***	**s***	E/	E/	E/	E/	W
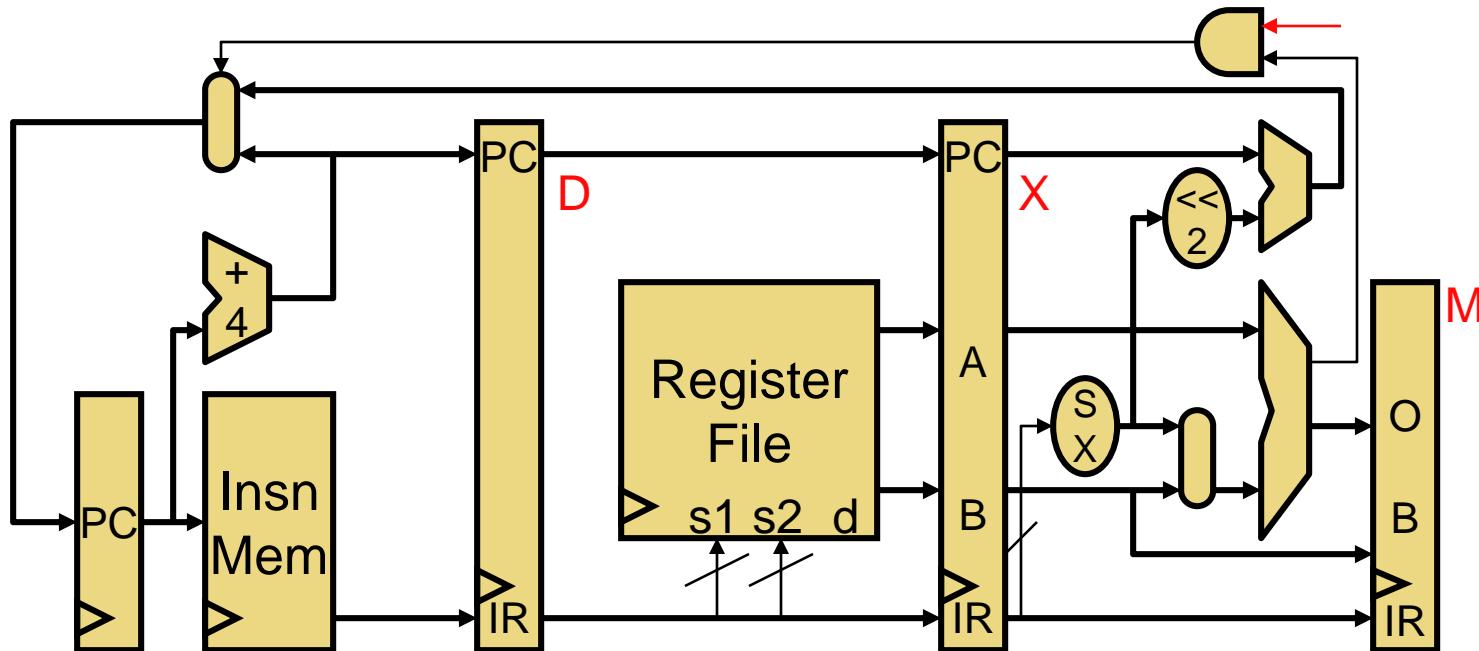
- **s*** = structural hazard, two insns need same structure
  - ISAs and pipelines designed to have few of these
  - Canonical example: all insns forced to go through M stage

Pipeline Depth

integer pipeline

floating point pipeline
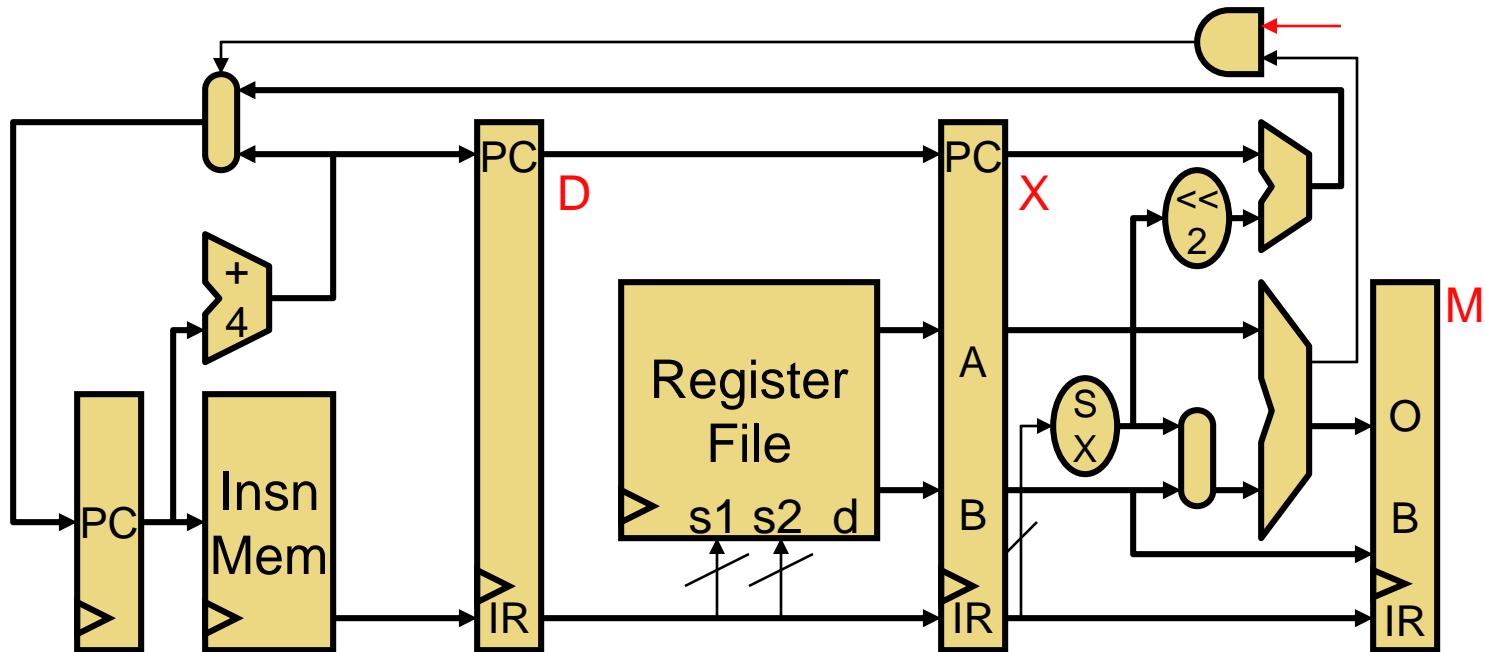
Branch Prediction
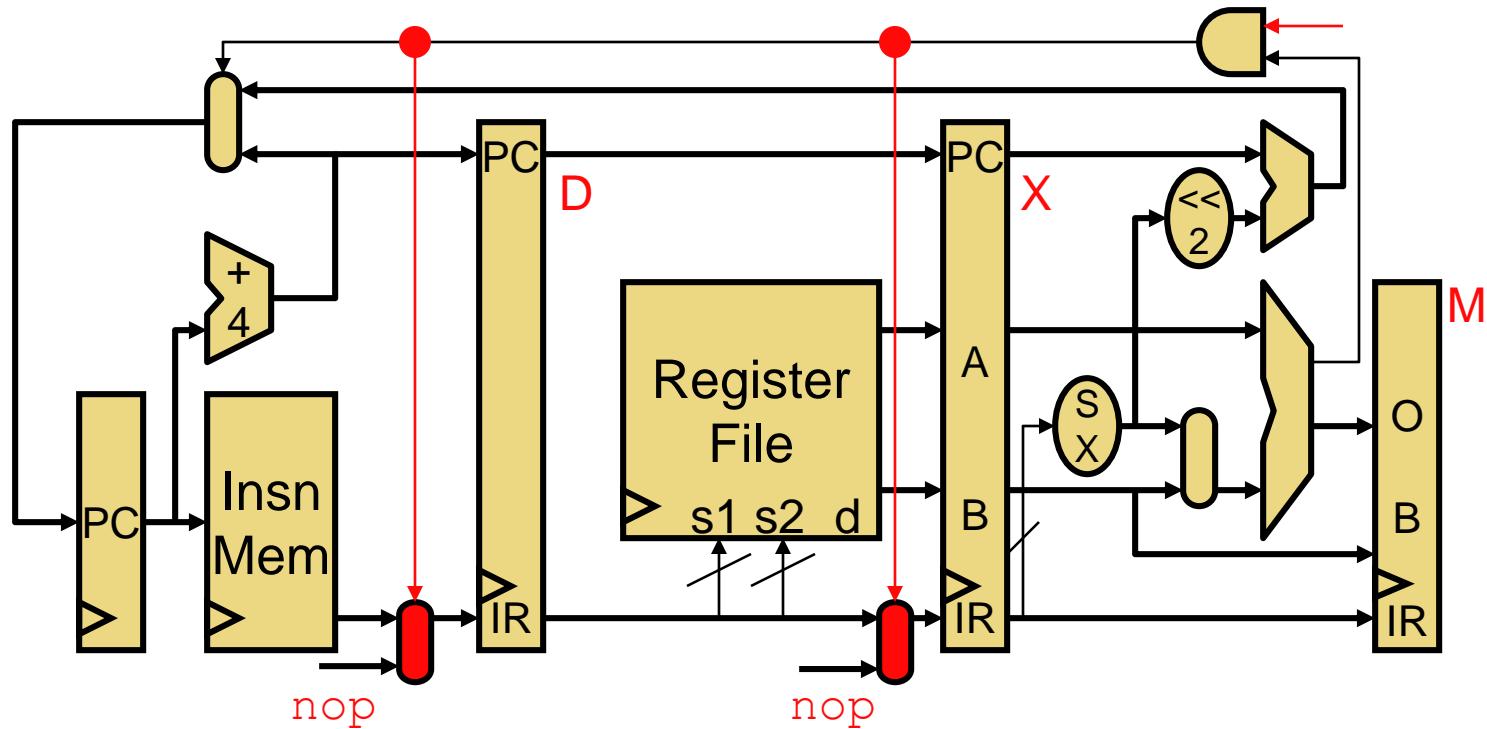
- **What is** a branch?

What About Branches?

- **When do we read from insn memory?**
- **When do we know branch outcome?**
  - When do we even know if we have a branch?

Branch Recovery

- **Branch recovery**: what to do when branch is actually taken
  - Insns that will be written into D and X are wrong
    - Flush them, i.e., replace them with `nops`
  - \+ They haven't written permanent state yet (regfile, DMem)
    - Two cycle penalty for taken branches

Big Idea: Speculative Execution

- **Speculative execution**
  - Execute before all parameters known with certainty
  - **Correct speculation** lets us avoid stalls & improve performance

- **Speculation requirements**
  - When should we speculate?
  - Was our speculation correct?
  - How do I cleanup a mis-speculation?
    - restore pre-speculation state

- **Control speculation**: speculation aimed at control hazards
  - Unknown parameter: are these the correct insns to execute next?

Control Speculation Mechanics

- Guess branch target, start fetching at guessed position
  - Doing nothing is implicitly guessing target is PC+4
  - Can actively guess other targets: **dynamic branch prediction**

- Execute branch to verify (check) guess
  - Correct speculation? keep going
  - Mis-speculation? Flush mis-speculated insns
    - Hopefully haven't modified permanent state (Regfile, DMem)
    + Happens naturally in in-order 5-stage pipeline

Branch Speculation and Recovery

		1	2	3	4	5	6	7	8	9
Correct:	`addi r1,1➔r3`	F	D	X	M	W				
	`bnez r3,targ`		F	D	X	M	W			
	`st r6➔[r7+4]`			**F**	**D**	X	M	W		
	`mul r8,r9➔r10`				**F**	D	X	M	W	

speculative

- **Mis-speculation recovery**: what to do on wrong guess
  - Not too painful in a short, in-order pipeline
  - Branch resolves in X
  - + Younger insns (in F, D) haven't changed permanent state
  - **On next cycle, flush** insns in D and X

		1	2	3	4	5	6	7	8	9
Recovery:	`addi r1,1➔r3`	F	D	X	M	W				
	`bnez r3,targ`		F	D	**X**	M	W			
	~~`st r6➔[r7+4]`~~			**F**	**D**	--	--	--		
	~~`mul r8,r9➔r10`~~				**F**	--	--	--	--	
	`targ:add r4,r5➔r4`					**F**	D	X	M	W

Modern Computer Architecture
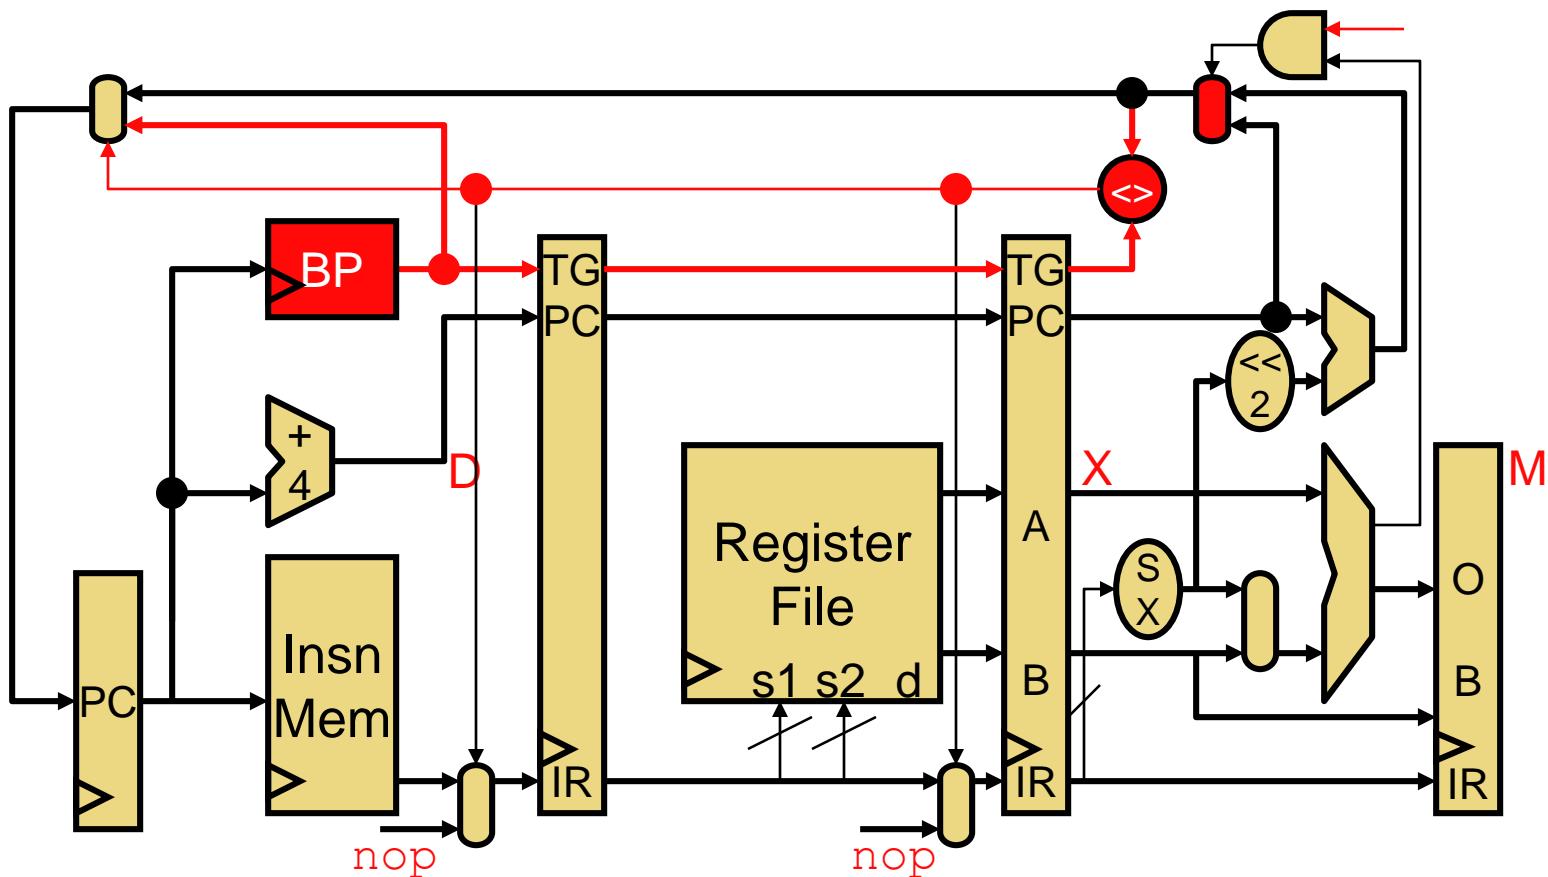
Branch Performance

- Back of the envelope calculation
  - **Branch: 20%**, load: 20%, store: 10%, other: 50%
  - Say, **75% of branches are taken**

- CPI = 1 + 20% * 75% * 2 =
  $$1 + \mathbf{0.20 * 0.75 * 2} = 1.3$$
  - **Branches cause 30% slowdown**
    - Worse with deeper pipelines (higher misprediction penalty)

- Can we do better than assuming branch is not taken?

Dynamic Branch Prediction
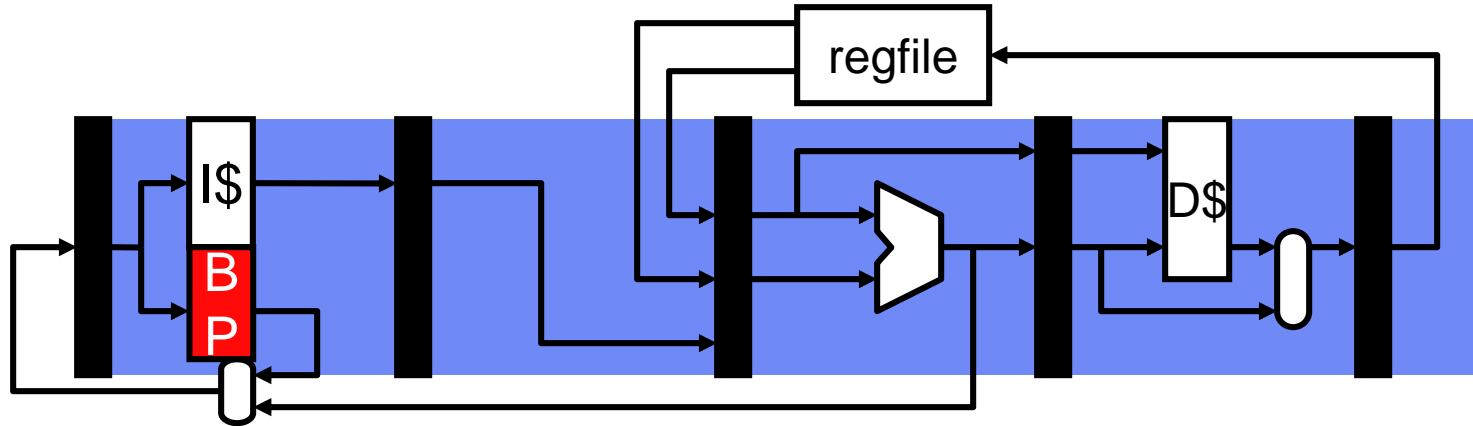
- **Dynamic branch prediction**: hardware guesses outcome
  - Start fetching from guessed address
  - Flush on **mis-prediction**

Branch Prediction Performance

- Parameters
  - **Branch: 20%**, load: 20%, store: 10%, other: 50%
  - 75% of branches are taken
- Dynamic branch prediction
  - Branches predicted with 95% accuracy
  - CPI = 1 + 20% * 5% * 2 = **1.02**
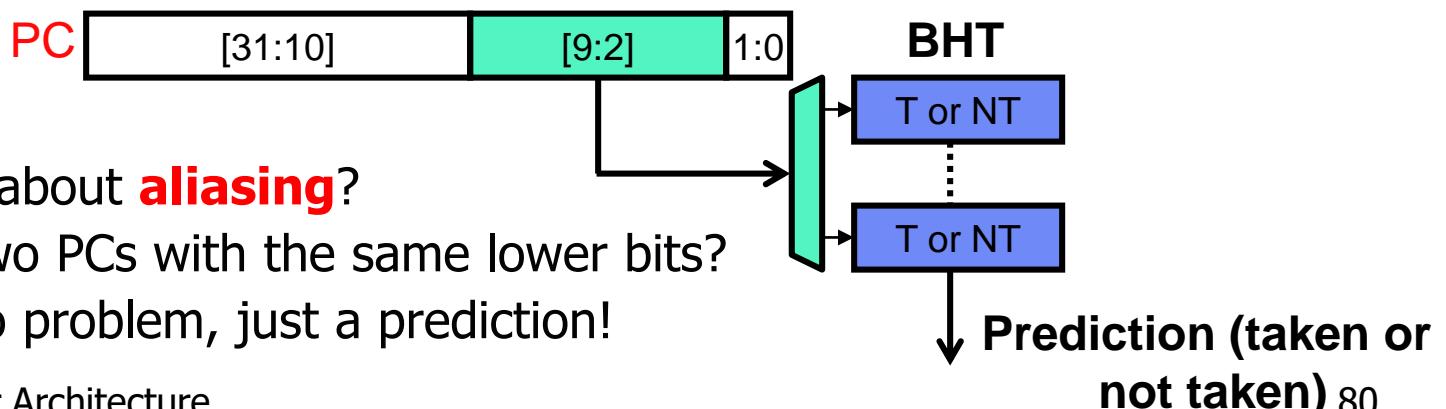
Dynamic Branch Prediction Components

- ## Step #1: is it a branch?
  - Easy after decode, but we're in fetch…
- ## Step #2: is the branch taken or not taken?
  - **Direction predictor** (applies to conditional branches only)
  - Predicts taken/not-taken
  - Easy after execute, but we're in fetch…
- ## Step #3: if the branch is taken, where does it go?
  - Easy after decode, but we're in fetch…

Branch Direction Prediction

- **Past performance may be indicative of future results**
  - Record the past in a hardware structure
- **Direction predictor (DIRP)**
  - Map conditional-branch PC to taken/not-taken (T/N) decision
  - Individual conditional branches often biased or weakly biased
    - 90%+ one way or the other considered **"biased"**
    - Why? Loop back edges, checking for uncommon conditions
- **Branch history table (BHT)**: simplest predictor
  - PC indexes table of bits (0 = N, 1 = T), no tags
  - Essentially: branch will go same way it went last time

PC [31:10] [9:2] 1:0   **BHT**

T or NT

T or NT

- What about **aliasing**?
  - Two PCs with the same lower bits?
  - No problem, just a prediction!

Prediction (taken or not taken)

Branch History Table (BHT)

- **Branch history table (BHT)**: simplest direction predictor
  - PC indexes table of bits (0 = N, 1 = T), no tags
  - Essentially: branch will go same way it went last time
  - Problem: **inner loop branch** below
    ```
    for (i=0;i<100;i++)
        for (j=0;j<3;j++)
            // loop body
    ```
    – Two "built-in" mis-predictions per inner loop iteration
    – Branch predictor "changes its mind too quickly"

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	T	T	T	Correct
3	T	T	T	Correct
4	T	T	N	Wrong
5	N	N	T	Wrong
6	T	T	T	Correct
7	T	T	T	Correct
8	T	T	N	Wrong
9	N	N	T	Wrong
10	T	T	T	Correct
11	T	T	T	Correct
12	T	T	N	Wrong

Two-Bit Saturating Counters (2bc)

- **Two-bit saturating counters (2bc)** [Smith 1981]
  - Replace each single-bit prediction
    - (0,1,2,3) = (N,n,t,T)
  - Adds "hysteresis"
    - Force predictor to mis-predict twice before "changing its mind"
  - One misprediction each loop execution (rather than two)
    - +Fixes this pathology (which is not contrived, by the way)
    - Can we do even better?

Time	State	Prediction	Outcome	Result?
1	N	N	T	Wrong
2	n	N	T	Wrong
3	t	T	T	Correct
4	T	T	N	Wrong
5	t	T	T	Correct
6	T	T	T	Correct
7	T	T	T	Correct
8	T	T	N	Wrong
9	t	T	T	Correct
10	T	T	T	Correct
11	T	T	T	Correct
12	T	T	N	Wrong

Correlated Predictor

- **Correlated (two-level) predictor** [Patt 1991]
  - Exploits observation that branch outcomes are correlated
  - Maintains separate prediction per (PC, BHR) pairs
    - **Branch history register (BHR)**: recent branch outcomes
  - Simple working example: assume program has one branch
    - BHT: one 1-bit DIRP entry
    - BHT+**2BHR**: 2^2 = **four** 1-bit DIRP entries
  - – Why didn't we do better?
    - BHT not long enough to capture pattern

Time	"Pattern"	State				Prediction	Outcome	Result?
		NN	NT	TN	TT			
1	NN	**N**	N	N	N	**N**	**T**	Wrong
2	NT	T	**N**	N	N	**N**	**T**	Wrong
3	TT	T	T	N	**N**	**N**	**T**	Wrong
4	TT	T	T	N	**T**	**T**	**N**	Wrong
5	TN	T	T	**N**	N	**N**	**T**	Wrong
6	NT	T	**T**	T	N	**T**	**T**	Correct
7	TT	T	T	T	**N**	**N**	**T**	Wrong
8	TT	T	T	T	**T**	**T**	**N**	Wrong
9	TN	T	T	**T**	N	**T**	**T**	Correct
10	NT	T	**T**	T	N	**T**	**T**	Correct
11	TT	T	T	T	**N**	**N**	**T**	Wrong
12	TT	T	T	T	**T**	**T**	**N**	Wrong

Correlated Predictor – 3 Bit Pattern

- Try 3 bits of history
- 2^3 DIRP entries per pattern

Time	"Pattern"	NNN	NNT	NTN	NTT	TNN	TNT	TTN	TTT	Prediction	Outcome	Result?
1	NNN	**N**	N	N	N	N	N	N	N	**N**	**T**	Wrong
2	NNT	T	**N**	N	N	N	N	N	N	**N**	**T**	Wrong
3	NTT	T	T	N	**N**	N	N	N	N	**N**	**T**	Wrong
4	TTT	T	T	N	T	N	N	N	**N**	**N**	**N**	Correct
5	TTN	T	T	N	T	N	N	**N**	N	**N**	**T**	Wrong
6	TNT	T	T	N	T	N	**N**	T	N	**N**	**T**	Wrong
7	NTT	T	T	N	**T**	N	T	T	N	**T**	**T**	Correct
8	TTT	T	T	N	T	N	T	T	**N**	**N**	**N**	Correct
9	TTN	T	T	N	T	N	T	**T**	N	**T**	**T**	Correct
10	TNT	T	T	N	T	N	**T**	T	N	**T**	**T**	Correct
11	NTT	T	T	N	**T**	N	T	T	N	**T**	**T**	Correct
12	TTT	T	T	N	T	N	T	T	**N**	**N**	**N**	Correct
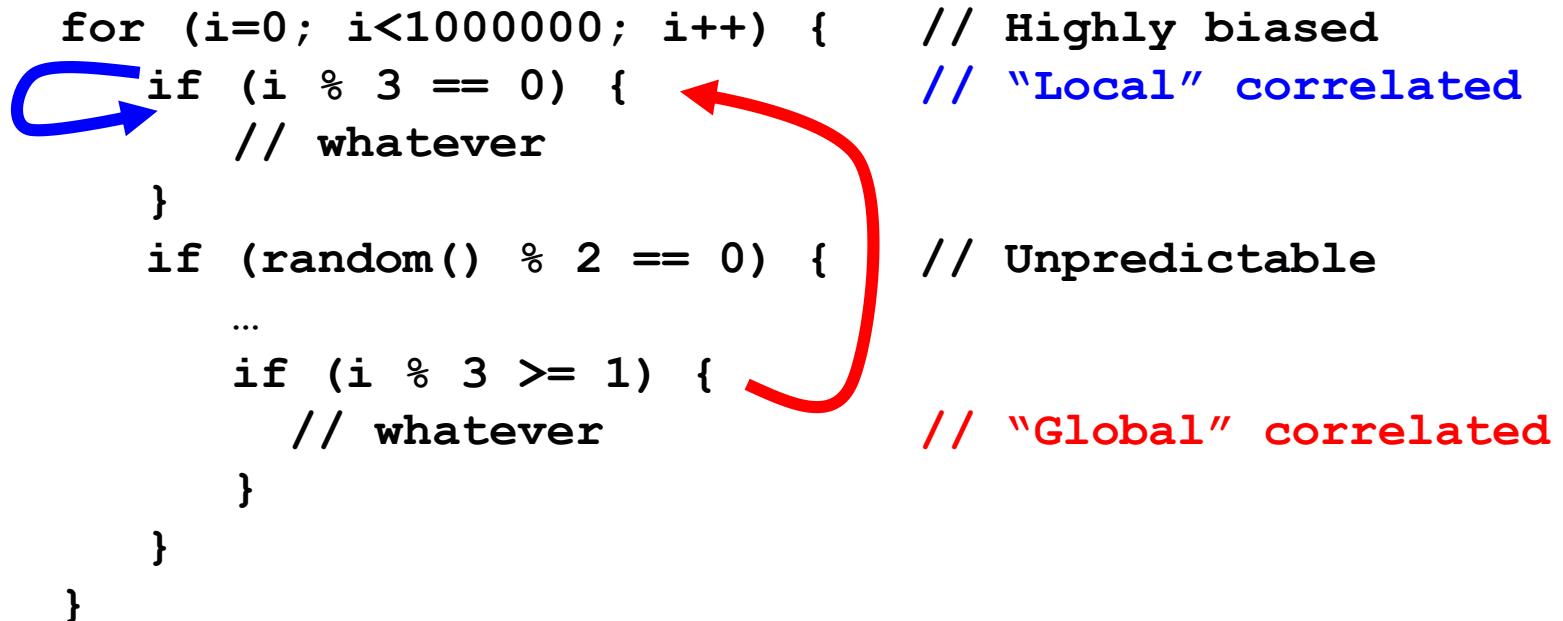
+ No mis-predictions after predictor learns all the relevant patterns!

Correlated Predictor Design I

- Design choice I: one **global** BHR or one per PC (**local**)?
  - Each one captures different kinds of patterns
    - Global history captures relationship among different branches
    - Local history captures "self" correlation
  - Local history requires another table to store the per-PC history
- Consider:

```
for (i=0; i<1000000; i++) {     // Highly biased
    if (i % 3 == 0) {           // "Local" correlated
        // whatever
    }
    if (random() % 2 == 0) {    // Unpredictable
        …
        if (i % 3 >= 1) {
            // whatever         // "Global" correlated
        }
    }
}
```

# Correlated Predictor Design II
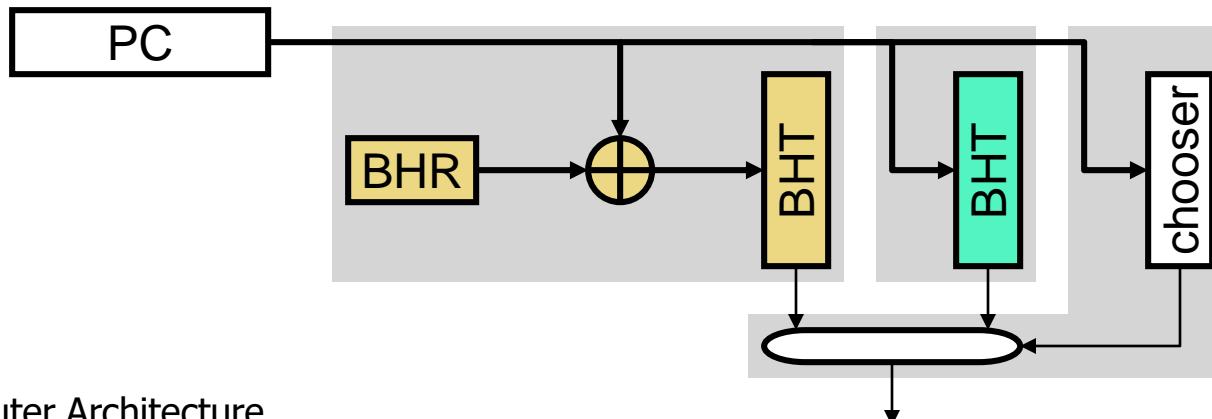
- Design choice II: how many history bits (BHR size)?
  - Tricky one
  - + Given unlimited resources, longer BHRs are better, but…
  - − BHT utilization decreases
    - − Many history patterns are never seen
    - − Many branches are history independent (don't care)
    - PC xor BHR allows multiple PCs to dynamically share BHT
    - BHR length $< \log_2$(BHT size)
  - − Predictor takes longer to train
  - Typical length: 8–12
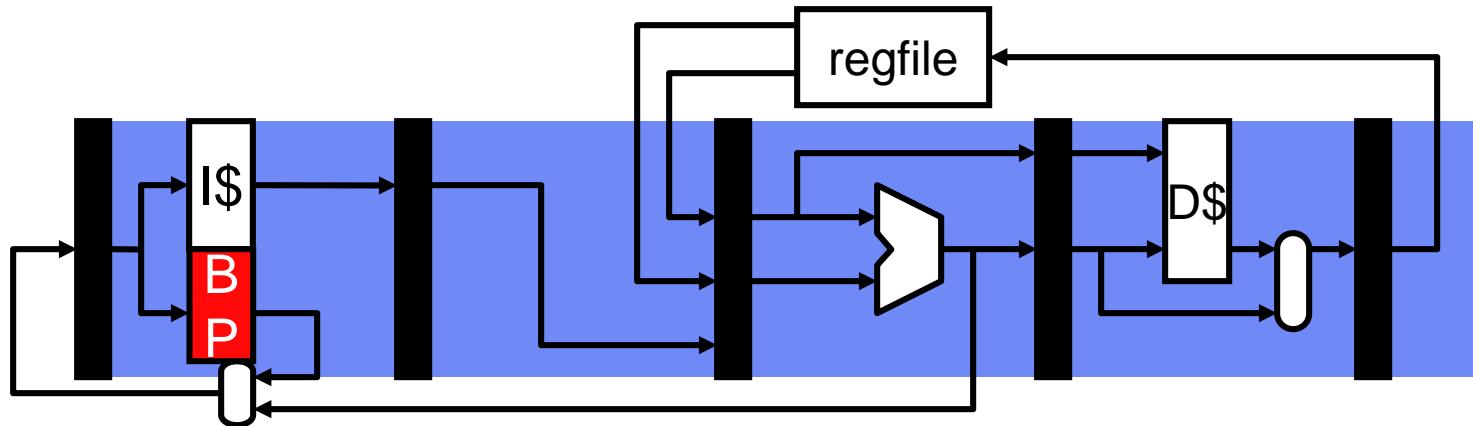
Hybrid Predictor

- **Hybrid (tournament) predictor** [McFarling 1993]
  - Attacks correlated predictor BHT capacity problem
  - Idea: combine two predictors
    - **Simple BHT** predicts history independent branches
    - **Correlated predictor** predicts only branches that need history
    - **Chooser** assigns branches to one predictor or the other
    - Branches start in simple BHT, move mis-prediction threshold
  - + Correlated predictor can be made **smaller**, handles fewer branches
 - + 90–95% accuracy
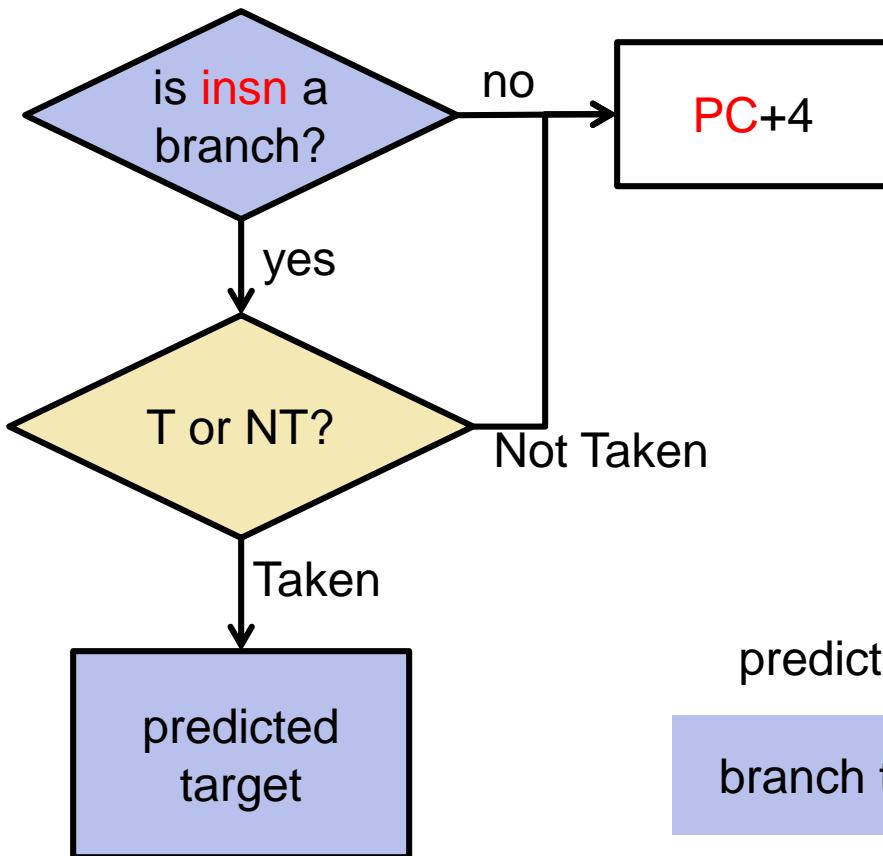
Revisiting Branch Prediction Components

- Step #1: is it a branch?
  - Easy after decode... during fetch need another **predictor**
- Step #2: is the branch taken or not taken?
  - **Direction predictor** (as before)
- Step #3: if the branch is taken, where does it go?
  - **Branch target buffer (BTB)**
  - Supplies target PC if branch is taken

Branch Prediction Steps

is insn a branch?

no → PC+4

yes

T or NT?

Not Taken

Taken

predicted target

- Which insn's behavior are we trying to predict?
- Where does PC come from?

prediction source:

branch target buffer

direction predictor

Is insn a branch?

- How do we know if an insn is a branch or not?
  - We don't until after decode
- Idea #1: record a list of insns that are **not** branches
  - what could go wrong?
  - must also consider aliasing issues
  - what happens if we think a non-branch insn is a branch?

Branch Target Buffer (BTB)

- As before: learn from past, predict the future
  - Record the past branches in a hardware structure

- **Branch target buffer (BTB)**:
  - A small hardware table
    - input = PC, data = 1 bit (0=not-branch,1=branch)
  - is-a-branch = BTB[hash(PC)]
  - Hash function is typically just extracting lower bits (as before)
  - Aliasing?
    - not a correctness issue, but can be big performance problem
    - there are **many** more non-branches than branches
    - what if…
      - BTB has 1K entries
      - programs have 20% branches
      - we have a 5KB program

BTB Entries

- Add "**tag**" field to BTB entries
- is-a-branch = (BTB[hash(PC)].**tag** == PC) ? BTB[PC].**iab** : 0
  - is-a-branch field is actually redundant now!
- reduces effects of aliasing considerably
  - only need storage ~(# branches)

index	is-a-branch	tag
0	0	0
1	1	0x4e3745
2	0	0
3	0	0
4	0	0
5	0	0
6	0	0
7	0	0

What is a taken branch's target?

- Add a "**target**" field to BTB entries
- is-a-branch = (BTB[hash(PC)].**tag** == PC) ? 1 : 0
- predicted-target = (BTB[hash(PC)].**tag** == PC) ? BTB[PC].**target** : 0
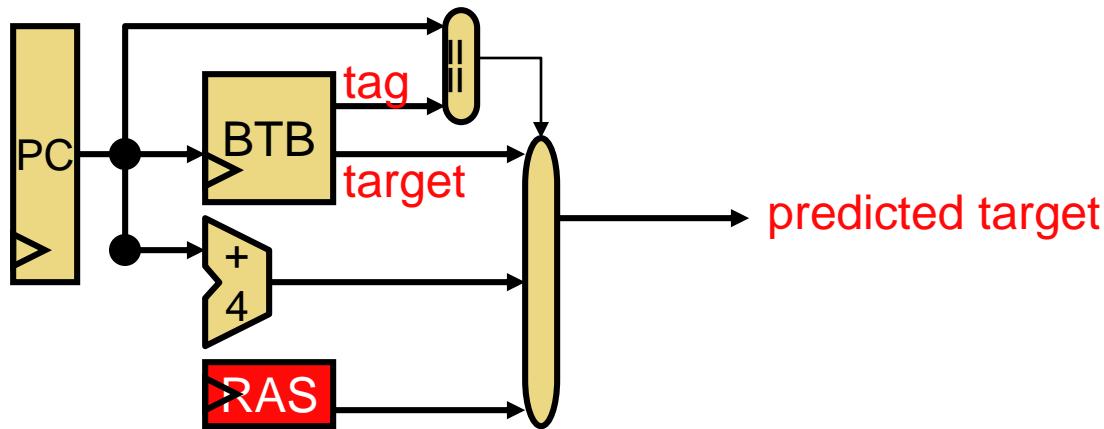
index	tag	target
0	0	0
1	0x4e3745	0x4738da
2	0	0
3	0	0
4	0	0
5	0	0
6	0	0
7	0	0

Why Does a BTB Work?

- Because most control insns use **direct targets**
  - Target encoded in insn itself $\rightarrow$ same "taken" target every time

- What about **indirect targets**?
  - Target held in a register $\rightarrow$ can be different each time
  - Two indirect call idioms
    + Dynamically linked functions (DLLs): target always the same
    - Dynamically dispatched (virtual) functions: hard but uncommon
  - Two indirect unconditional jump idioms
    - Switches: hard but uncommon
    - Function returns: **hard and common**

Return Address Stack (RAS)
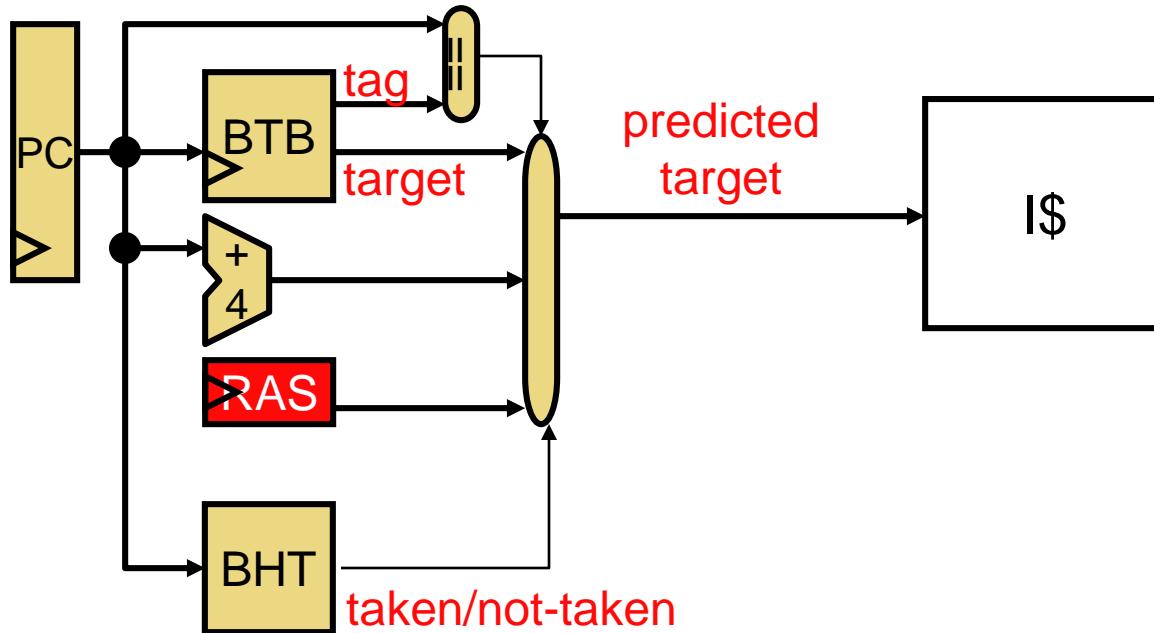
- **Return address stack (RAS)**
  - Call instruction? RAS[TopOfStack++] = PC+4
  - Return instruction? Predicted-target = RAS[--TopOfStack]
  - Q: how can you tell if an insn is a call/return before decoding it?
    - Accessing RAS on every insn BTB-style doesn't work
  - Answer: another predictor (or put them in BTB marked as "return")
    - Or, **pre-decode bits** in insn mem, written when first executed

Putting It All Together

- BTB & branch direction predictor during fetch

- If branch prediction correct, no taken branch penalty

Branch Prediction Performance
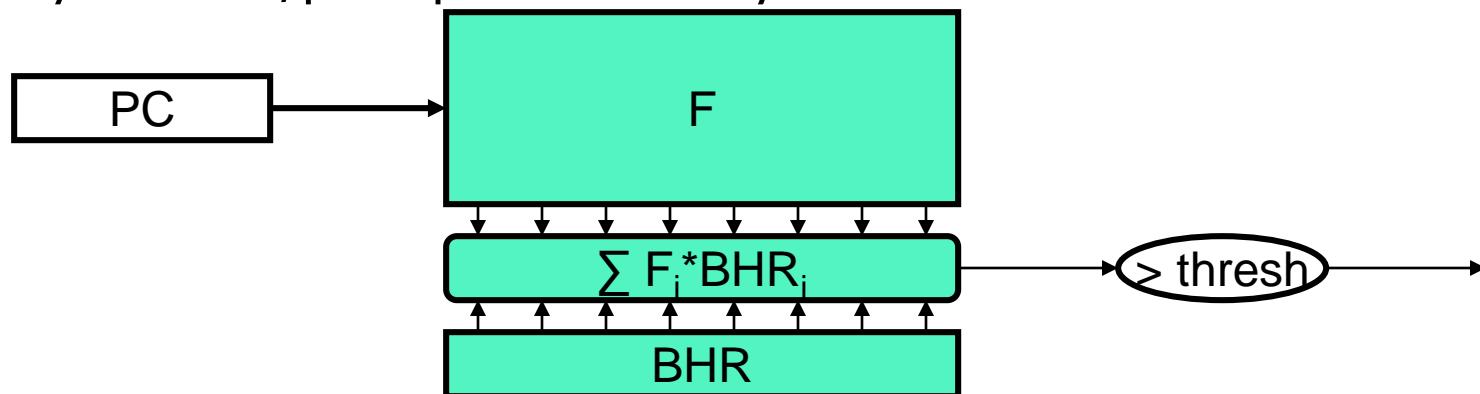
- Dynamic branch prediction
  - 20% of instruction branches
  - Simple predictor: branches predicted with 75% accuracy
    - CPI = 1 + (20% * **25%** * 2) = **1.1**
  - More advanced predictor: 95% accuracy
    - CPI = 1 + (20% * **5%** * 2) = **1.02**

- Branch mis-predictions still a big problem though
  - Pipelines are long: typical mis-prediction penalty is 10+ cycles
 - For cores that do more per cycle, predictions more costly (later)

Research: Perceptron Predictor
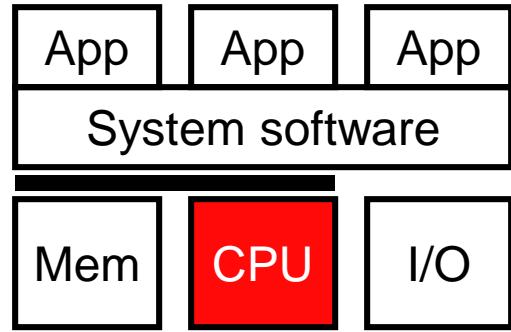
- **Perceptron predictor** [Jimenez]
  - Attacks predictor size problem using machine learning approach
  - History table replaced by table of function coefficients F_i (signed)
  - Predict taken if $\sum(BHR_i * F_i) >$ threshold
  - \+ Table size $\#PC * |BHR| * |F|$ (can use long BHR: ~60 bits)
    - – Equivalent correlated predictor would be $\#PC * 2^{|BHR|}$
  - How does it learn? Update F_i when branch is taken
    - BHR$_i$ == 1 ? F_i++ : F_i− −;
    - "don't care" F_i bits stay near 0, important F_i bits saturate
  - \+ Hybrid BHT/perceptron accuracy: 95–98%

Summary

App	App	App

System software

Mem	CPU	I/O

- **Single-cycle datapaths**
- **Latency vs. throughput & performance**
- **Basic pipelining**
- **Data hazards**
  - Bypassing
  - Load-use stalling
- **Pipelined multi-cycle operations**
- **Control hazards**
  - Branch prediction