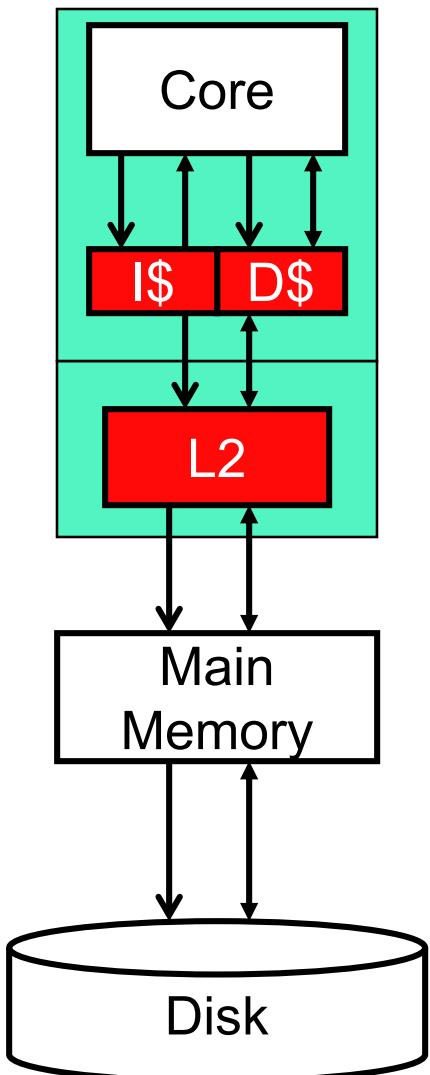


# Caches

---



- “Cache”: hardware managed
  - Hardware automatically retrieves missing data
  - Built from fast on-chip SRAM
  - In contrast to off-chip, DRAM “main memory”
- **Average access time** of a memory component
  - $latency_{avg} = latency_{hit} + (\%_{miss} * latency_{miss})$
  - Hard to get low  $latency_{hit}$  and  $\%_{miss}$  in one structure  
→ memory hierarchy
- Cache ABCs (**associativity, block size, capacity**)
- **Performance optimizations**
  - Prefetching & data restructuring
- **Handling writes**
  - Write-back vs. write-through
- **Memory hierarchy**
  - Smaller, faster, expensive → bigger, slower, cheaper

# Motivation

---

- Processor can compute only as fast as memory
  - A 3Ghz processor can execute an “add” operation in 0.33ns
  - Today’s “main memory” latency is more than 33ns
  - Naïve implementation:
    - loads/stores can be 100x slower than other operations
- Unobtainable goal:
  - Memory that operates at processor speeds
  - Memory as large as needed for all running programs
  - Memory that is cost effective
- Can’t achieve all of these goals at once
  - Example: latency of an SRAM is at least:  $\sqrt{\text{number of bits}}$

# **Memories (SRAM & DRAM)**

# Types of Memory

---

- **Static RAM (SRAM)**
  - 6 or 8 transistors per bit
    - Two inverters (4 transistors) + transistors for reading/writing
  - Optimized for speed (first) and density (second)
  - Fast (sub-nanosecond latencies for small SRAM)
    - Speed roughly proportional to its area ( $\sim \text{sqrt}(\text{number of bits})$ )
  - Mixes well with standard processor logic
- **Dynamic RAM (DRAM)**
  - 1 transistor + 1 capacitor per bit
  - Optimized for density (in terms of cost per bit)
  - Slow (>30ns internal access, ~50ns pin-to-pin)
  - Different fabrication steps (does not mix well with logic)
- Nonvolatile storage: Magnetic disk, Flash RAM

# Memory & Storage Technologies

---

- **Cost** - what can \$200 buy (2009)?
  - SRAM: 16MB
  - DRAM: 4,000MB (4GB) – 250x cheaper than SRAM
  - Flash: 64,000MB (64GB) – 16x cheaper than DRAM
  - Disk: 2,000,000MB (2TB) – 32x vs. Flash (512x vs. DRAM)
- **Latency**
  - SRAM: <1 to 2ns (on chip)
  - DRAM: ~50ns – 100x or more slower than SRAM
  - Flash: 75,000ns (75 microseconds) – 1500x vs. DRAM
  - Disk: 10,000,000ns (10ms) – 133x vs Flash (200,000x vs DRAM)
- **Bandwidth**
  - SRAM: 300GB/sec (e.g., 12-port 8-byte register file @ 3Ghz)
  - DRAM: ~25GB/s
  - Flash: 0.25GB/s (250MB/s), 100x less than DRAM
  - Disk: 0.1 GB/s (100MB/s), 250x vs DRAM, **sequential** access only

# The Memory Hierarchy

# Known From the Beginning

---

“Ideally, one would desire an infinitely large memory capacity such that any particular word would be immediately available ... We are forced to recognize the possibility of constructing a hierarchy of memories, each of which has a greater capacity than the preceding but which is less quickly accessible.”

Burks, Goldstine, VonNeumann

“Preliminary discussion of the logical design of an electronic computing instrument”

IAS memo 1946

# Big Observation: Locality & Caching

---

- **Locality of memory references**
  - Empirical property of real-world programs, few exceptions
- **Temporal locality**
  - Recently referenced data is likely to be referenced again soon
  - **Reactive**: “cache” recently used data in small, fast memory
- **Spatial locality**
  - More likely to reference data near recently referenced data
  - **Proactive**: “cache” large chunks of data to include nearby data
- Both properties hold for both data and instructions
- Cache: “Hash table” of recently used blocks of data
  - In hardware, finite-sized, transparent to software

# Spatial and Temporal Locality Example

---

- Which memory accesses demonstrate spatial locality?
- Which memory accesses demonstrate temporal locality?

```
int sum = 0;  
int x[1000];  
  
for(int c = 0; c < 1000; c++) {  
    sum += c;  
  
    x[c] = 0;  
}
```

# Library Analogy

---

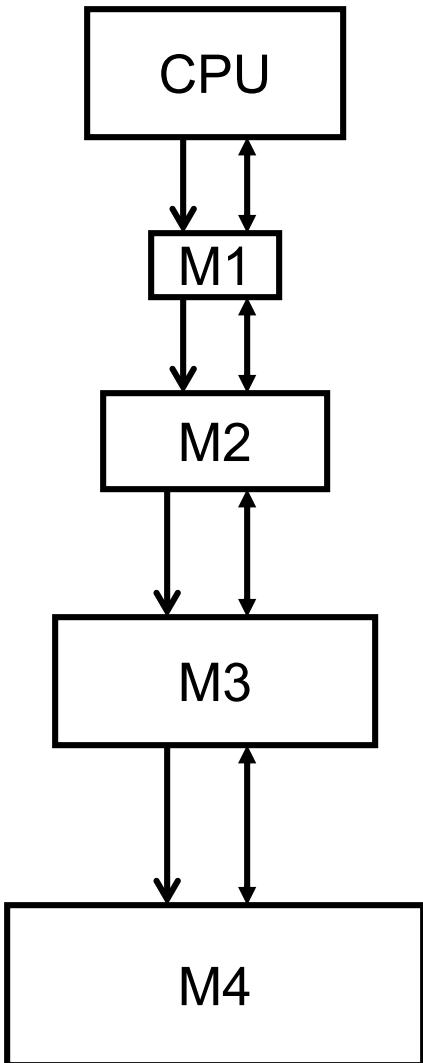
- Consider books in a library
- Library has lots of books, but it is slow to access
  - Far away (time to walk to the library)
  - Big (time to walk within the library)
- How can you avoid these latencies?
  - Check out books, take them home with you
    - Put them on desk, on bookshelf, etc.
  - But desks & bookshelves have limited capacity
    - Keep recently used books around (**temporal locality**)
    - Grab books on related topic at the same time (**spatial locality**)
    - Guess what books you'll need in the future (prefetching)

# Library Analogy Explained

---

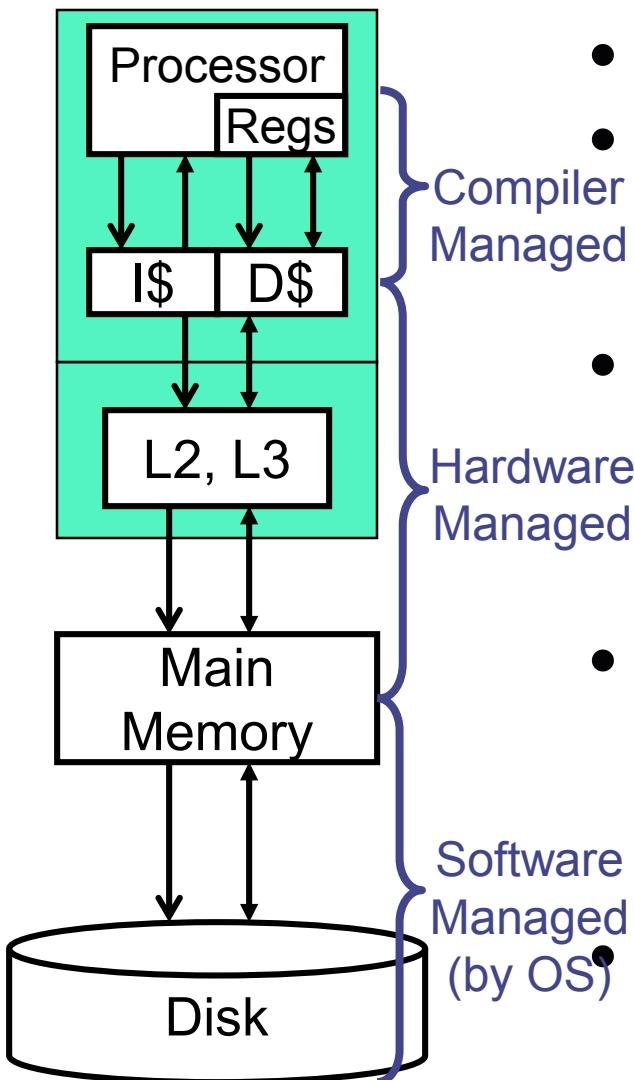
- Registers  $\leftrightarrow$  books on your desk
  - Actively being used, small capacity
- Caches  $\leftrightarrow$  bookshelves
  - Moderate capacity, pretty fast to access
- Main memory  $\leftrightarrow$  library
  - Big, holds almost all data, but slow
- Disk (virtual memory)  $\leftrightarrow$  inter-library loan
  - Very slow, but hopefully really uncommon

# Exploiting Locality: Memory Hierarchy



- Hierarchy of memory components
  - Upper components
    - Fast  $\leftrightarrow$  Small  $\leftrightarrow$  Expensive
  - Lower components
    - Slow  $\leftrightarrow$  Big  $\leftrightarrow$  Cheap
- Connected by “buses”
  - Which also have latency and bandwidth issues
- Most frequently accessed data in M1
  - M1 + next most frequently accessed in M2, etc.
  - Move data up-down hierarchy
- Optimize average access time
  - $latency_{avg} = latency_{hit} + (%_{miss} * latency_{miss})$
  - Attack each component

# Concrete Memory Hierarchy



- 0th level: **Registers**
- 1st level: **Primary caches**
  - Split instruction (I\$) and data (D\$)
  - Typically 8KB to 64KB each
- 2nd level: **2<sup>nd</sup> and 3<sup>rd</sup> cache (L2, L3)**
  - On-chip, typically made of SRAM
  - 2<sup>nd</sup> level typically ~256KB to 512KB
  - “Last level cache” typically 4MB to 16MB
- 3rd level: **main memory**
  - Made of DRAM (“Dynamic” RAM)
  - Typically 2-16GB for desktops/laptops
    - Servers can have >1 TB
- 4th level: **disk (swap and files)**
  - Uses magnetic disks or flash drives

# Caches

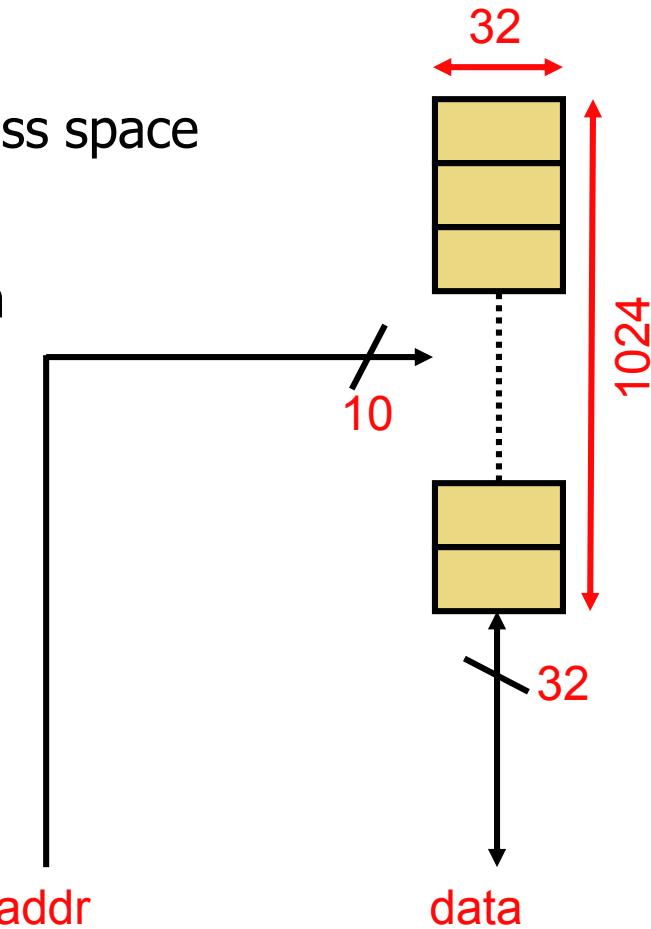
# Analogy to a Software Hashtable

---

- What is a “hash table”?
  - What is it used for?
  - How does it work?
- Short answer:
  - Maps a “key” to a “value”
    - Constant time lookup/insert
  - Have a table of some size, say  $N$ , of “buckets”
  - Take a “key” value, apply a hash function to it
  - Insert and lookup a “key” at “ $\text{hash(key)} \bmod N$ ”
    - Need to store the “key” and “value” in each bucket
    - Need to check to make sure the “key” matches
  - Need to handle conflicts/overflows somehow (chaining, re-hashing)

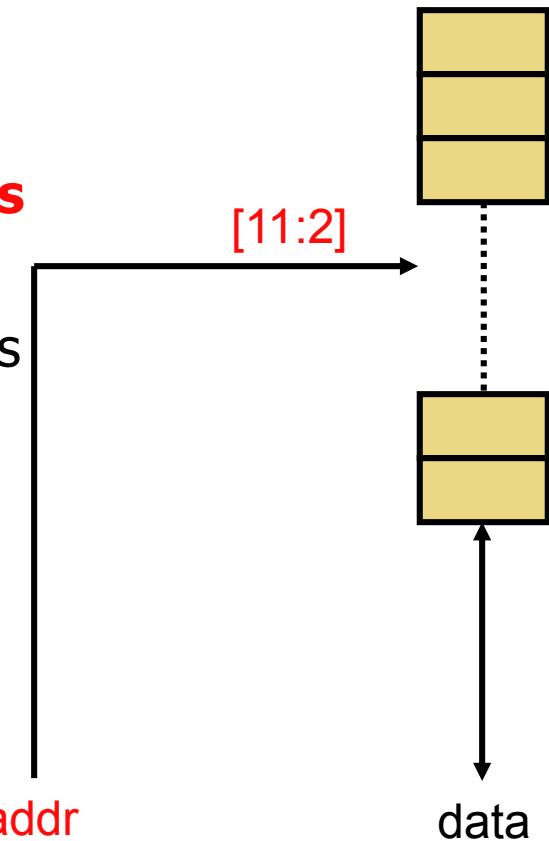
# Hardware Cache Organization

- **Cache is a hardware hashtable**
  - with one exception: data not always present
- The setup
  - 32-bit ISA → 4B addresses,  $2^{32}$  B address space
- Logical cache organization
  - 4KB, organized as 1K **blocks** of 4B each
  - Each block can hold a 4B word
- Physical cache implementation
  - 1K (1024 bit) by 4B **SRAM**
  - Called **data array**
  - 10-bit address input
  - 32-bit data input/output



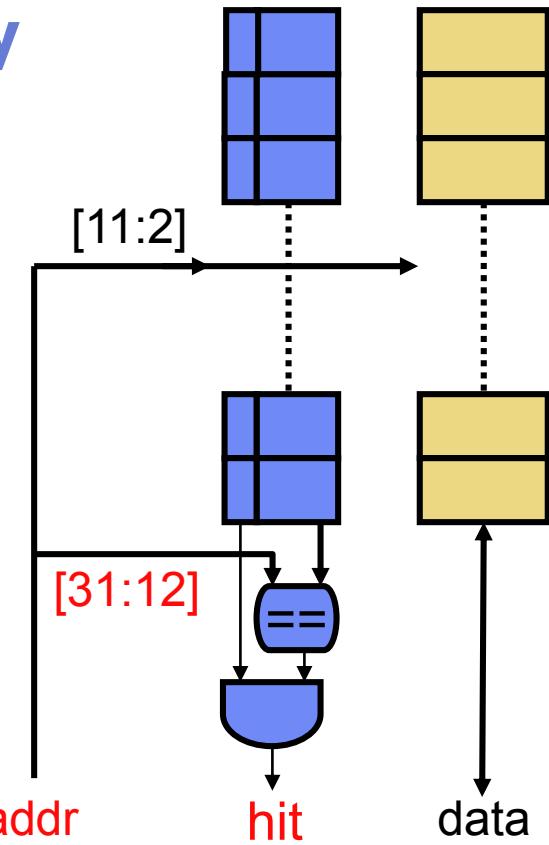
# Looking Up A Block

- Q: which 10 of the 32 address bits to use?
- A: bits [11:2]
  - 2 least significant (LS) bits [1:0] are the **offset bits**
    - Locate byte within word
    - Don't need these to locate word
  - Next 10 LS bits [11:2] are the **index bits**
    - These locate the word
    - Nothing says index must be these bits
    - But these work best in practice
      - Why? (think about it)



# Knowing that You Found It

- $2^{20}$  blocks map to each cache row (**set**)
  - How to know which if any is currently there?
  - Tag each cache word with remaining address bits [31:12]
- Build separate and parallel **tag array**
  - 1K by 21-bit SRAM
  - 20-bit (next slide) **tag** + 1 **valid bit**
- Lookup algorithm
  - Read tag indicated by index bits
  - If tag matches & valid bit set:
    - then: Hit  $\rightarrow$  data is good
    - else: Miss  $\rightarrow$  data is garbage, wait...

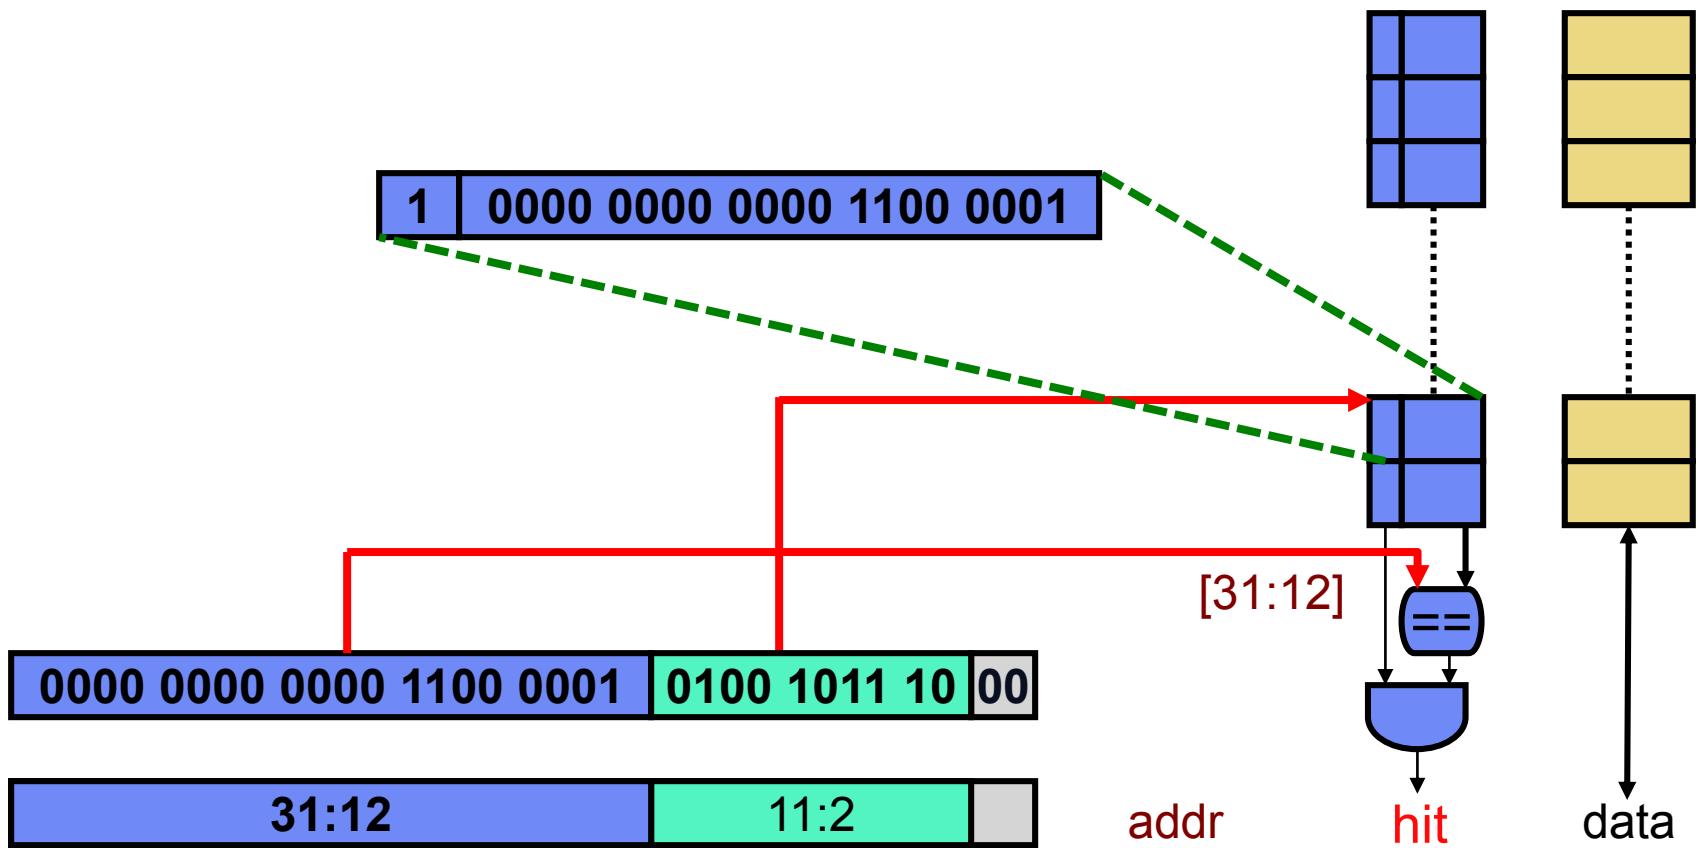


31:12

11:2

# A Concrete Example

- Lookup address x000C14B8
  - Index = addr [11:2] = (addr >> 2) & x3FF = x12E
  - Tag = addr [31:12] = (addr >> 12) = x000C1



# Calculating Tag Overhead

---

- “32KB cache” means cache holds 32KB of data
  - Called **capacity**
  - Tag storage is considered overhead
- Tag overhead of 32KB cache with 1024 32B **frames**
  - frame is physical structure holding a block, frame size==block size
  - 32B frame → 5-bit offset
  - 1024 frames → 10-bit index
  - 32-bit address – 5-bit offset – 10-bit index = 17-bit tag
  - (17-bit tag + 1-bit valid) \* 1024 frames = 18Kb tags = 2.2KB tags
  - ~6% overhead
- What about 64-bit addresses?
  - Tag increases to 49 bits, ~20% overhead (worst case)

# Handling a Cache Miss

---

- What if requested data isn't in the cache?
  - How does it get in there?
- **Cache controller**: finite state machine
  - Remembers miss address
  - Accesses next level of memory
  - Waits for response
  - Writes data/tag into proper locations
- All of this happens on the **fill path**
- Sometimes called **backside**

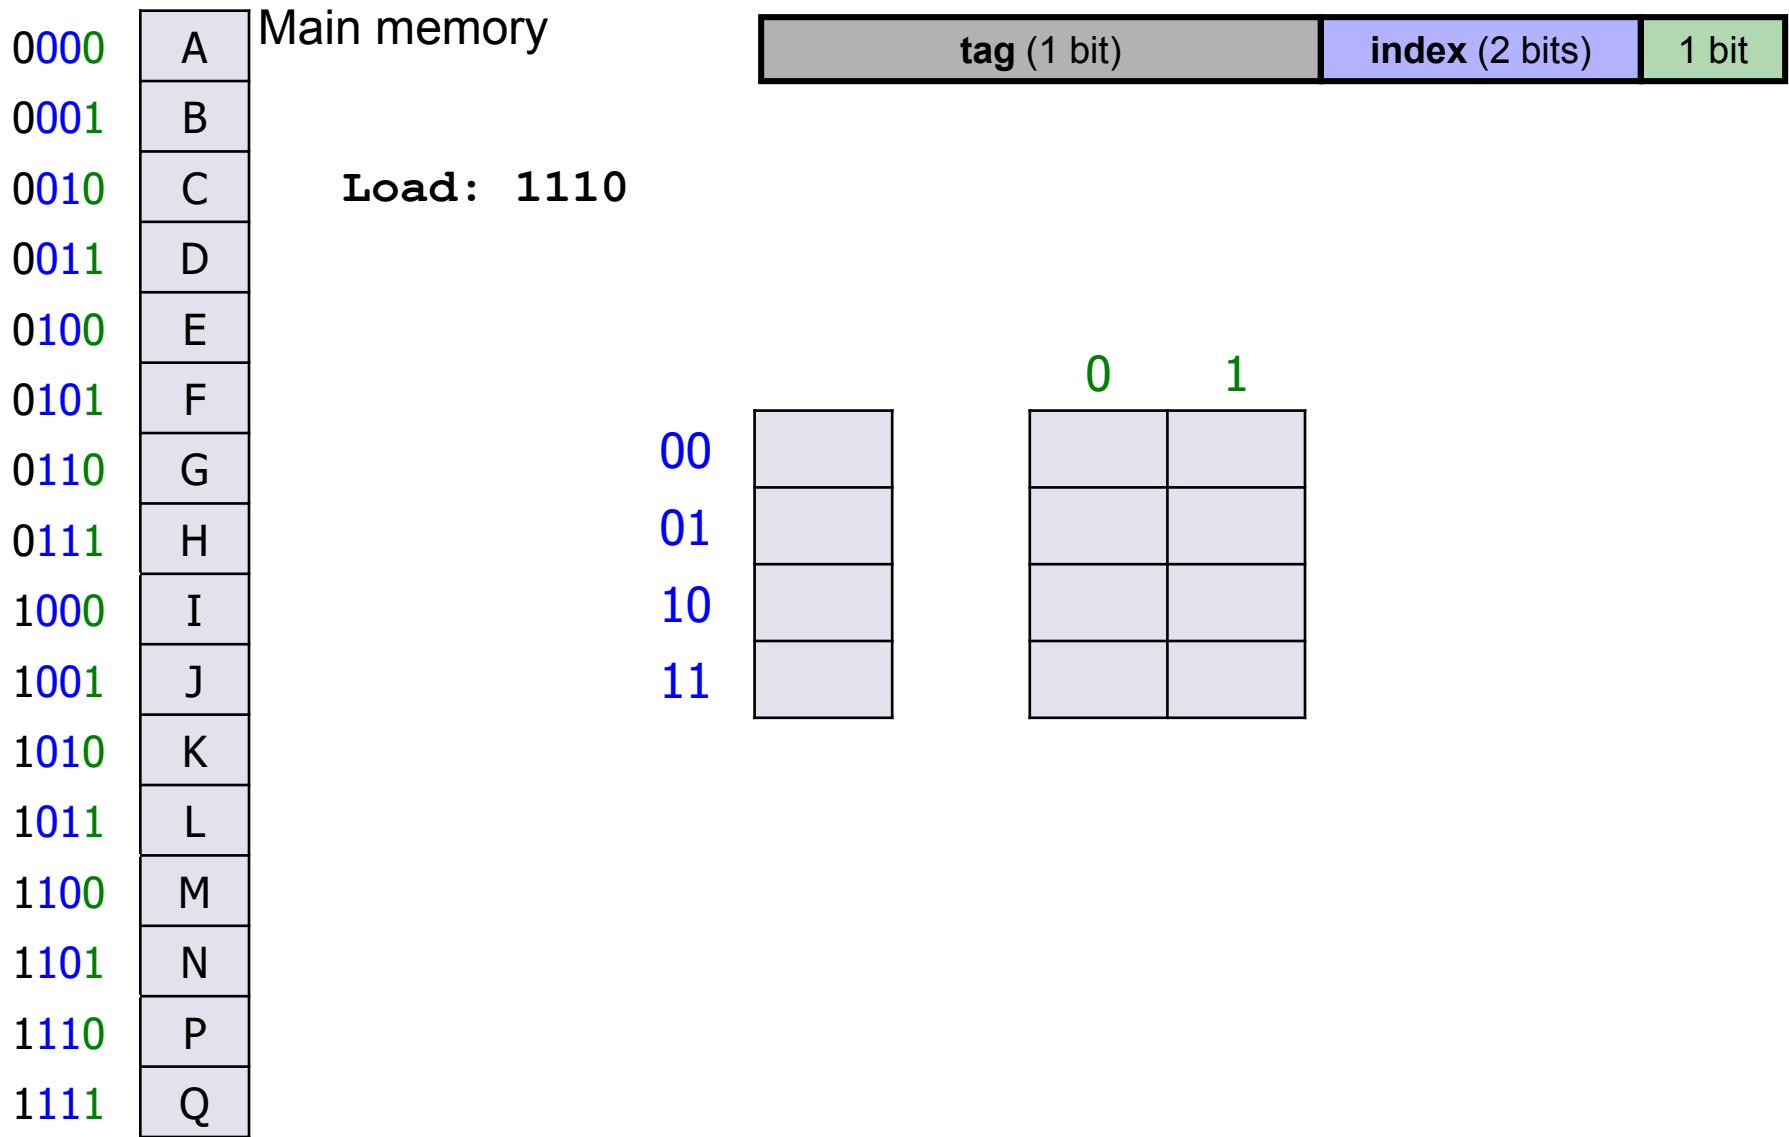
# Cache Examples

---

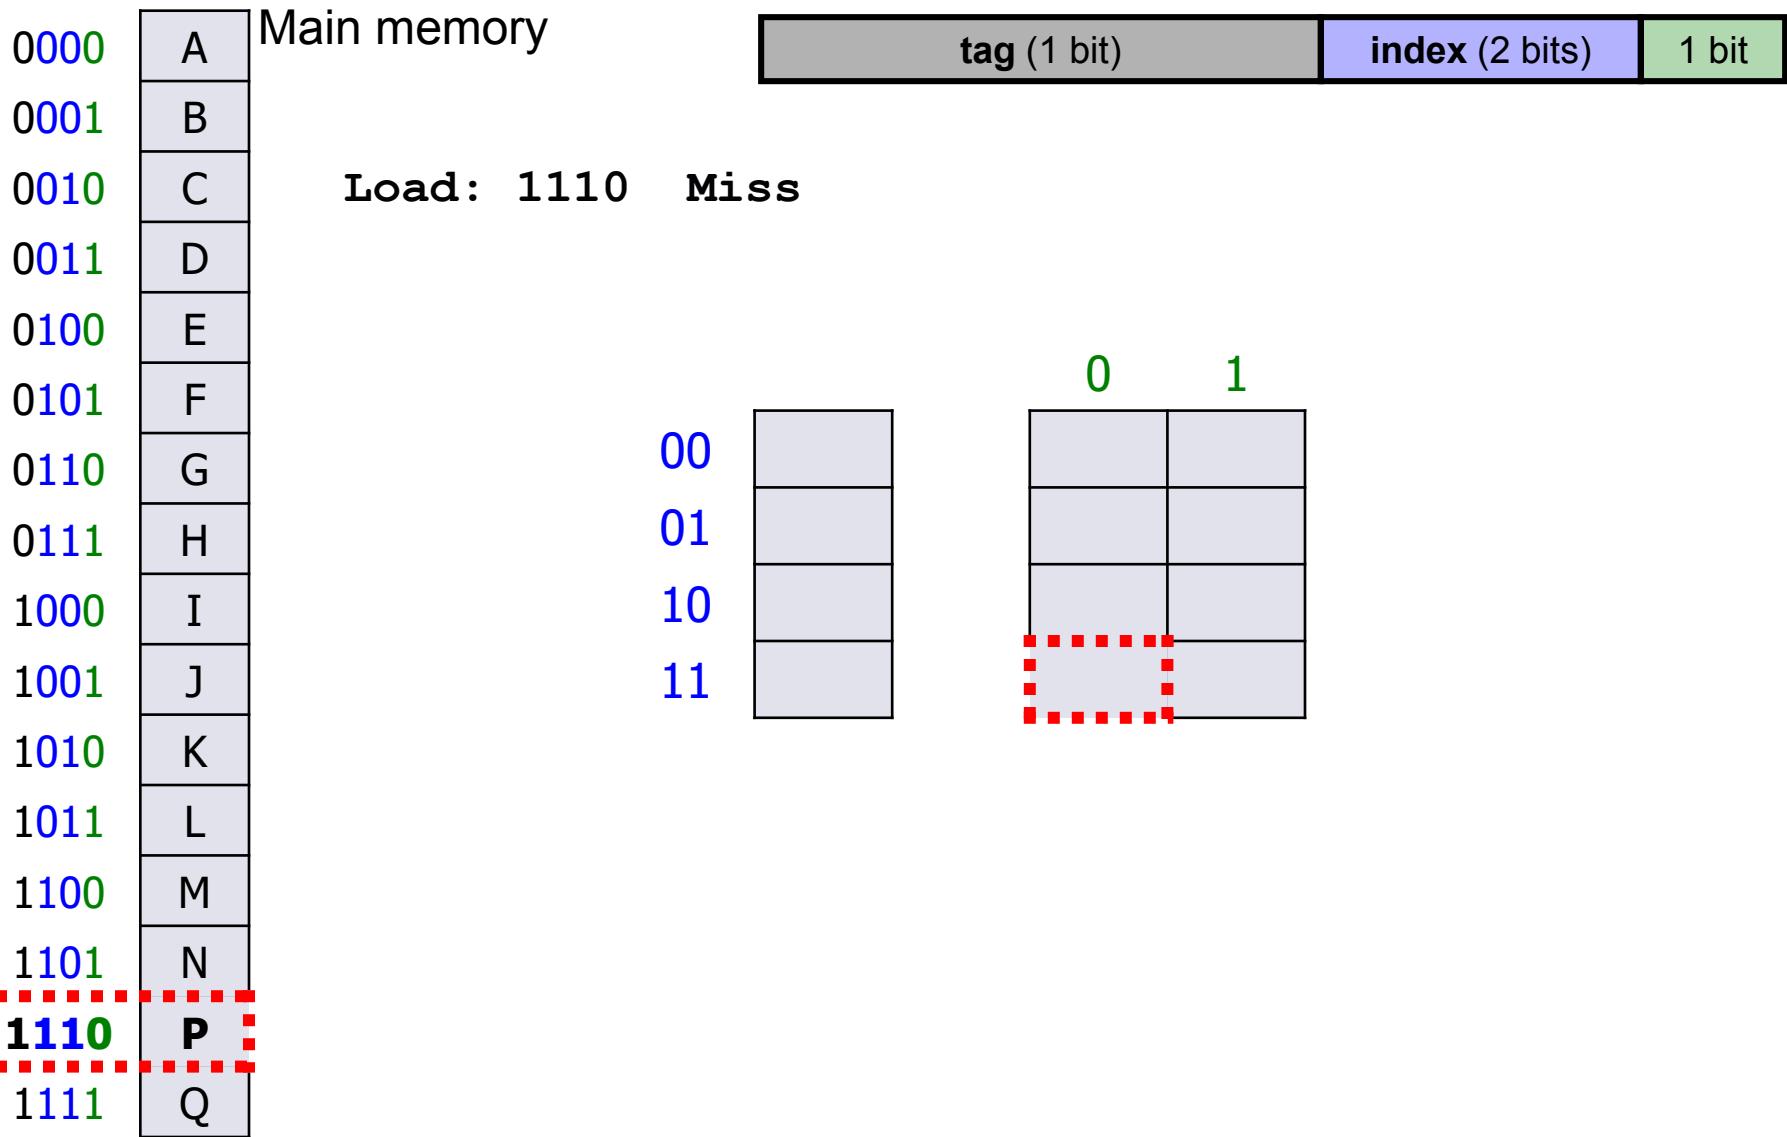
- 4-bit addresses → 16B memory
  - Simpler cache diagrams than 32-bits
- 8B cache, 2B blocks
  - Figure out number of rows/sets
    - 4 (capacity / block-size)
  - Figure out how address splits into offset/index/tag bits
    - Offset: least-significant  $\log_2(\text{block-size}) = \log_2(2) = 1 \rightarrow 000\text{0}$
    - Index: next  $\log_2(\text{number-of-sets}) = \log_2(4) = 2 \rightarrow 0\text{00}0$
    - Tag: rest =  $4 - 1 - 2 = 1 \rightarrow \text{0000}$



# 4-bit Address, 8B Cache, 2B Blocks



# 4-bit Address, 8B Cache, 2B Blocks



# 4-bit Address, 8B Cache, 2B Blocks

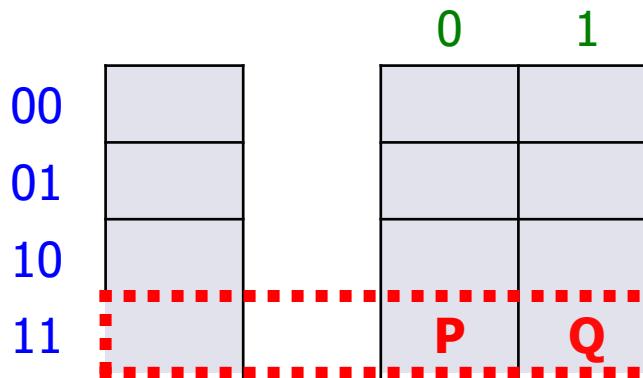
Main memory

0000	A
0001	B
0010	C
0011	D
0100	E
0101	F
0110	G
0111	H
1000	I
1001	J
1010	K
1011	L
1100	M
1101	N
1110	P
1111	Q

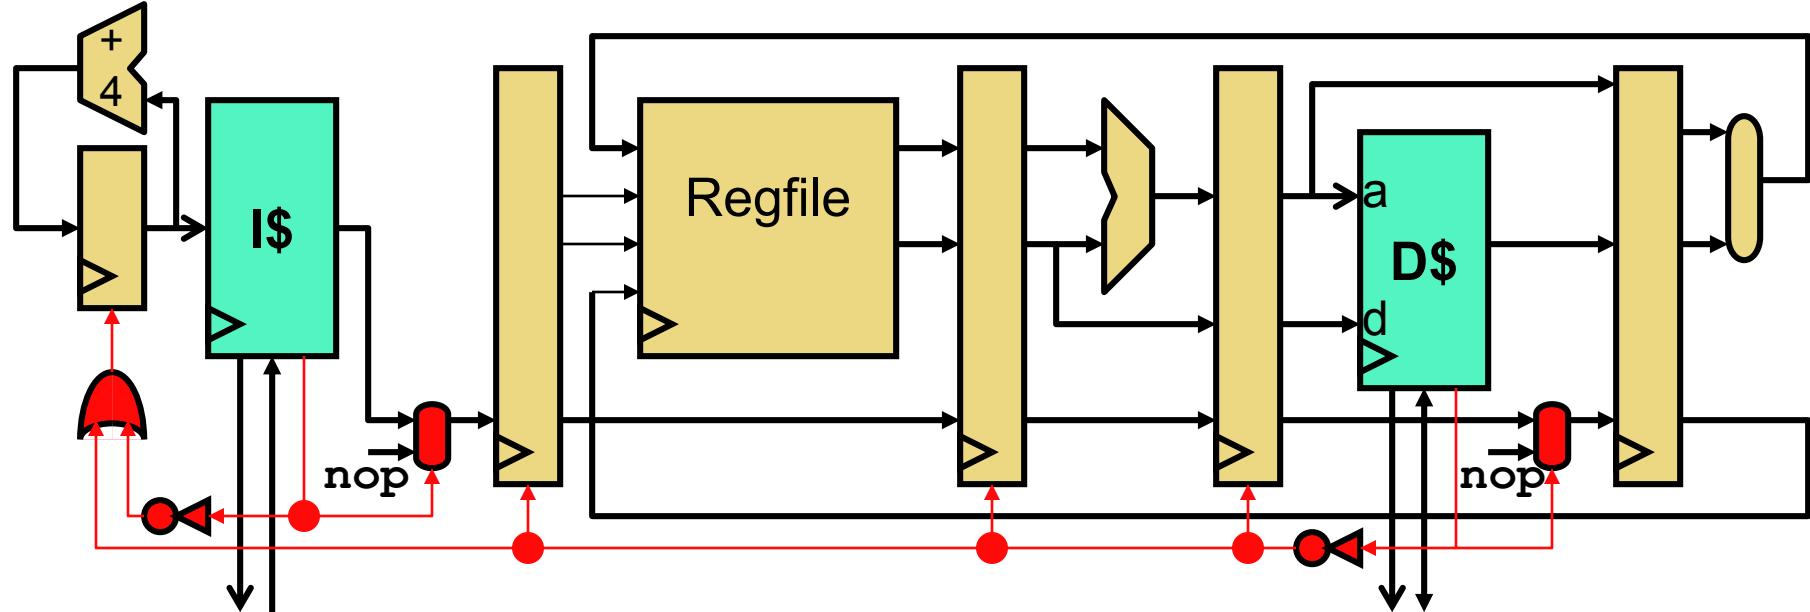
Load: 1110



Load: 1110 Miss

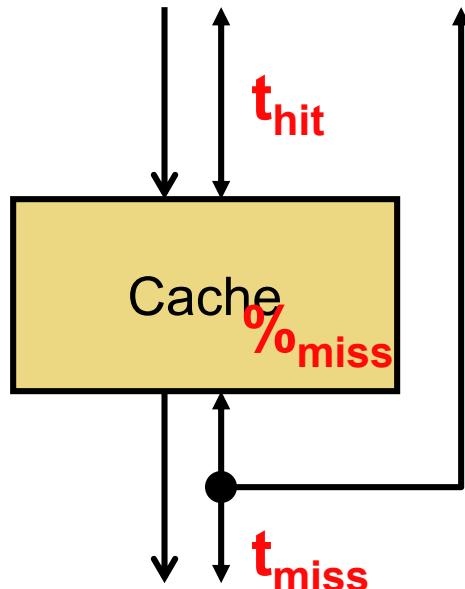


# Cache Misses and Pipeline Stalls



- I\$ and D\$ misses stall pipeline just like data hazards
  - Stall logic driven by miss signal
    - Cache “logically” re-evaluates hit/miss every cycle
    - Block is filled → miss signal de-asserts → pipeline restarts

# Cache Performance Equation



- For a cache
  - **Access**: read or write to cache
  - **Hit**: desired data found in cache
  - **Miss**: desired data not found in cache
    - Must get from another component
    - No notion of “miss” in register file
  - **Fill**: action of placing data into cache
- **%miss** (miss-rate): #misses / #accesses
- **$t_{hit}$** : time to read data from (write data to) cache
- **$t_{miss}$** : time to read data into cache
- Performance metric: average access time

$$t_{avg} = t_{hit} + (%miss * t_{miss})$$

# CPI Calculation with Cache Misses

---

- Parameters
  - Simple pipeline with base CPI of 1
  - Instruction mix: 30% loads/stores
  - I\$:  $\%_{\text{miss}} = 2\%$ ,  $t_{\text{miss}} = 10$  cycles
  - D\$:  $\%_{\text{miss}} = 10\%$ ,  $t_{\text{miss}} = 10$  cycles
- What is new CPI?
  - $\text{CPI}_{I\$} = \%_{\text{missI\$}} * t_{\text{miss}} = 0.02 * 10 \text{ cycles} = 0.2 \text{ cycle}$
  - $\text{CPI}_{D\$} = \%_{\text{load/store}} * \%_{\text{missD\$}} * t_{\text{missD\$}} = 0.3 * 0.1 * 10 \text{ cycles} = 0.3 \text{ cycle}$
  - $\text{CPI}_{\text{new}} = \text{CPI} + \text{CPI}_{I\$} + \text{CPI}_{D\$} = 1 + 0.2 + 0.3 = 1.5$

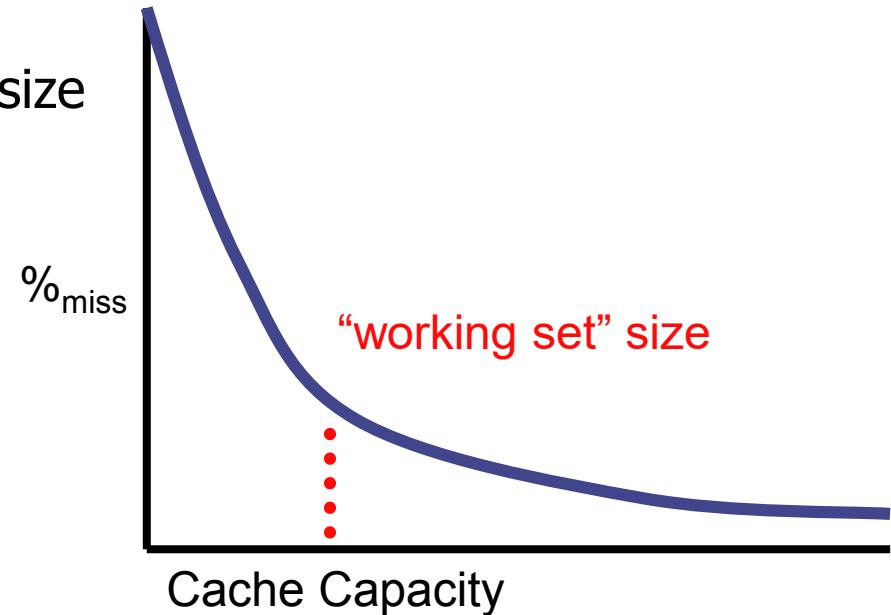
# Measuring Cache Performance

---

- Ultimate metric is  $t_{avg}$ 
  - Cache capacity and circuits roughly determines  $t_{hit}$
  - Lower-level memory structures determine  $t_{miss}$
  - Measure  $\%_{miss}$ 
    - Hardware performance counters
    - Simulation

# Capacity and Performance

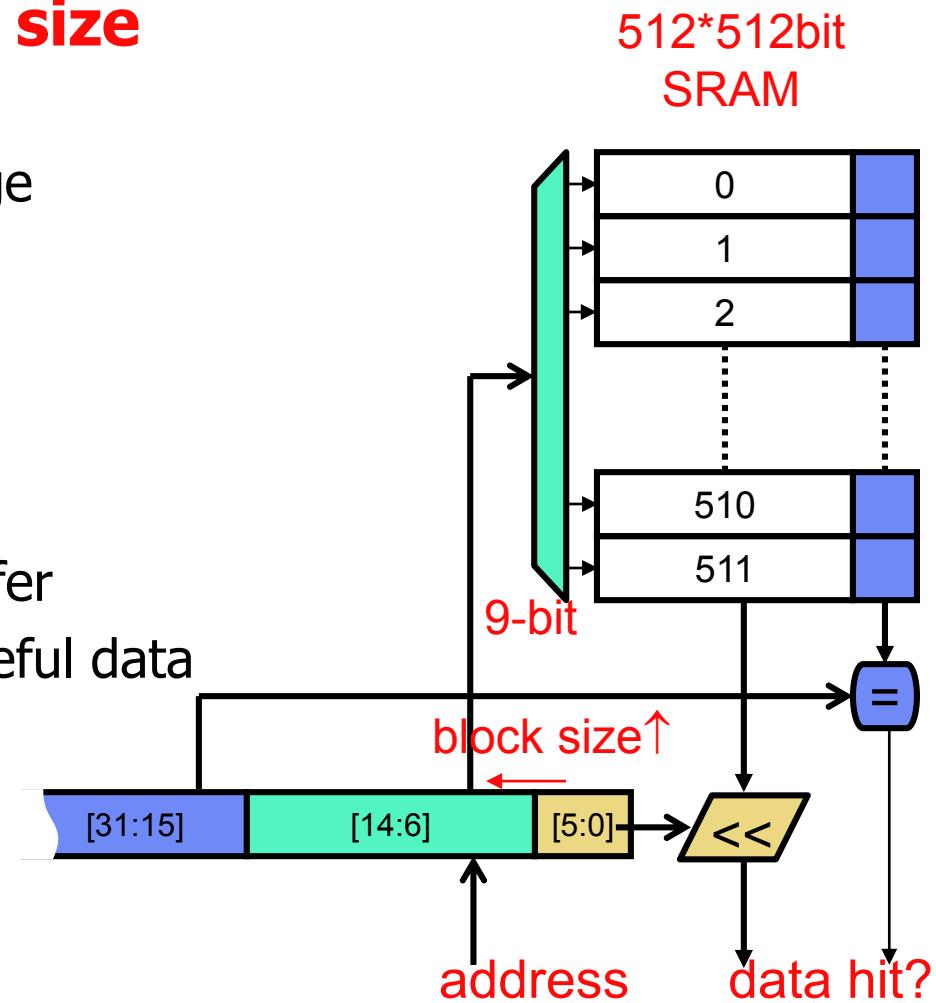
- Simplest way to reduce  $\%_{\text{miss}}$ : increase capacity
  - + Miss rate decreases monotonically
    - **"Working set"**: instructions/data program is actively using
    - Diminishing returns
  - However  $t_{\text{hit}}$  increases
    - Latency grows with cache size
  - what happens to  $t_{\text{avg}}$  ?



- For fixed capacity, reduce  $\%_{\text{miss}}$  by changing **organization**

# Block Size

- Given capacity, manipulate  $\%_{\text{miss}}$  by changing organization
- One option: increase **block size**
  - Exploit **spatial locality**
  - Notice index/offset bits change
  - Tag remains the same
- Ramifications
  - + Reduce  $\%_{\text{miss}}$  (up to a point)
  - + Reduce tag overhead (why?)
  - Potentially useless data transfer
  - Premature replacement of useful data

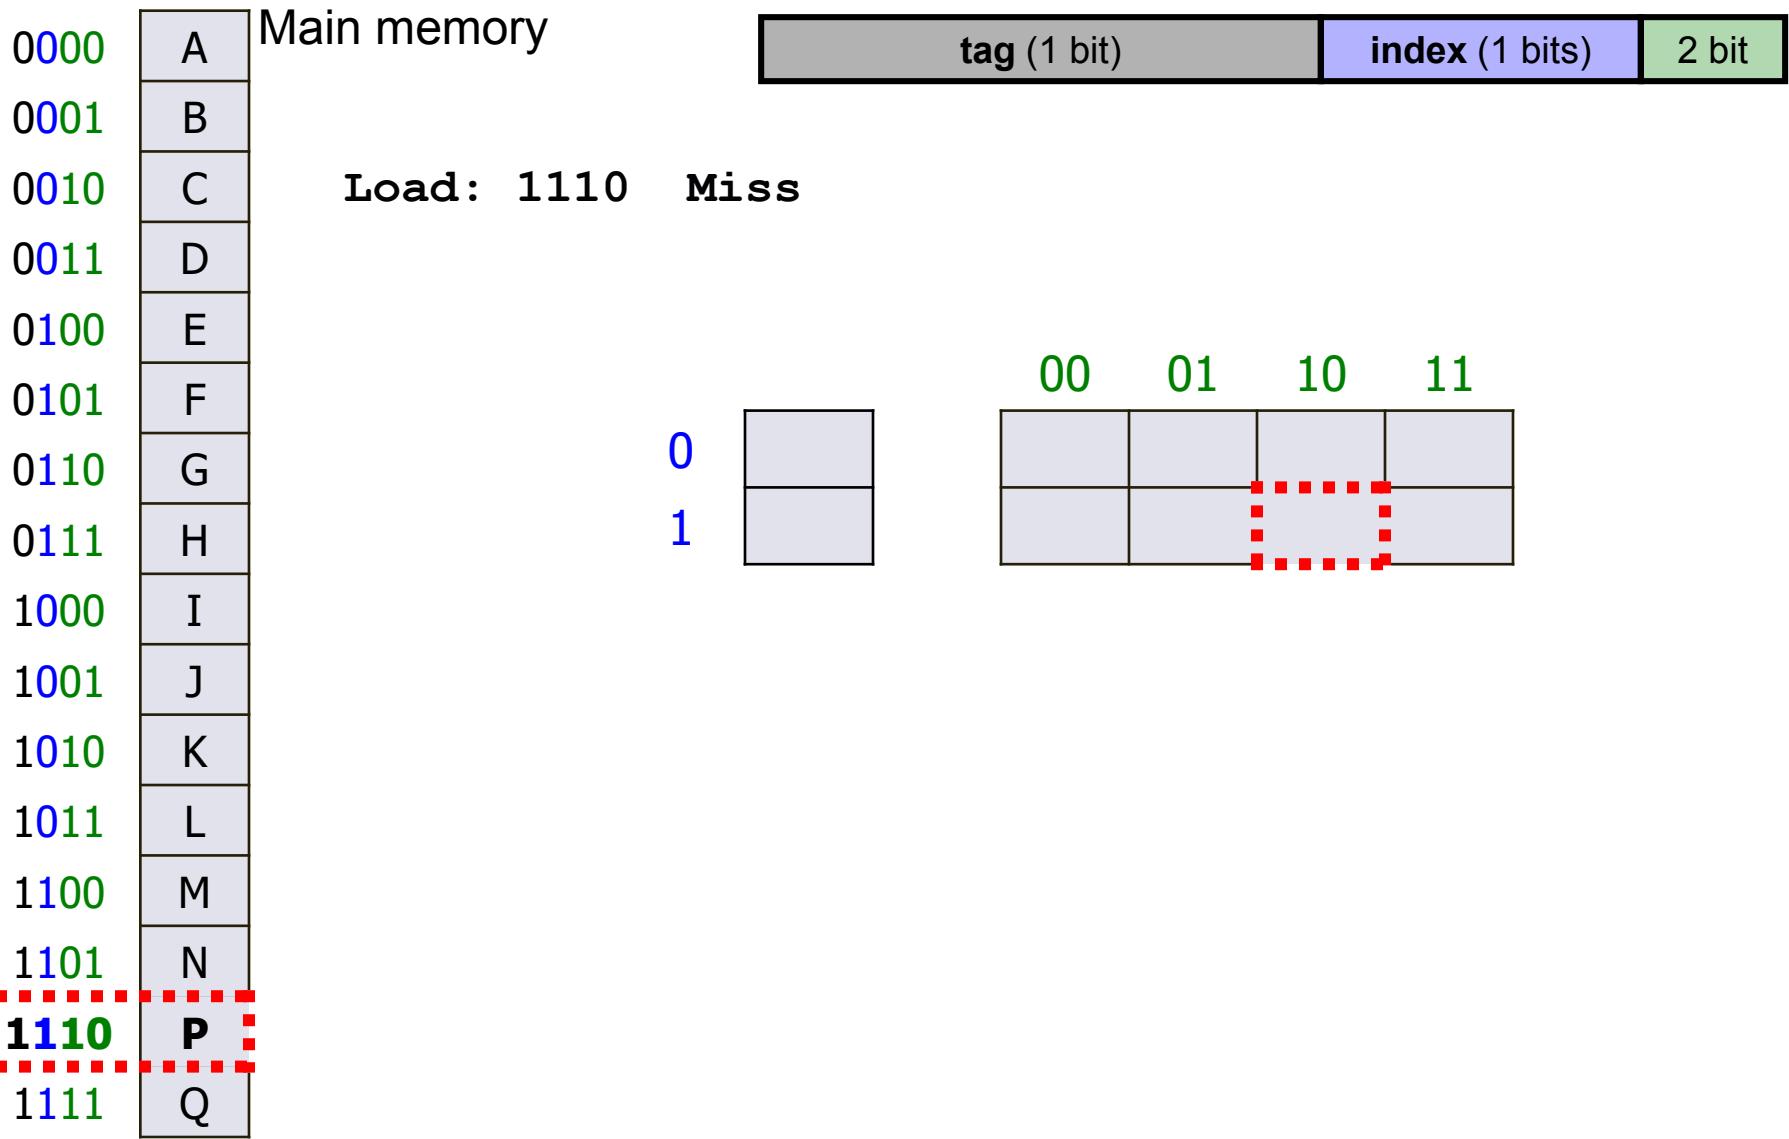


# Larger Blocks to Lower Tag Overhead

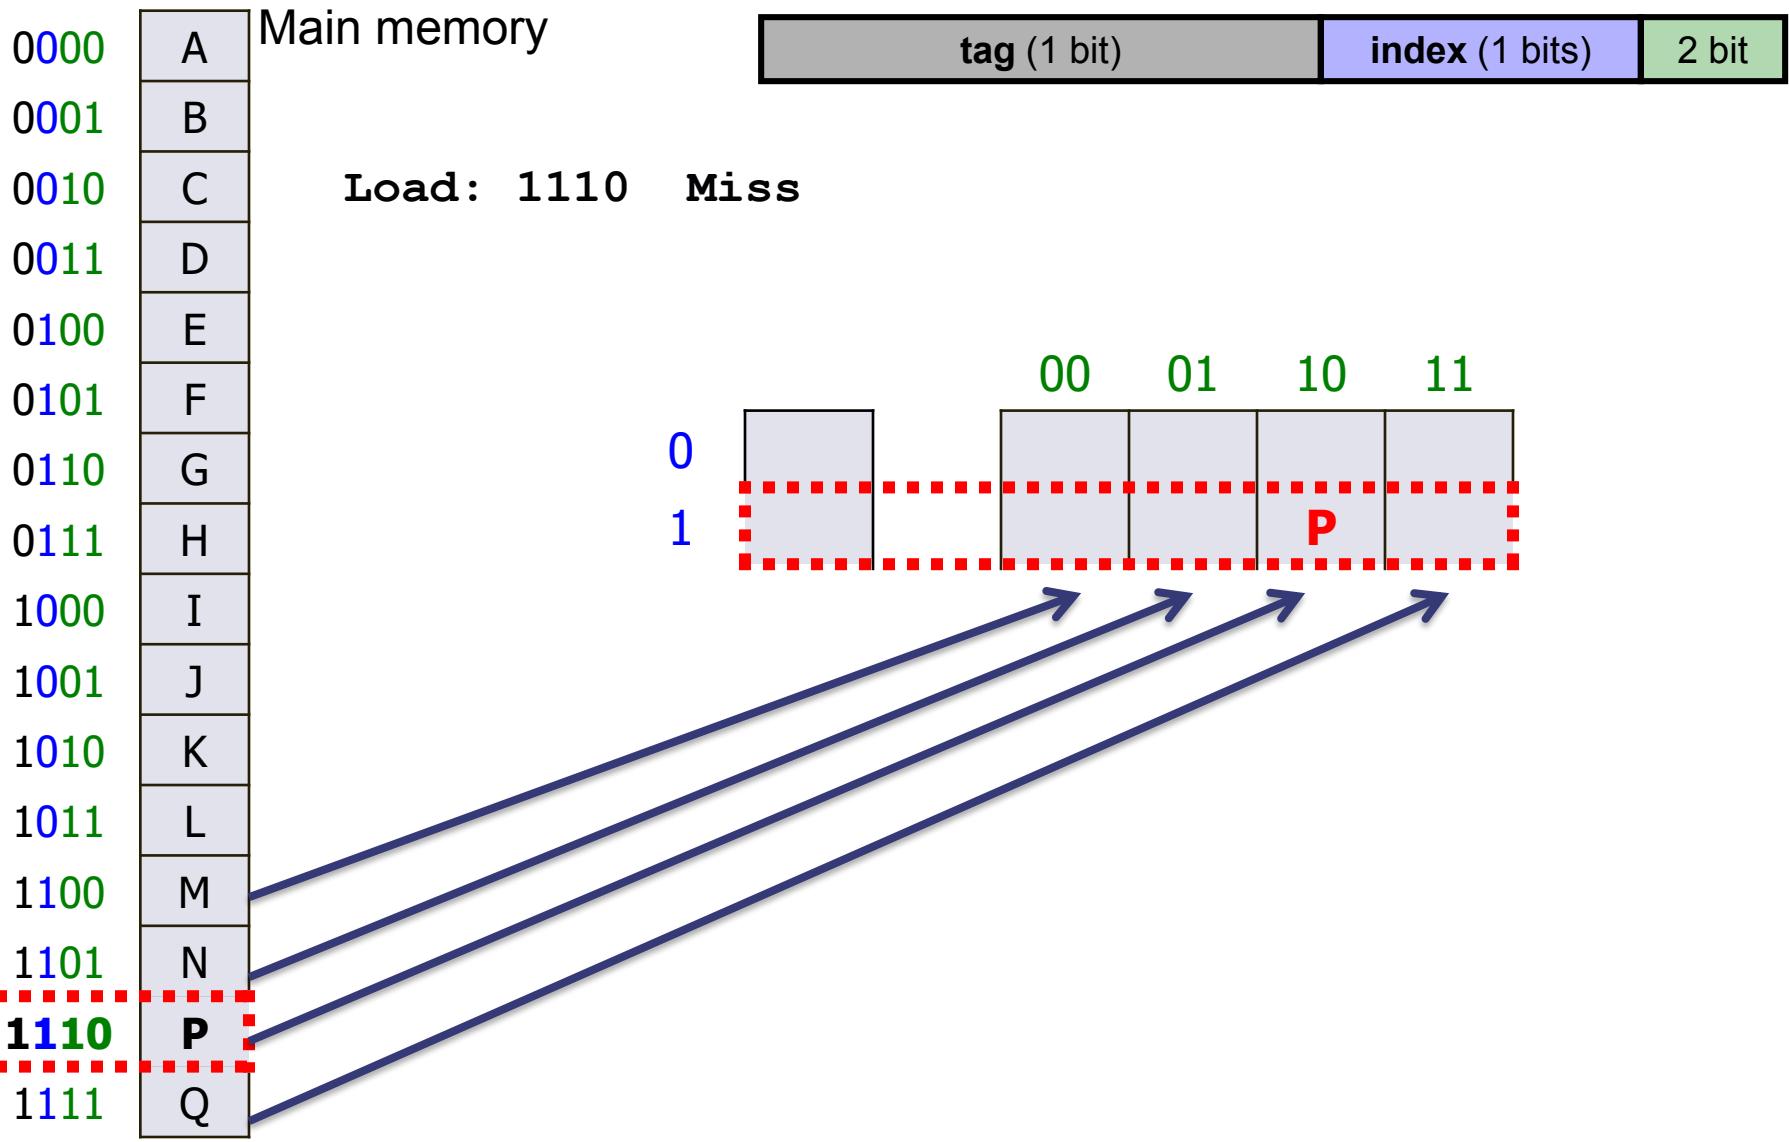
---

- Tag overhead of 32KB cache with 1024 32B frames
  - 32B frames → 5-bit offset
  - 1024 frames → 10-bit index
  - 32-bit address – 5-bit offset – 10-bit index = 17-bit tag
  - (17-bit tag + 1-bit valid) \* 1024 frames = 18Kb tags = 2.2KB tags
  - ~6% overhead
- Tag overhead of 32KB cache with 512 64B frames
  - 64B frames → 6-bit offset
  - 512 frames → 9-bit index
  - 32-bit address – 6-bit offset – 9-bit index = 17-bit tag
  - (17-bit tag + 1-bit valid) \* 512 frames = 9Kb tags = 1.1KB tags
  - + ~3% overhead

# 4-bit Address, 8B Cache, 4B Blocks

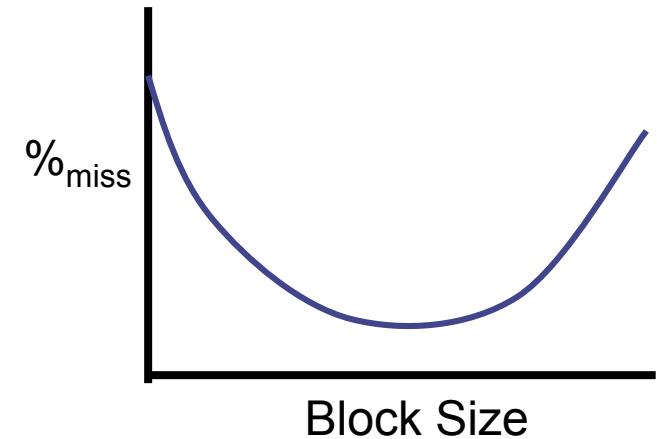


# 4-bit Address, 8B Cache, 4B Blocks



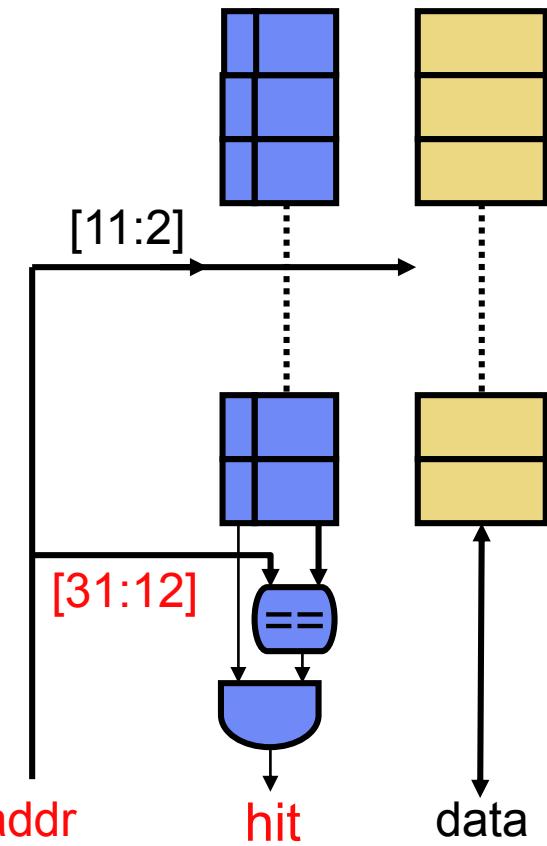
# Effect of Block Size on Miss Rate

- Two effects on miss rate
  - + **Spatial prefetching (good)**
    - For adjacent locations
    - Turns miss/miss into miss/hit pairs
  - **Interference (bad)**
    - For non-adjacent locations that map to adjacent frames
    - Turns hits into misses by disallowing simultaneous residence
    - Consider entire cache as one big block
- Both effects always present
  - Spatial “prefetching” dominates initially
    - Depends on size of the cache
  - Reasonable block sizes are 32B–128B
- But also increases traffic
  - More data moved, not all used



# Cache Conflicts

- Consider two frequently-accessed variables...
- What if their addresses have the same “index” bits?
  - Such addresses “conflict” in the cache
  - Can’t hold both in the cache at once...
  - Can result in lots of misses (bad!)
- Conflicts increase cache miss rate
  - Worse, result in non-robust performance
  - Small program change -> changes memory layout -> changes cache mapping of variables -> dramatically increase/decrease conflicts
- How can we mitigate conflicts?



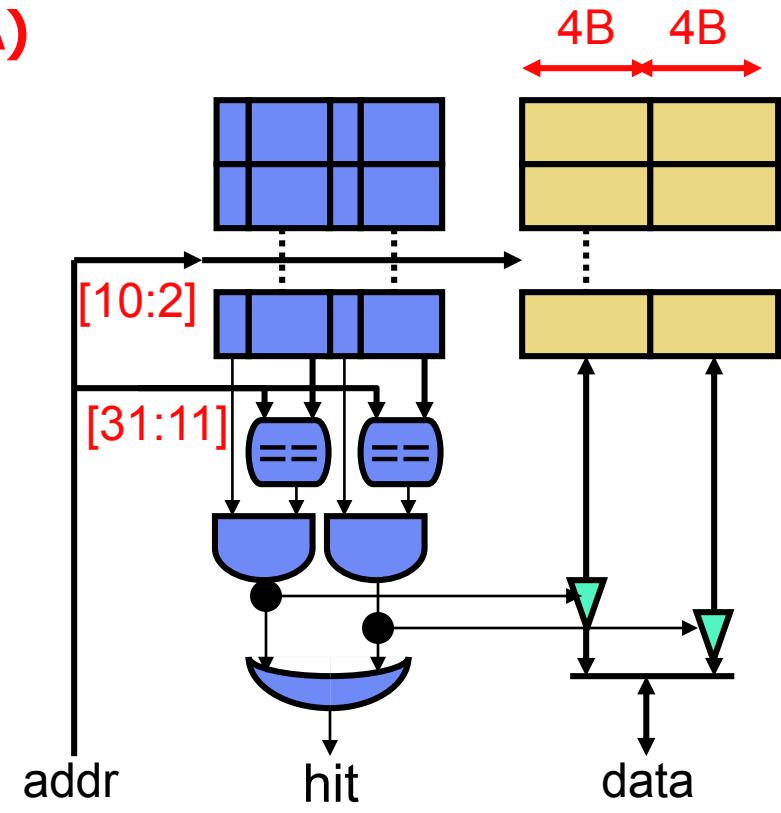
# Associativity

- **Set-associativity**

- Block can reside in one of few frames
- Frame groups called *sets*
- Each frame in a set called a *way*
- This is **2-way set-associative (SA)**
- 1-way → **direct-mapped (DM)**
- 1-set → **fully-associative (FA)**

- + Reduces conflicts
- Increases latency<sub>hit</sub>:
  - additional tag match & muxing

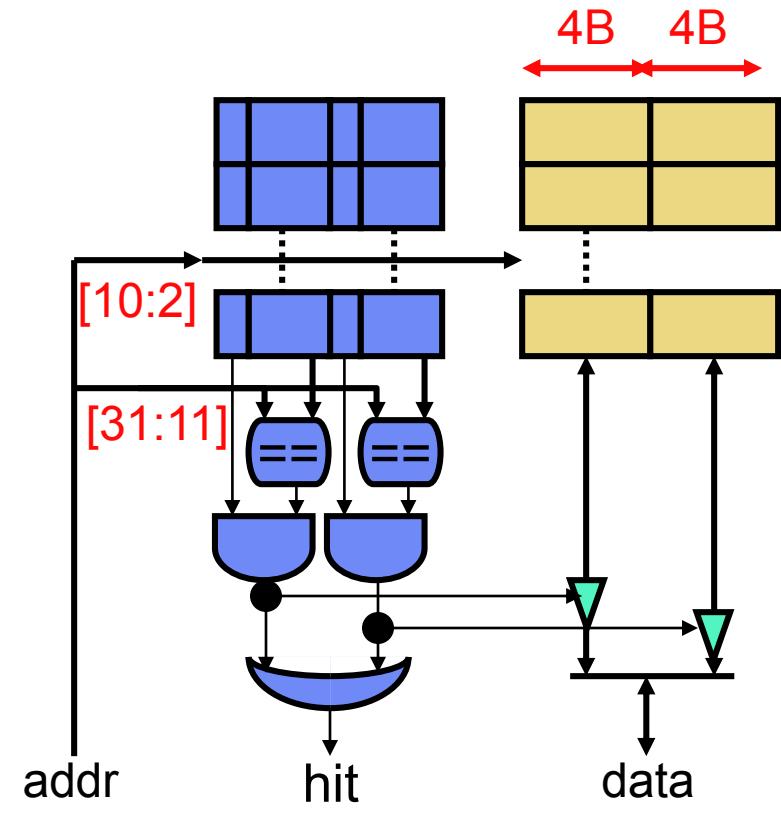
more associativity  
→



# Associativity

- Lookup algorithm
  - Use index bits to find set
  - Read data/tags in all frames in parallel
  - **Any** (match and valid bit), Hit
- Notice tag/index/offset bits
  - Only 9-bit index (versus 10-bit for direct mapped)

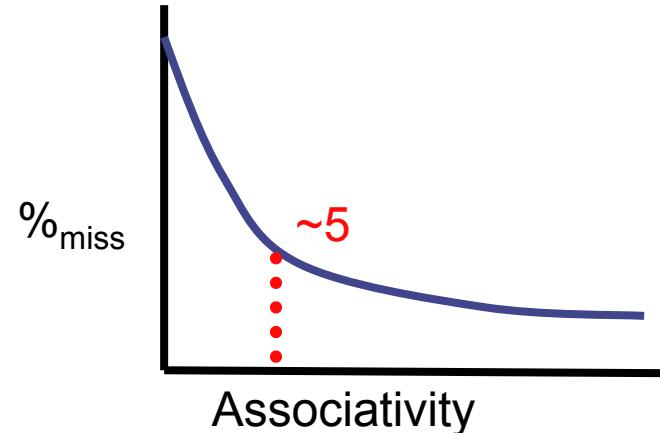
more associativity  
→



# Associativity and Performance

---

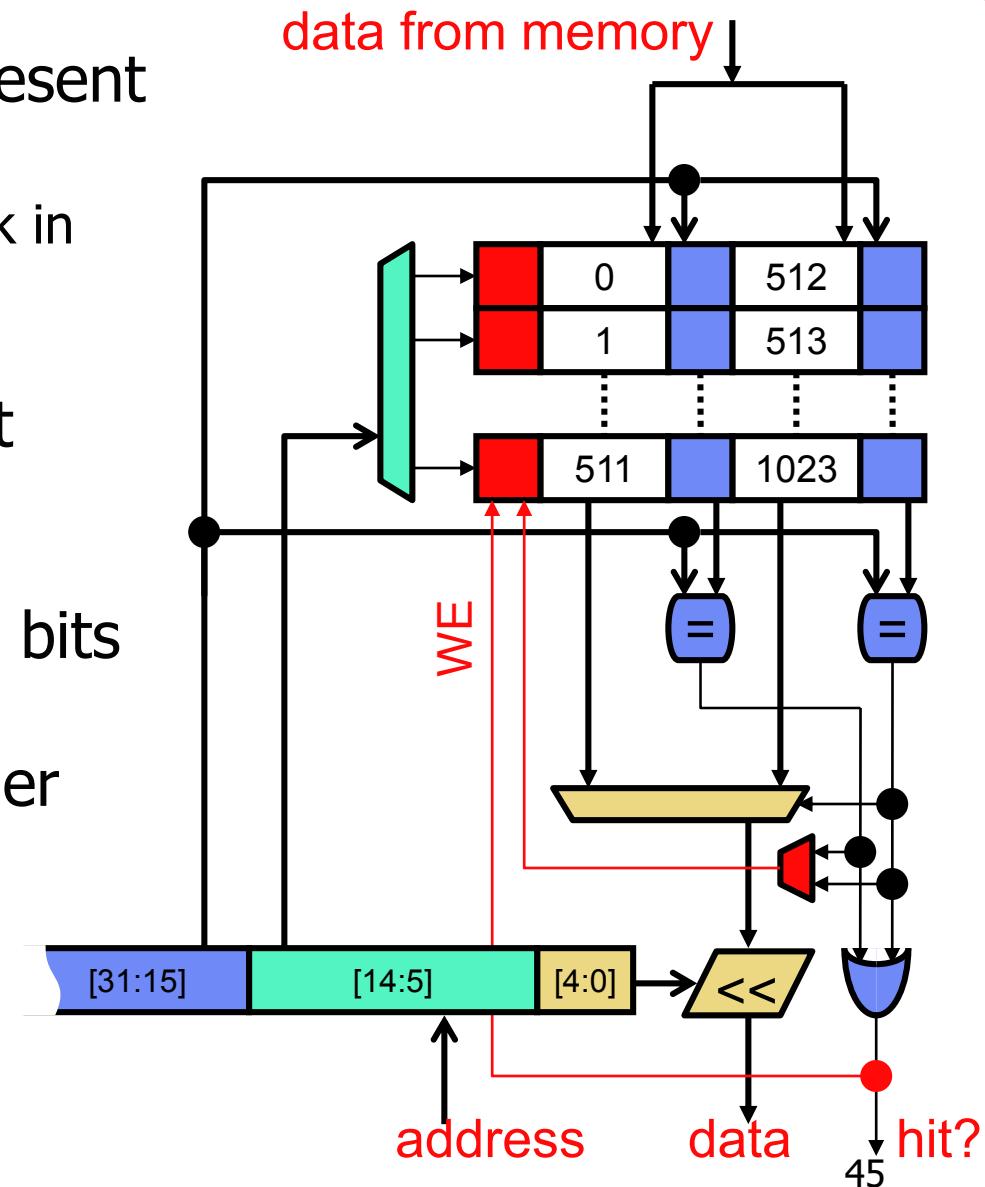
- Higher associative caches
  - + Have better (lower)  $\%_{\text{miss}}$ 
    - Diminishing returns
  - However  $t_{\text{hit}}$  increases
    - The more associative, the slower
    - What about  $t_{\text{avg}}$ ?



- Block-size and number of sets should be powers of two
  - Makes indexing easier (just rip bits out of the address)
- 3-way set-associativity? No problem

# Miss Handling & Replacement Policies

- Set-associative caches present a new design choice
  - On cache miss, which block in set to replace (kick out)?
- Add **LRU** field to each set
  - “Least recently used”
- Each access updates LRU bits
- Pseudo-LRU used for larger associativity caches

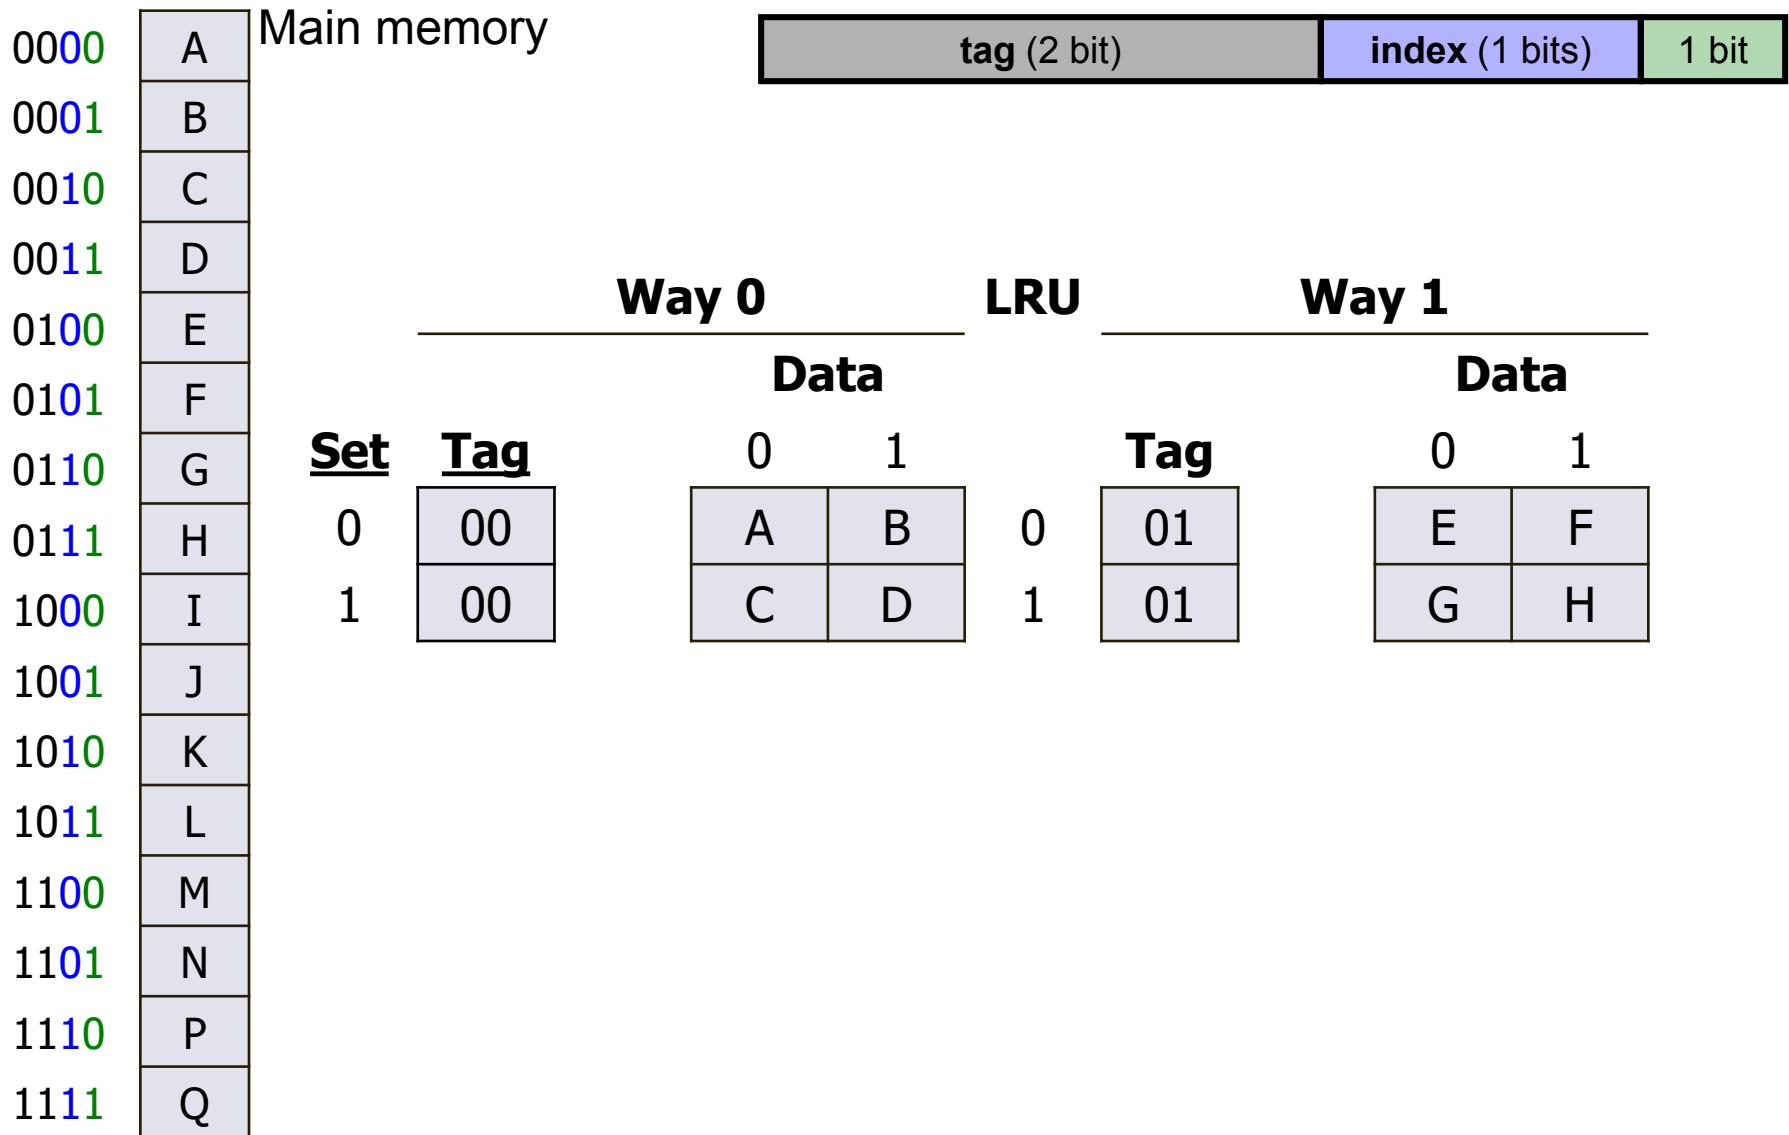


# Replacement Policies

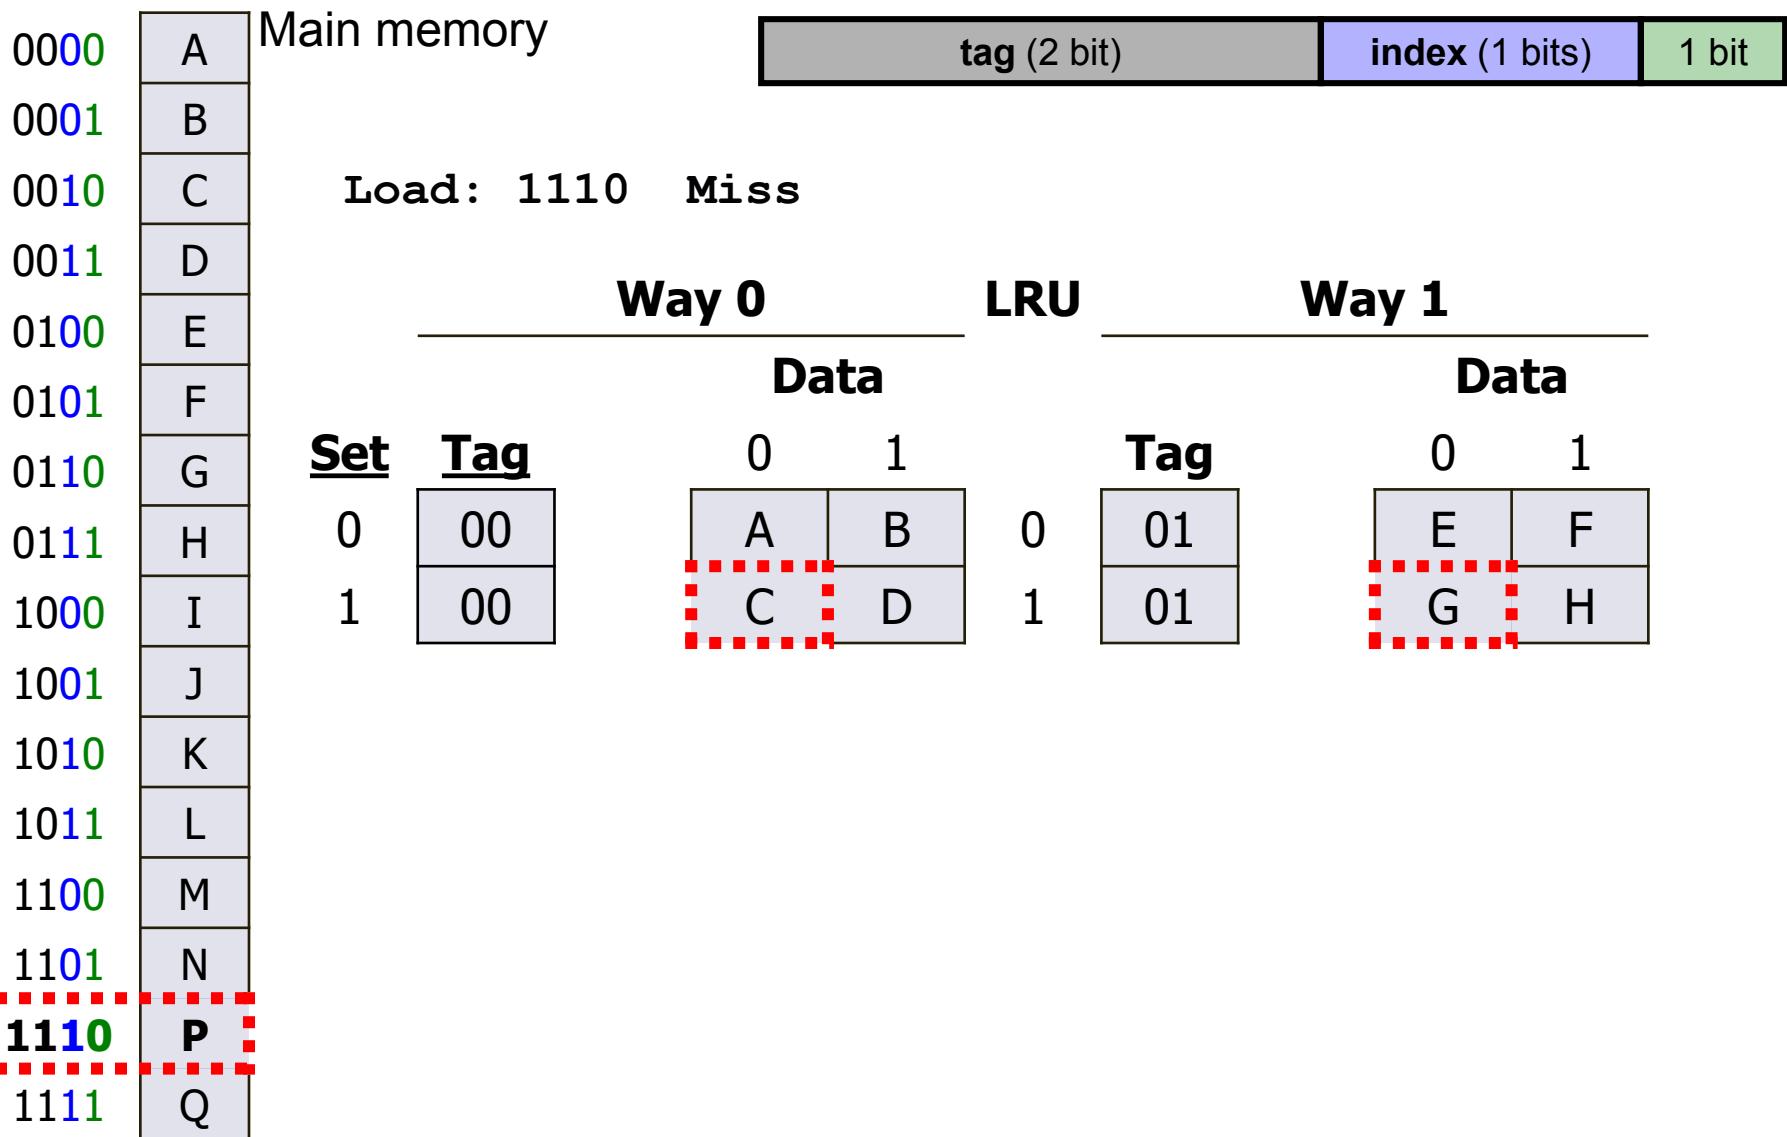
---

- Set-associative caches present a new design choice
  - On cache miss, which block in set to replace (kick out)?
- Some options
  - **Random**
  - **FIFO (first-in first-out)**
  - **LRU (least recently used)**
    - Fits with temporal locality, LRU = least likely to be used in future
  - **NMRU (not most recently used)**
    - An easier to implement approximation of LRU
    - Same as LRU for 2-way set-associative caches
  - **Belady's**: replace block that will be used furthest in future
    - Unachievable optimum

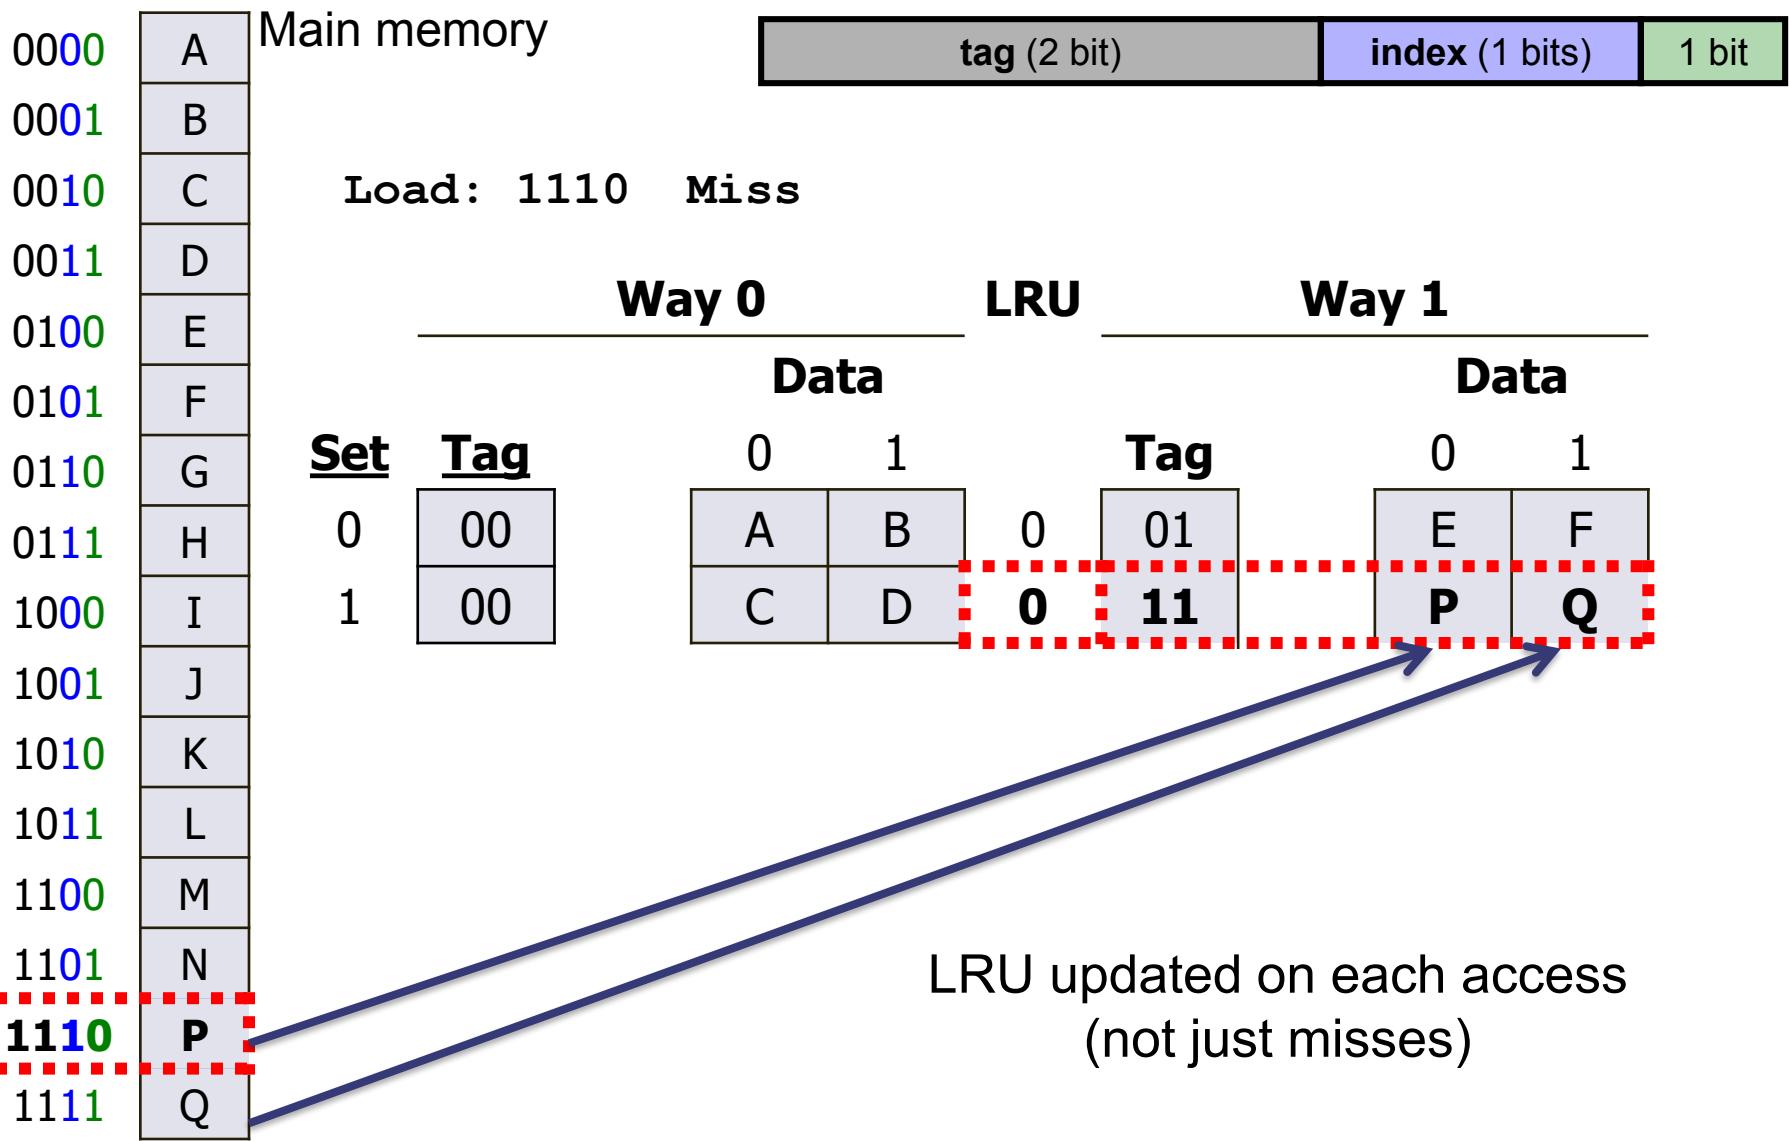
# 4-bit Address, 8B Cache, 2B Blocks, 2-way



# 4-bit Address, 8B Cache, 2B Blocks, 2-way



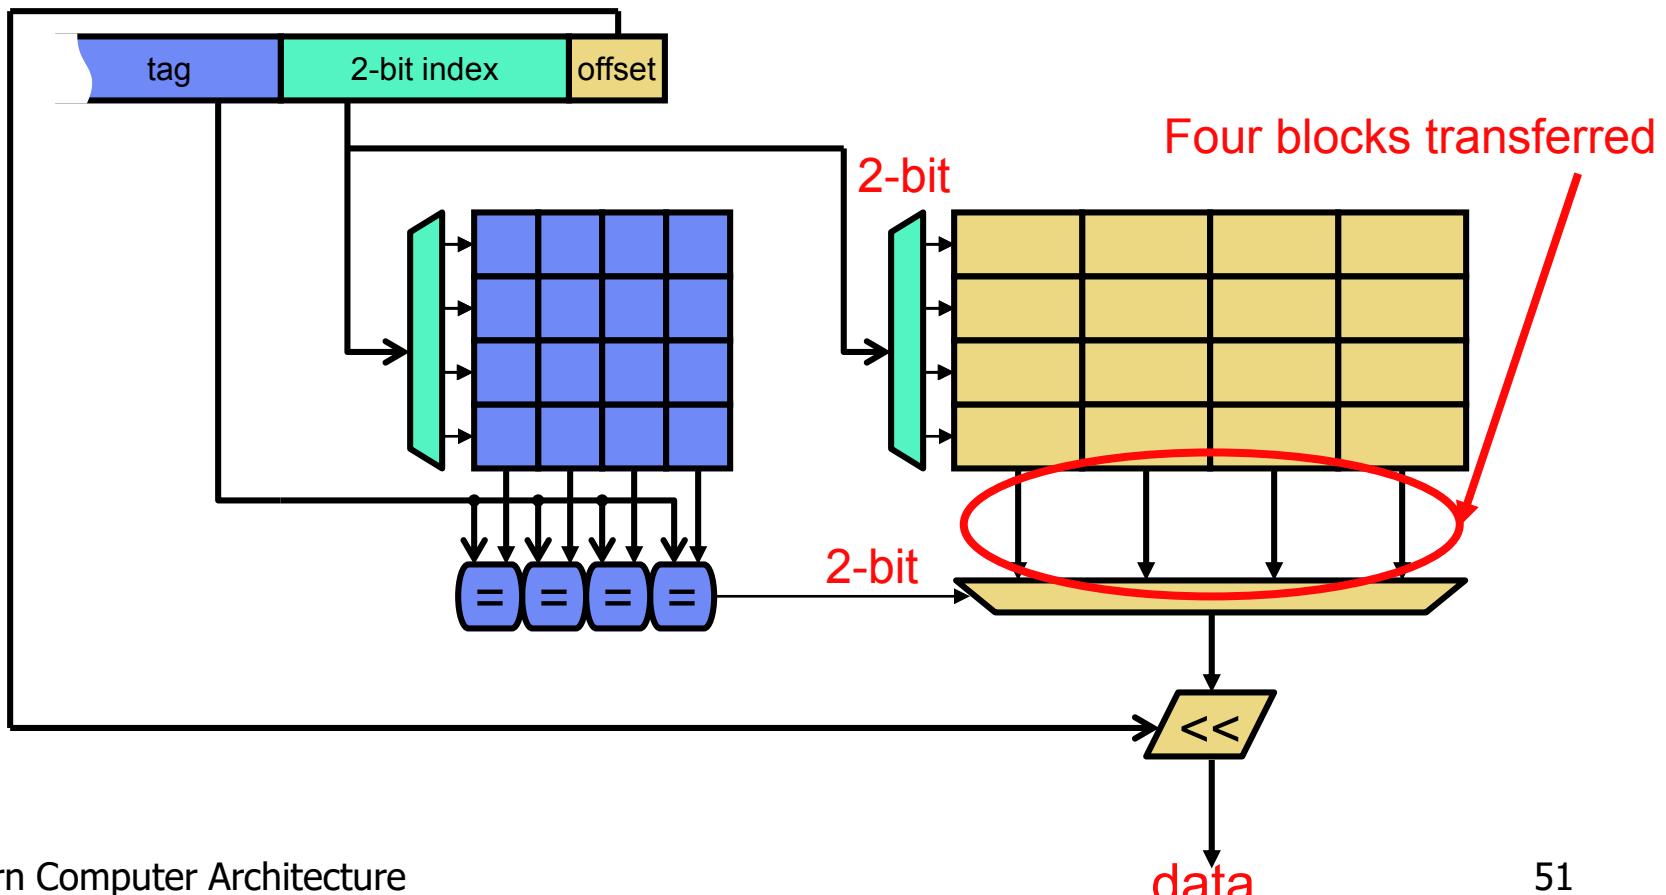
# 4-bit Address, 8B Cache, 2B Blocks, 2-way



# **Implementing Set- Associative Caches**

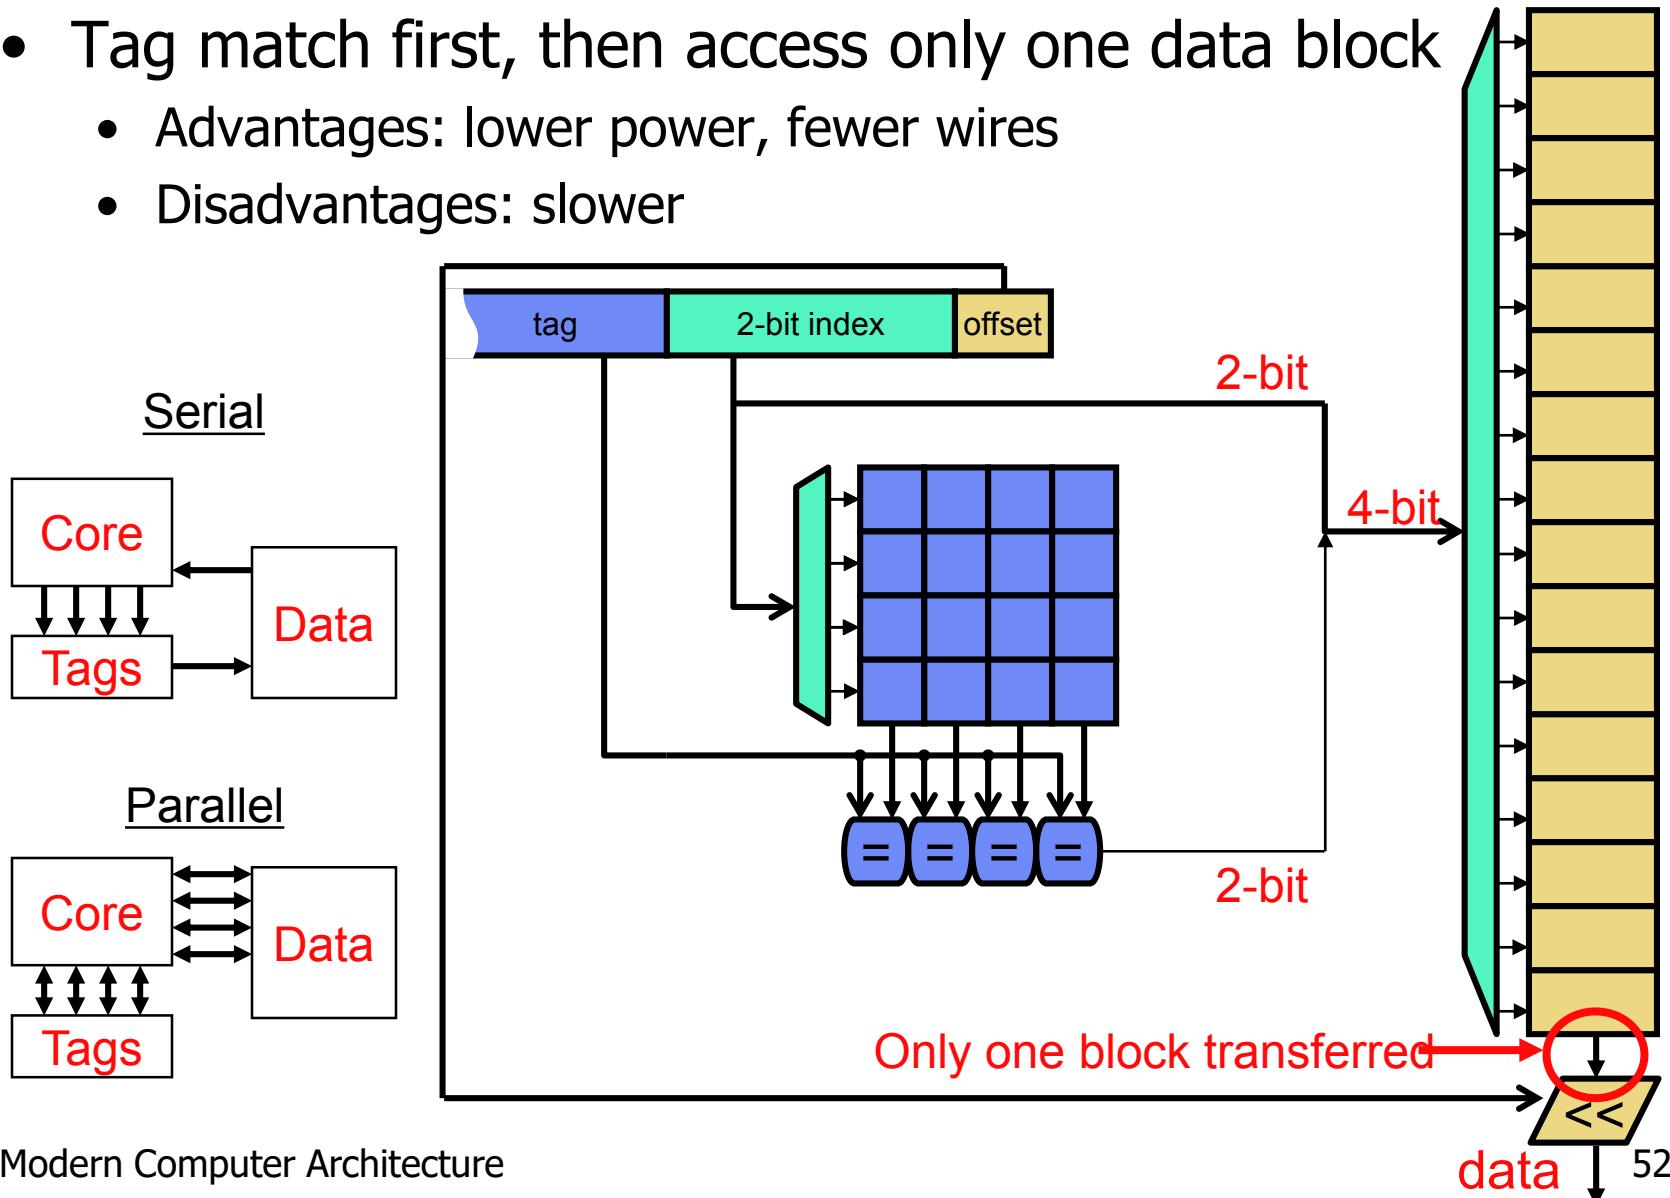
# Option#1: Parallel Tag Access

- Data and tags actually physically separate
  - Split into two different memory structures
- Option#1: read both structures in parallel:



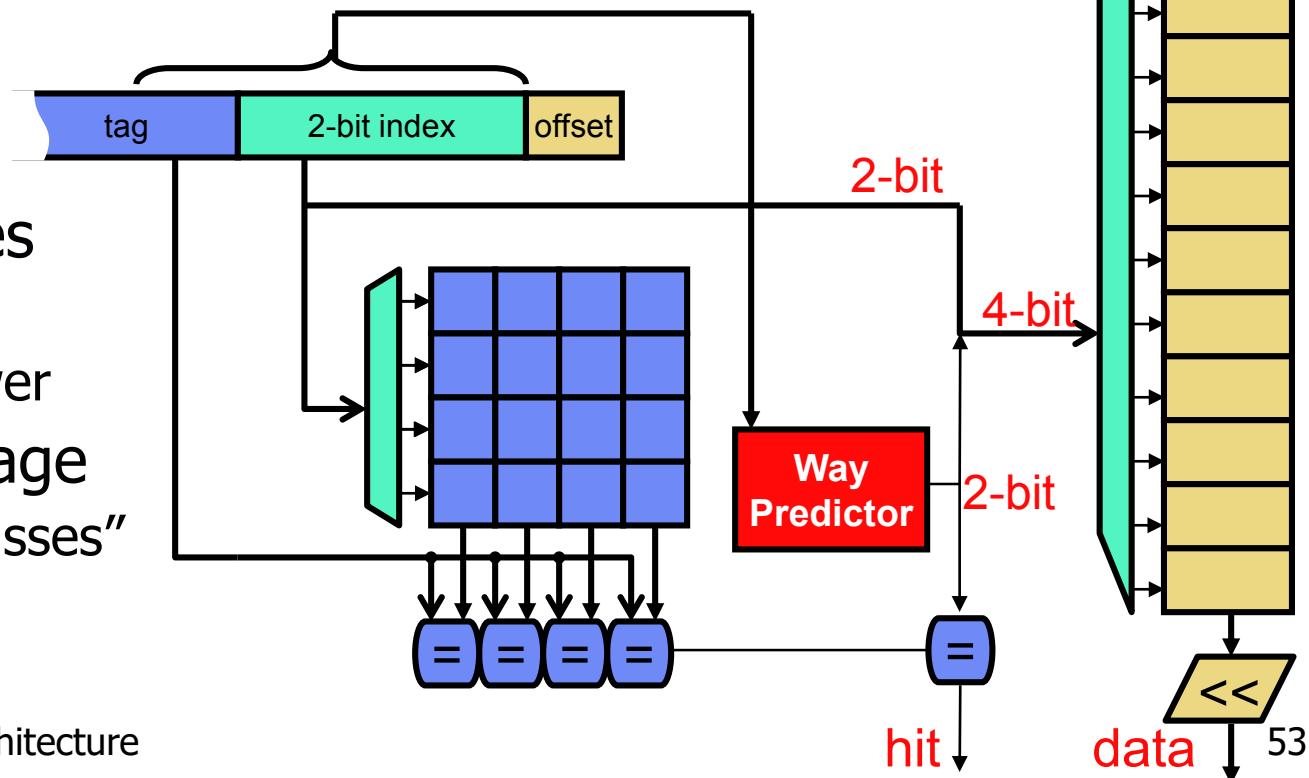
# Option#2: Serial Tag Access

- Tag match first, then access only one data block
  - Advantages: lower power, fewer wires
  - Disadvantages: slower



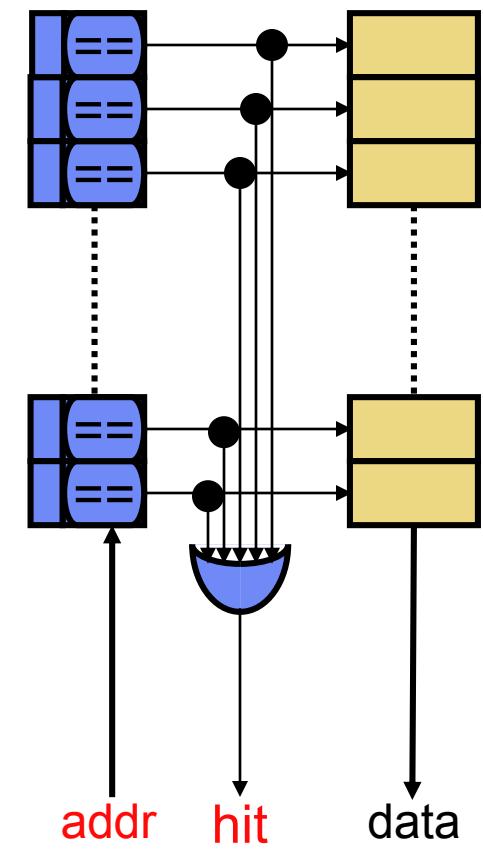
# Best of Both? Way Prediction

- Predict “way” of block
  - Just a “hint”
  - Use the index plus some tag bits
  - Table of  $n$ -bit entries for  $2^n$  associative cache
  - Update on mis-prediction or replacement



# Highly Associative Caches with “CAMs”

- **CAM: content addressable memory**
  - Array of words with **built-in comparators**
  - **No separate “decoder” logic**
  - Input is value to match (tag)
  - Generates 1-hot encoding of matching slot
- Fully associative cache
  - Tags as CAM, data as RAM
  - Effective but somewhat expensive
    - But cheaper than any other method
  - Used for high (16-/32-way) associativity
    - No good way to build 1024-way associativity
    - + No real need for it, either
- CAMs are used elsewhere, too...



# **Cache Optimizations**

# Classifying Misses: 3C Model

---

- Divide cache misses into three categories
  - **Compulsory (cold)**: never seen this address before
    - **Would miss even in infinite cache**
  - **Capacity**: miss because cache is too small
    - **Would miss even in fully associative cache**
    - Identify? Consecutive accesses to block separated by access to at least  $N$  other distinct blocks ( $N$  is number of frames in cache)
  - **Conflict**: miss because cache associativity is too low
    - Identify? **All other misses**
  - **(Coherence)**: miss due to external invalidations
    - Only in shared memory multiprocessors (later)
- Calculated by multiple simulations
  - Simulate infinite cache, fully-associative cache, normal cache
  - Subtract to find each count

# Miss Rate: ABC

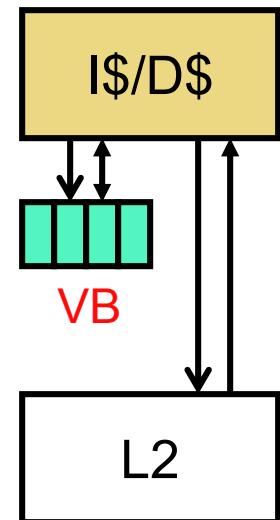
---

- Why do we care about 3C miss model?
  - So that we know what to do to eliminate misses
  - e.g. if have no conflict misses, increasing associativity won't help
- **Associativity**
  - + Decreases conflict misses
  - Increases latency<sub>hit</sub>
- **Block size**
  - Increases conflict/capacity misses
  - + Decreases compulsory misses
  - No significant effect on latency<sub>hit</sub>
- **Capacity**
  - + Decreases capacity misses
  - Increases latency<sub>hit</sub>

# Reducing Conflict Misses: Victim Buffer

---

- Conflict misses: not enough associativity
  - High-associativity is expensive, but also rarely needed
    - 3 blocks mapping to same 2-way set
- **Victim buffer (VB)**: small fully-associative cache
  - Sits on I\$/D\$ miss path
  - Small so very fast (e.g., 8 entries)
  - Blocks evicted from I\$/D\$ are placed in VB
  - On miss, check VB: hit? Place block back in I\$/D\$
  - 8 extra ways, shared among all sets
    - + Only a few sets will need it at any given time
  - + Very effective in practice



# Overlapping Misses: Lockup Free Cache

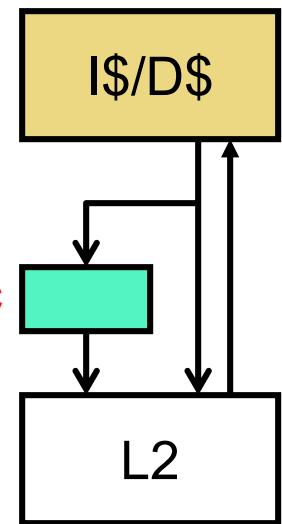
---

- **Lockup free:** allows other accesses while miss is pending
  - Consider: Load [r1] -> r2; Load [r3] -> r4; Add r2, r4 -> r5
- **Handle misses in parallel**
  - Allows “overlapping” misses
  - “memory-level parallelism”
- Implementation: **miss status holding register (MSHR)**
  - Remember: miss address, chosen frame, requesting instruction
  - When miss returns know where to put block, who to inform

# Prefetching

---

- Bring data into cache proactively/**speculatively**
  - If successful, reduces number of cache misses
- Key: anticipate upcoming miss addresses accurately
  - Can do in software or hardware
- Simple hardware prefetching: **next block prefetching**
  - Miss on address **X** → anticipate miss on **X+block-size**
    - + Works for insns: sequential execution
    - + Works for data: arrays
- Table-driven hardware prefetching
  - Use **predictor** to detect strides, common patterns
- Effectiveness determined by:
  - **Timeliness**: initiate prefetches sufficiently in advance
  - **Coverage**: prefetch for as many misses as possible
  - **Accuracy**: don't pollute with unnecessary data



# Software Prefetching

---

- Use a special “prefetch” instruction
  - Tells the hardware to bring in data, doesn’t actually read it
  - Just a hint
- Inserted by programmer or compiler
- Example

```
int tree_add(tree_t* t) {  
    if (t == NULL) return 0;  
    __builtin_prefetch(t->left);  
    return t->val + tree_add(t->right) + tree_add(t->left);  
}
```

- Multiple prefetches bring multiple blocks in parallel
  - More “memory-level” parallelism (MLP)

# Software Restructuring: Data

---

- Capacity misses: poor spatial or temporal locality
  - Several code restructuring techniques to improve both
    - Loop blocking (break into cache-sized chunks), loop fusion
  - Compiler must know that restructuring preserves semantics
- **Loop interchange**: spatial locality
  - Example: for a row-major matrix, `x[i][j]` should be followed by `x[i][j+1]`
  - Poor code: `x[i][j]` followed by `x[i+1][j]`

```
for (j = 0; j<NCOLS; j++)
    for (i = 0; i<NROWS; i++)
        sum += x[i][j];
```
  - Better code
    - `for (i = 0; i<NROWS; i++)`
    - `for (j = 0; j<NCOLS; j++)`
    - `sum += x[i][j];`

# Software Restructuring: Data

---

- **Loop blocking**: temporal locality

- Poor code

```
for (k=0; k<NUM_ITERATIONS; k++)
    for (i=0; i<NUM_ELEMS; i++)
        X[i] = f(X[i]); // say
```

- Better code

- Cut array into CACHE\_SIZE chunks
- Run all phases on one chunk, proceed to next chunk

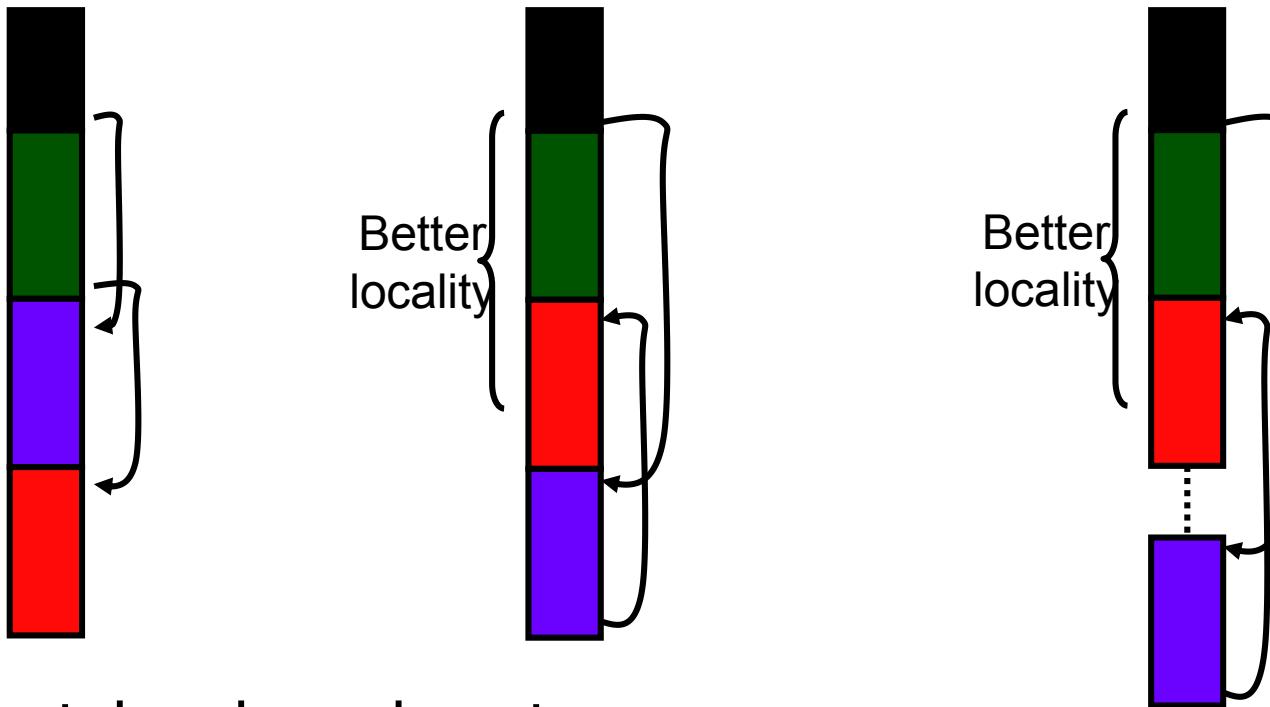
```
for (i=0; i<NUM_ELEMS; i+=CACHE_SIZE)
    for (k=0; k<NUM_ITERATIONS; k++)
        for (j=0; j<CACHE_SIZE; j++)
            X[i+j] = f(X[i+j]);
```

- Assumes you know CACHE\_SIZE, do you?

# Software Restructuring: Code

---

- Compiler can layout code for better locality
  - If (a) { **code1;** } else { **code2;** } **code3;**
  - But, code2 case never happens (say, error condition)

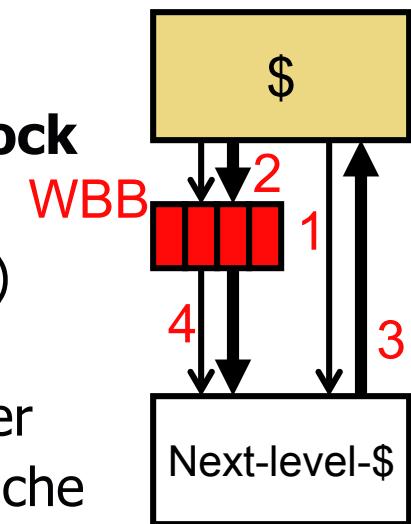


- Fewer taken branches, too

# **What About Stores?**

# Handling Cache Writes

- When to propagate new value to (lower level) memory?
- **Option #1: Write-through:** immediately
  - On hit, update cache
  - Immediately send the write to the next level
- **Option #2: Write-back:** when block is replaced
  - Requires additional “**dirty**” bit per block
    - Replace **clean** block: **no extra traffic**
    - Replace **dirty** block: **extra “writeback” of block**
  - + **Writeback-buffer (WBB):**
    - Hide latency of writeback (keep off critical path)
    - Step#1: Send “fill” request to next-level
    - Step#2: While waiting, write dirty block to buffer
    - Step#3: When new blocks arrives, put it into cache
    - Step#4: Write buffer contents to next-level



# Determining WBB Size

---

- How big should we make the L1's WBB?
- $L = \lambda W$  (Little's Law)
- $L = (\%miss_{L1}) * (\text{latency}_{L2})$
- configuration:
  - 1GHz processor
  - 20% of insns are loads
  - $\%miss_{L1} = 5\%$
  - L2 access = 20 cycles
  - $L = .01 * 20 \rightarrow 1 \text{ entry suffices}$

# Write Propagation Comparison

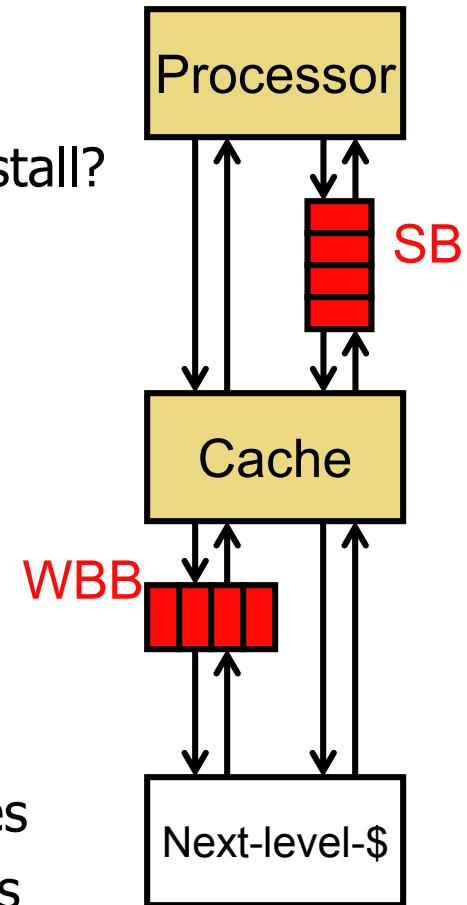
---

- **Write-through**
  - Creates additional traffic
    - Consider repeated write hits
  - Next level must handle small writes (1, 2, 4, 8-bytes)
  - + No need for dirty bits in cache
  - + No need to handle “writeback” operations
    - Simplifies miss handling (no write-back buffer)
  - Sometimes used for L1 caches (IBM, GPUs)
  - Usually **write-non-allocate**: on write miss, just write to next level
- **Write-back**
  - + Key advantage: uses less bandwidth
  - Reverse of other pros/cons above
  - Used by Intel, AMD, and many ARM cores
  - Second-level and beyond are generally write-back caches
  - Usually **write-allocate**: on write miss, fill block from next level

# Write Misses and Store Buffers

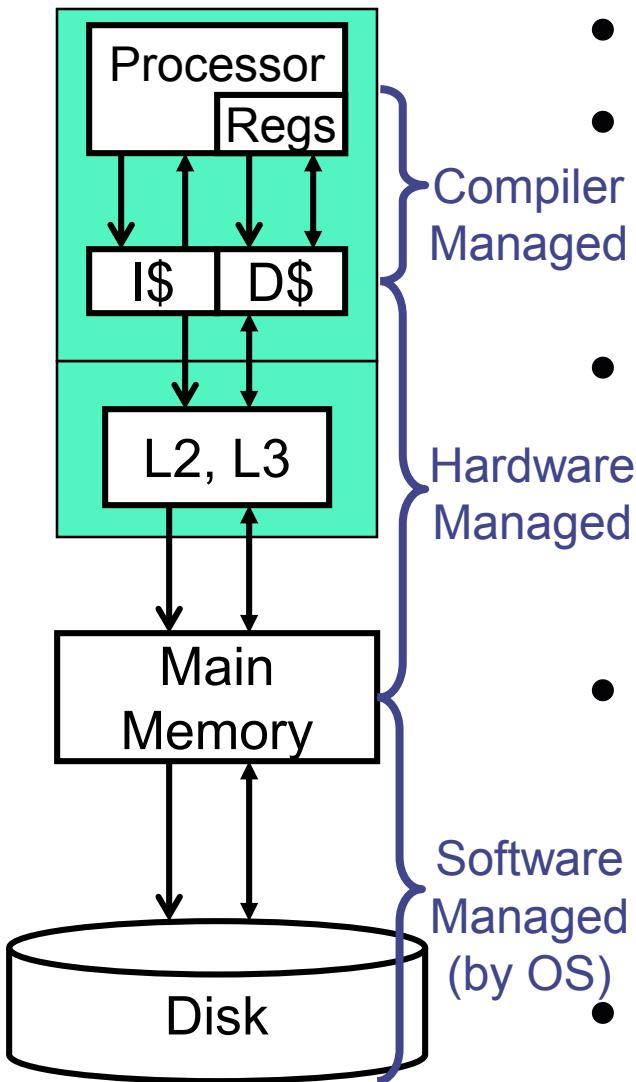
---

- Read miss?
  - Load can't go on without the data, it must stall
- Write miss?
  - Technically, no instruction is waiting for data, why stall?
- **Store buffer:** a small buffer for store misses
  - Stores put address/value into SB, **keep going**
  - SB writes to D\$ in the background
  - Loads must search SB (in addition to D\$)
  - (mostly) eliminates stalls on write misses
  - creates some problems in multiprocessors (later)
- Store buffer vs. writeback buffer
  - Store buffer: “in front” of D\$, for hiding store misses
  - Writeback buffer: “behind” D\$, for hiding writebacks



# **Cache Hierarchies**

# Concrete Memory Hierarchy



- 0th level: **Registers**
- 1st level: **Primary caches**
  - Split instruction (I\$) and data (D\$)
  - Typically 8KB to 64KB each
- 2nd level: **2<sup>nd</sup> and 3<sup>rd</sup> cache (L2, L3)**
  - On-chip, typically made of SRAM
  - 2<sup>nd</sup> level typically ~256KB to 512KB
  - “Last level cache” typically 4MB to 16MB
- 3rd level: **main memory**
  - Made of DRAM (“Dynamic” RAM)
  - Typically 1GB to 4GB for desktops/laptops
    - Servers can have >1 TB
- 4th level: **disk (swap and files)**
  - Uses magnetic disks or flash drives

# Designing a Cache Hierarchy

---

- For any memory component:  $t_{hit}$  vs.  $\%_{miss}$  tradeoff
- Upper components (I\$, D\$) emphasize low  $t_{hit}$ 
  - **Frequent access  $\rightarrow t_{hit}$  important**
  - $t_{miss}$  is low  $\rightarrow \%_{miss}$  less important
  - Lower capacity and lower associativity (to reduce  $t_{hit}$ )
  - Small-medium block-size (to reduce conflicts)
  - **Split instruction & data cache to allow simultaneous access**
- Moving down (L2, L3) emphasis turns to  $\%_{miss}$ 
  - **Infrequent access  $\rightarrow t_{hit}$  less important**
  - $t_{miss}$  is high  $\rightarrow \%_{miss}$  important
  - High capacity, associativity, and block size (to reduce  $\%_{miss}$ )
  - **Unified insn+data caches to dynamically allocate capacity**

# Split vs. Unified Caches

---

- **Split I\$/D\$:** insns and data in different caches
  - To minimize structural hazards and  $t_{hit}$
  - Larger unified I\$/D\$ would be slow, 2nd port even slower
  - Optimize I\$ to fetch multiple instructions, no writes
- **Unified L2, L3:** insns and data together
  - To minimize  $\%_{miss}$
  - + Fewer capacity misses: unused insn capacity can be used for data
  - More conflict misses: insn/data conflicts
    - A much smaller effect in large caches
  - Insn/data structural hazards are rare: simultaneous I\$/D\$ miss
  - Go even further: unify L2/L3 of multiple cores in a multi-core

# Hierarchy: Inclusion versus Exclusion

---

- **Inclusion**

- **Bring block from memory into L2 then L1**
  - A block in the L1 is always in the L2 (but not vice-versa)
  - If block evicted from L2, must also evict it from L1
  - Why? more on this when we talk about multicore

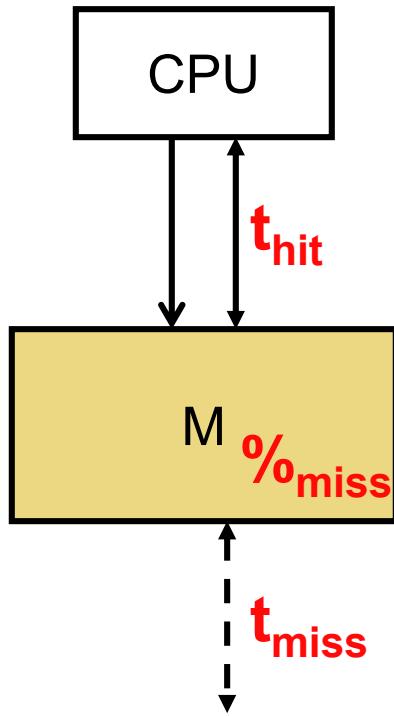
- **Exclusion**

- **Bring block from memory into L1 but not L2**
  - Move block to L2 on L1 eviction
    - L2 becomes a large victim cache
  - Block is either in L1 or L2 (never both)
- Good if L2 is small relative to L1
  - Example: AMD's Duron 64KB L1s, 64KB L2

- **Non-inclusion**

- No guarantees

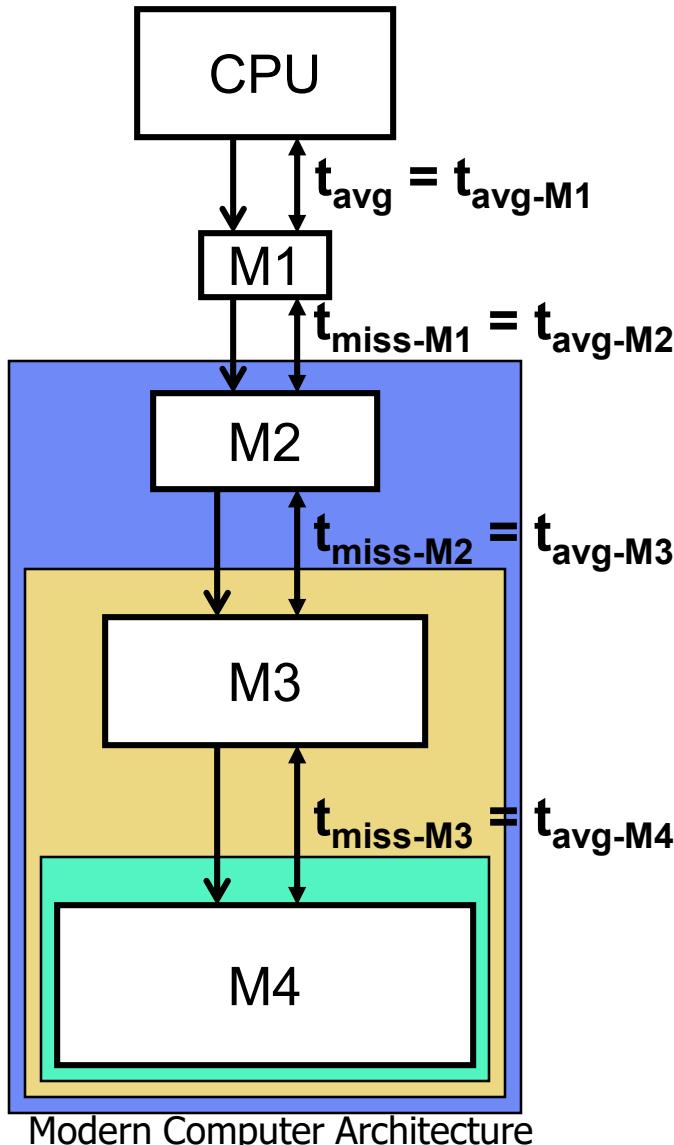
# Memory Performance Equation



- For memory component M
  - **Access**: read or write to M
  - **Hit**: desired data found in M
  - **Miss**: desired data not found in M
    - Must get from another (slower) component
  - **Fill**: action of placing data in M
- $\%_{\text{miss}}$  (miss-rate): #misses / #accesses
- $t_{\text{hit}}$ : time to read data from (write data to) M
- $t_{\text{miss}}$ : time to read data into M
- Performance metric
  - $t_{\text{avg}}$ : average access time

$$t_{\text{avg}} = t_{\text{hit}} + (\%_{\text{miss}} * t_{\text{miss}})$$

# Hierarchy Performance



$t_{avg}$

$t_{avg-M1}$

$t_{hit-M1} + (%_{miss-M1} * t_{miss-M1})$

$t_{hit-M1} + (%_{miss-M1} * t_{avg-M2})$

$t_{hit-M1} + (%_{miss-M1} * (t_{hit-M2} + (%_{miss-M2} * t_{miss-M2})))$

$t_{hit-M1} + (%_{miss-M1} * (t_{hit-M2} + (%_{miss-M2} * t_{avg-M3})))$

...

# Performance Calculation

---

- In a pipelined processor, I\$/D\$  $t_{hit}$  is “built in” (effectively 0)
- Parameters
  - Base pipeline CPI = 1
  - Instruction mix: 30% loads/stores
  - I\$:  $\%_{miss} = 2\%$ ,  $t_{miss} = 10$  cycles
  - D\$:  $\%_{miss} = 10\%$ ,  $t_{miss} = 10$  cycles
- What is CPI with cache misses?
  - $CPI_{I\$} = \%_{missI\$} * t_{miss} = 0.02 * 10 \text{ cycles} = 0.2 \text{ cycle}$
  - $CPI_{D\$} = \%_{memory} * \%_{missD\$} * t_{missD\$} = 0.30 * 0.10 * 10 \text{ cycles} = 0.3 \text{ cycle}$
  - $CPI_{new} = CPI + CPI_{I\$} + CPI_{D\$} = 1 + 0.2 + 0.3 = 1.5$

# Miss Rates: per “access” vs “instruction”

---

- Miss rates can be expressed two ways:
  - Misses per “instruction” (or instructions per miss), -or-
  - Misses per “cache access” (or accesses per miss)
- For first-level caches, use instruction mix to convert
  - If memory ops are 1/3<sup>rd</sup> of instructions..
  - 2% of instructions miss (1 in 50) is 6% of “accesses” miss (1 in 17)
- What about second-level caches?
  - Misses per “instruction” still straight-forward (“global” miss rate)
  - Misses per “access” is trickier (“local” miss rate)
    - Depends on number of accesses (which depends on L1 miss rate)

# Performance Calculation (Revisited)

---

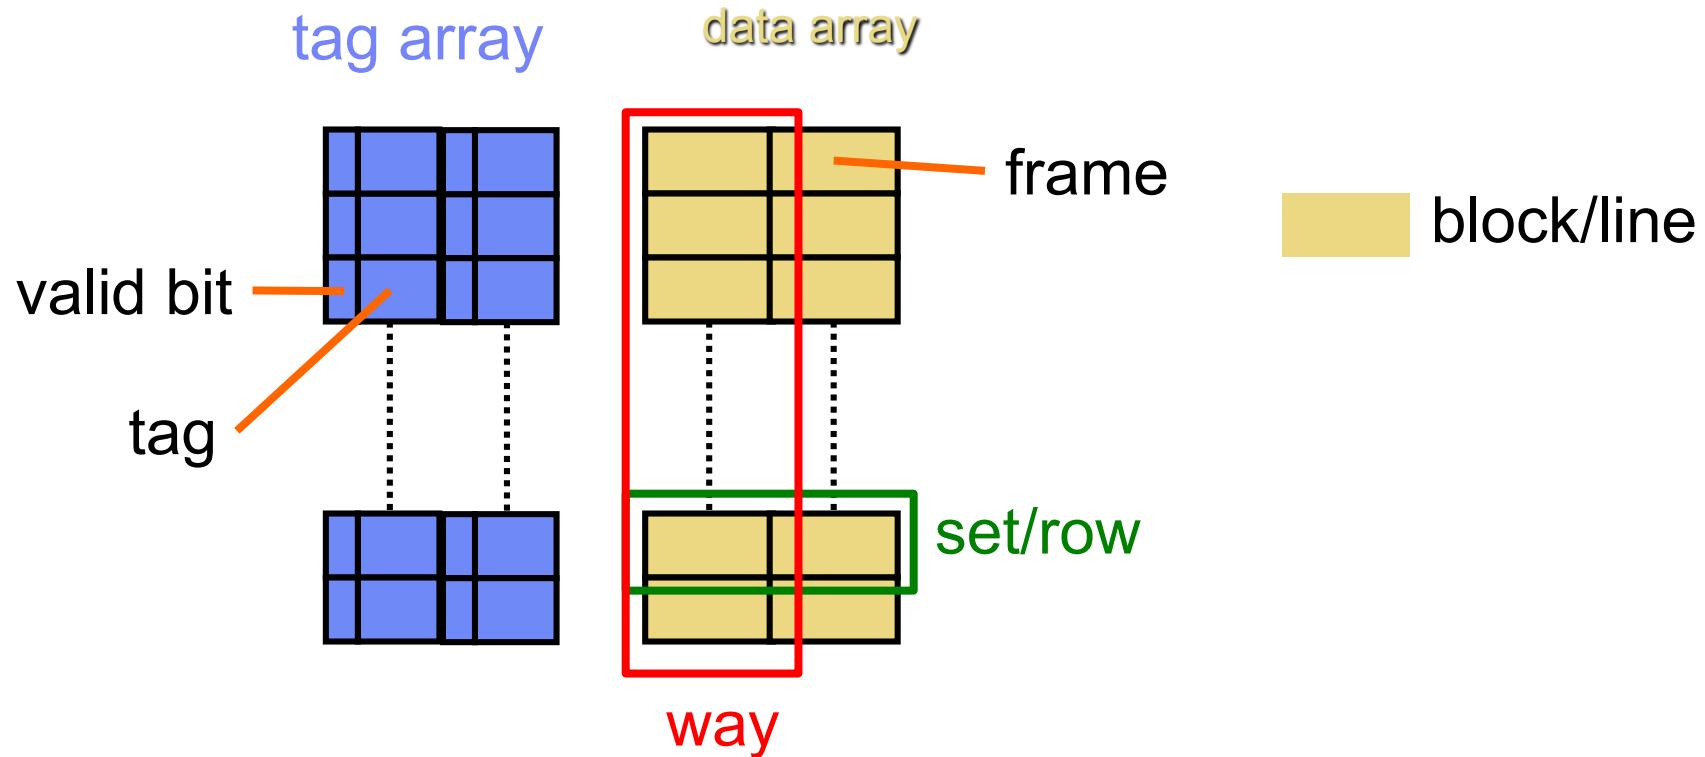
- Parameters
  - Base pipeline CPI = 1
    - In this case, already incorporates  $t_{hit}$
  - I\$:  $\%_{miss} = 2\%$  of instructions,  $t_{miss} = 10$  cycles
  - D\$:  $\%_{miss} = 3\%$  of instructions,  $t_{miss} = 10$  cycles
- What is new CPI?
  - $CPI_{I\$} = \%_{missI\$} * t_{miss} = 0.02 * 10 \text{ cycles} = 0.2 \text{ cycle}$
  - $CPI_{D\$} = \%_{missD\$} * t_{missD\$} = 0.03 * 10 \text{ cycles} = 0.3 \text{ cycle}$
  - $CPI_{new} = CPI + CPI_{I\$} + CPI_{D\$} = 1 + 0.2 + 0.3 = 1.5$

# Multilevel Performance Calculation

---

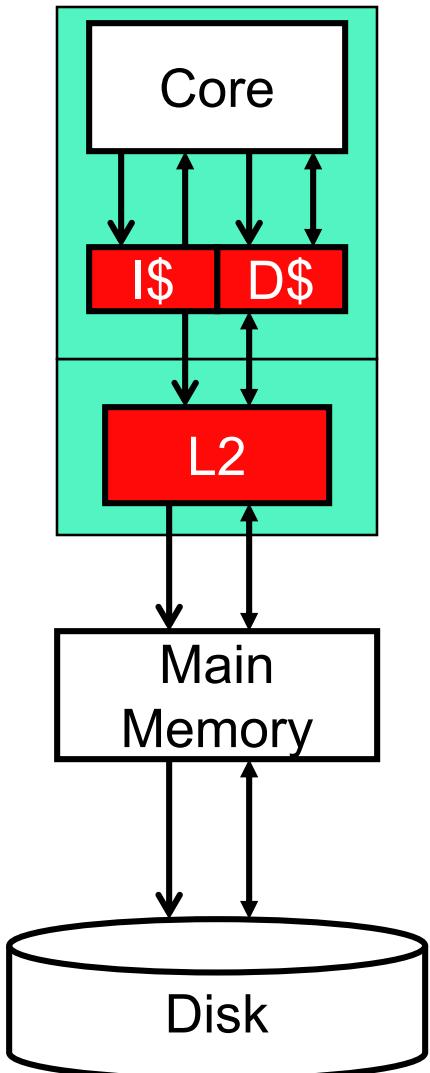
- Parameters
  - 30% of instructions are memory operations
  - L1:  $t_{hit} = 1$  cycles (included in CPI of 1),  $\%_{miss} = 5\%$  of accesses
  - L2:  $t_{hit} = 10$  cycles,  $\%_{miss} = 20\%$  **of L2 accesses**
  - Main memory:  $t_{hit} = 50$  cycles
- Calculate CPI
  - $CPI = 1 + 30\% * 5\% * t_{missD\$}$
  - $t_{missD\$} = t_{avgL2} = t_{hitL2} + (\%_{missL2} * t_{hitMem}) = 10 + (20\% * 50) = 20$  cycles
  - Thus,  $CPI = 1 + 30\% * 5\% * 20 = 1.3$  CPI
- Alternate CPI calculation:
  - What % of instructions miss in L1 cache?  $30\% * 5\% = 1.5\%$
  - What % of instructions miss in L2 cache?  $20\% * 1.5\% = 0.3\%$  of insn
  - $CPI = 1 + (1.5\% * 10) + (0.3\% * 50) = 1 + 0.15 + 0.15 = 1.3$  CPI

# Cache Glossary



# Summary

---



- “Cache”: hardware managed
  - Hardware automatically retrieves missing data
  - Built from fast on-chip SRAM
  - In contrast to off-chip, DRAM “main memory”
- **Average access time** of a memory component
  - $latency_{avg} = latency_{hit} + (\%_{miss} * latency_{miss})$
  - Hard to get low  $latency_{hit}$  and  $\%_{miss}$  in one structure  
→ memory hierarchy
- Cache ABCs (**associativity, block size, capacity**)
- **Performance optimizations**
  - Prefetching & data restructuring
- **Handling writes**
  - Write-back vs. write-through
- **Memory hierarchy**
  - Smaller, faster, expensive → bigger, slower, cheaper