

# Modern Computer Architecture

with sources that included UPenn & University of Wisconsin slides  
by Joe Devietti, Milo Martin & Amir Roth, Mark Hill, Guri Sohi, Jim Smith, and David Wood


# Schedule

---

- Introduction and Transistors
- Parallel computing (Isaac D. Scherson, University of California, Irvine, in October)
- ISAs
- Performance
- Pipelining Basic
- Branch Prediction
- Caches
- Virtual Memory
- Out-of-Order Execution
- Multicore multi-thread
- Vectors/GPUs for data parallelism.

# This Unit: Static & Dynamic Scheduling

---



- Code scheduling
  - To reduce pipeline stalls
  - To increase ILP (insn level parallelism)
- Static scheduling by the compiler
  - Approach & limitations
- Dynamic scheduling in hardware
  - Register renaming
  - Instruction selection
  - Handling memory operations

# **Code Scheduling & Limitations**

# Code Scheduling

---

- Scheduling: act of finding independent instructions
  - “Static” done at compile time by the compiler (software)
  - “Dynamic” done at runtime by the processor (hardware)
- Why schedule code?
  - Scalar pipelines: fill in load-to-use delay slots to improve CPI
  - Superscalar: place independent instructions together
    - As above, load-to-use delay slots
    - Allow multiple-issue decode logic to let them execute at the same time

# Compiler Scheduling

- Compiler can schedule (move) instructions to reduce stalls
  - Basic pipeline scheduling**: eliminate back-to-back load-use pairs
  - Example code sequence: **a = b + c; d = f - e;**
    - sp stack pointer, sp+0 is “a”, sp+4 is “b”, etc...

Before

```
ld [sp+4]→r2
ld [sp+8]→r3
add r2,r3→r1 //stall
st r1→[sp+0]
ld [sp+16]→r5
ld [sp+20]→r6
sub r6,r5→r4 //stall
st r4→[sp+12]
```

After

```
ld [sp+4]→r2
ld [sp+8]→r3
ld [sp+16]→r5
add r2,r3→r1 //no stall
ld [sp+20]→r6
st r1→[sp+0]
sub r6,r5→r4 //no stall
st r4→[sp+12]
```

# Compiler Scheduling Requires

---

- **Large scheduling scope**

- Independent instruction to put between load-use pairs
  - + Original example: large scope, two independent computations
  - This example: small scope, one computation

Before

```
ld [sp+4]→r2
ld [sp+8]→r3
add r2,r3→r1 //stall
st r1→[sp+0]
```

After (same!)

```
ld [sp+4]→r2
ld [sp+8]→r3
add r2,r3→r1 //stall
st r1→[sp+0]
```

- Compiler can create **larger** scheduling scopes
  - For example: loop unrolling & function inlining

# Scheduling Scope Limited by Branches

---

**r1 and r2 are inputs**

loop:

  jz r1, not\_found

  ld [r1+0] → r3

  sub r2, r3 → r4

  jz r4, found

  ld [r1+4] → r1

  jmp loop



Aside: what does this code do?

Legal to move load up past branch?

# Compiler Scheduling Requires

---

- **Enough registers**

- To hold additional “live” values
- Example code contains 7 different values (including `sp`)
- Before: max 3 values live at any time → 3 registers enough
- After: max 4 values live → 3 registers not enough

Original

```
ld [sp+4]→r2
ld [sp+8]→r1
add r1,r2→r1 //stall
st r1→[sp+0]
ld [sp+16]→r2
ld [sp+20]→r1
sub r2,r1→r1 //stall
st r1→[sp+12]
```

Wrong!

```
ld [sp+4]→r2
ld [sp+8]→r1
ld [sp+16]→r2
add r1,r2→r1 // wrong r2
ld [sp+20]→r1
st r1→[sp+0] // wrong r1
sub r2,r1→r1
st r1→[sp+12]
```

# Compiler Scheduling Requires

- **Alias analysis**
  - Ability to tell whether load/store reference same memory locations
    - Effectively, whether load/store can be rearranged
  - Previous example: easy, loads/stores use same base register (`sp`)
  - New example: can compiler tell that `r8 != r9`?
  - Must be **conservative**

Before

```
ld [r9+4]→r2
ld [r9+8]→r3
add r3,r2→r1 //stall
st r1→[r9+0]
ld [r8+0]→r5
ld [r8+4]→r6
sub r5,r6→r4 //stall
st r4→[r8+8]
```

Wrong(?)

```
ld [r9+4]→r2
ld [r9+8]→r3
ld [r8+0]→r5 //does r8==r9?
add r3,r2→r1
ld [r8+4]→r6 //does r8+4==r9?
st r1→[r9+0]
sub r5,r6→r4
st r4→[r8+8]
```

# A Good Case: Static Scheduling of SAXPY

---

- **SAXPY** (Single-precision A X Plus Y)
  - Linear algebra routine (used in solving systems of equations)

```
for (i=0;i<N;i++)  
    Z[i]=(A*X[i])+Y[i];  
  
0: ldf [X+r1]→f1      // loop  
1: mulf f0,f1→f2      // A in f0  
2: ldf [Y+r1]→f3      // X,Y,Z are constant addresses  
3: addf f2,f3→f4  
4: stf f4→[Z+r1]  
5: addi r1,4→r1        // i in r1  
6: blt r1,r2,0          // N*4 in r2
```

- Static scheduling works great for SAXPY
  - All loop iterations independent
  - Use loop unrolling to increase scheduling scope
  - Aliasing analysis is tractable (just ensure X, Y, Z are independent)
  - Still limited by number of registers

# Unrolling & Scheduling SAXPY

- Fuse two (in general, K) iterations of loop
  - Fuse loop control: induction variable (*i*) increment + branch
  - Adjust register names & induction uses (constants → constants+4)
  - Reorder operations to reduce stalls

```
ldf [X+r1]→f1  
mulf f0,f1→f2  
ldf [Y+r1]→f3  
addf f2,f3→f4  
stf f4→[Z+r1]  
addi r1,4→r1  
blt r1,r2,0
```



```
ldf [X+r1]→f1  
mulf f0,f1→f2  
ldf [Y+r1]→f3  
addf f2,f3→f4  
stf f4→[Z+r1]  
  
ldf [X+r1+4]→f5  
mulf f0,f5→f6  
ldf [Y+r1+4]→f7  
addf f6,f7→f8  
stf f8→[Z+r1+4]  
addi r1,8→r1  
blt r1,r2,0
```



```
ldf [X+r1]→f1  
ldf [X+r2+4]→f5  
mulf f0,f1→f2  
mulf f0,f5→f6  
ldf [Y+r1]→f3  
ldf [Y+r1+4]→f7  
addf f2,f3→f4  
addf f6,f7→f8  
stf f4→[Z+r1]  
stf f8→[Z+r1+4]  
addi r1,8→r1  
blt r1,r2,0
```

# Compiler Scheduling Limitations

---

- Scheduling scope
  - Example: can't generally move memory operations past branches
- Limited number of registers (set by ISA)
- Inexact “memory aliasing” information
  - Often prevents reordering of loads above stores by compiler
- Caches misses (or any runtime event) confound scheduling
  - How can the compiler know which loads will miss vs hit?
  - Can impact the compiler’s scheduling decisions

# **Dynamic (Hardware) Scheduling**

# Can Hardware Overcome These Limits?

---

- **Dynamically-scheduled processors**
  - Also called “out-of-order” processors
  - Hardware re-schedules insns...
  - ...within a sliding window of VonNeumann insns
  - As with pipelining and superscalar, ISA unchanged
    - Same hardware/software interface, appearance of in-order
- Increases scheduling scope
  - Does loop unrolling transparently!
  - Uses branch prediction to “unroll” branches
- Examples:
  - Pentium Pro/II/III (3-wide), Core 2 (4-wide), Alpha 21264 (4-wide), MIPS R10000 (4-wide), Power5 (5-wide)

# Example: In-Order Limitations #1

	0	1	2	3	4	5	6	7	8	9	10	11	12
Ld [r1] → r2	F	D	X	M <sub>1</sub>	M <sub>2</sub>	W							
add r2 + r3 → r4	F	D	<b>d*</b>	<b>d*</b>	<b>d*</b>	X	M <sub>1</sub>	M <sub>2</sub>	W				
xor r4 ^ r5 → r6		F	D	<b>d*</b>	<b>d*</b>	<b>d*</b>	X	M <sub>1</sub>	M <sub>2</sub>	W			
ld [r7] → r4		F	D	<b>p*</b>	<b>p*</b>	<b>p*</b>	X	M <sub>1</sub>	M <sub>2</sub>	W			

- In-order pipeline, three-cycle load-use penalty
  - 2-wide
- Why not the following?

	0	1	2	3	4	5	6	7	8	9	10	11	12
Ld [r1] → r2	F	D	X	M <sub>1</sub>	M <sub>2</sub>	W							
add r2 + r3 → <b>r4</b>	F	D	<b>d*</b>	<b>d*</b>	<b>d*</b>	X	M <sub>1</sub>	M <sub>2</sub>	W				
xor <b>r4</b> ^ r5 → r6		F	D	<b>d*</b>	<b>d*</b>	<b>d*</b>	X	M <sub>1</sub>	M <sub>2</sub>	W			
ld [r7] → <b>r4</b>		F	D	<b>X</b>	<b>M<sub>1</sub></b>	<b>M<sub>2</sub></b>	<b>W</b>						

# Example: In-Order Limitations #2

	0	1	2	3	4	5	6	7	8	9	10	11	12
Ld [p1] → p2	F	D	X	M <sub>1</sub>	M <sub>2</sub>	W							
add p2 + p3 → p4	F	D	<b>d*</b>	<b>d*</b>	<b>d*</b>	X	M <sub>1</sub>	M <sub>2</sub>	W				
xor p4 ^ p5 → p6		F	D	<b>d*</b>	<b>d*</b>	<b>d*</b>	X	M <sub>1</sub>	M <sub>2</sub>	W			
ld [p7] → p8		F	D	<b>p*</b>	<b>p*</b>	<b>p*</b>	X	M <sub>1</sub>	M <sub>2</sub>	W			

- In-order pipeline, three-cycle load-use penalty
  - 2-wide
- Why not the following:


	0	1	2	3	4	5	6	7	8	9	10	11	12
Ld [p1] → p2	F	D	X	M <sub>1</sub>	M <sub>2</sub>	W							
add p2 + p3 → p4	F	D	<b>d*</b>	<b>d*</b>	<b>d*</b>	X	M <sub>1</sub>	M <sub>2</sub>	W				
xor p4 ^ p5 → p6		F	D	<b>d*</b>	<b>d*</b>	<b>d*</b>	X	M <sub>1</sub>	M <sub>2</sub>	W			
ld [p7] → p8		F	D	<b>X</b>	<b>M<sub>1</sub></b>	<b>M<sub>2</sub></b>	<b>W</b>						

# Out-of-Order to the Rescue

	0	1	2	3	4	5	6	7	8	9	10	11	12
Ld [p1] → p2	F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W	C				
add p2 + p3 → p4	F	Di				I	RR	X	W	C			
xor p4 ^ p5 → p6		F	Di			I	RR	X	W	C			
ld [p7] → p8		F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W	C			

- “Dynamic scheduling” done by the hardware
- Still 2-wide superscalar, but now out-of-order, too
  - Allows instructions to issue when dependences are ready
- Longer pipeline
  - In-order front end: Fetch, “**Dispatch**”
  - Out-of-order execution core:
    - “**Issue**”, “**RegisterRead**”, Execute, Memory, Writeback
  - In-order retirement: “**Commit**”

# Out-of-Order Pipeline



# Out-of-Order Execution


---

- Also called “Dynamic scheduling”
  - Done by the hardware on-the-fly during execution
- Looks at a “window” of instructions waiting to execute
  - Each cycle, picks the next ready instruction(s)
- Two steps to enable out-of-order execution:
  - Step #1: Register renaming – to avoid “false” dependencies
  - Step #2: Dynamically schedule – to enforce “true” dependencies
- Key to understanding out-of-order execution:
  - **Data dependencies**


# Dependence types

---


- **RAW** (Read After Write) = “true dependence” (true)



- **WAW** (Write After Write) = “output dependence” (false)



- **WAR** (Write After Read) = “anti-dependence” (false)



- WAW & WAR are “false”, Can be **totally eliminated** by “renaming”

# Step #1: Register Renaming

- To eliminate register conflicts/hazards
- “Architected” vs “Physical” registers – level of indirection
  - Names:  $r_1, r_2, r_3$
  - Locations:  $p_1, p_2, p_3, p_4, p_5, p_6, p_7$
  - Original mapping:  $r_1 \rightarrow p_1, r_2 \rightarrow p_2, r_3 \rightarrow p_3, p_4 \rightarrow p_7$  are “available”

	MapTable	FreeList	Original insns	Renamed insns																			
Time	<table border="1"><tr><td><math>r_1</math></td><td><math>r_2</math></td><td><math>r_3</math></td></tr><tr><td><math>p_1</math></td><td><math>p_2</math></td><td><math>p_3</math></td></tr><tr><td><math>p_4</math></td><td><math>p_2</math></td><td><math>p_3</math></td></tr><tr><td><math>p_4</math></td><td><math>p_2</math></td><td><math>p_5</math></td></tr><tr><td><math>p_4</math></td><td><math>p_2</math></td><td><math>p_6</math></td></tr></table>	$r_1$	$r_2$	$r_3$	$p_1$	$p_2$	$p_3$	$p_4$	$p_2$	$p_3$	$p_4$	$p_2$	$p_5$	$p_4$	$p_2$	$p_6$	<table border="1"><tr><td><math>p_4, p_5, p_6, p_7</math></td></tr><tr><td><math>p_5, p_6, p_7</math></td></tr><tr><td><math>p_6, p_7</math></td></tr><tr><td><math>p_7</math></td></tr></table>	$p_4, p_5, p_6, p_7$	$p_5, p_6, p_7$	$p_6, p_7$	$p_7$	<hr/> <p>add <math>r_2, r_3 \rightarrow r_1</math></p> <p>sub <math>r_2, r_1 \rightarrow r_3</math></p> <p>mul <math>r_2, r_3 \rightarrow r_3</math></p> <p>div <math>r_1, 4 \rightarrow r_1</math></p> <hr/>	<hr/> <p>add <math>p_2, p_3 \rightarrow p_4</math></p> <p>sub <math>p_2, p_4 \rightarrow p_5</math></p> <p>mul <math>p_2, p_5 \rightarrow p_6</math></p> <p>div <math>p_4, 4 \rightarrow p_7</math></p> <hr/>
$r_1$	$r_2$	$r_3$																					
$p_1$	$p_2$	$p_3$																					
$p_4$	$p_2$	$p_3$																					
$p_4$	$p_2$	$p_5$																					
$p_4$	$p_2$	$p_6$																					
$p_4, p_5, p_6, p_7$																							
$p_5, p_6, p_7$																							
$p_6, p_7$																							
$p_7$																							

- Renaming – conceptually write each register once
  - + Removes **false** dependences
  - + Leaves **true** dependences intact!
- When to reuse a physical register? After overwriting insn done

# Register Renaming Algorithm

---

- Two key data structures:
  - `maptable[architectural_reg] → physical_reg`
  - Free list: allocate (new) & free registers (implemented as a queue)
- Algorithm: at “decode” stage for each instruction:

```
insn.phys_input1 = maptable[insn.arch_input1]
```

```
insn.phys_input2 = maptable[insn.arch_input2]
```

```
insn.old_phys_output = maptable[insn.arch_output]
```

```
new_reg = new_phys_reg()
```


```
maptable[insn.arch_output] = new_reg
```

```
insn.phys_output = new_reg
```


- At “commit”

- Once all older instructions have committed, free register  
`free_phys_reg(insn.old_phys_output)`

# Out-of-order Pipeline



# Step #2: Dynamic Scheduling



- Instructions fetch/decoded/renamed into *Instruction Buffer*
  - Also called “instruction window” or “instruction scheduler”
- Instructions (conceptually) check ready bits every cycle
  - Execute oldest “ready” instruction, set output as “ready”

# Dynamic Scheduling/Issue Algorithm

---

- Data structures:
  - Ready table[phys\_reg] → yes/no (part of “issue queue”)
- Algorithm at “issue” stage (prior to read registers):

foreach instruction:

```
if table[insn.phys_input1] == ready &&
    table[insn.phys_input2] == ready then
        insn is “ready”
select the oldest “ready” instruction
table[insn.phys_output] = ready
```
- Multiple-cycle instructions? (such as loads)
  - For an insn with latency of N, set “ready” bit N-1 cycles in future

# **Register Renaming**

# Register Renaming Algorithm (Simplified)

- Two key data structures:
  - `maptable[architectural_reg] → physical_reg`
  - Free list: allocate (new) & free registers (implemented as a queue)
    - ignore freeing of registers for now
- Algorithm: at “decode” stage for each instruction:

```
insn.phys_input1 = maptable[insn.arch_input1]  
insn.phys_input2 = maptable[insn.arch_input2]
```

```
new_reg = new_phys_reg()  
maptable[insn.arch_output] = new_reg  
insn.phys_output = new_reg
```

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

→ xor p1 ^ p2 →

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

# Renaming example

$\text{xor r1} \wedge \text{r2} \rightarrow \text{r3}$        $\longrightarrow$        $\text{xor p1} \wedge \text{p2} \rightarrow \text{p6}$   
 $\text{add r3} + \text{r4} \rightarrow \text{r4}$   
 $\text{sub r5} - \text{r2} \rightarrow \text{r3}$   
 $\text{addi r3} + 1 \rightarrow \text{r1}$

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

# Renaming example

$\text{xor r1} \wedge \text{r2} \rightarrow \text{r3}$        $\longrightarrow$        $\text{xor p1} \wedge \text{p2} \rightarrow \text{p6}$   
 $\text{add r3} + \text{r4} \rightarrow \text{r4}$   
 $\text{sub r5} - \text{r2} \rightarrow \text{r3}$   
 $\text{addi r3} + 1 \rightarrow \text{r1}$

r1	p1
r2	p2
<b>r3</b>	<b>p6</b>
r4	p4
r5	p5

Map table

p7
p8
p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 →

r1	p1
r2	p2
r3	<b>p6</b>
r4	<b>p4</b>
r5	p5

Map table

p7
p8
p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → **p7**  
sub p8 - p2 → p9  
addi p9 + 1 → p10

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table

<b>p7</b>
p8
p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table

p8
p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub **r5** - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub **p5** - p2 →

r1	p1
<b>r2</b>	<b>p2</b>
r3	p6
r4	p7
<b>r5</b>	<b>p5</b>

Map table

p8
p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → **p8**  
addi p3 + 1 → p1

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table

<b>p8</b>
p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8

r1	p1
r2	p2
<b>r3</b>	<b>p8</b>
r4	p7
r5	p5

Map table

p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi **r3** + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi **p8** + 1 →

r1	p1
r2	p2
<b>r3</b>	<b>p8</b>
r4	p7
r5	p5

Map table

p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → **p9**

r1	p1
r2	p2
r3	p8
r4	p7
r5	p5

Map table

<b>p9</b>
p10

Free-list


# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9


r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table



Free-list

# Out-of-order Pipeline




# **Dynamic Scheduling Mechanisms**

# Dispatch

---

- Put renamed instructions into out-of-order structures
- Re-order buffer (ROB)
  - Holds instructions from Fetch through Commit
- Issue Queue
  - Central piece of scheduling logic
  - Holds instructions from Dispatch through Issue
  - Tracks ready inputs
    - Physical register names + ready bit
    - “AND” the bits to tell if ready



# Dispatch Steps

---

- Allocate Issue Queue (IQ) slot
  - Full? Stall
- Read **ready bits** of inputs
  - 1-bit per physical reg
- Clear **ready bit** of output in table
  - Instruction has not produced value yet
- Write data into Issue Queue (IQ) slot

# Dispatch Example

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

**Ready bits**

p1	y
p2	y
p3	y
p4	y
p5	y
p6	y
p7	y
p8	y
p9	y

**Issue Queue**

Insn	Inp1	R	Inp2	R	Dst	Bday

# Dispatch Example

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

**Ready bits**

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	y
p8	y
p9	y

**Issue Queue**

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	p6	0

# Dispatch Example

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

**Ready bits**

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	n
p8	y
p9	y

**Issue Queue**

<b>Insn</b>	<b>Inp1</b>	<b>R</b>	<b>Inp2</b>	<b>R</b>	<b>Dst</b>	<b>Bday</b>
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1

# Dispatch Example

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

**Ready bits**

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	n
<b>p8</b>	<b>n</b>
p9	y

**Issue Queue**

<b>Insn</b>	<b>Inp1</b>	<b>R</b>	<b>Inp2</b>	<b>R</b>	<b>Dst</b>	<b>Bday</b>
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1
sub	p5	y	p2	y	p8	2

# Dispatch Example

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

**Ready bits**

p1	y
p2	y
p3	y
p4	y
p5	y
p6	n
p7	n
p8	n
p9	n


**Issue Queue**

<b>Insn</b>	<b>Inp1</b>	<b>R</b>	<b>Inp2</b>	<b>R</b>	<b>Dst</b>	<b>Bday</b>
xor	p1	y	p2	y	p6	0
add	p6	n	p4	y	p7	1
sub	p5	y	p2	y	p8	2
addi	p8	n	---	y	p9	3

# Out-of-order pipeline

---

- Execution (out-of-order) stages
- **Select** ready instructions
  - Send for execution
- **Wakeup** dependents



# Dynamic Scheduling/Issue Algorithm

---

- Data structures:
  - Ready table[phys\_reg] → yes/no (part of issue queue)
- Algorithm at “schedule” stage (prior to read registers):

foreach instruction:

```
if table[insn.phys_input1] == ready &&
    table[insn.phys_input2] == ready then
        insn is "ready"
```

select the oldest “ready” instruction

```
table[insn.phys_output] = ready
```

# Issue = Select + Wakeup

---

- **Select** oldest of “ready” instructions
  - “xor” is the oldest ready instruction below
  - “xor” and “sub” are the two oldest ready instructions below
  - Note: may have resource constraints: i.e. load/store/floating point

Insn	Inp1	R	Inp2	R	Dst	Bday	
xor	p1	y	p2	y	p6	0	Ready!
add	p6	n	p4	y	p7	1	
sub	p5	y	p2	y	p8	2	Ready!
addi	p8	n	---	y	p9	3	

# Issue = Select + Wakeup

- Wakeup dependent instructions
  - Search for destination (Dst) in inputs & set “ready” bit
    - Implemented with a special memory array circuit called a Content Addressable Memory (CAM)
  - Also update ready-bit table for future instructions

Insn	Inp1	R	Inp2	R	Dst	Bday
xor	p1	y	p2	y	<b>p6</b>	0
add	<b>p6</b>	<b>y</b>	p4	y	p7	1
sub	p5	y	p2	y	<b>p8</b>	2
addi	<b>p8</b>	<b>y</b>	---	y	p9	3

Ready bits

p1	y
p2	y
p3	y
p4	y
p5	y
<b>p6</b>	<b>y</b>
p7	n
<b>p8</b>	<b>y</b>
p9	n

- For multi-cycle operations (loads, floating point)
  - Wakeup deferred a few cycles
  - Include checks to avoid structural hazards

# Issue


---

- **Select/Wakeup** one cycle
- Dependent instructions execute on back-to-back cycles
  - Next cycle: add/addi are ready:


Insn	Inp1	R	Inp2	R	Dst	Bday
add	p6	y	p4	y	p7	1
addi	p8	y	---	y	p9	3

- Issued instructions are removed from issue queue
  - Free up space for subsequent instructions


# 000 execution (2-wide)




# 000 execution (2-wide)




# 000 execution (2-wide)




# 000 execution (2-wide)



# 000 execution (2-wide)




# 000 execution (2-wide)



# OOO execution (2-wide)

Note similarity  
to in-order



# When Does Register Read Occur?

---

- Our approach: after select, right before execute
  - **Not during in-order part of pipeline, in out-of-order part**
  - Read **physical** register (renamed)
  - Or get value via bypassing (based on physical register name)
  - This is Pentium 4, MIPS R10k, Alpha 21264, IBM Power4, Intel's "Sandy Bridge" (2011)
  - Physical register file may be large
    - Multi-cycle read
- Alternative approach:
  - Read as part of "issue" stage, keep values in Issue Queue
    - At commit, write them back to "architectural register file"
  - Pentium Pro, Core 2, Core i7
  - Simpler, but may be less energy efficient (more data movement)

# **Renaming Revisited**

# Re-order Buffer (ROB)

---

- ROB entry holds all info for recovery/commit
  - **All instructions** & in order
  - Architectural register names, physical register names, insn type
  - Not removed until very last thing ("commit")
- Operation
  - Fetch: insert at tail (if full, stall)
  - Commit: remove from head (if not yet done, stall)
- Purpose: tracking for in-order commit
  - Maintain appearance of in-order execution
  - Needed to support:
    - **Misprediction recovery**
    - **Freeing of physical registers**

# Renaming revisited

---

- Track (or “log”) the “overwritten register” in ROB
  - Free this register at commit
  - Also used to restore the map table on “recovery”
    - Used for branch misprediction recovery

# Register Renaming Algorithm (Full)

---

- Two key data structures:
  - `maptable[architectural_reg] → physical_reg`
  - Free list: allocate (new) & free registers (implemented as a queue)
- Algorithm: at “decode” stage for each instruction:

```
insn.phys_input1 = maptable[insn.arch_input1]
insn.phys_input2 = maptable[insn.arch_input2]
insn.old_phys_output = maptable[insn.arch_output]
new_reg = new_phys_reg()
maptable[insn.arch_output] = new_reg
insn.phys_output = new_reg
```

- At “commit”
  - **Once all older instructions have committed, free register `free_phys_reg(insn.old_phys_output)`**

# Recovery

---

- Completely remove wrong path instructions
  - Flush from IQ
  - Remove from ROB
  - Restore map table to before misprediction
  - Free destination registers
- How to restore map table?
  - Option #1: log-based reverse renaming to recover each instruction
    - Tracks the old mapping to allow it to be reversed
    - Done sequentially for each instruction (slow)
    - See next slides
  - Option #2: checkpoint-based recovery
    - Checkpoint state of maptable and free list each cycle
    - Faster recovery, but requires more state
  - Option #3: hybrid (checkpoint for branches, unwind for others)

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3      →      xor p1 ^ p2 → [ p3 ]  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

r1	p1
r2	p2
r3	<b>p3</b>
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

# Renaming example

$\text{xor r1} \wedge \text{r2} \rightarrow \text{r3}$      $\longrightarrow$      $\text{xor p1} \wedge \text{p2} \rightarrow \text{p6}$     [ p3 ]  
 $\text{add r3} + \text{r4} \rightarrow \text{r4}$   
 $\text{sub r5} - \text{r2} \rightarrow \text{r3}$   
 $\text{addi r3} + 1 \rightarrow \text{r1}$

r1	p1
r2	p2
r3	<b>p6</b>
r4	p4
r5	p5

Map table

p7
p8
p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

xor p1 ^ p2 → p6  
add p6 + p4 →

[ p3 ]  
[ p4 ]

r1	p1
r2	p2
r3	p6
r4	<b>p4</b>
r5	p5

Map table

p7
p8
p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7

[ p3 ]  
[ p4 ]

r1	p1
r2	p2
r3	p6
r4	p7
r5	p5

Map table

p8
p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 →

[ p3 ]  
[ p4 ]  
[ p6 ]

r1	p1
r2	p2
r3	<b>p6</b>
r4	p7
r5	p5

Map table

p8
p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8

[ p3 ]  
[ p4 ]  
[ p6 ]

r1	p1
r2	p2
r3	<b>p8</b>
r4	p7
r5	p5

Map table

p9
p10

Free-list

# Renaming example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1



xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 →

[ p3 ]  
[ p4 ]  
[ p6 ]  
[ p1 ]

r1	<b>p1</b>
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p9
p10

Free-list

# Renaming example


xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

[ p3 ]  
[ p4 ]  
[ p6 ]  
[ p1 ]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table



Free-list

# Recovery Example

**Now, let's use this info. to recover from a branch misprediction**

**bnz r1 loop**

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

**bnz p1, loop**

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

[ ]

[ p3 ]

[ p4 ]

[ p6 ]

[ p1 ]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table



Free-list

# Recovery Example

bnz r1 loop  
xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

bnz p1, loop  
xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → **p9**

[ ]  
[ p3 ]  
[ p4 ]  
[ p6 ]  
[ p1 ]

r1	<b>p1</b>
r2	p2
r3	p8
r4	p7
r5	p5

Map table

<b>p9</b>
p10

Free-list

# Recovery Example

bnz r1 loop  
xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3

bnz p1, loop  
xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → **p8**

[ ]  
[ p3 ]  
[ p4 ]  
[ p6 ]

r1	p1
r2	p2
r3	<b>p6</b>
r4	p7
r5	p5

Map table

<b>p8</b>
p9
p10

Free-list

# Recovery Example

bnz r1 loop  
xor r1 ^ r2 → r3  
add r3 + r4 → r4

bnz p1, loop  
xor p1 ^ p2 → p6  
add p6 + p4 → p7

[ ]  
[ p3 ]  
[ p4 ]

r1	p1
r2	p2
r3	p6
r4	p4
r5	p5

Map table

p7
p8
p9
p10

Free-list

# Recovery Example

bnz r1 loop  
xor r1 ^ r2 → r3

bnz p1, loop  
xor p1 ^ p2 → p6

[ ]  
[ p3 ]

r1	p1
r2	p2
r3	<b>p3</b>
r4	p4
r5	p5

Map table

<b>p6</b>
p7
p8
p9
p10

Free-list

# Recovery Example

bnz r1, loop

bnz p1, loop

[ ]

r1	p1
r2	p2
r3	p3
r4	p4
r5	p5

Map table

p6
p7
p8
p9
p10

Free-list

# Commit

$\text{xor r1} \wedge \text{r2} \rightarrow \text{r3}$   
 $\text{add r3} + \text{r4} \rightarrow \text{r4}$   
 $\text{sub r5} - \text{r2} \rightarrow \text{r3}$   
 $\text{addi r3} + 1 \rightarrow \text{r1}$

$\text{xor p1} \wedge \text{p2} \rightarrow \text{p6}$  [ p3 ]  
 $\text{add p6} + \text{p4} \rightarrow \text{p7}$  [ p4 ]  
 $\text{sub p5} - \text{p2} \rightarrow \text{p8}$  [ p6 ]  
 $\text{addi p8} + 1 \rightarrow \text{p9}$  [ p1 ]

- Commit: instruction becomes **architected state**
  - In-order, only when instructions are finished
    - Free overwritten register (why?)

# Freeing over-written register

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

[ p3 ]  
[ p4 ]  
[ p6 ]  
[ p1 ]

- P3 was r3 **before** xor
- P6 is r3 **after** xor
  - Anything older than xor should read p3
  - Anything younger than xor should read p6 (until another insn writes r3)
- At commit of xor, **no older instructions exist**

# Commit Example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

[ p3 ]  
[ p4 ]  
[ p6 ]  
[ p1 ]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table



Free-list

# Commit Example

xor r1 ^ r2 → r3  
add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

xor p1 ^ p2 → p6  
add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

[ p3 ]  
[ p4 ]  
[ p6 ]  
[ p1 ]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3

Free-list

# Commit Example

add r3 + r4 → r4  
sub r5 - r2 → r3  
addi r3 + 1 → r1

add p6 + p4 → p7  
sub p5 - p2 → p8  
addi p8 + 1 → p9

[ p4 ]  
[ p6 ]  
[ p1 ]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4

Free-list

# Commit Example

sub r5 - r2 → r3  
addi r3 + 1 → r1

sub p5 - p2 → p8  
addi p8 + 1 → p9

[ p6 ]  
[ p1 ]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4
p6

Free-list

# Commit Example

addi r3 + 1 → r1

addi p8 + 1 → p9

[ p1 ]

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4
p6
p1

Free-list

# Commit Example

r1	p9
r2	p2
r3	p8
r4	p7
r5	p5

Map table

p10
p3
p4
p6
p1

Free-list

# Textbook OoO Terminology

Table 3.4. Instruction flow and resources involved in an out-of-order processor

	Step	Resources read	Resources written or utilized
Front-end	Fetch	PC	PC
		Branch predictor	Instruction buffer
		I-cache	
	Decode-rename	Instruction buffer	Decode buffer
		Register map	Register map
			ROB
	Dispatch	Decode buffer	Reservation stations
		Register map	ROB
		Register file (logical and physical)	
Back-end	Issue	Reservation stations	Functional units
			D-cache
	Execute	Functional units	Reservation stations
		D-cache	ROB
			Physical register file
			Branch predictor
Commit	Commit	Store buffer, and so on.	Store buffer, and so on.
		ROB	ROB
		Physical register file	Logical register file
		Store buffer	Register map
			D-cache

# Textbook:Lecture “Map Table”

Textbook	Lecture
instruction buffer	-
decode buffer	-
register map	map table
reservation station	issue queue entry
ROB	ROB
logical register file	register file
physical register file	register file

# Lecture version of Textbook Table 3.4

<b>Step/Stage</b>	<b>resources read</b>	<b>resources written/utilized</b>
Fetch	PC, branch predictor, I\$	PC, ROB
Decode-rename	map table	map table, ROB
Dispatch	ready table	ROB, issue queue
Issue	issue queue, regfile	functional units
Execute	D\$	functional units, issue queue, ROB, D\$, branch predictor, regfile
Commit	ROB	ROB, map table, D\$

# **Dynamic Scheduling Example**

# Dynamic Scheduling Example

---

- The following slides are a detailed but concrete example
- Yet, it contains enough detail to be overwhelming
  - Try not to worry about the details
- Focus on the big picture:

**Hardware can reorder instructions  
to extract instruction-level parallelism**

# Recall: Motivating Example

	0	1	2	3	4	5	6	7	8	9	10	11	12
ld [p1] → p2	F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W	C				
add p2 + p3 → p4	F	Di				I	RR	X	W	C			
xor p4 ^ p5 → p6		F	Di			I	RR	X	W	C			
ld [p7] → p8		F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W	C			

- How would this execution occur cycle-by-cycle?
- Execution latencies assumed in this example:
  - Loads have two-cycle load-to-use penalty
    - Three cycle total execution latency
  - All other instructions have single-cycle execution latency
- “Issue queue”: hold all waiting (un-executed) instructions
  - Holds ready/not-ready status
  - Faster than looking up in ready table each cycle

# Out-of-Order Pipeline – Cycle 0

Map	Table
r1	p8
r2	p7
r3	p6
r4	p5
r5	p4
r6	p3
r7	p2
r8	p1

Ready Table	
p1	yes
p2	yes
p3	yes
p4	yes
p5	yes
p6	yes
p7	yes
p8	yes
p9	---
p10	---
p11	---
p12	---

Reorder Buffer	<b>Insn</b>	<b>To Free</b>	<b>Done?</b>
	ld		no
	add		no

<b>Insn</b>	<b>Src1</b>	<b>R?</b>	<b>Src2</b>	<b>R?</b>	<b>Dest</b>	<b>Bdy</b>

# Out-of-Order Pipeline – Cycle 1a

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di											
add r2 + r3 → r4	F												
xor r4 ^ r5 → r6													
Id [r7] → r4													

Map Table
r1 p8
r2 p9
r3 p6
r4 p5
r5 p4
r6 p3
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 yes
p4 yes
p5 yes
p6 yes
p7 yes
p8 yes
p9 no
p10 ---
p11 ---
p12 ---

Reorder Buffer	Insn	To Free	Done?				
	Id	p7	no				
add			no				
Issue Queue	Insn	Src1	R?	Src2	R?	Dest	Bdy
	Id	p8	yes	---	yes	p9	0

# Out-of-Order Pipeline – Cycle 1b

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di											
add r2 + r3 → r4	F	Di											
xor r4 ^ r5 → r6													
Id [r7] → r4													

Map Table
r1 p8
r2 p9
r3 p6
r4 p10
r5 p4
r6 p3
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 yes
p4 yes
p5 yes
p6 yes
p7 yes
p8 yes
p9 no
p10 no
p11 ---
p12 ---

Reorder Buffer		Insn	To Free	Done?
		Id	p7	no
		add	p5	no

Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes	---	yes	p9	0
add	p9	no	p6	yes	p10	1

# Out-of-Order Pipeline – Cycle 1c

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di											
add r2 + r3 → r4	F	Di											
xor r4 ^ r5 → r6			F										
Id [r7] → r4			F										

Map Table
r1 p8
r2 p9
r3 p6
r4 p10
r5 p4
r6 p3
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 yes
p4 yes
p5 yes
p6 yes
p7 yes
p8 yes
p9 no
p10 no
p11 ---
p12 ---

Reorder Buffer

Insn	To Free	Done?
Id	p7	no
add	p5	no
xor		no
Id		no

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes	---	yes	p9	0
add	p9	no	p6	yes	p10	1

# Out-of-Order Pipeline – Cycle 2a

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I										
add r2 + r3 → r4	F	Di											
xor r4 ^ r5 → r6			F										
Id [r7] → r4			F										

Map Table
r1 p8
r2 p9
r3 p6
r4 p10
r5 p4
r6 p3
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 yes
p4 yes
p5 yes
p6 yes
p7 yes
p8 yes
p9 no
p10 no
p11 ---
p12 ---

Reorder Buffer						
Insn	To Free	Done?				
Id	p7	no				
add	p5	no				
xor		no				
Id		no				

Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes	---	yes	p9	0
add	p9	no	p6	yes	p10	1

# Out-of-Order Pipeline – Cycle 2b

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I										
add r2 + r3 → r4	F	Di											
xor r4 ^ r5 → r6		F	Di										
Id [r7] → r4		F											

Map Table
r1 p8
r2 p9
r3 p6
r4 p10
r5 p4
r6 p11
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 yes
p4 yes
p5 yes
p6 yes
p7 yes
p8 yes
p9 no
p10 no
p11 no
---

Reorder Buffer

Insn	To Free	Done?
Id	p7	no
add	p5	no
xor	p3	no
Id		no

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes		yes	p9	0
add	p9	no	p6	yes	p10	1
xor	p10	no	p4	yes	p11	2

# Out-of-Order Pipeline – Cycle 2c

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I										
add r2 + r3 → r4	F	Di											
xor r4 ^ r5 → r6		F	Di										
Id [r7] → r4		F	Di										

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Reorder Buffer	Insn	To Free	Done?
	Id	p7	no
	add	p5	no
	xor	p3	no
	Id	p10	no

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes		yes	p9	0
add	p9	no	p6	yes	p10	1
xor	p10	no	p4	yes	p11	2
Id	p2	yes	---	yes	p12	3

# Out-of-Order Pipeline – Cycle 3

	0	1	2	3	4	5	6	7	8	9	10	11	12
Id [r1] → r2	F	Di	I	RR									
add r2 + r3 → r4	F	Di											
xor r4 ^ r5 → r6		F	Di										
Id [r7] → r4		F	Di	I									

Map Table

r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Ready Table

p1	yes
p2	yes
p3	yes
p4	yes
p5	yes
p6	yes
p7	yes
p8	yes
p9	no
p10	no
p11	no
p12	no

Reorder Buffer

Insn	To Free	Done?
Id	p7	no
add	p5	no
xor	p3	no
Id	p10	no

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
Id	p8	yes		yes	p9	0
add	p9	no	p6	yes	p10	1
xor	p10	no	p4	yes	p11	2
Id	p2	yes	---	yes	p12	3

# Out-of-Order Pipeline – Cycle 4

	0	1	2	3	4	5	6	7	8	9	10	11	12
ld [r1] → r2	F	Di	I	RR	X								
add r2 + r3 → r4	F	Di											
xor r4 ^ r5 → r6		F	Di										
ld [r7] → r4		F	Di	I	RR								

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Reorder Buffer						
Insn	To Free	Done?				
ld	p7	no				
add	p5	no				
xor	p3	no				
ld	p10	no				

Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
ld	p8	yes		yes	p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	no	p4	yes	p11	2
ld	p2	yes		yes	p12	3

# Out-of-Order Pipeline – Cycle 5a

	0	1	2	3	4	5	6	7	8	9	10	11	12
ld [r1] → r2	F	Di	I	RR	X	M <sub>1</sub>							
add r2 + r3 → r4	F	Di				I							
xor r4 × r5 → r6		F	Di										
ld [r7] → r4		F	Di	I	RR	X							

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Reorder Buffer	Insn	To Free	Done?
	ld	p7	no
	add	p5	no
	xor	p3	no
	ld	p10	no

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
ld	p8	yes		yes	p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
ld	p2	yes		yes	p12	3

# Out-of-Order Pipeline – Cycle 5b

	0	1	2	3	4	5	6	7	8	9	10	11	12
ld [r1] → r2	F	Di	I	RR	X	M <sub>1</sub>							
add r2 + r3 → r4	F	Di				I							
xor r4 × r5 → r6		F	Di										
ld [r7] → r4		F	Di	I	RR	X							

Map Table	Ready Table
r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Reorder Buffer	Insn	To Free	Done?
	ld	p7	no
	add	p5	no
	xor	p3	no
	ld	p10	no

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
ld	p8	yes		yes	p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
ld	p2	yes		yes	p12	3

# Out-of-Order Pipeline – Cycle 6

	0	1	2	3	4	5	6	7	8	9	10	11	12
ld [r1] → r2	F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>						
add r2 + r3 → r4	F	Di				I	RR						
xor r4 × r5 → r6		F	Di					I					
ld [r7] → r4		F	Di	I	RR	X	M <sub>1</sub>						

Map Table
r1 p8
r2 p9
r3 p6
r4 p12
r5 p4
r6 p11
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 yes
p4 yes
p5 yes
p6 yes
p7 yes
p8 yes
p9 yes
p10 yes
p11 yes
p12 yes

Reorder Buffer						
Insn	To Free	Done?				
ld	p7					no
add	p5					no
xor	p3					no
ld	p10					no

Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
ld	p8	yes		yes	p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
ld	p2	yes		yes	p12	3

# Out-of-Order Pipeline – Cycle 7

	0	1	2	3	4	5	6	7	8	9	10	11	12
ld [r1] → r2	F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W					
add r2 + r3 → r4	F	Di				I	RR	X					
xor r4 × r5 → r6		F	Di				I	RR					
ld [r7] → r4		F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>					

Map Table
r1 p8
r2 p9
r3 p6
r4 p12
r5 p4
r6 p11
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 yes
p4 yes
p5 yes
p6 yes
p7 yes
p8 yes
p9 yes
p10 yes
p11 yes
p12 yes

Reorder Buffer	Insn	To Free	Done?
	ld	p7	yes
add	p5		no
xor	p3		no
ld	p10		no

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
ld	p8	yes		yes	p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
ld	p2	yes		yes	p12	3

# Out-of-Order Pipeline – Cycle 8a

	0	1	2	3	4	5	6	7	8	9	10	11	12
ld [r1] → r2	F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W	C				
add r2 + r3 → r4	F	Di				I	RR	X					
xor r4 × r5 → r6		F	Di				I	RR					
ld [r7] → r4		F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>					

Map Table

r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Ready Table

p1	yes
p2	yes
p3	yes
p4	yes
p5	yes
p6	yes
p7	---
p8	yes
p9	yes
p10	yes
p11	yes
p12	yes

Reorder Buffer

Insn	To Free	Done?
ld	p7	yes
add	p5	no
xor	p3	no
ld	p10	no

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
ld	p8	yes		yes	p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
ld	p2	yes		yes	p12	3

# Out-of-Order Pipeline – Cycle 8b

	0	1	2	3	4	5	6	7	8	9	10	11	12
ld [r1] → r2	F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W	C				
add r2 + r3 → r4	F	Di				I	RR	X	W				
xor r4 × r5 → r6		F	Di				I	RR	X				
ld [r7] → r4		F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W				

Map Table

r1	p8
r2	p9
r3	p6
r4	p12
r5	p4
r6	p11
r7	p2
r8	p1

Ready Table

p1	yes
p2	yes
p3	yes
p4	yes
p5	yes
p6	yes
p7	---
p8	yes
p9	yes
p10	yes
p11	yes
p12	yes

Reorder Buffer

Insn	To Free	Done?
ld	p7	yes
add	p5	yes
xor	p3	no
ld	p10	yes

Issue Queue

Insn	Src1	R?	Src2	R?	Dest	Bdy
ld	p8	yes		yes	p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
ld	p2	yes		yes	p12	3

# Out-of-Order Pipeline – Cycle 9a

	0	1	2	3	4	5	6	7	8	9	10	11	12
ld [r1] → r2	F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W	C				
add r2 + r3 → r4	F	Di				I	RR	X	W	C			
xor r4 × r5 → r6		F	Di			I	RR	X					
ld [r7] → r4		F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W				

Map Table
r1 p8
r2 p9
r3 p6
r4 p12
r5 p4
r6 p11
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 yes
p4 yes
p5 ---
p6 yes
p7 ---
p8 yes
p9 yes
p10 yes
p11 yes
p12 yes

Reorder Buffer						
Insn	To Free	Done?				
ld	p7	yes				
add	p5	yes				
xor	p3	no				
ld	p10	yes				

Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
ld	p8	yes			p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
ld	p2	yes			p12	3

# Out-of-Order Pipeline – Cycle 9b

	0	1	2	3	4	5	6	7	8	9	10	11	12
ld [r1] → r2	F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W	C				
add r2 + r3 → r4	F	Di				I	RR	X	W	C			
xor r4 × r5 → r6		F	Di			I	RR	X	W				
ld [r7] → r4		F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W				

Map Table
r1 p8
r2 p9
r3 p6
r4 p12
r5 p4
r6 p11
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 yes
p4 yes
p5 ...
p6 yes
p7 ...
p8 yes
p9 yes
p10 yes
p11 yes
p12 yes

Reorder Buffer						
Insn	To Free	Done?				
ld	p7					yes
add	p5					yes
xor	p3					yes
ld	p10					yes

Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
ld	p8	yes			p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
ld	p2	yes			p12	3

# Out-of-Order Pipeline – Cycle 10

	0	1	2	3	4	5	6	7	8	9	10	11	12
ld [r1] → r2	F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W	C				
add r2 + r3 → r4	F	Di				I	RR	X	W	C			
xor r4 × r5 → r6		F	Di			I	RR	X	W	C			
ld [r7] → r4		F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W	C			

Map Table
r1 p8
r2 p9
r3 p6
r4 p12
r5 p4
r6 p11
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 ---
p4 yes
p5 ---
p6 yes
p7 ---
p8 yes
p9 yes
p10 ---
p11 yes
p12 yes

Reorder Buffer						
Insn	To Free	Done?				
ld	p7					yes
add	p5					yes
xor	p3					yes
ld	p10					yes

Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
ld	p8	yes			p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
ld	p2	yes			p12	3

# Out-of-Order Pipeline – Done!

	0	1	2	3	4	5	6	7	8	9	10	11	12
ld [r1] → r2	F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W	C				
add r2 + r3 → r4	F	Di				I	RR	X	W	C			
xor r4 × r5 → r6		F	Di			I	RR	X	W	C			
ld [r7] → r4		F	Di	I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W		C		

Map Table
r1 p8
r2 p9
r3 p6
r4 p12
r5 p4
r6 p11
r7 p2
r8 p1

Ready Table
p1 yes
p2 yes
p3 ...
p4 yes
p5 ...
p6 yes
p7 ...
p8 yes
p9 yes
p10 ...
p11 yes
p12 yes

Reorder Buffer						
Insn	To Free	Done?				
ld	p7					yes
add	p5					yes
xor	p3					yes
ld	p10					yes


Issue Queue						
Insn	Src1	R?	Src2	R?	Dest	Bdy
ld	p8	yes			p9	0
add	p9	yes	p6	yes	p10	1
xor	p10	yes	p4	yes	p11	2
ld	p2	yes			p12	3

# Handling Memory Operations


# Recall: Types of Dependencies

---


- RAW (Read After Write) = “true dependence”



- WAW (Write After Write) = “output dependence”



- WAR (Write After Read) = “anti-dependence”



- WAW & WAR are “false”, Can be **totally eliminated** by “renaming”

# Also Have Dependencies via Memory

---

- If value in “r2” and “r3” is the same...
- RAW (Read After Write) – True dependency

st r1 → [r2]

...  
ld [r3] → r4

- WAW (Write After Write)

st r1 → [r2]

...  
st r4 → [r3]

- WAR (Write After Read)

ld [r2] → r1

...  
st r4 → [r3]

WAR/WAW are “false dependencies”

- But can't rename memory in same way as registers
  - Why? Addresses are not known at rename
  - Need to use other tricks

# Let's Start with Just Stores

---

- Stores: Write data cache, not registers
  - Can we rename memory?
- No (at least not easily)
  - Cache writes unrecoverable
- Solution: write stores into cache only when certain
  - When are we certain? At "commit"

# Handling Stores

	0	1	2	3	4	5	6	7	8	9	10	11	12
mul p1 * p2 → p3	F	Di	I	RR	X <sub>1</sub>	X <sub>2</sub>	X <sub>3</sub>	X <sub>4</sub>	W	C			
jump-not-zero p3	F	Di					I	RR	X	W	C		
st p5 → [p3+4]		F	Di				I	RR	X	M	W	C	
st p4 → [p6+8]		F	Di	I?									

- Can “st p4 → [p6+8]” issue in cycle 3?
  - Its register inputs are ready...

# Problem #1: Out-of-Order Stores

	0	1	2	3	4	5	6	7	8	9	10	11	12
mul p1 * p2 → p3	F	Di	I	RR	X <sub>1</sub>	X <sub>2</sub>	X <sub>3</sub>	X <sub>4</sub>	W	C			
jump-not-zero p3	F	Di					I	RR	X	W	C		
st p5 → [p3+4]		F	Di				I	RR	X	M	W	C	
st p4 → [p6+8]		F	Di	I?	RR	X	M	W				C	

- Can “st p4 → [p6+8]” write the cache in cycle 6?
  - “st p5 → [p3+4]” has not yet executed
- What if p3+4 == p6+8?
  - The two stores write the same address! WAW dependency!
  - Not known until their “X” stages (cycle 5 & 8)
- Unappealing solution: all stores execute in-order
- We can do better...

# Problem #2: Speculative Stores

	0	1	2	3	4	5	6	7	8	9	10	11	12
mul p1 * p2 → p3	F	Di	I	RR	X <sub>1</sub>	X <sub>2</sub>	X <sub>3</sub>	X <sub>4</sub>	W	C			
jump-not-zero p3	F	Di					I	RR	X	W	C		
st p5 → [p3+4]		F	Di				I	RR	X	M	W	C	
st p4 → [p6+8]		F	Di	I?	RR	X	M	W					C

- Can “st p4 → [p6+8]” write the cache in cycle 6?
  - Store is still “speculative” at this point
- What if “jump-not-zero” is mis-predicted?
  - Not known until its “X” stage (cycle 8)
- How does it “undo” the store once it hits the cache?
  - Answer: it can’t; stores write the cache only at **commit**
  - Guaranteed to be non-speculative at that point

# Store Queue (SQ)

---

- Solves two problems
  - Allows for recovery of speculative stores
  - Allows out-of-order stores
- Store Queue (SQ)
  - **At dispatch, each store is given a slot in the Store Queue**
  - First-in-first-out (FIFO) queue
  - Each entry contains: “address”, “value”, and “bday”
- Operation:
  - Dispatch (in-order): allocate entry in SQ (stall if full)
  - Execute (out-of-order): write store value into store queue
  - Commit (in-order): read value from SQ and write into data cache
  - Branch recovery: remove entries from the store queue
- Also solves problems with loads

# Store Queue Operation

	0	1	2	3	4	5	6	7	8	9	10	11	12
fdiv p1 / p2 → p9	F	Di	I	RR	X <sub>1</sub>	X <sub>2</sub>	X <sub>3</sub>	X <sub>4</sub>	X <sub>5</sub>	X <sub>6</sub>	W	C	
st p4 → [p5+4]	F	Di	I	RR	X	SQ						C	
st p3 → [p6+8]		F	Di	I	R R	X	SQ						C

- Stores write to SQ, not M
  - similar to register renaming, where we allocated a new physical register for each insn
- What if the fdiv triggers a /0 exception?

# Memory Forwarding

	0	1	2	3	4	5	6	7	8	9	10	11	12
fdiv p1 / p2 → p9	F	Di	I	RR	X <sub>1</sub>	X <sub>2</sub>	X <sub>3</sub>	X <sub>4</sub>	X <sub>5</sub>	X <sub>6</sub>	W	C	
st p4 → [p5+4]	F	Di	I	RR	X	SQ						C	
st p3 → [p6+8]		F	Di	I	RR	X	SQ						C
ld [p7] → p8		F	Di	I?	RR	X	M <sub>1</sub>	M <sub>2</sub>	W				C

- Can “ld [p7] → p8” issue and begin execution?
  - Why or why not?

# Memory Forwarding

	0	1	2	3	4	5	6	7	8	9	10	11	12
fdiv p1 / p2 → p9	F	Di	I	RR	X <sub>1</sub>	X <sub>2</sub>	X <sub>3</sub>	X <sub>4</sub>	X <sub>5</sub>	X <sub>6</sub>	W	C	
st p4 → [p5+4]	F	Di	I	RR	X	SQ						C	
st p3 → [p6+8]		F	Di	I	RR	X	SQ						C
ld [p7] → p8		F	Di	I?	RR	X	M <sub>1</sub>	M <sub>2</sub>	W				C

- Can “ld [p7] → p8” issue and begin execution?
  - Why or why not?
- If the load reads from either of the stores’ addresses...
  - Load must get correct value, but stores don’t write cache until commit...
- Solution: “memory forwarding”
  - Load also searches the Store Queue (in parallel with cache access)
  - Conceptually like register bypassing, but different implementation
    - Why? Addresses unknown until execute


# Problem #3: WAR Hazards

	0	1	2	3	4	5	6	7	8	9	10	11	12
mul p1 * p2 → p3	F	Di	I	RR	X <sub>1</sub>	X <sub>2</sub>	X <sub>3</sub>	X <sub>4</sub>	W	C			
jump-not-zero p3	F	Di						I	RR	X	W	C	
ld [p3+4] → p5		F	Di					I	RR	X	M <sub>1</sub>	M <sub>2</sub>	W C
st p4 → [p6+8]		F	Di	I	RR	X	SQ						C

- What if “ $p3+4 == p6 + 8$ ”?
  - WAR: need to make sure that load doesn’t read store’s result
  - Need to get values based on “program order” not “execution order”
- Bad solution: require all stores/loads to execute in-order
- Good solution: add “age” fields to store queue (SQ)
  - Loads read from **youngest older matching** store
  - Another reason the SQ is a FIFO queue

# Memory Forwarding via Store Queue

- Store Queue (SQ)
  - Holds all in-flight stores
  - CAM: searchable by address
  - Age logic: determine youngest matching store older than load
- Store rename/dispatch
  - Allocate entry in SQ
- Store execution
  - Update SQ
    - Address + Data
- Load execution
  - Search SQ identify youngest older matching store
    - Match? Read SQ
    - No Match? Read cache



# Store Queue (SQ)

---

- On load execution, select the store that is:
  - To same address as load
  - Older than the load (before the load in program order)
- Of these matching stores, select the youngest
  - The store to the same address that immediately precedes the load

# When Can Loads Execute?

	0	1	2	3	4	5	6	7	8	9	10	11	12
mul p1 * p2 → p3	F	Di	I	RR	X <sub>1</sub>	X <sub>2</sub>	X <sub>3</sub>	X <sub>4</sub>	W	C			
jump-not-zero p3	F	Di					I	RR	X	W	C		
st p5 → [p3+4]		F	Di				I	RR	X	SQ	C		
ld [p6+8] → p7		F	Di	I?	RR	X	M <sub>1</sub>	M <sub>2</sub>	W				C

- Can “ld [p6+8] → p7” issue in cycle 3
  - Why or why not?

# When Can Loads Execute?

	0	1	2	3	4	5	6	7	8	9	10	11	12
mul p1 * p2 → p3	F	Di	I	RR	X <sub>1</sub>	X <sub>2</sub>	X <sub>3</sub>	X <sub>4</sub>	W	C			
jump-not-zero p3	F	Di					I	RR	X	W	C		
st p5 → [p3+4]		F	Di				I	RR	X	SQ	C		
ld [p6+8] → p7		F	Di	I?	RR	X	M <sub>1</sub>	M <sub>2</sub>	W				C

- Aliasing! Does  $p3+4 == p6+8$ ?
  - If no, load should get value from memory
    - **Can it start to execute?**
  - If yes, load should get value from store
    - By reading the store queue?
    - **But the value isn't put into the store queue until cycle 9**
- **Key challenge:** don't know addresses until execution!
  - One solution: require all loads to wait for all earlier (prior) stores

# Conservative Load Scheduling

---

- Conservative load scheduling:
  - All older stores have executed
    - Some architectures: split store address / store data
      - Only requires knowing addresses (not the store values)
  - Advantage: always safe
  - Disadvantage: performance (limits ILP)

# Conservative Load Scheduling

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
ld [p1] → p4	F	Di	I	Rr	X	M <sub>1</sub>	M <sub>2</sub>	W	C							
ld [p2] → p5	F	Di	I	Rr	X	M <sub>1</sub>	M <sub>2</sub>	W	C							
add p4, p5 → p6		F	Di			I	Rr	X	W	C						
st p6 → [p3]		F	Di			I	Rr	X	SQ	C						
ld [p1+4] → p7			F	Di			I	Rr	X	M <sub>1</sub>	M <sub>2</sub>	W	C			
ld [p2+4] → p8			F	Di			I	Rr	X	M <sub>1</sub>	M <sub>2</sub>	W	C			
add p7, p8 → p9				F	Di					I	Rr	X	W	C		
st p9 → [p3+4]				F	Di						I	Rr	X	SQ	C	

Conservative load scheduling: can't issue ld [p1+4] until cycle 7!

Might as well be an in-order machine on this example

Can we do better? How?

# Optimistic Load Scheduling

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
ld [p1] → p4	F	Di	I	Rr	X	M <sub>1</sub>	M <sub>2</sub>	W	C							
ld [p2] → p5	F	Di	I	Rr	X	M <sub>1</sub>	M <sub>2</sub>	W	C							
add p4, p5 → p6		F	Di			I	Rr	X	W	C						
st p6 → [p3]		F	Di				I	Rr	X	SQ	C					
ld [p1+4] → p7			F	Di	I	Rr	X	M <sub>1</sub>	M <sub>2</sub>	W	C					
ld [p2+4] → p8			F	Di	I	Rr	X	M <sub>1</sub>	M <sub>2</sub>	W		C				
add p7, p8 → p9				F	Di			I	Rr	X	W	C				
st p9 → [p3+4]				F	Di				I	Rr	X	SQ	C			

**Optimistic load scheduling: can actually benefit from out-of-order!**  
**Let's speculate!**


# Load Speculation

---

- Speculation requires three things.....
  - 1. When do we speculate?
  - 2. How do we detect a mis-speculation?
  - 3. How do we recover from mis-speculations?
    - Squash offending load and all newer insns
    - Similar to branch mis-prediction recovery

# Load Queue

- Detects load ordering violations
- Load execution: Write address into LQ
  - Also note any store forwarded from
- Store execution: Search LQ
  - Younger load with same addr?
    - Did younger load forward from younger store? [See slide 149]



# Store Queue + Load Queue

---

- Store Queue: handles forwarding, allows OoO stores
  - Entry per store (allocated @ dispatch, deallocated @ commit)
  - Written by stores (@ execute)
  - Searched by loads (@ execute)
  - Read from SQ to write data cache (@ commit)
- Load Queue: detects ordering violations
  - Entry per load (allocated @ dispatch, deallocated @ commit)
  - Written by loads (@ execute)
  - Searched by stores (@ execute)
- Both together
  - Allows aggressive load scheduling
  - Stores don't constrain load execution

# Optimistic Load Scheduling Problem

---

- Allows loads to issue before older stores
  - Increases ILP
    - + Good: When no conflict, increases performance
    - Bad: Conflict => squash => worse performance than waiting
- Can we have our cake AND eat it too?

# Predictive Load Scheduling

---

- Predict which loads must wait for stores
- Fool me once, shame on you-- fool me twice?
  - Loads default to aggressive
  - Keep table of load PCs that have been caused squashes
    - Schedule these conservatively
  - + Simple predictor
  - Makes “bad” loads wait for **all** older stores
- More complex predictors used in practice
  - Predict which stores loads should wait for
  - “Store Sets” paper

# **Load/Store Queue Examples**

# Initial State

(Stores to different addresses)

1. St p1 → [p2]
2. St p3 → [p4]
3. Ld [p5] → p6

RegFile		Load Queue		Store Queue	
		Bdy	Addr	Bdy	Addr
p1	5				
p2	10				
p3	0				
p4	9				
p5	20				
p6	0				
p5	10				
p6	0				
p6	---				
p7	---				
p8	Cache				
Cache		Addr		Val	
100		13			
200		17			

RegFile		Load Queue		Store Queue	
		Bdy	Addr	Bdy	Addr
p1	5				
p2	10				
p3	0				
p4	9				
p5	20				
p6	0				
p5	10				
p6	0				
p6	---				
p7	---				
p8	Cache				
Cache		Addr		Val	
100		13			
200		17			

RegFile		Load Queue		Store Queue	
		Bdy	Addr	Bdy	Addr
p1	5				
p2	10				
p3	0				
p4	9				
p5	20				
p6	0				
p5	10				
p6	0				
p6	---				
p7	---				
p8	Cache				
Cache		Addr		Val	
100		13			
200		17			

# Good Interleaving

(Shows importance of address check)

1. St p1 → [p2]
2. St p3 → [p4]
3. Ld [p5] → p6

1. St p1 → [p2]

RegFile		Load Queue		Store Queue	
		Bdy	Addr	Bdy	Addr
p1	5				
p2	10				
p3	0				
p4	9				
p5	20				
p6	0				
p5	10				
p6	0				
p5	1	100	5		
p6	---				
p7	---				
p8	Cache				
Cache		Addr		Val	
100	13				
200	17				

2. St p3 → [p4]

RegFile		Load Queue		Store Queue	
		Bdy	Addr	Bdy	Addr
p1	5				
p2	10				
p3	0				
p4	9				
p5	20				
p6	0				
p5	10				
p6	0				
p5	1	100	5		
p6	---				
p7	---				
p8	Cache				
Cache		Addr		Val	
100	13				
200	17				

3. Ld [p5] → p6

RegFile		Load Queue		Store Queue	
		Bdy	Addr	Bdy	Addr
p1	5				
p2	10				
p3	0				
p4	9				
p5	20				
p6	0				
p5	10				
p6	0				
p5	1	100	5		
p6	5				
p7	---				
p8	Cache				
Cache		Addr		Val	
100	13				
200	17				

# Different Initial State

(All to same address)

1. St p1 → [p2]
2. St p3 → [p4]
3. Ld [p5] → p6

RegFile		Load Queue		Store Queue		
		Bdy	Addr	Bdy	Addr	Val
p1	5					
p2	10					
p3	0					
p4	9					
p5	10					
p6	0					
p7	---					
p8	Cache					
Cache		Addr	Val	Addr	Val	
100		13		100	13	
200		17		200	17	

RegFile		Load Queue		Store Queue		
		Bdy	Addr	Bdy	Addr	Val
p1	5					
p2	10					
p3	0					
p4	9					
p5	10					
p6	0					
p7	---					
p8	Cache					
Cache		Addr	Val	Addr	Val	
100		13		100	13	
200		17		200	17	

RegFile		Load Queue		Store Queue		
		Bdy	Addr	Bdy	Addr	Val
p1	5					
p2	10					
p3	0					
p4	9					
p5	10					
p6	0					
p7	---					
p8	Cache					
Cache		Addr	Val	Addr	Val	
100		13		100	13	
200		17		200	17	

# Good Interleaving #1

(Program Order)

1. St p1 → [p2]

RegFile		Load Queue		Store Queue		Cache	
		Bdy	Addr	Bdy	Addr	Addr	Val
p1	5						
p2	10						
p3	0						
p4	9						
p5	10						
p6	0						
p5	10						
p6	0						
p5	1	100	5				
p6	---						
p7							
p8	Cache						

2. St p3 → [p4]

RegFile		Load Queue		Store Queue		Cache	
		Bdy	Addr	Bdy	Addr	Addr	Val
p1	5						
p2	10						
p3	0						
p4	9						
p5	10						
p6	0						
p5	10						
p6	0						
p5	1	100	5				
p6	---						
p7							
p8	Cache						

3. Ld [p5] → p6

RegFile		Load Queue		Store Queue		Cache	
		Bdy	Addr	Bdy	Addr	Addr	Val
p1	5						
p2	10						
p3	0						
p4	9						
p5	10						
p6	0						
p5	10						
p6	0						
p5	1	100	5				
p6	9						
p7							
p8	Cache						

1. St p1 → [p2]
2. St p3 → [p4]
3. Ld [p5] → p6

# Good Interleaving #2

(Stores reordered)

2. St p3 → [p4]

	RegFile	Load Queue		
	Bdy	Addr		
p1	5			
p2	10			
p3	9			
p4	10			
p5	10			
p6	0			
p7	---			
p8	Cache			
	Bdy	Addr	Val	
p5	1			
p6	2	100	9	
	Addr	Val		
p8	100	13		
	200	17		


1. St p1 → [p2]

	RegFile	Load Queue		
	Bdy	Addr		
p1	5			
p2	10			
p3	9			
p4	10			
p5	10			
p6	0			
p7	---			
p8	Cache			
	Bdy	Addr	Val	
p5	1			
p6	2	100	9	
	Addr	Val		
p8	100	13		
	200	17		

3. Ld [p5] → p6

	RegFile	Load Queue		
	Bdy	Addr		
p1	5			
p2	10			
p3	9			
p4	10			
p5	10			
p6	0			
p7	---			
p8	Cache			
	Bdy	Addr	Val	
p5	10			
p6	0			
p7	9			
p8	2	100	9	
	Addr	Val		
p8	100	13		
	200	17		

1. St p1 → [p2]
2. St p3 → [p4]
3. Ld [p5] → p6




# Bad Interleaving #1


(Load reads the cache)

1. St p1 → [p2]
2. St p3 → [p4]
3. Ld [p5] → p6

3. Ld [p5] → p6



2. St p3 → [p4]



# Bad Interleaving #2

(Load gets value from wrong store)

1. St p1 → [p2]
2. St p3 → [p4]
3. Ld [p5] → p6

1. St p1 → [p2]

RegFile		Load Queue			Store Queue	
		Bdy	Addr	Bdy of Fwd. Store		
p1	5					
p2	10					
p3	9					
p4	10					
p5	10					
p6	---					
p7	---					
p8	Cache					
Addr		Val				
100	13	100	13			
200	17	200	17			

3. Ld [p5] → p6

RegFile		Load Queue			Store Queue	
		Bdy	Addr	Bdy of Fwd. Store		
p1	5					
p2	10					
p3	9					
p4	10					
p5	10					
p6	5					
p7	---					
p8	Cache					
Bdy		Addr			Val	
100	100	1	100	1	5	
200	2					

2. St p3 → [p4]

RegFile		Load Queue			Store Queue	
		Bdy	Addr	Bdy of Fwd. Store		
p1	5					
p2	10					
p3	9					
p4	10					
p5	10					
p6	5					
p7	---					
p8	Cache					
Bdy		Addr			Val	
100	100	1	100	1	5	
200	2					

# Bad/Good Interleaving

(Load gets value from correct store, but does it work?)

1. St p1 → [p2]
2. St p3 → [p4]
3. Ld [p5] → p6

2. St p3 → [p4]

RegFile		Load Queue		
		Bdy	Addr	Bdy of Fwd. Store
p1	5			
p2	10			
p3	9			
p4	10			
p5	10			
p6	0			
p7	---			
p8	Cache			
Store Queue				
	Bdy	Addr	Val	
p5	10			
p6	0			
p6	2	100	9	
Cache				
	Addr	Val		
p8	100	13		
p8	200	17		

3. Ld [p5] → p6

RegFile		Load Queue		
		Bdy	Addr	Bdy of Fwd. Store
p1	5			
p2	10			
p3	9			
p4	10			
p5	10			
p6	0			
p6	10	100	2	
Store Queue				
	Bdy	Addr	Val	
p5	10			
p6	0			
p6	1			
Cache				
	Addr	Val		
p8	100	13		
p8	200	17		

1. St p1 → [p2]

RegFile		Load Queue		
		Bdy	Addr	Bdy of Fwd. Store
p1	5			
p2	10			
p3	9			
p4	10			
p5	10			
p6	0			
p6	3	100	2	
Store Queue				
	Bdy	Addr	Val	
p5	10			
p6	0			
p6	1			
Cache				
	Addr	Val		
p8	100	5		
p8	200	9		
Cache				
	Addr	Val		
p8	100	13		
p8	200	17		

# **Out-of-Order: Benefits & Challenges**

# Dynamic Scheduling Operation (Recap)

---


- Dynamic scheduling
  - Totally in the hardware (not visible to software)
  - Also called “out-of-order execution” (OoO)
- Fetch many instructions into instruction window
  - Use branch prediction to speculate past (multiple) branches
  - Flush pipeline on branch misprediction
- Rename registers to avoid false dependencies
- Execute instructions as soon as possible
  - Register dependencies are known
  - Handling memory dependencies is harder
- “Commit” instructions in order
  - Anything strange happens before commit, just flush the pipeline

# Haswell Buffer Sizes

Extract more parallelism in every generation

	Nehalem	Sandy Bridge	Haswell
Out-of-order Window	128	168	192
In-flight Loads	48	64	72
In-flight Stores	32	36	42
Scheduler Entries	36	54	60
Integer Register File	N/A	160	168
FP Register File	N/A	144	168
Allocation Queue	28/thread	28/thread	56

# Haswell Core at a Glance



## Next generation branch prediction

- Improves performance *and* saves wasted work

## Improved front-end

- Initiate TLB and cache misses speculatively
- Handle cache misses in parallel to hide latency
- Leverages improved branch prediction

## Deeper buffers

- Extract more instruction parallelism
- More resources when running a single thread

## More execution units, shorter latencies

- Power down when not in use

## More load/store bandwidth


- Better prefetching, better cache line split latency & throughput, double L2 bandwidth
- New modes save power without losing performance

## No pipeline growth

- Same branch misprediction latency
- Same L1/L2 cache latency

# Haswell Execution Unit Overview

## Unified Reservation Station



# Core Cache Size/Latency/Bandwidth

Metric	Nehalem	Sandy Bridge	Haswell
L1 Instruction Cache	32K, 4-way	32K, 8-way	32K, 8-way
L1 Data Cache	32K, 8-way	32K, 8-way	32K, 8-way
Fastest Load-to-use	4 cycles	4 cycles	4 cycles
Load bandwidth	16 Bytes/cycle	32 Bytes/cycle (banked)	64 Bytes/cycle
Store bandwidth	16 Bytes/cycle	16 Bytes/cycle	32 Bytes/cycle
L2 Unified Cache	256K, 8-way	256K, 8-way	256K, 8-way
Fastest load-to-use	10 cycles	11 cycles	11 cycles
Bandwidth to L1	32 Bytes/cycle	32 Bytes/cycle	64 Bytes/cycle
L1 Instruction TLB	4K: 128, 4-way 2M/4M: 7/thread	4K: 128, 4-way 2M/4M: 8/thread	4K: 128, 4-way 2M/4M: 8/thread
L1 Data TLB	4K: 64, 4-way 2M/4M: 32, 4-way 1G: fractured	4K: 64, 4-way 2M/4M: 32, 4-way 1G: 4, 4-way	4K: 64, 4-way 2M/4M: 32, 4-way 1G: 4, 4-way
L2 Unified TLB	4K: 512, 4-way	4K: 512, 4-way	4K+2M shared: 1024, 8-way

All caches use 64-byte lines

# OoO Execution is all around us

---

- Joe's phone's CPU: Qualcomm Krait 400 processor
  - based on ARM Cortex A15 processor
  - **out-of-order** 2.5GHz quad-core
  - 3-wide fetch/decode
  - 4-wide issue
  - 11-stage integer pipeline
  - 28nm process technology
  - 4/4KB DM L1\$, 16/16KB 4-way SA L2\$, 2MB 8-way SA L3\$

# Out of Order: Benefits

---

- Allows speculative re-ordering
  - Loads / stores
  - Branch prediction to look past branches
- Done by hardware
  - Compiler may want different schedule for different hw configs
  - Hardware has only its own configuration to deal with
- Schedule can change due to cache misses
- **Memory-level parallelism**
  - Executes “around” cache misses to find independent instructions
  - Finds and initiates independent misses, reducing memory latency
    - Especially good at hiding L2 hits (~12 cycles in Core i7)

# Challenges for Out-of-Order Cores

---

- Design complexity
  - More complicated than in-order? Certainly!
  - But, we have managed to overcome the design complexity
- Clock frequency
  - Can we build a “high ILP” machine at high clock frequency?
  - Yep, with some additional pipe stages, clever design
- Limits to (efficiently) scaling the window and ILP
  - Large physical register file
  - Fast register renaming/wakeup/select/load queue/store queue
    - Active areas of micro-architectural research
  - Branch & memory depend. prediction (limits effective window size)
    - 95% branch mis-prediction: 1 in 20 branches, or 1 in 100 insn.
  - Plus all the issues of building “wide” in-order superscalar
- Power efficiency
  - Today, even mobile phone chips are out-of-order cores


# Meltdown and Spectre

Vulnerabilities in modern computers leak passwords and sensitive data.

---


- Meltdown
  - It breaks the most fundamental isolation between user applications and the operating system. This attack allows a program to access the memory, and thus also the secrets, of other programs and the operating system.
- Spectre
  - It breaks the isolation between different applications. It allows an attacker to trick error-free programs, which follow best practices, into leaking their secrets. In fact, the safety checks of said best practices actually increase the attack surface and may make applications more susceptible to Spectre

# Meltdown



```
1 raise_exception();
2 // the line below is never reached
3 access(probe_array[data * 4096]);
```

Listing 1: A toy example to illustrate side-effects of out-of-order execution.



```
1 ; rcx = kernel address, rbx = probe array
2 xor rax, rax
3 retry:
4 mov al, byte [rcx]
5 shl rax, 0xc
6 jz retry
7 mov rbx, qword [rbx + rax]
```

Listing 2: The core of Meltdown. An inaccessible kernel address is moved to a register, raising an exception. Subsequent instructions are executed out of order before the exception is raised, leaking the data from the kernel address through the indirect memory access.


# Redux: HW vs. SW Scheduling

---

- Static scheduling
  - Performed by compiler, limited in several ways
- Dynamic scheduling
  - Performed by the hardware, overcomes limitations
- Static limitation → dynamic mitigation
  - Number of registers in the ISA → register renaming
  - Scheduling scope → branch prediction & speculation
  - Inexact memory aliasing information → speculative memory ops
  - Unknown latencies of cache misses → execute when ready
- Which to do? **Compiler does what it can, hardware the rest**
  - Why? dynamic scheduling needed to sustain more than 2-way issue
  - **Helps with hiding memory latency** (execute around misses)
  - Intel Core i7 is four-wide execute w/ scheduling window of 100+
  - Even mobile phones have dynamically scheduled cores (ARM A9, A15)

# Summary: Scheduling

---



- Code scheduling
  - To reduce pipeline stalls
  - To increase ILP (insn-level parallelism)
- Static scheduling by the compiler
  - Approach & limitations
- Dynamic scheduling in hardware
  - Register renaming
  - Instruction selection
  - Handling memory operations
- Up next: multicore