

Modern Computer Architecture

with sources that included UPenn & University of Wisconsin slides
by Joe Devietti, Milo Martin & Amir Roth, Mark Hill, Guri Sohi, Jim Smith, and David Wood

Schedule

- Introduction and Transistors
- Parallel computing (Isaac D. Scherson, University of California, Irvine, in October)
- ISAs
- Performance
- Pipelining Basic
- Branch Prediction
- Caches
- Virtual Memory
- Out-of-Order Execution
- Multicore multi-thread
- Vectors/GPUs for data parallelism.

How to Compute This Fast?

- Performing the **same** operations on **many** data items
 - Example: SAXPY

```
for (I = 0; I < 1024; I++) {           L1: ldf [X+r1]->f1 // I is in r1
    Z[I] = A*X[I] + Y[I];             mulf f0,f1->f2 // A is in f0
}                                         ldf [Y+r1]->f3
                                         addf f2,f3->f4
                                         stf f4->[Z+r1]
                                         addi r1,4->r1
                                         blti r1,4096,L1
```

- Instruction-level parallelism (ILP) - fine grained
 - Loop unrolling with static scheduling –or– dynamic scheduling
 - Wide-issue superscalar (non-)scaling limits benefits
- Thread-level parallelism (TLP) - coarse grained
 - Multicore
 - Can we do some “medium grained” parallelism?

Data-Level Parallelism

- **Data-level parallelism (DLP)**
 - Single operation repeated on multiple data elements
 - SIMD (**S**ingle-**I**nstruction, **M**ultiple-**D**ata)
 - Less general than ILP: parallel insns are all same operation
 - Exploit with **vectors**
- Old idea: Cray-1 supercomputer from late 1970s
 - Eight 64-entry x 64-bit floating point “vector registers”
 - 4096 bits (0.5KB) in each register! 4KB for vector register file
 - Special vector instructions to perform vector operations
 - Load vector, store vector (wide memory operation)
 - Vector+Vector or Vector+Scalar
 - addition, subtraction, multiply, etc.
 - In Cray-1, each instruction specifies 64 operations!
 - ALUs were expensive, so one operation per cycle (not parallel)

Example Vector ISA Extensions (SIMD)

- Extend ISA with vector storage ...
 - **Vector register**: fixed-size array of FP/int elements
 - **Vector length**: For example: 4, 8, 16, 64, ...
- ... and example operations for vector length of 4
 - Load vector: `ldf.v [X+r1]->v1`
`ldf [X+r1+0]->v10`
`ldf [X+r1+1]->v11`
`ldf [X+r1+2]->v12`
`ldf [X+r1+3]->v13`
 - Add two vectors: `addf.vv v1,v2->v3`
`addf v1i,v2i->v3i (where i is 0,1,2,3)`
 - Add vector to scalar: `addf-vs v1,f2,v3`
`addf v1i,f2->v3i (where i is 0,1,2,3)`
- Today's vectors: short (128-512 bits), but fully parallel

Example Use of Vectors – 4-wide

```
ldf [X+r1]->f1  
mulf f0,f1->f2  
ldf [Y+r1]->f3  
addf f2,f3->f4  
stf f4->[Z+r1]  
  
addi r1,4->r1  
blti r1,4096,L1
```

7x1024 instructions

```
ldf.v [X+r1]->v1  
mulf-vs v1,f0->v2  
ldf.v [Y+r1]->v3  
addf.vv v2,v3->v4  
stf.v v4,[Z+r1]  
  
addi r1,16->r1  
blti r1,4096,L1
```

7x256 instructions
(4x fewer instructions)

- Operations
 - Load vector: **ldf.v [X+r1]->v1**
 - Multiply vector to scalar: **mulf-vs v1,f0->v2**
 - Add two vectors: **addf.vv v1,v2->v3**
 - Store vector: **stf.v v1->[X+r1]**
- Performance?
 - Best case: 4x speedup
 - But, vector instructions don't always have single-cycle throughput
 - Execution width (implementation) vs vector width (ISA)

Vector Datapath & Implementation

- Vector insn. are just like normal insn... only “wider”
 - Single instruction fetch (no extra N^2 checks)
 - Wide register read & write (not multiple ports)
 - Wide execute: replicate floating point unit (same as superscalar)
 - Wide bypass (avoid N^2 bypass problem)
 - Wide cache read & write (single cache tag check)
- Execution width (implementation) vs vector width (ISA)
 - Example: Pentium 4 and “Core 1” executes vector ops at half width
 - “Core 2” executes them at full width
- Because they are just instructions...
 - ...superscalar execution of vector instructions
 - Multiple n-wide vector instructions per cycle

Vector Insn Sets for Different ISAs

- x86
 - Intel and AMD: MMX, SSE, SSE2, SSE3, SSE4, AVX, AVX2
 - currently: AVX 512 offers 512b vectors
- PowerPC
 - AltiVEC/VMX: 128b
- ARM
 - NEON: 128b
 - Scalable Vector Extensions (SVE): up to 2048b
 - implementation is narrower than this!

Other Vector Instructions

- These target specific domains: e.g., image processing, crypto
 - Vector reduction (sum all elements of a vector)
 - Geometry processing: 4x4 translation/rotation matrices
 - Saturating (non-overflowing) subword add/sub: image processing
 - Byte asymmetric operations: blending and composition in graphics
 - Byte shuffle/permute: crypto
 - Population (bit) count: crypto
 - Max/min/argmax/argmin: video codec
 - Absolute differences: video codec
 - Multiply-accumulate: digital-signal processing
 - Special instructions for AES encryption
- More advanced (but in Intel's Xeon Phi)
 - Scatter/gather loads: indirect store (or load) from a vector of pointers
 - Vector mask: predication (conditional execution) of specific elements

Vector Masks (Predication)

- **Vector Masks**: 1 bit per vector element
 - Implicit predicate in all vector operations

```
for (I=0; I<N; I++) if (mask_I) { vop... }
```
 - Usually stored in a “scalar” register (up to 64-bits)
 - Used to vectorize loops with conditionals in them

```
cmp_eq.v, cmp_lt.v, etc.: sets vector predicates
```

```
for (I=0; I<32; I++)
    if (X[I] != 0.0) Z[I] = A/X[I];
```

```
ldf.v [X+r1] -> v1
cmp_ne.v v1,f0 -> r2          // 0.0 is in f0
divf.sv {r2} v1,f1 -> v2      // A is in f1
stf.v {r2} v2 -> [Z+r1]
```

Scatter Stores & Gather Loads

- How to vectorize:

```
for(int i = 1, i<N, i++) {  
    int bucket = val[i] / scalefactor;  
    found[bucket] = 1;
```

- Easy to vectorize the divide, but what about the store?

- Solution: hardware support for vector “scatter stores”

- `stf.v v2->[r1+v1]`
- Each address calculated from $r1+v1_i$
 $stf\ v2_0->[r1+v1_0],\ stf\ v2_1->[r1+v1_1],$
 $stf\ v2_2->[r1+v1_2],\ stf\ v2_3->[r1+v1_3]$

- Vector “gather loads” defined analogously
 - `ldf.v [r1+v1]->v2`
- Scatter/gathers slower than regular vector load/store ops
 - Still provides a throughput advantage over non-vector version

Using Vectors in Your Code

- Write in assembly
 - Ugh
- Use “intrinsic” functions and data types
 - For example: `_mm_mul_ps()` and “`__m128`” datatype
- Use vector data types
 - `typedef double v2df __attribute__ ((vector_size (16)));`
- Use a library someone else wrote
 - Let them do the hard work
 - Matrix and linear algebra packages
- Let the compiler do it (automatic vectorization, with feedback)
 - GCC’s “`-ftree-vectorize`” option, `-ftree-vectorizer-verbose=n`
 - Limited impact for C/C++ code (old, hard problem)

By the numbers: CPUs vs GPUs

	Intel Xeon Platinum 8168 “Skylake”	Nvidia Tesla P100	Intel Xeon Phi 7290F
frequency	2.7 GHz	1.3 GHz	1.5 GHz
cores / threads	24 / 48	56 (“3584”) / 10Ks	72 / 288
RAM	768 GB	16 GB	384 GB
DP TFLOPS	1.0	4.7	3.5
Transistors	>5B ?	15.3B	>5B ?
Price	\$5,900	\$6,000	\$3,400

GPUs and SIMD/Vector Data Parallelism

- How do GPUs have such high peak FLOPS & FLOPS/Joule?
 - Exploit massive data parallelism – focus on total throughput
 - Remove hardware structures that accelerate single threads
 - Specialized for graphics: e.g., data-types & dedicated texture units
- “SIMT” execution model
 - Single instruction multiple threads
 - Similar to both “vectors” and “SIMD”
 - A key difference: better support for conditional control flow
- Program it with CUDA or OpenCL
 - Extensions to C/C++
 - Perform a “shader task” (a snippet of scalar computation) over many elements
 - Internally, GPU uses scatter/gather and vector mask operations

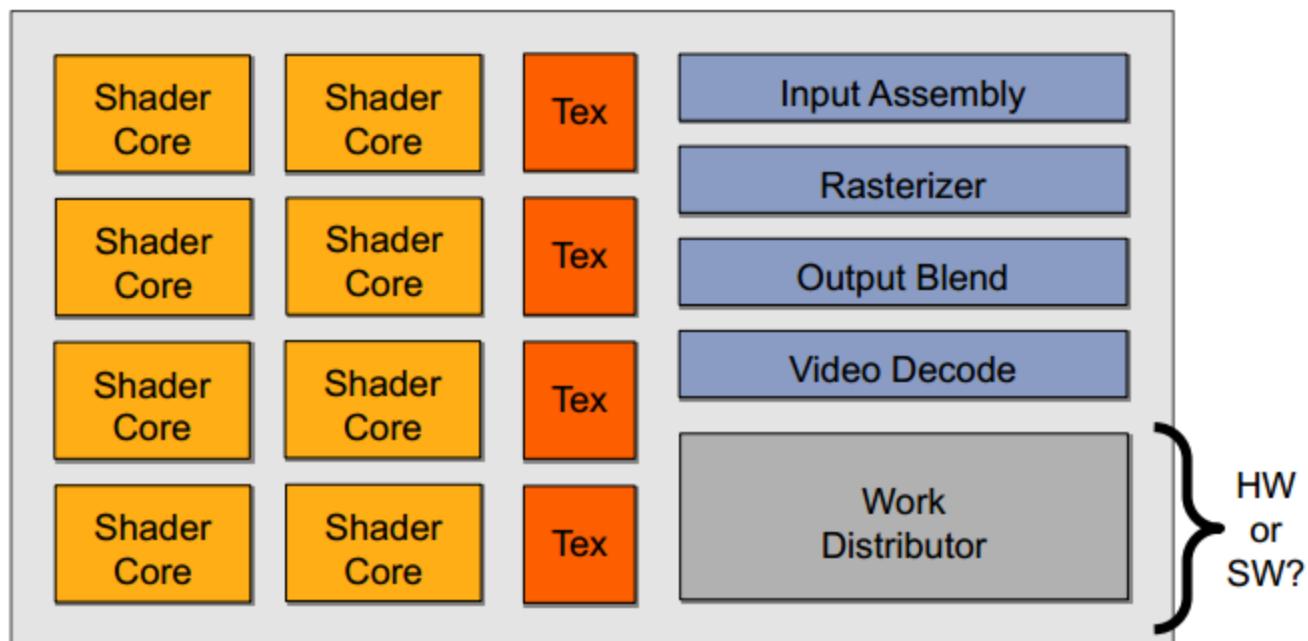
- following slides “Beyond Programmable Shading” course
“Real-Time Rendering Architectures”

<http://bps11.idav.ucdavis.edu/>

Mike Houston, AMD

What's in a GPU?

A GPU is a heterogeneous chip multi-processor (highly tuned for graphics)



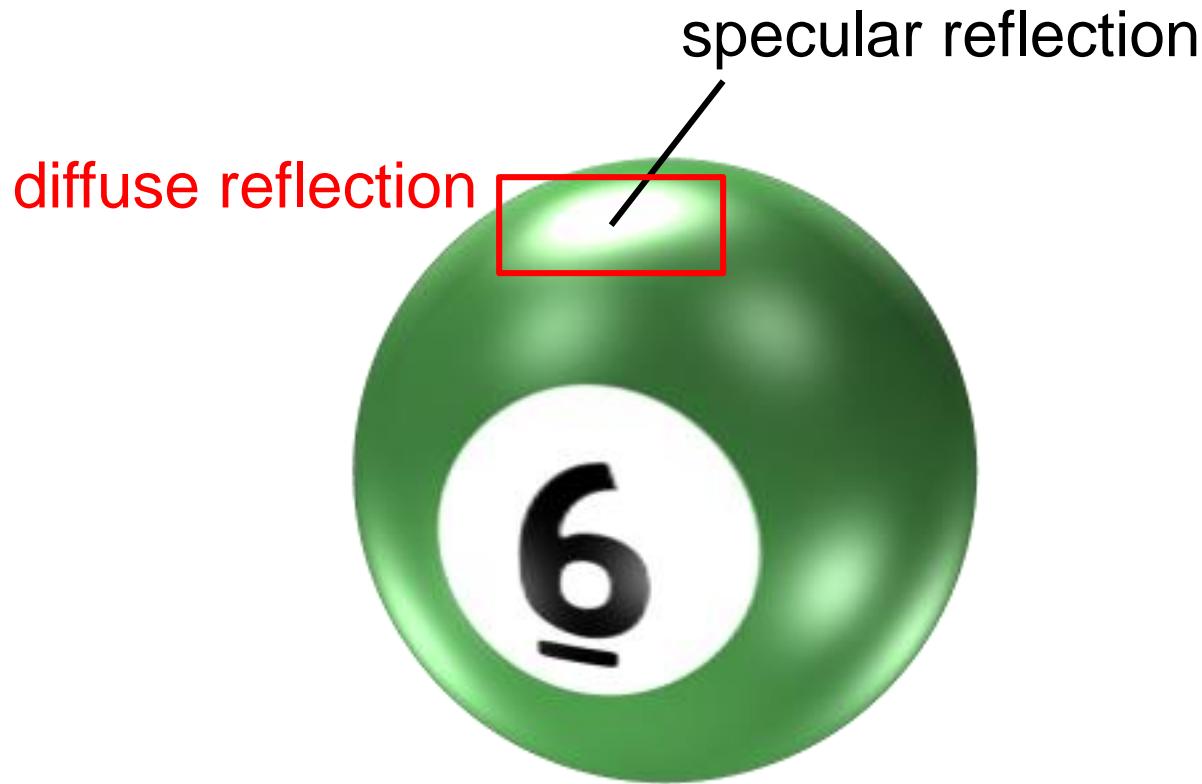
A diffuse reflectance shader

```
sampler mySamp;  
Texture2D<float3> myTex;  
float3 lightDir;  
  
float4 diffuseShader(float3 norm, float2 uv)  
{  
    float3 kd;  
    kd = myTex.Sample(mySamp, uv);  
    kd *= clamp( dot(lightDir, norm), 0.0, 1.0);  
    return float4(kd, 1.0);  
}
```

Shader programming model:

Fragments are processed **independently**,
but there is no explicit parallel
programming

diffuse reflection example



c/o <http://www.iconarchive.com/show/pool-ball-icons-by-barkerbaggies/Ball-6-icon.html>

Compile shader

```
sampler mySamp;  
Texture2D<float3> myTex;  
float3 lightDir;  
  
float4 diffuseShader(float3 norm, float2 uv)  
{  
    float3 kd;  
    kd = myTex.Sample(mySamp, uv);  
    kd *= clamp( dot(lightDir, norm), 0.0, 1.0);  
    return float4(kd, 1.0);  
}
```

1 unshaded fragment input record



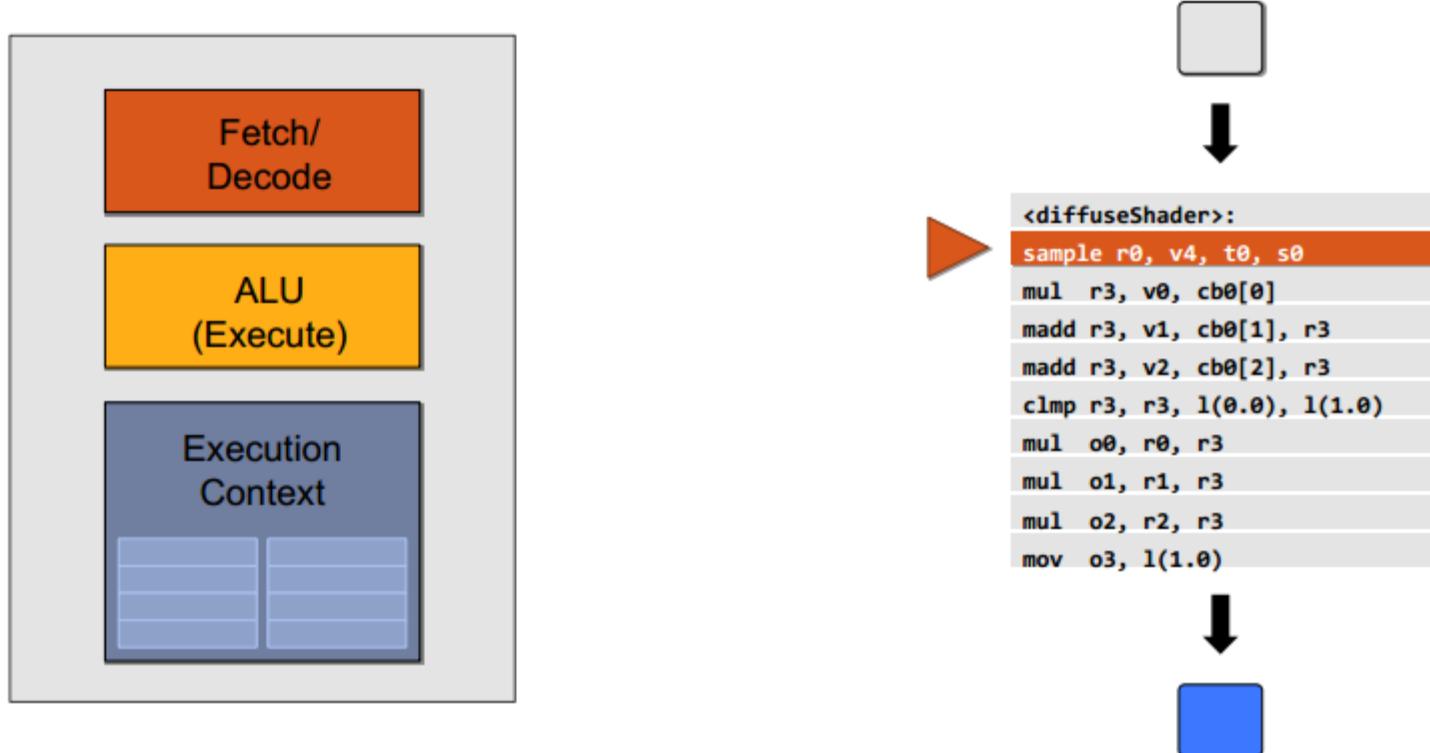
```
<diffuseShader>:  
sample r0, v4, t0, s0  
mul r3, v0, cb0[0]  
madd r3, v1, cb0[1], r3  
madd r3, v2, cb0[2], r3  
clmp r3, r3, 1(0.0), 1(1.0)  
mul o0, r0, r3  
mul o1, r1, r3  
mul o2, r2, r3  
mov o3, l(1.0)
```



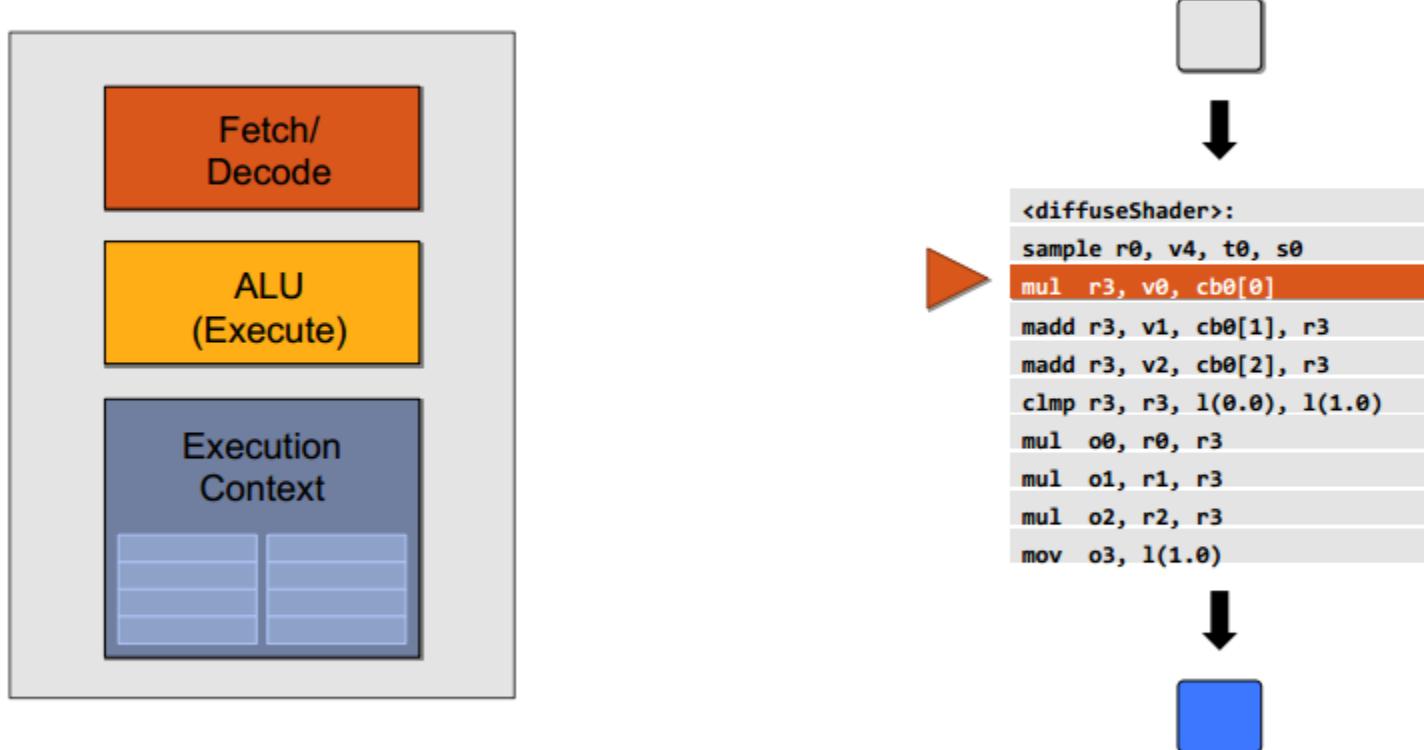
1 shaded fragment output record



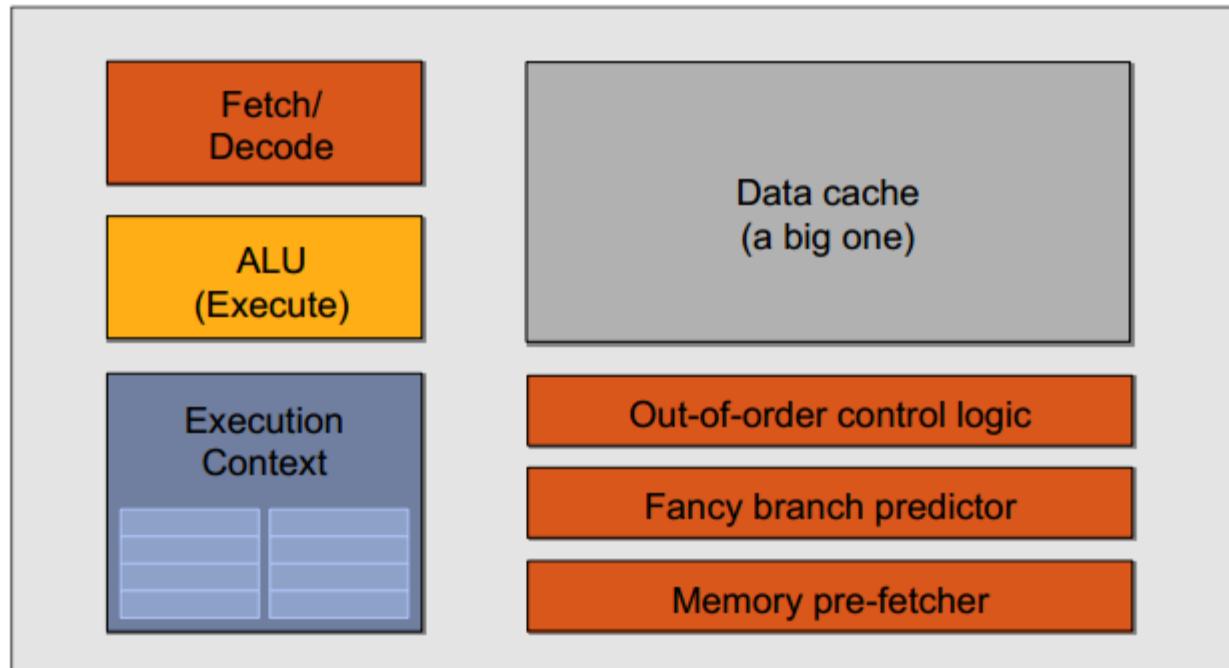
Execute shader



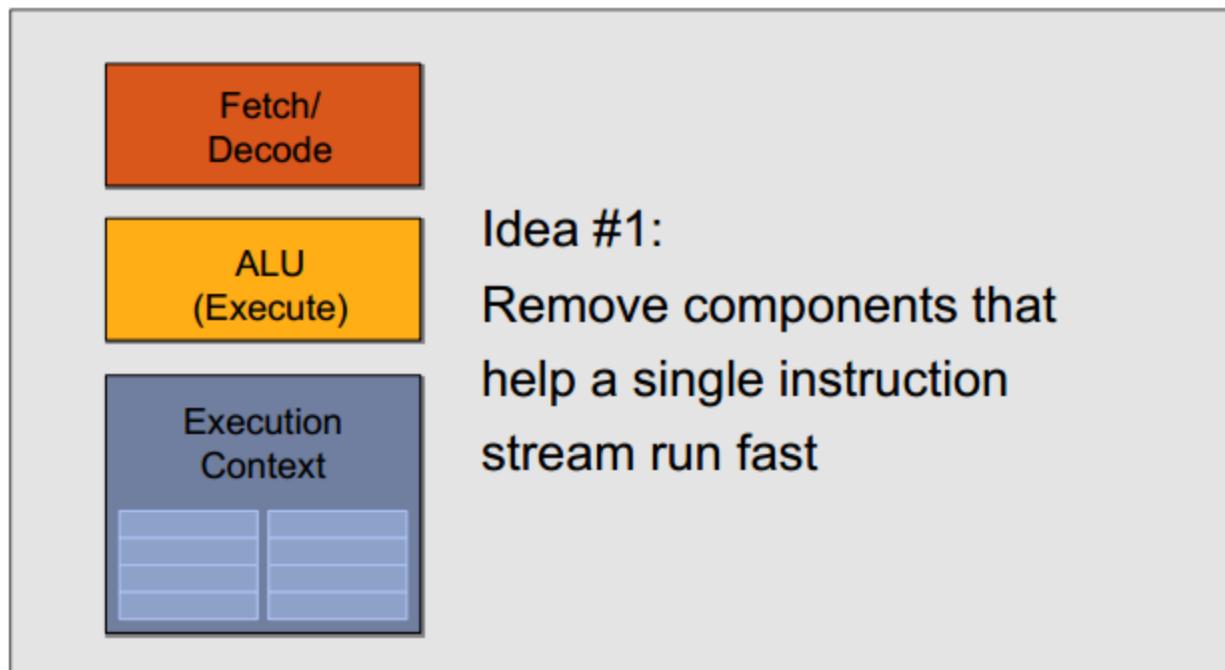
Execute shader



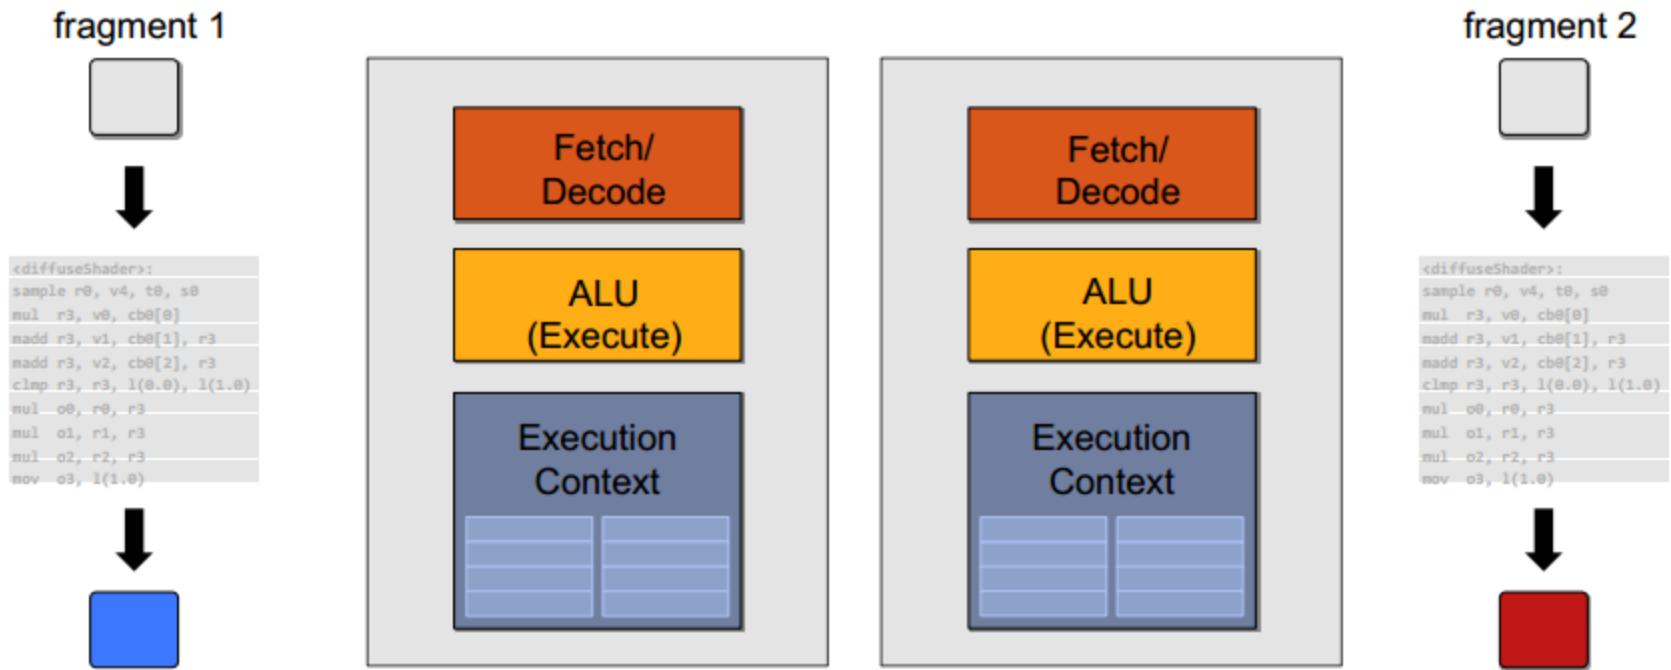
“CPU-style” cores



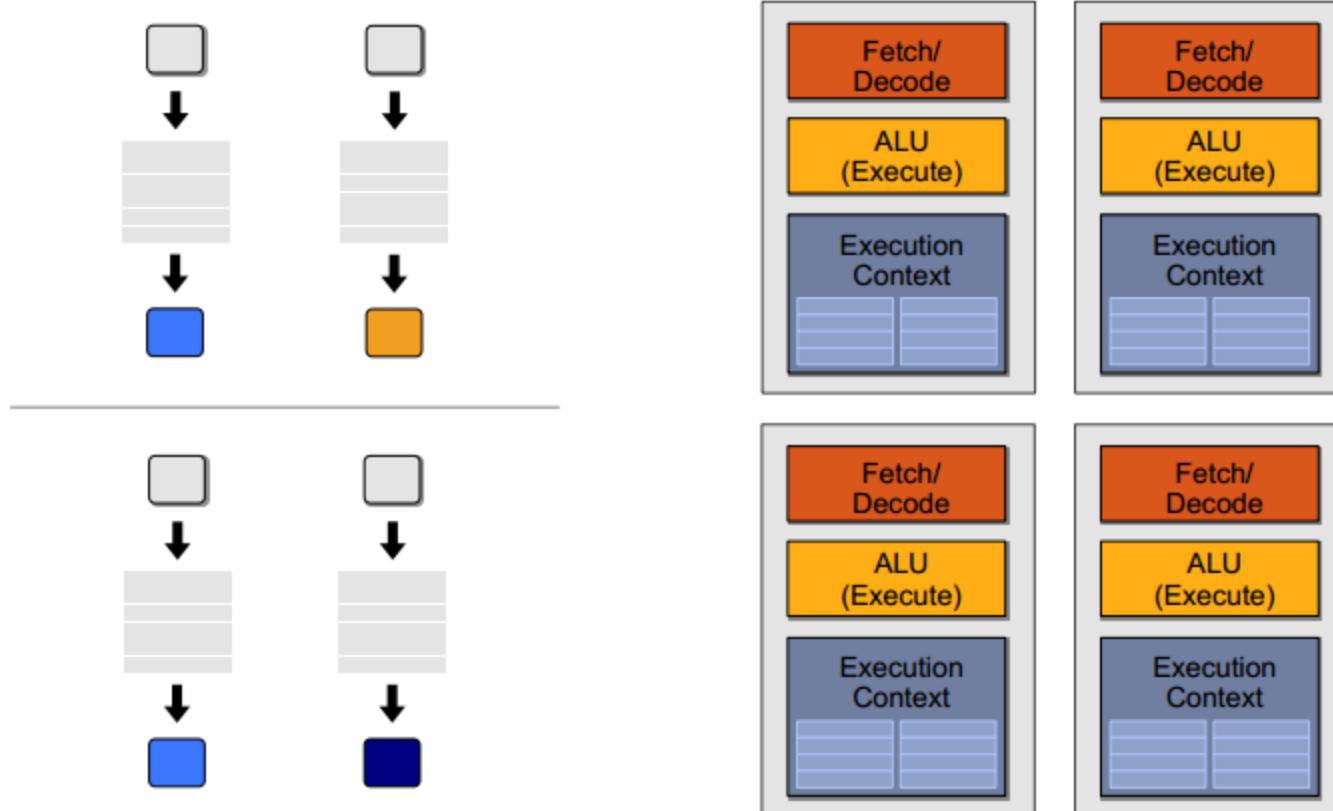
Slimming down



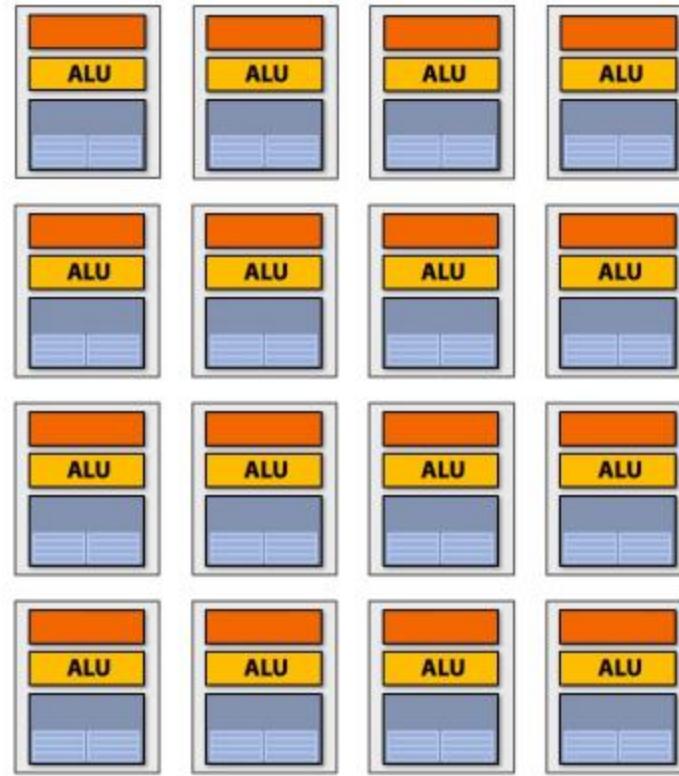
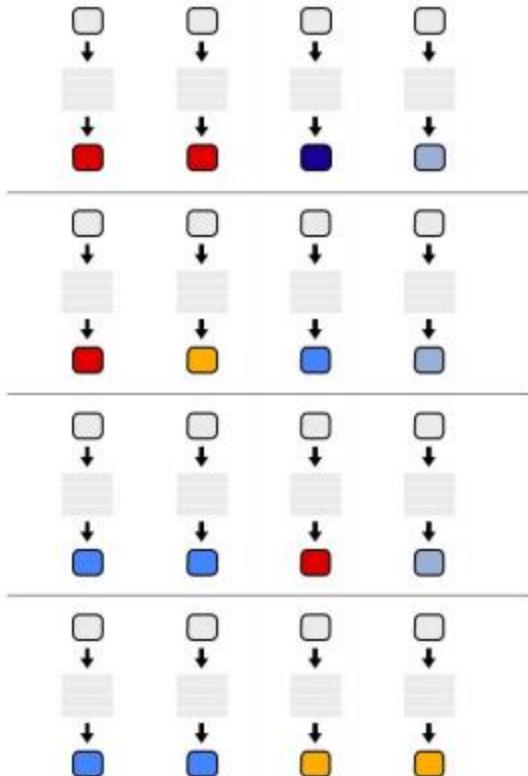
Two cores (two fragments in parallel)



Four cores (four fragments in parallel)

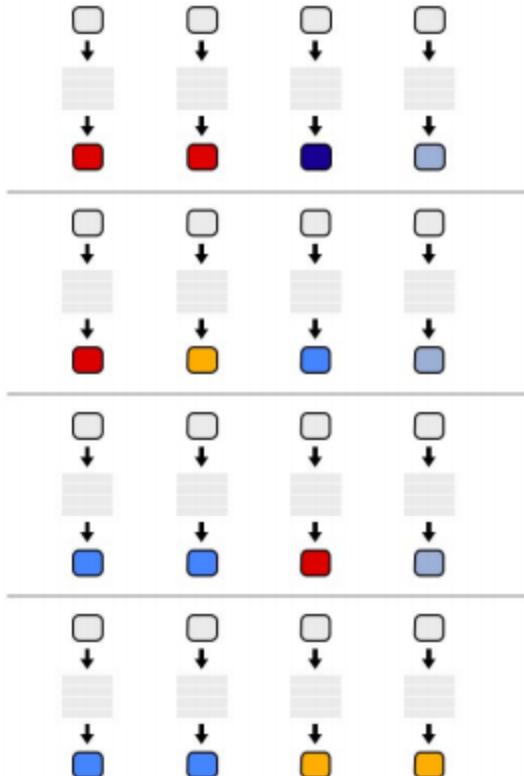


Sixteen cores (sixteen fragments in parallel)



16 cores = 16 simultaneous instruction streams
Beyond Programmable Shading Course, ACM SIGGRAPH 2011

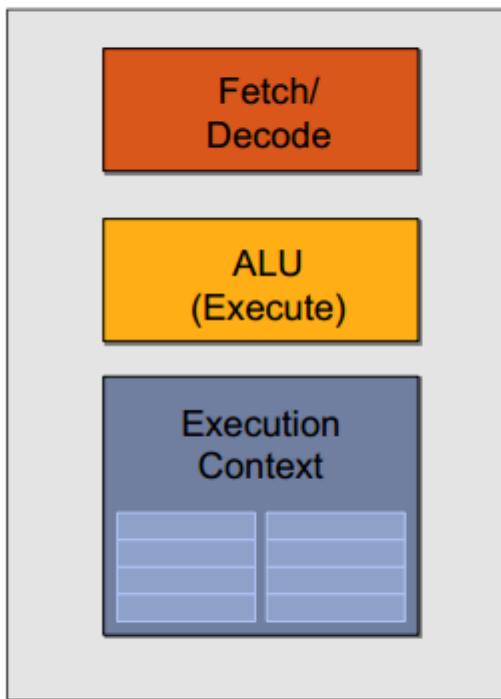
Instruction stream sharing



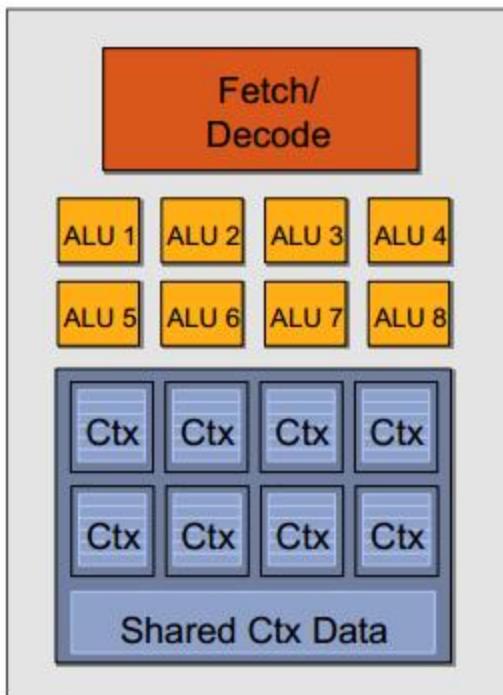
But ... many fragments
should be able to share an
instruction stream!

```
<diffuseShader>:  
sample r0, v4, t0, s0  
mul r3, v0, cb0[0]  
madd r3, v1, cb0[1], r3  
madd r3, v2, cb0[2], r3  
clmp r3, r3, 1(0.0), 1(1.0)  
mul o0, r0, r3  
mul o1, r1, r3  
mul o2, r2, r3  
mov o3, 1(1.0)
```

Recall: simple processing core



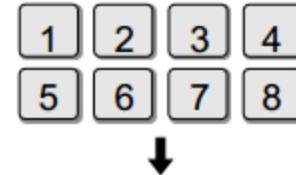
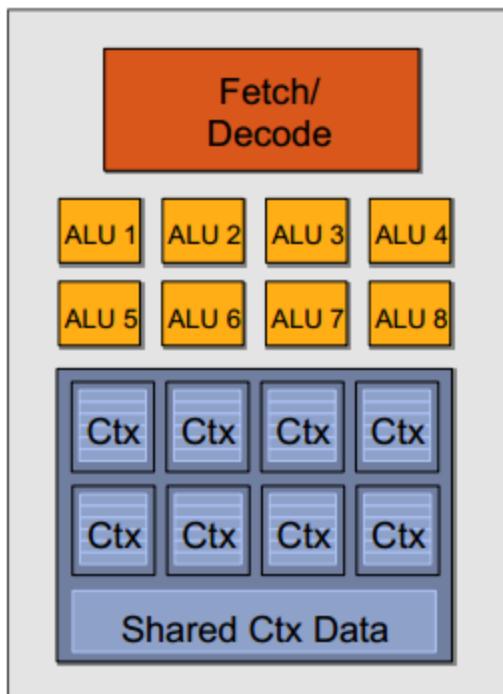
Add ALUs



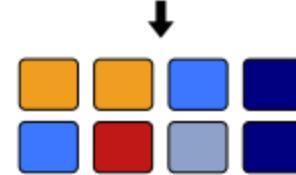
Idea #2:
Amortize cost/complexity of
managing an instruction
stream across many ALUs

SIMD processing

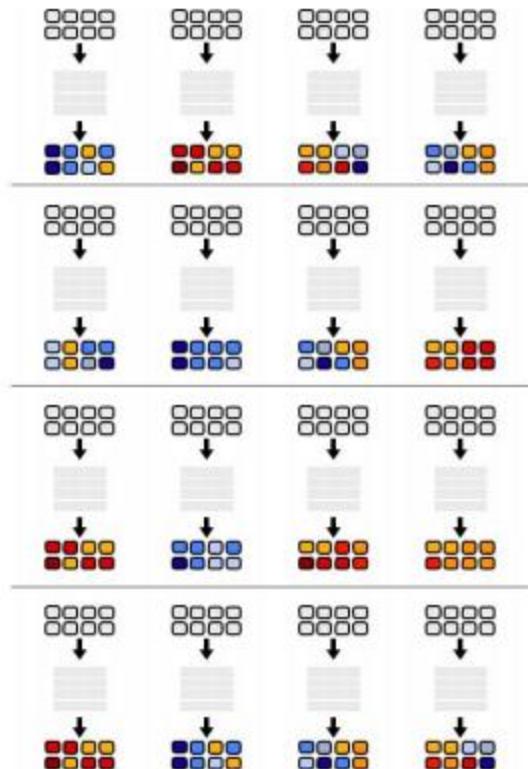
Modifying the shader



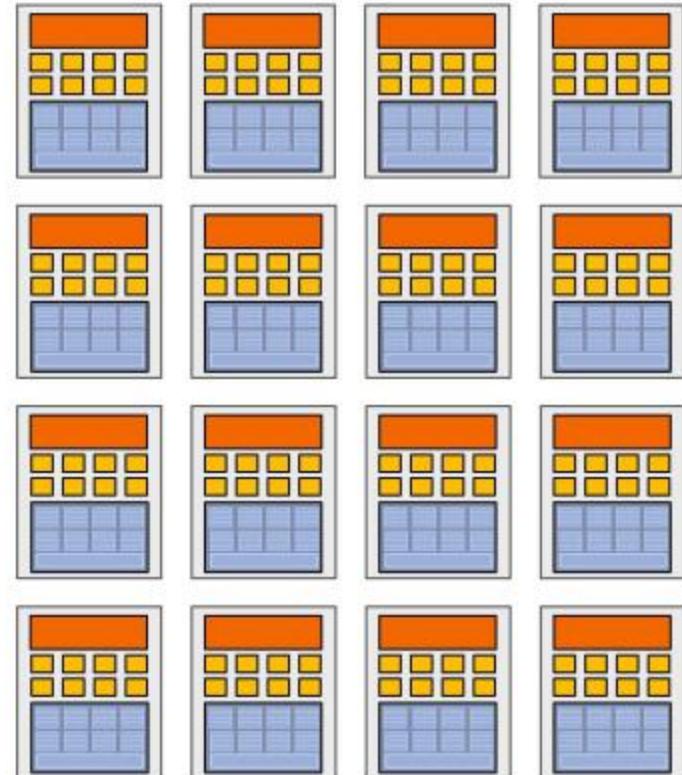
```
<VEC8_diffuseShader>:  
VEC8_sample vec_r0, vec_v4, t0, vec_s0  
VEC8_mul  vec_r3, vec_v0, cb0[0]  
VEC8_madd vec_r3, vec_v1, cb0[1], vec_r3  
VEC8_madd vec_r3, vec_v2, cb0[2], vec_r3  
VEC8_clmp vec_r3, vec_r3, 1(0.0), 1(1.0)  
VEC8_mul  vec_o0, vec_r0, vec_r3  
VEC8_mul  vec_o1, vec_r1, vec_r3  
VEC8_mul  vec_o2, vec_r2, vec_r3  
VEC8_mov  o3, 1(1.0)
```



128 fragments in parallel



16 cores = 128 ALUs



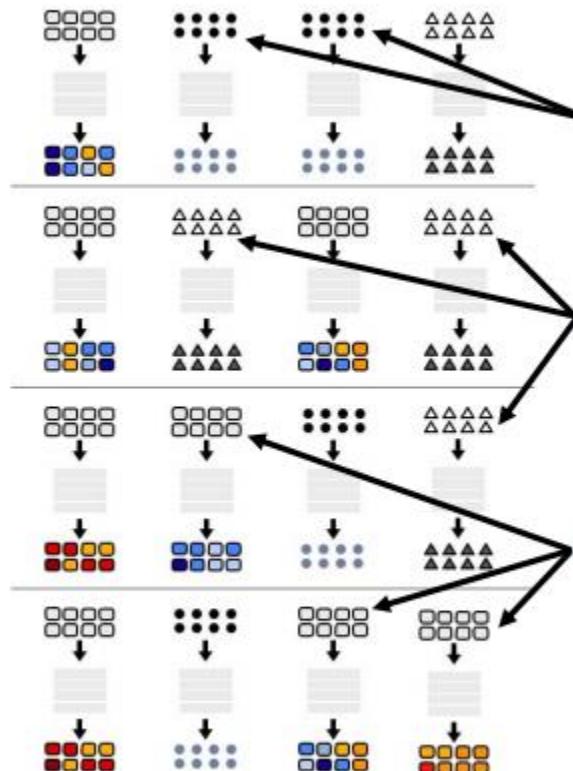
, 16 simultaneous instruction streams

Beyond Programmable Shading Course, ACM SIGGRAPH 2011

128 [

vertices/fragments
primitives
OpenCL work items

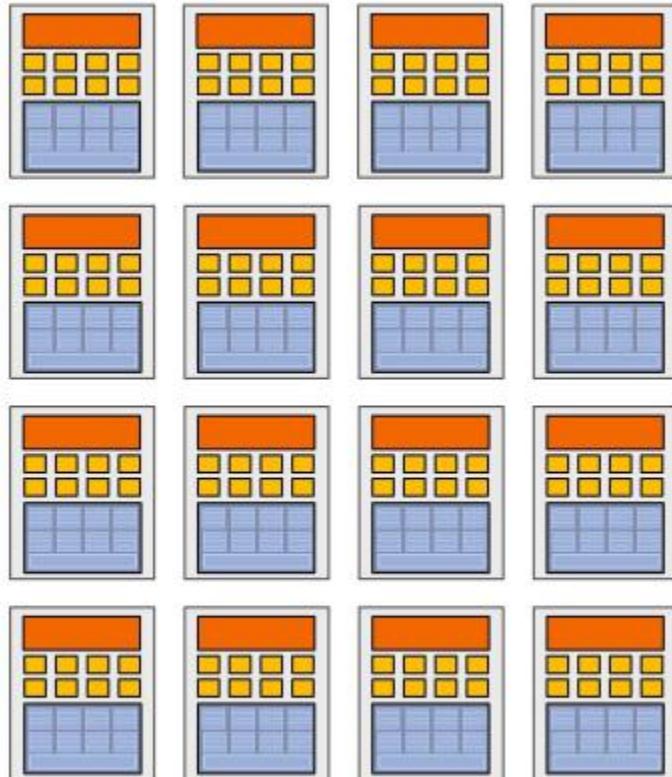
] in parallel



vertices

primitives

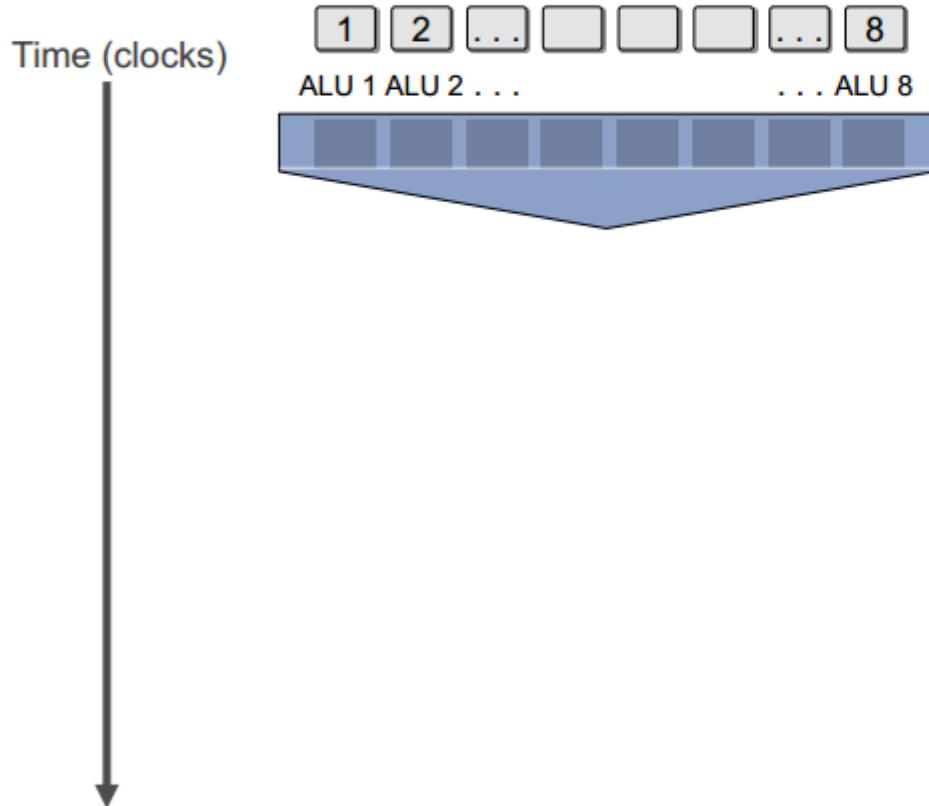
fragments



SIGGRAPH2011

Beyond Programmable Shading Course, ACM SIGGRAPH 2011

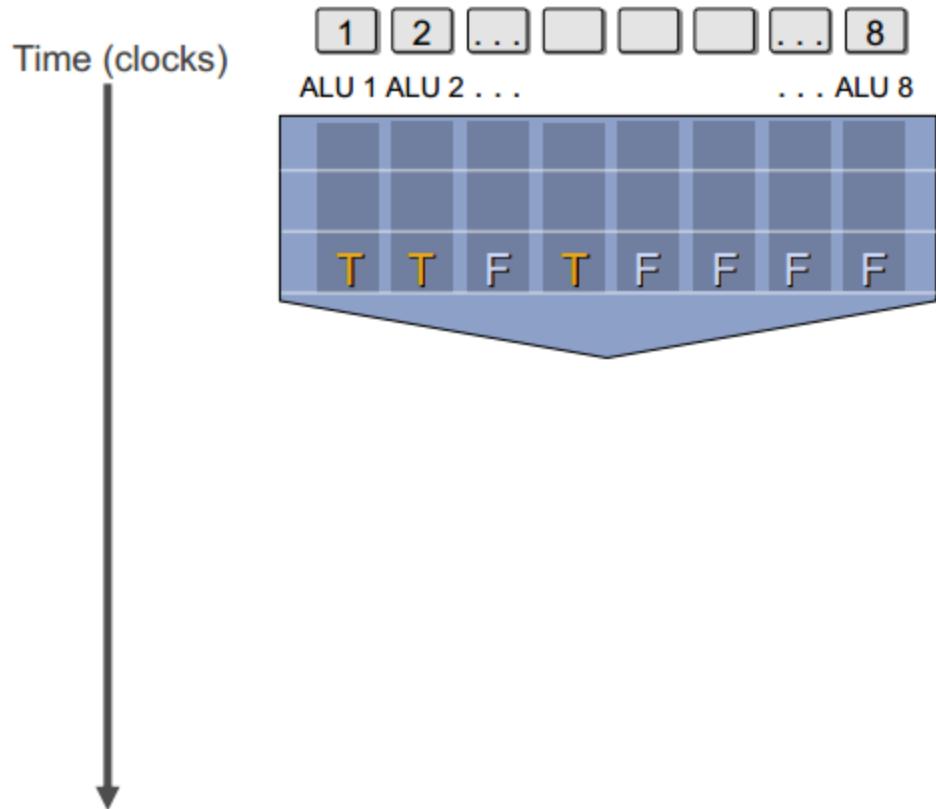
But what about branches?



```
<unconditional  
shader code>

if (x > 0) {  
    y = pow(x, exp);  
    y *= Ks;  
    refl = y + Ka;  
} else {  
    x = 0;  
    refl = Ka;  
}  
<resume unconditional  
shader code>
```

But what about branches?

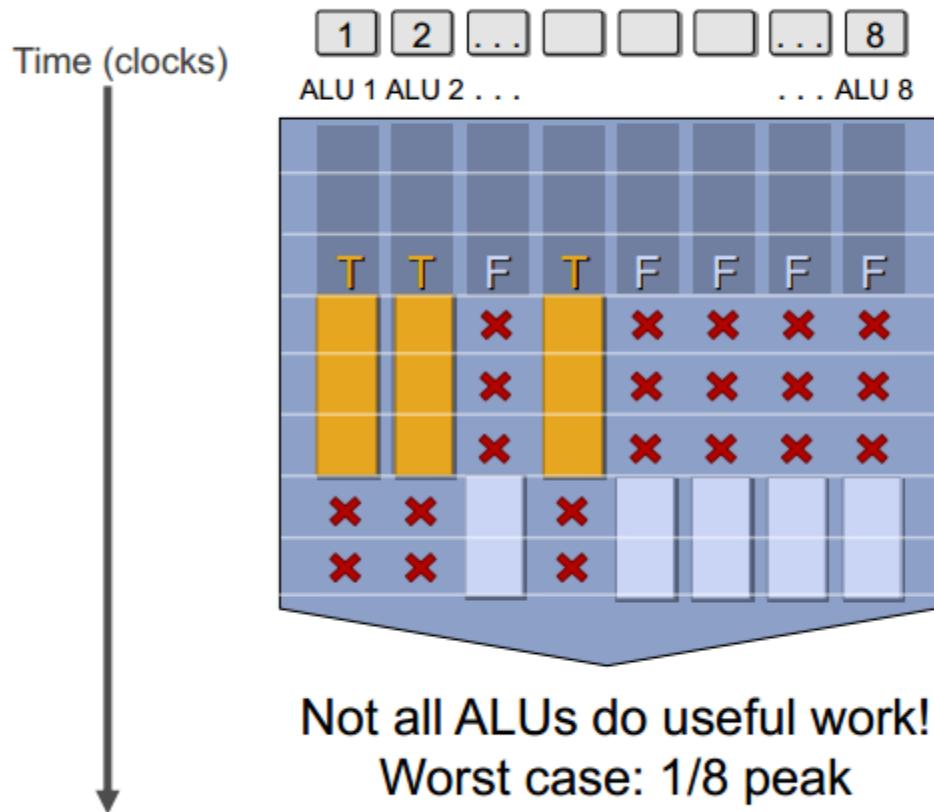


<unconditional
shader code>

```
if (x > 0) {  
    y = pow(x, exp);  
    y *= Ks;  
    refl = y + Ka;  
} else {  
    x = 0;  
    refl = Ka;  
}
```

<resume unconditional
shader code>

But what about branches?



```
<unconditional  
shader code>  
  
if (x > 0) {  
    y = pow(x, exp);  
    y *= Ks;  
    refl = y + Ka;  
} else {  
    x = 0;  
    refl = Ka;  
}  
  
<resume unconditional  
shader code>
```

Clarification

SIMD processing does not imply SIMD instructions

- Option 1: explicit vector instructions
 - x86 SSE, AVX, Intel Larrabee
- Option 2: scalar instructions, implicit HW vectorization
 - HW determines instruction stream sharing across ALUs (amount of sharing hidden from software)
 - NVIDIA GeForce (“SIMT” warps), ATI Radeon architectures (“wavefronts”)



In practice: 16 to 64 fragments share an instruction stream.

SIMD vs SIMT

- SIMD: single insn multiple **data**
 - write 1 insn that operates on a vector of data
 - handle control flow via explicit masking operations
- SIMT: single insn multiple **thread**
 - write 1 insn that operates on scalar data
 - each of many threads runs this insn
 - compiler+hw aggregate threads into groups that execute on SIMD hardware
 - compiler+hw handle masking for control flow

Stalls!

Stalls occur when a core cannot run the next instruction because of a dependency on a previous operation.

Texture access latency = 100's to 1000's of cycles

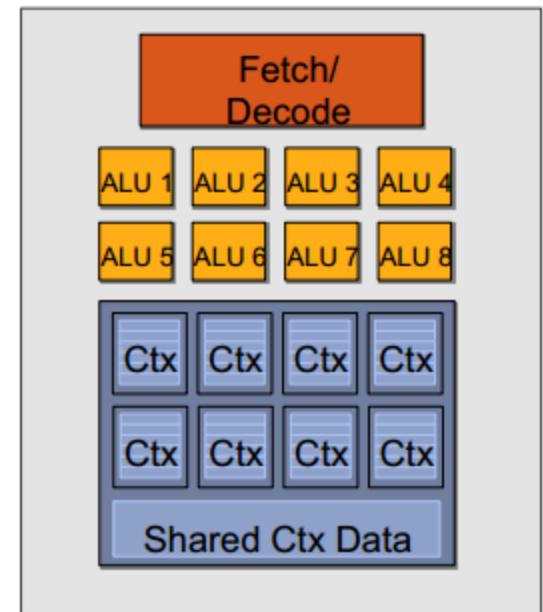
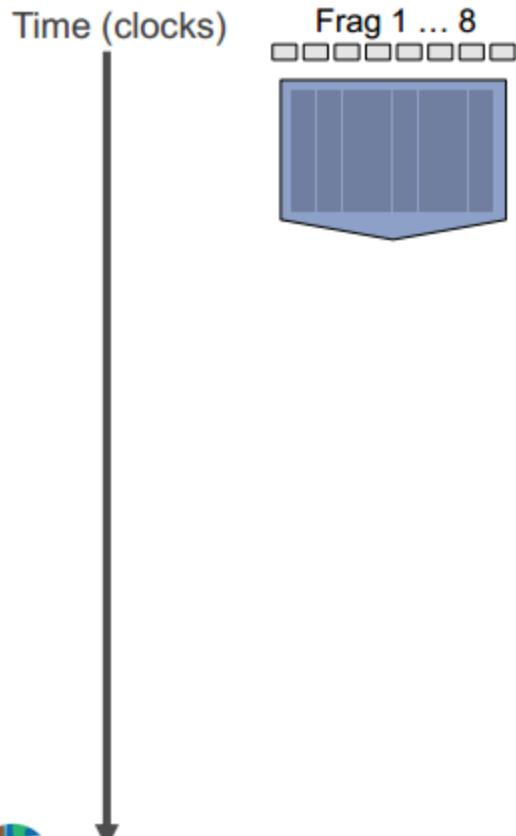
We've removed the fancy caches and logic that helps avoid stalls.

But we have **LOTS** of independent fragments.

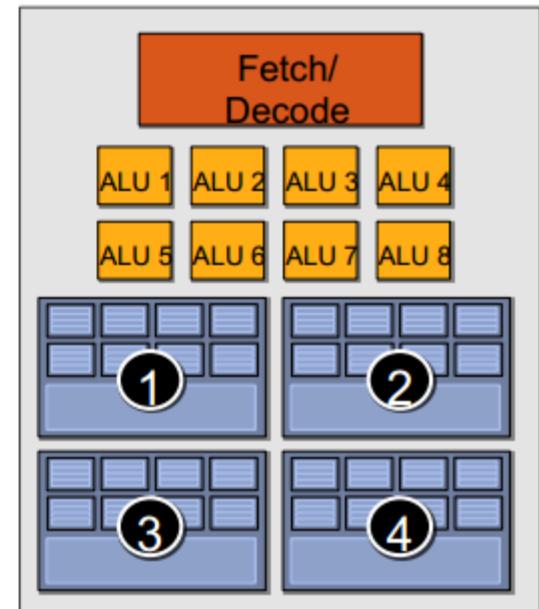
Idea #3:

Interleave processing of many fragments on a single core to avoid stalls caused by high latency operations.

Hiding shader stalls

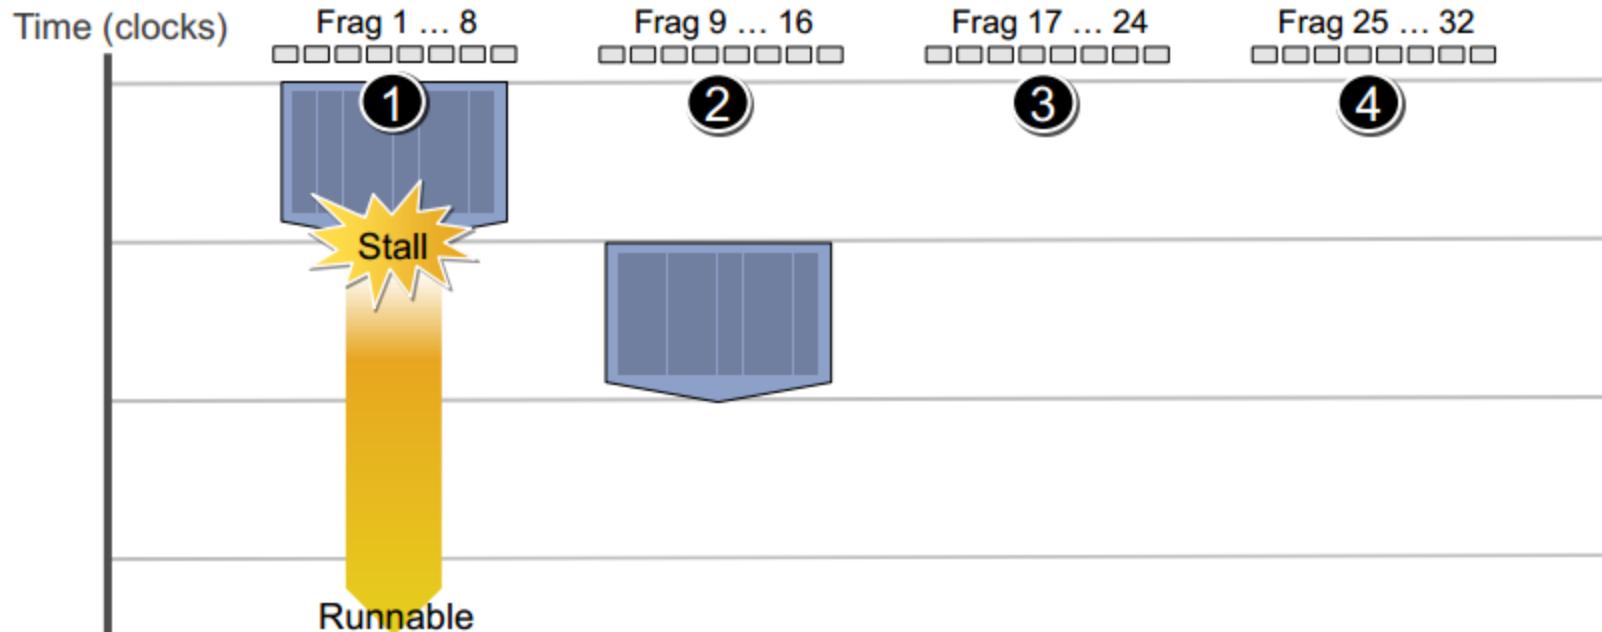


Hiding shader stalls

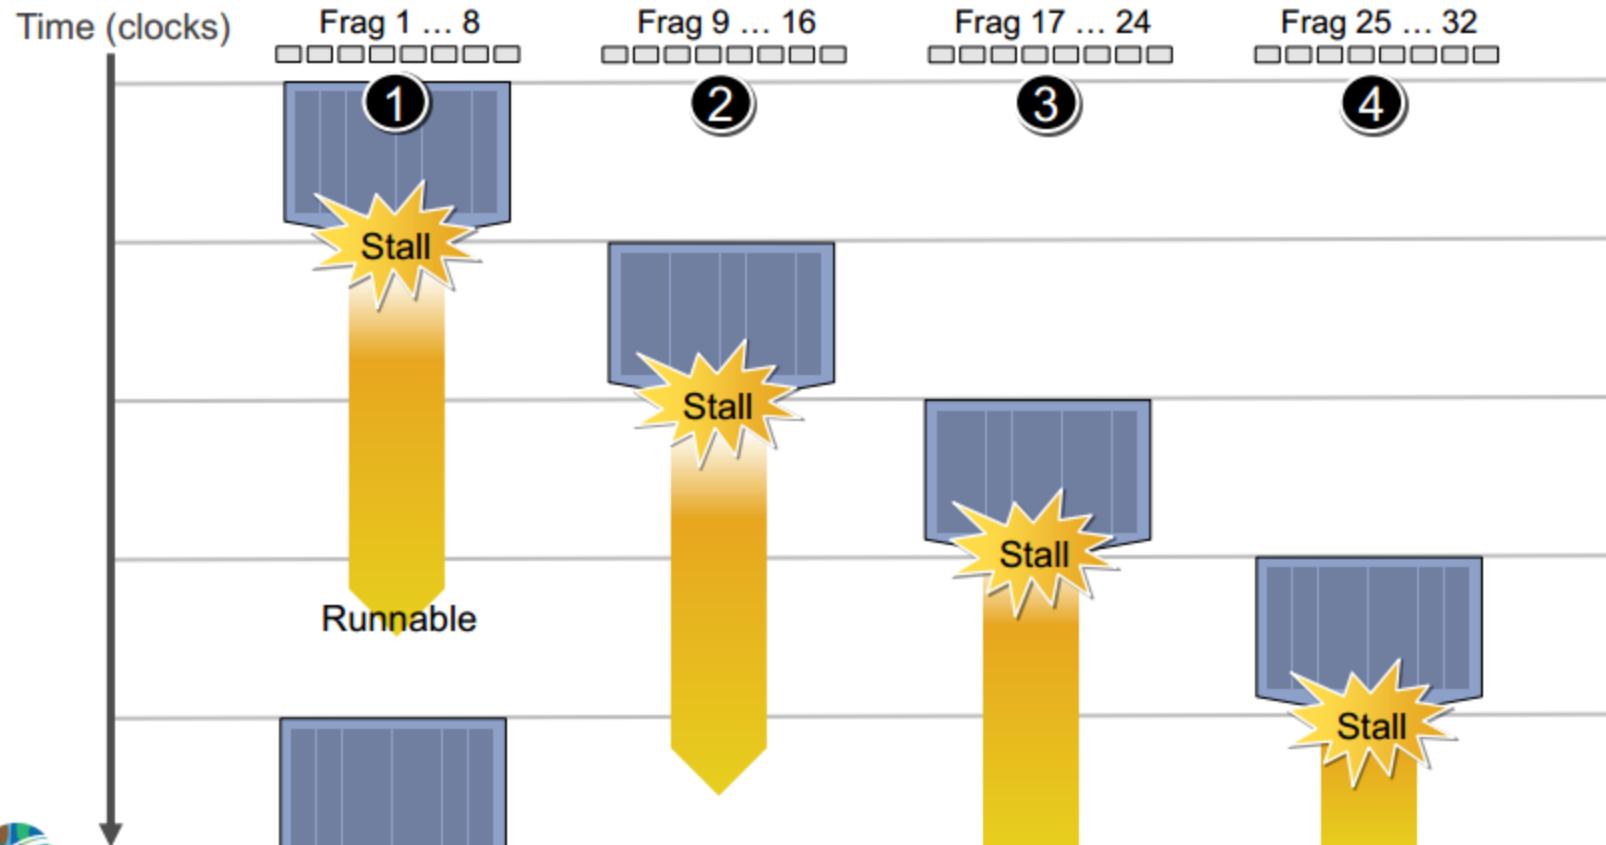


Beyond Programmable Shading Course, ACM SIGGRAPH 2011

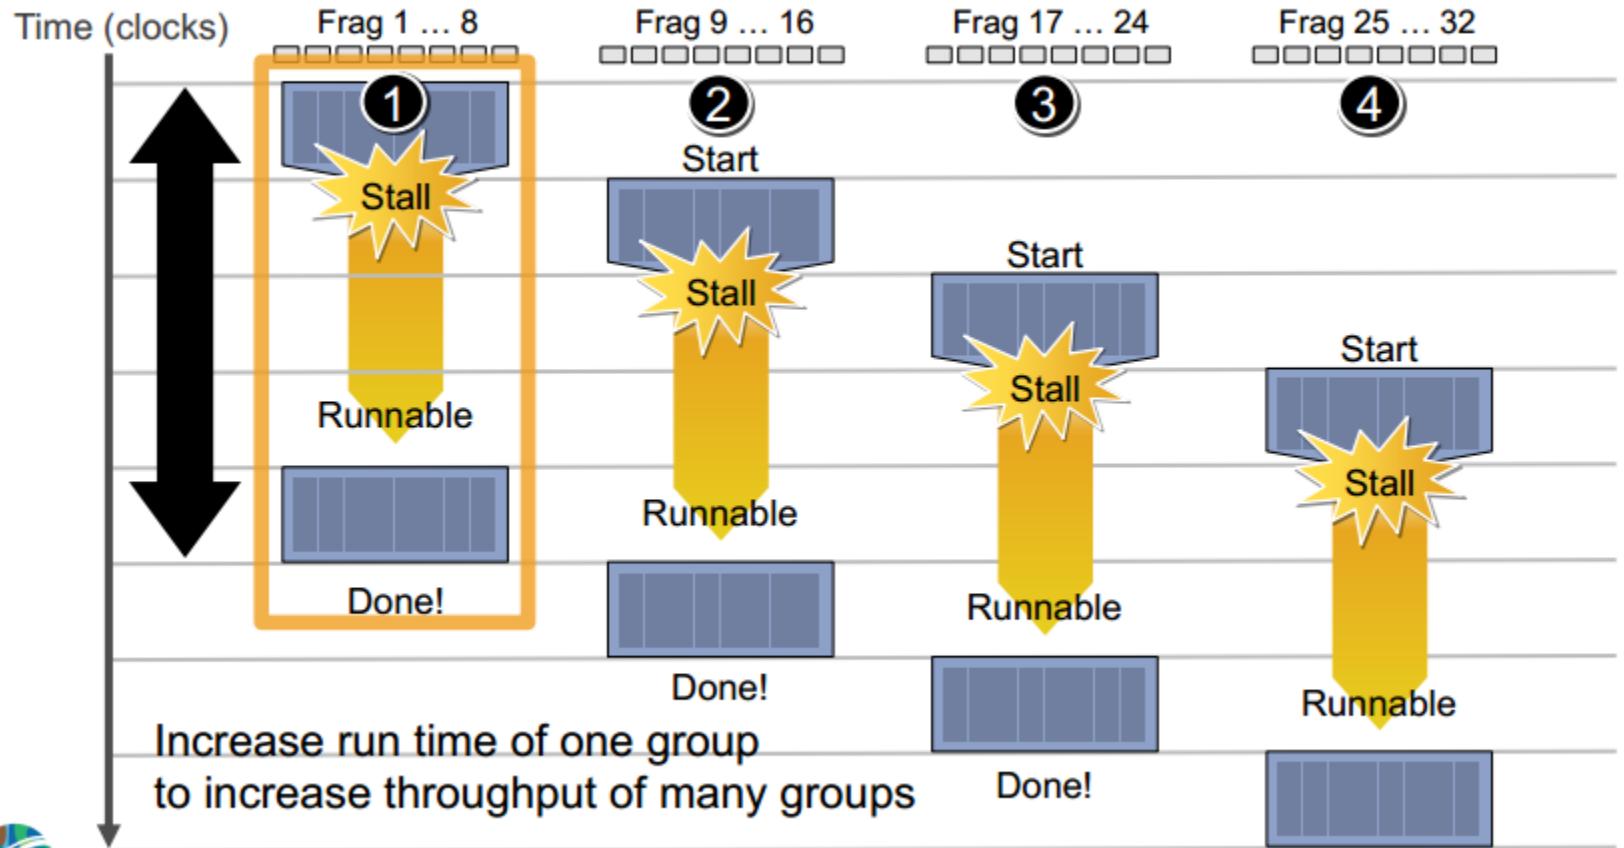
Hiding shader stalls



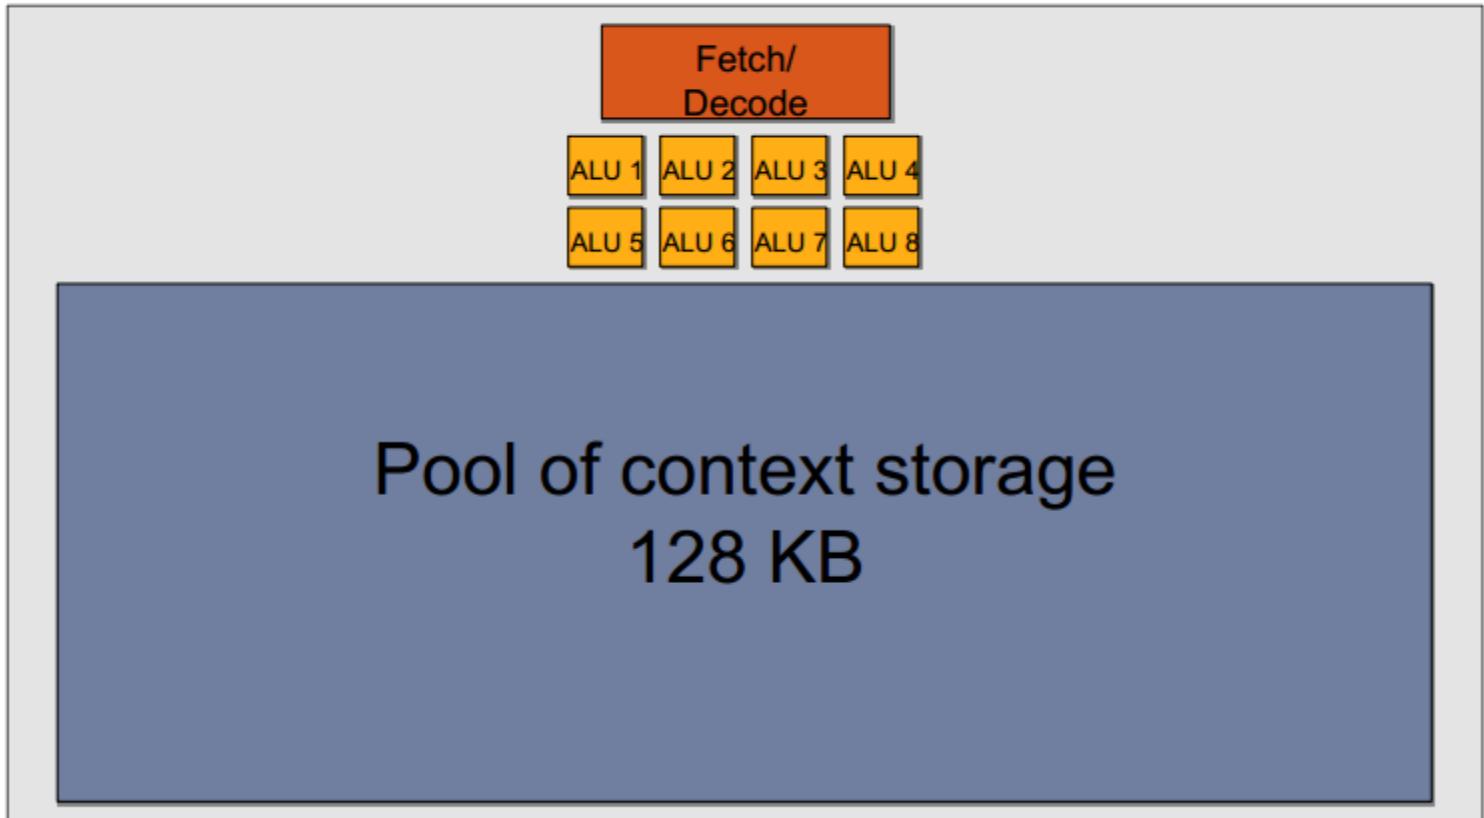
Hiding shader stalls



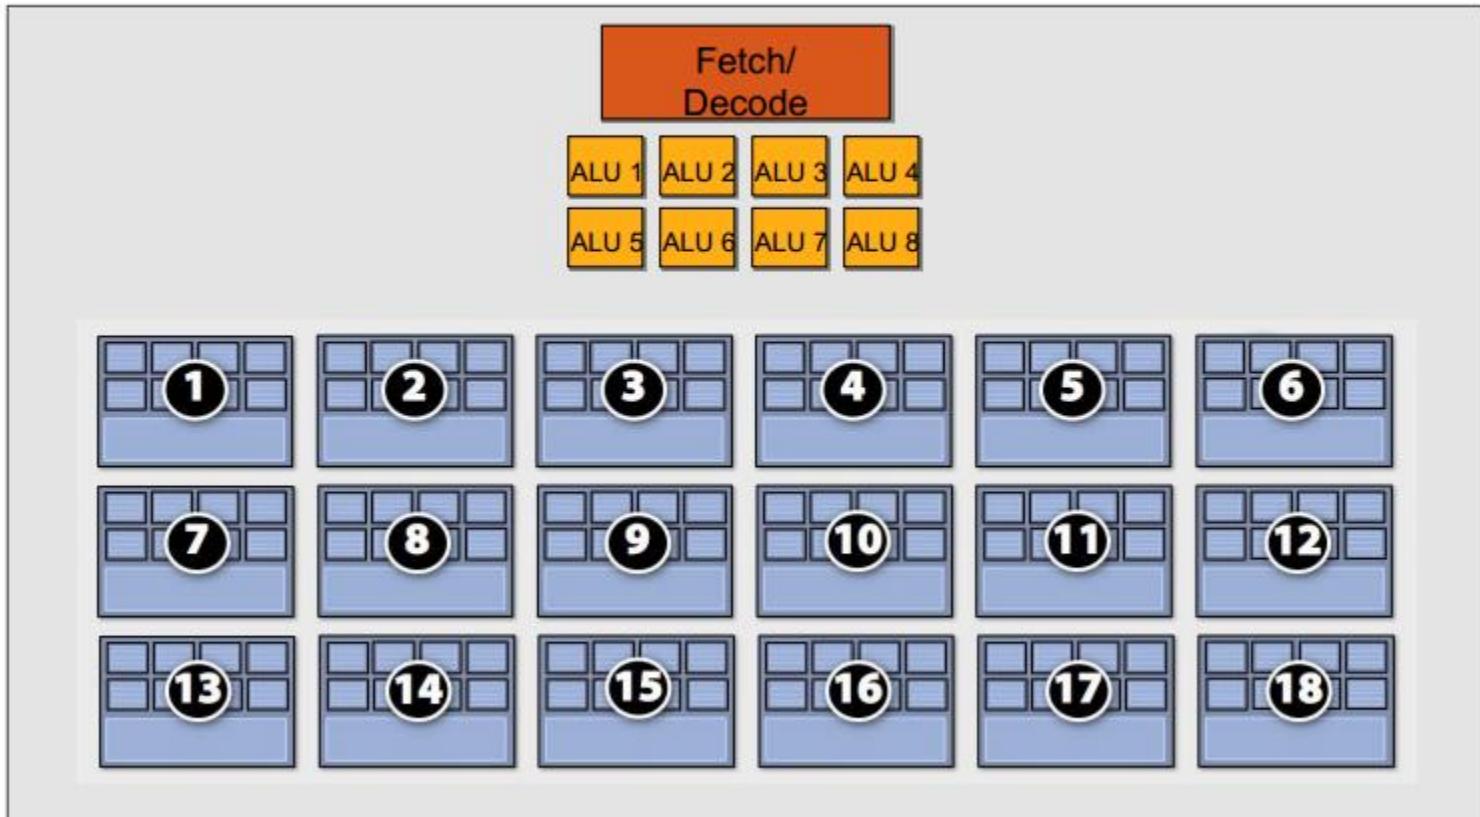
Throughput!



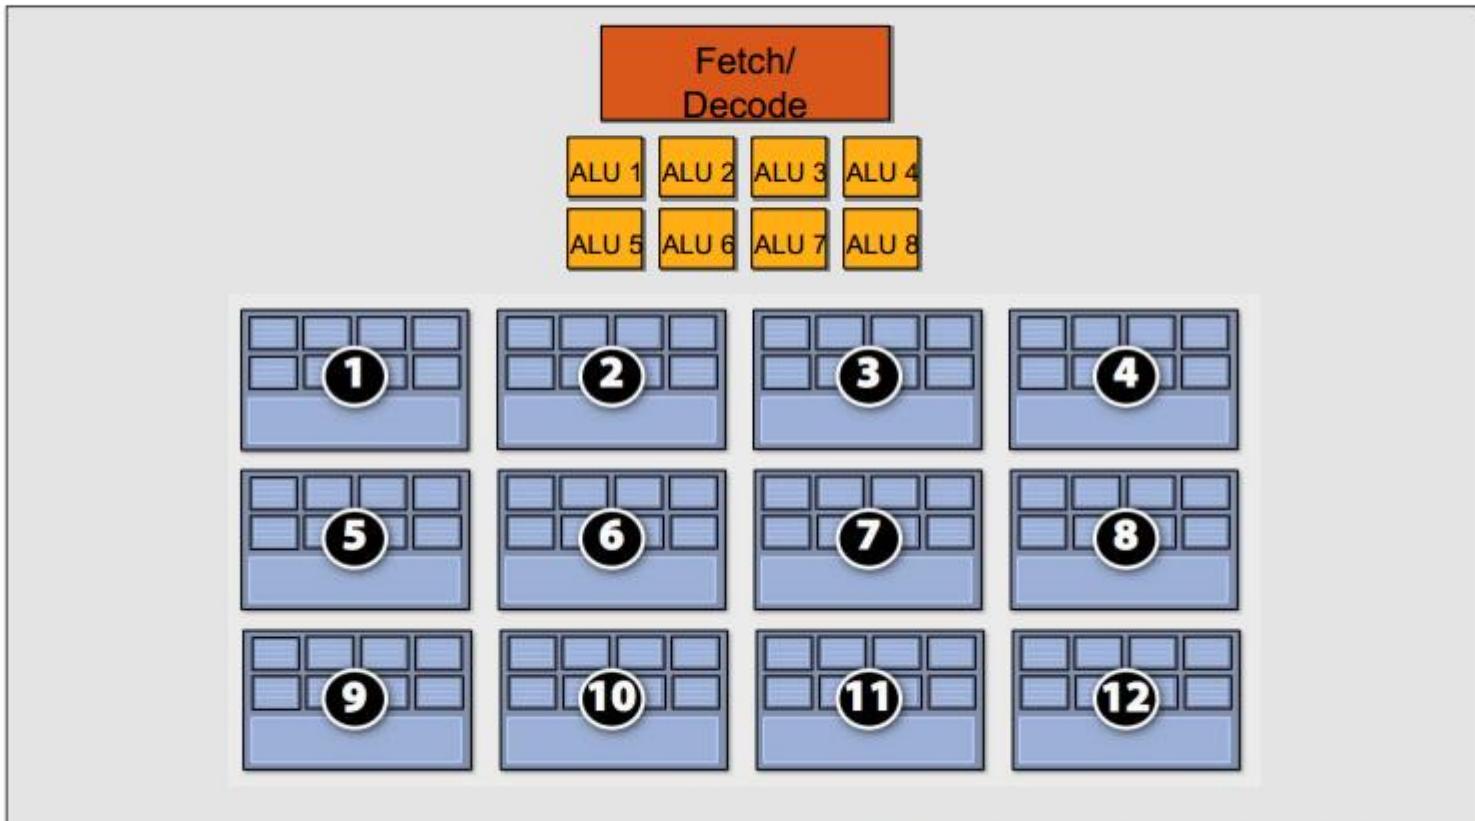
Storing contexts



Eighteen small contexts (maximal latency hiding)

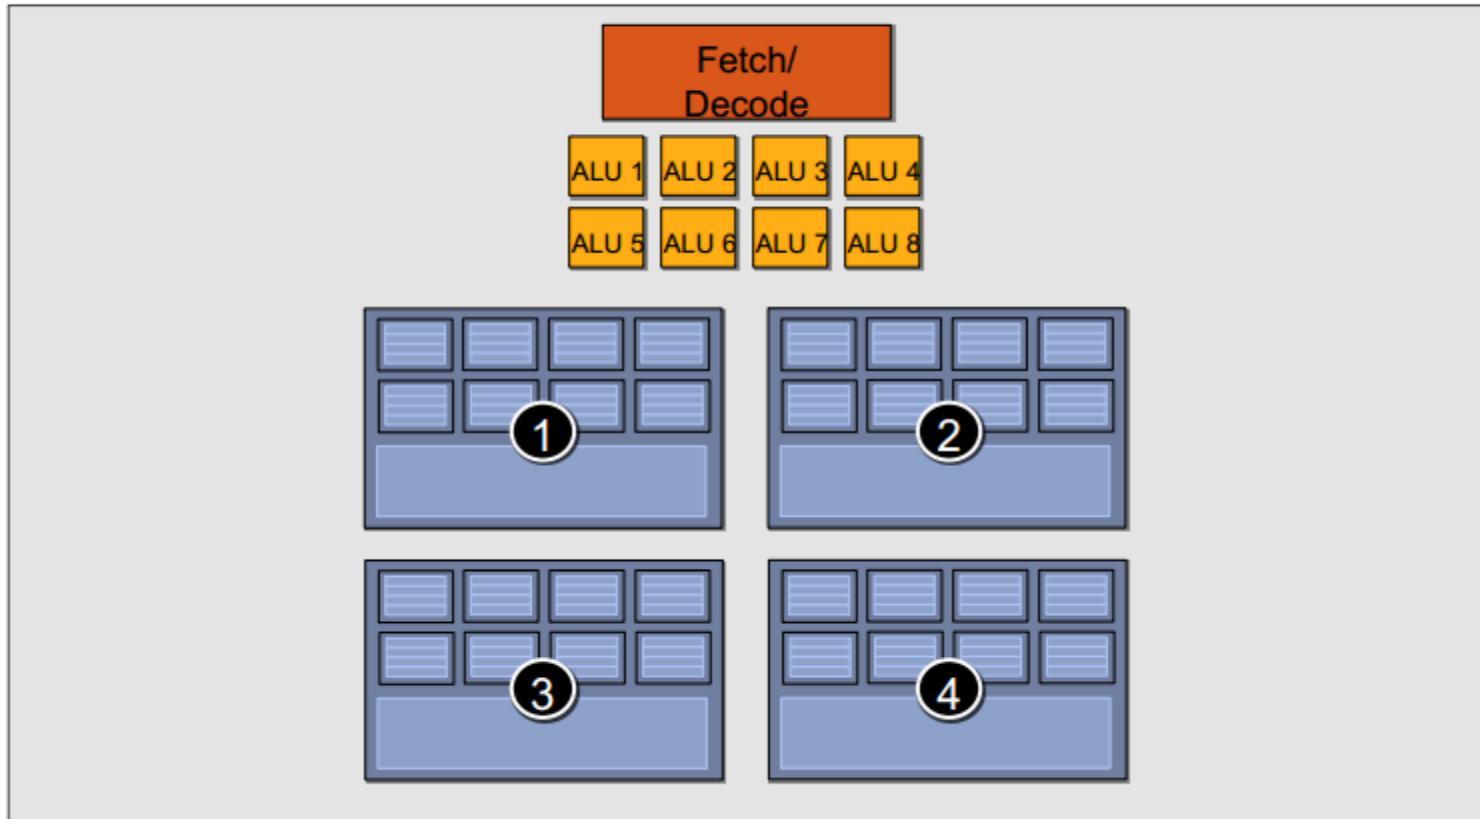


Twelve medium contexts



Four large contexts

(low latency hiding ability)



Clarification

Interleaving between contexts can be managed by hardware or software (or both!)

- NVIDIA / ATI Radeon GPUs
 - HW schedules / manages all contexts (lots of them)
 - Special on-chip storage holds fragment state
- Intel Larrabee
 - HW manages four x86 (big) contexts at fine granularity
 - SW scheduling interleaves many groups of fragments on each HW context
 - L1-L2 cache holds fragment state (as determined by SW)

Example chip

16 cores

8 mul-add ALUs per core
(128 total)

16 simultaneous
instruction streams

64 concurrent (but interleaved)
instruction streams

512 concurrent fragments

= 256 GFLOPs (@ 1GHz)



Summary: three key ideas

1. Use many “slimmed down cores” to run in parallel
2. Pack cores full of ALUs (by sharing instruction stream across groups of fragments)
 - Option 1: Explicit SIMD vector instructions
 - Option 2: Implicit sharing managed by hardware
3. Avoid latency stalls by interleaving execution of many groups of fragments
 - When one group stalls, work on another group

Data Parallelism Summary

- Data Level Parallelism
 - “medium-grained” parallelism between ILP and TLP
 - Still one flow of execution (unlike TLP)
 - Compiler/programmer must explicitly express it (unlike ILP)
- Hardware support: new “wide” instructions (SIMD)
 - Wide registers, perform multiple operations in parallel
- Trends
 - Wider: 64-bit (MMX, 1996), 128-bit (SSE2, 2000), 256-bit (AVX, 2011), 512-bit (Xeon Phi, 2013)
 - More advanced and specialized instructions
- GPUs
 - Embrace data parallelism via “SIMT” execution model
 - Becoming more programmable all the time
- Today’s chips exploit parallelism at **all** levels: ILP, DLP, TLP